

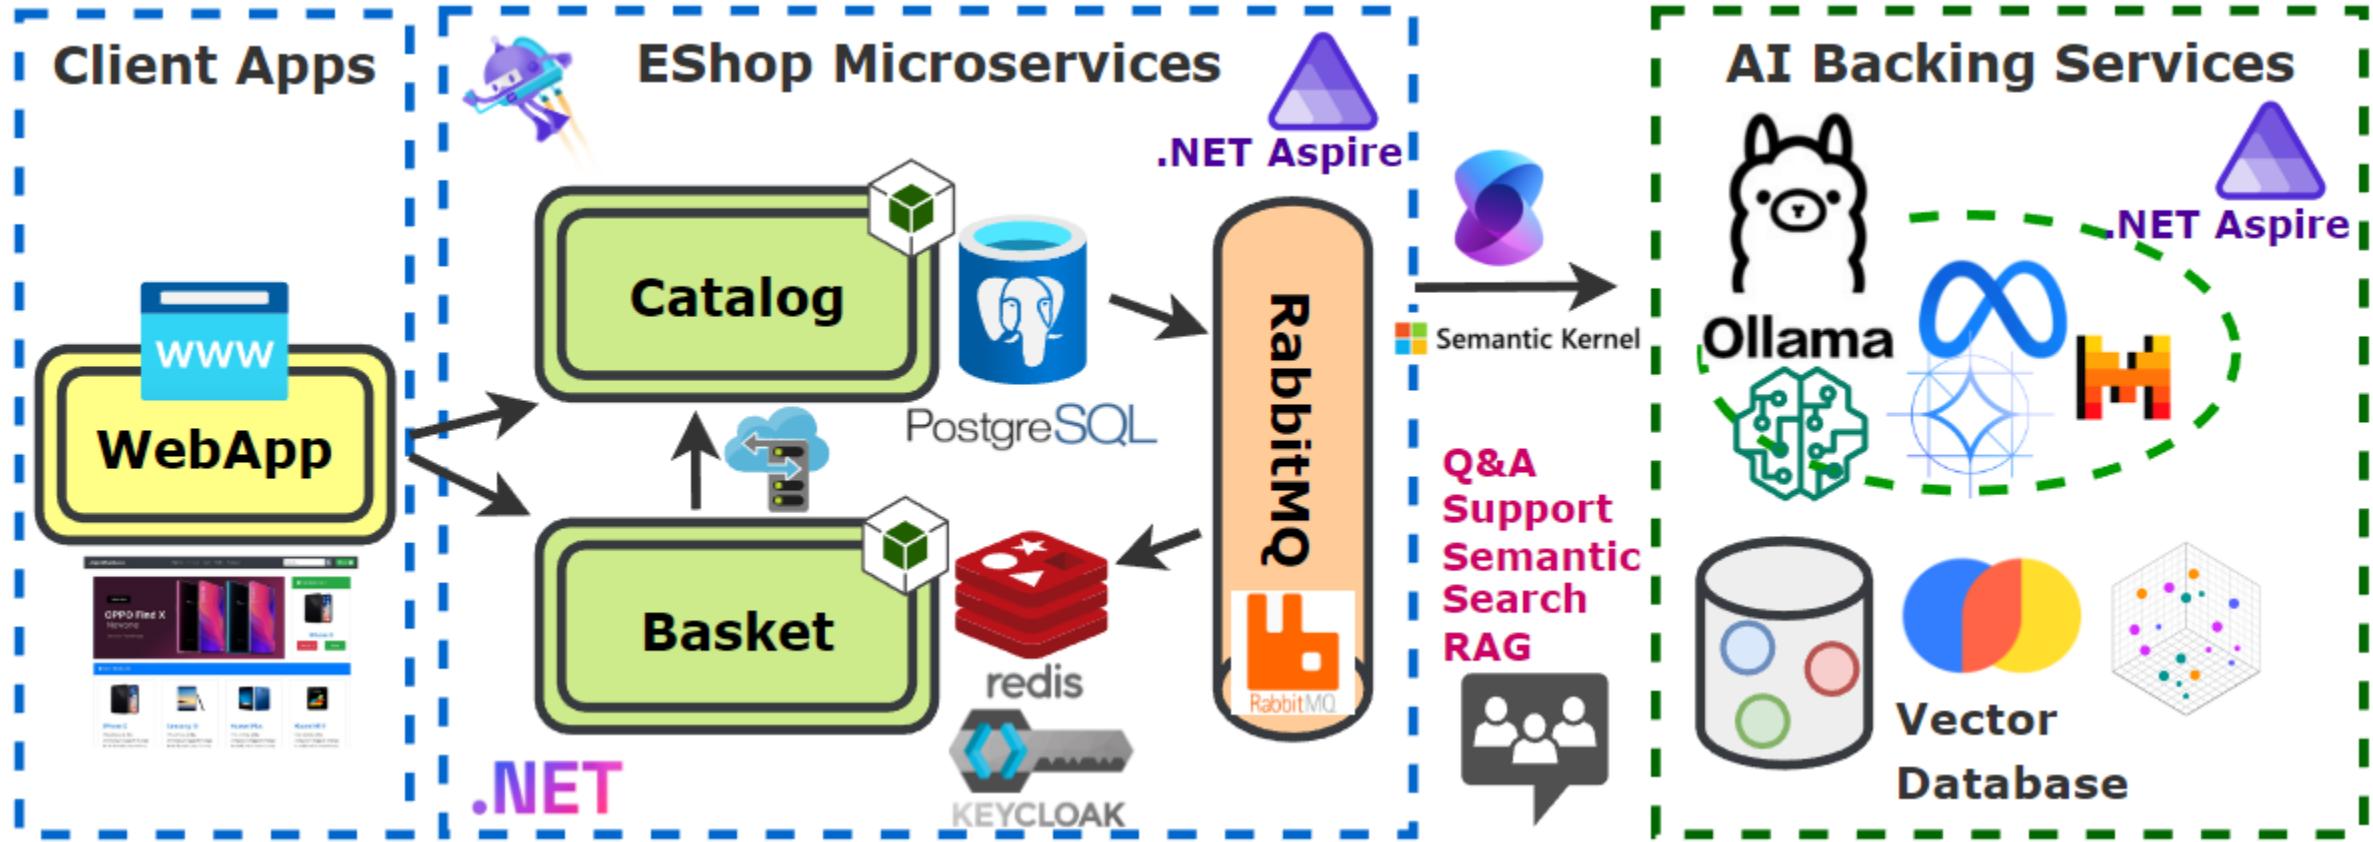
.NET Aspire and GenAI Develop Distributed Architectures

Develop EShop Catalog and Basket microservices integrate with PostgreSQL, Redis, RabbitMQ, Keycloak, Ollama and Semantic Kernel to Create Intelligent E-Shop Solutions

Mehmet Ozkaya



.NET Distributed Architectures w/ Aspire & GenAI



.NET Aspire & GenAI Distributed Architectures

Cloud-Native Distributed Architecture

.NET Aspire Framework

E-Shop Microservices

Deploy to Azure Container Apps

.NET GenAI Development

Use Cases: Customer Support Q&A using Ollama Product Semantic Search leveraging a VectorDB



.NET Aspire

A cloud ready stack for building observable, production ready, distributed applications

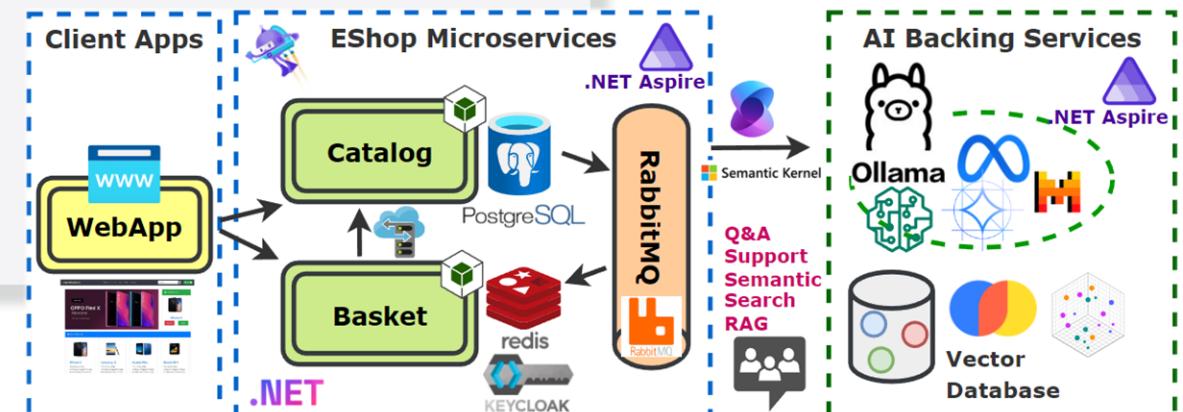
Developer Dashboard

Orchestration

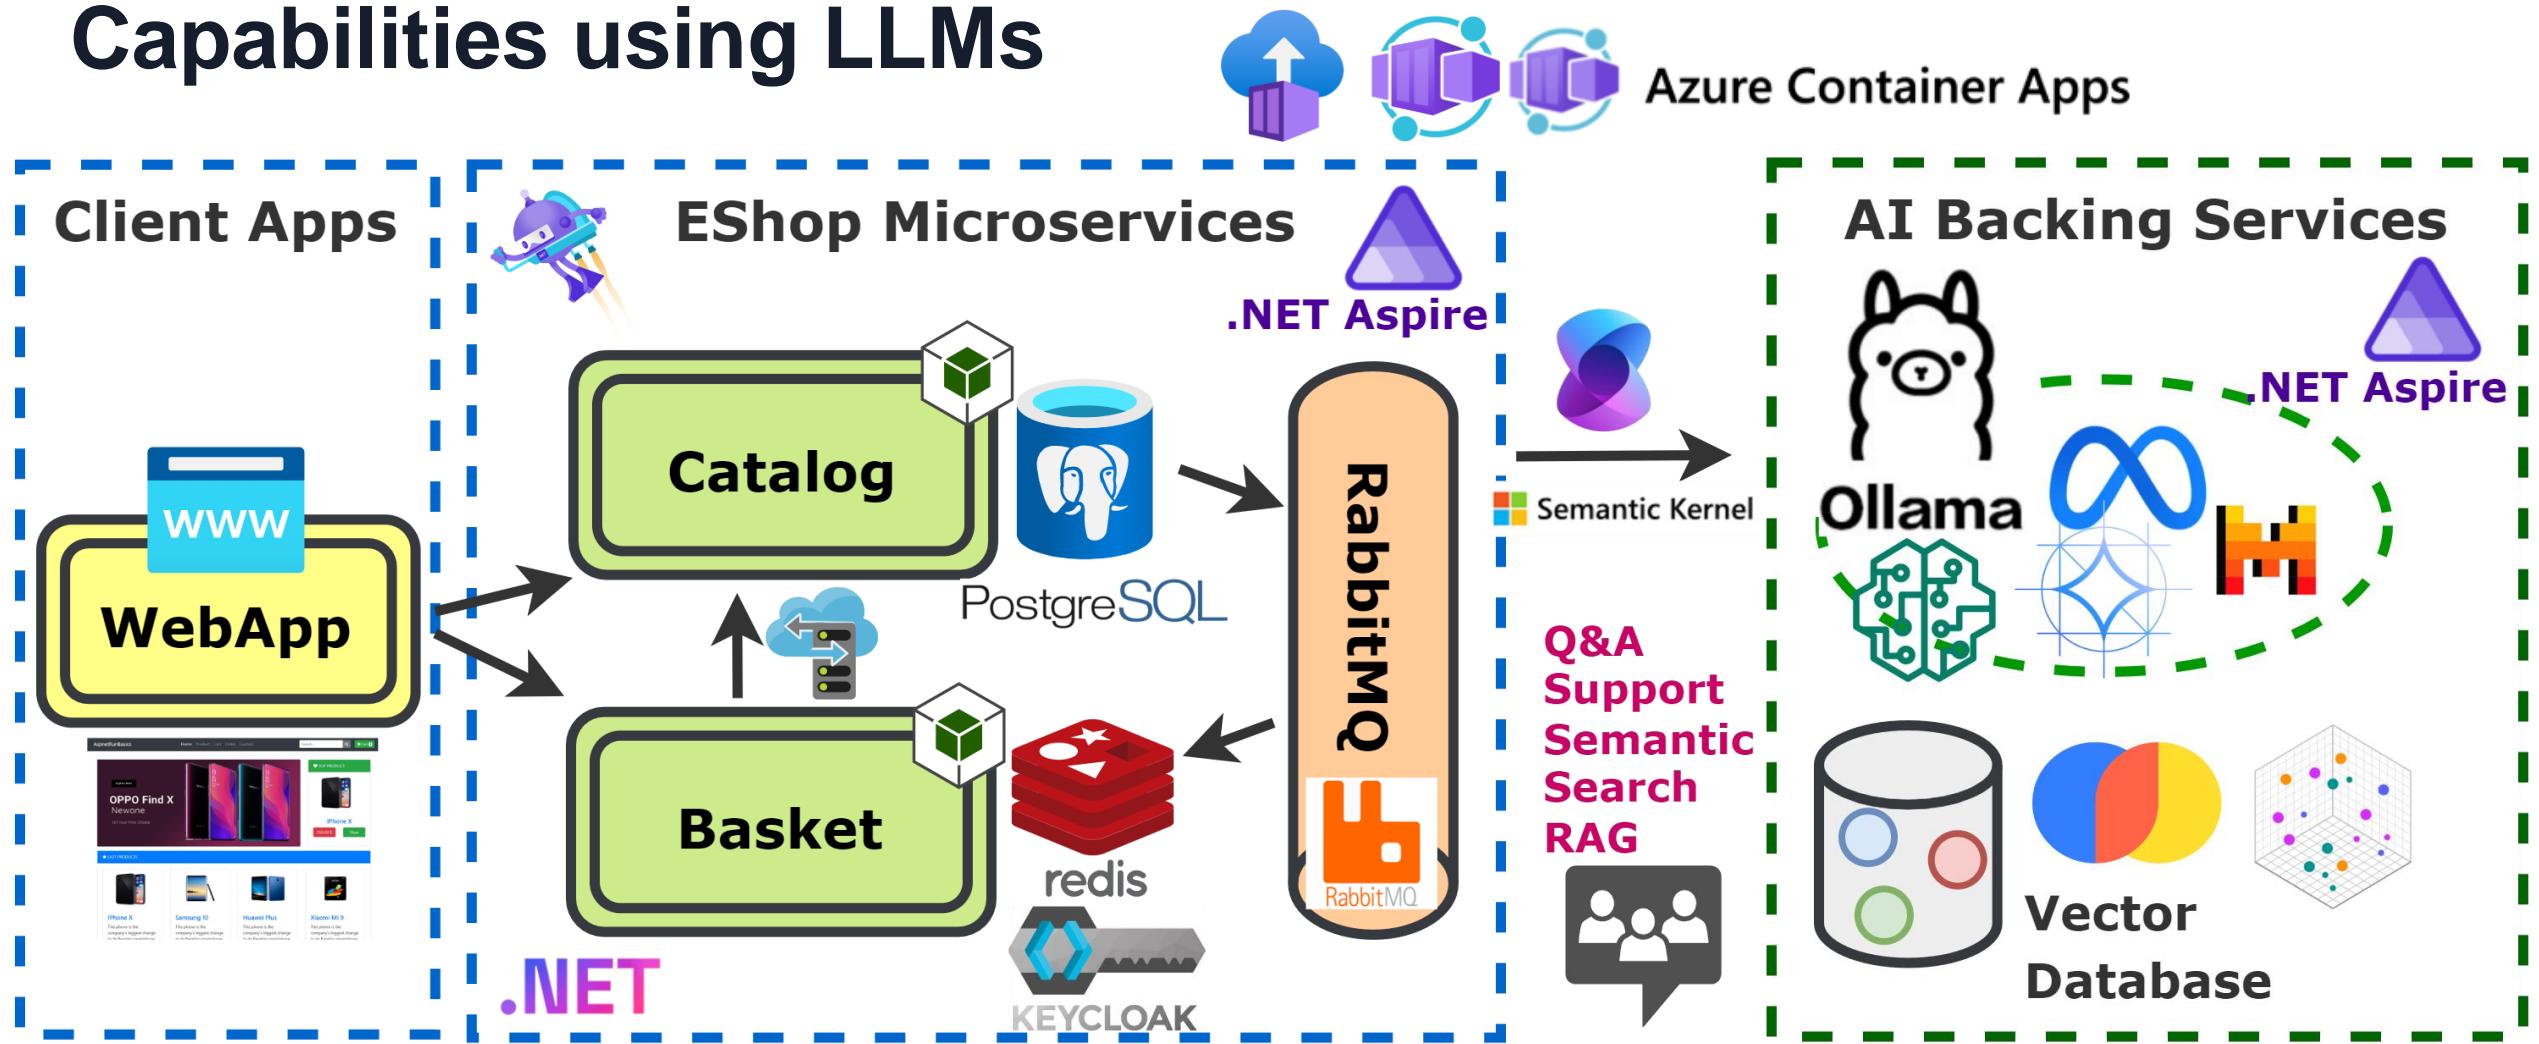
Components

Service Discovery

Deployment



Course Project: EShop Microservices with AI-Powered Capabilities using LLMs



Why Learn AI-Powered Distributed Architectures ?

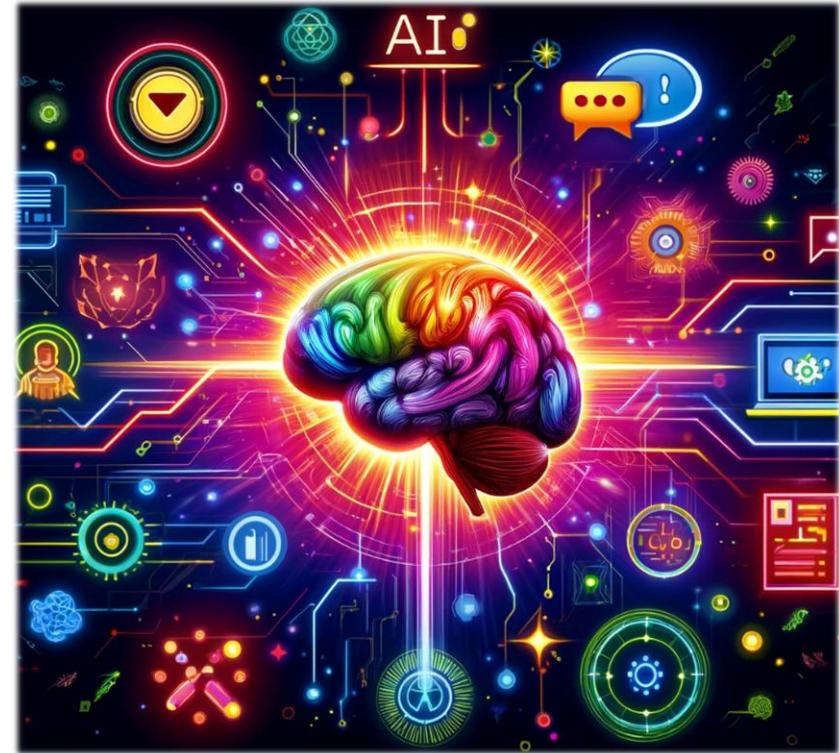
Growing demand for scalable, intelligent applications

Modern apps aren't just about serving static data

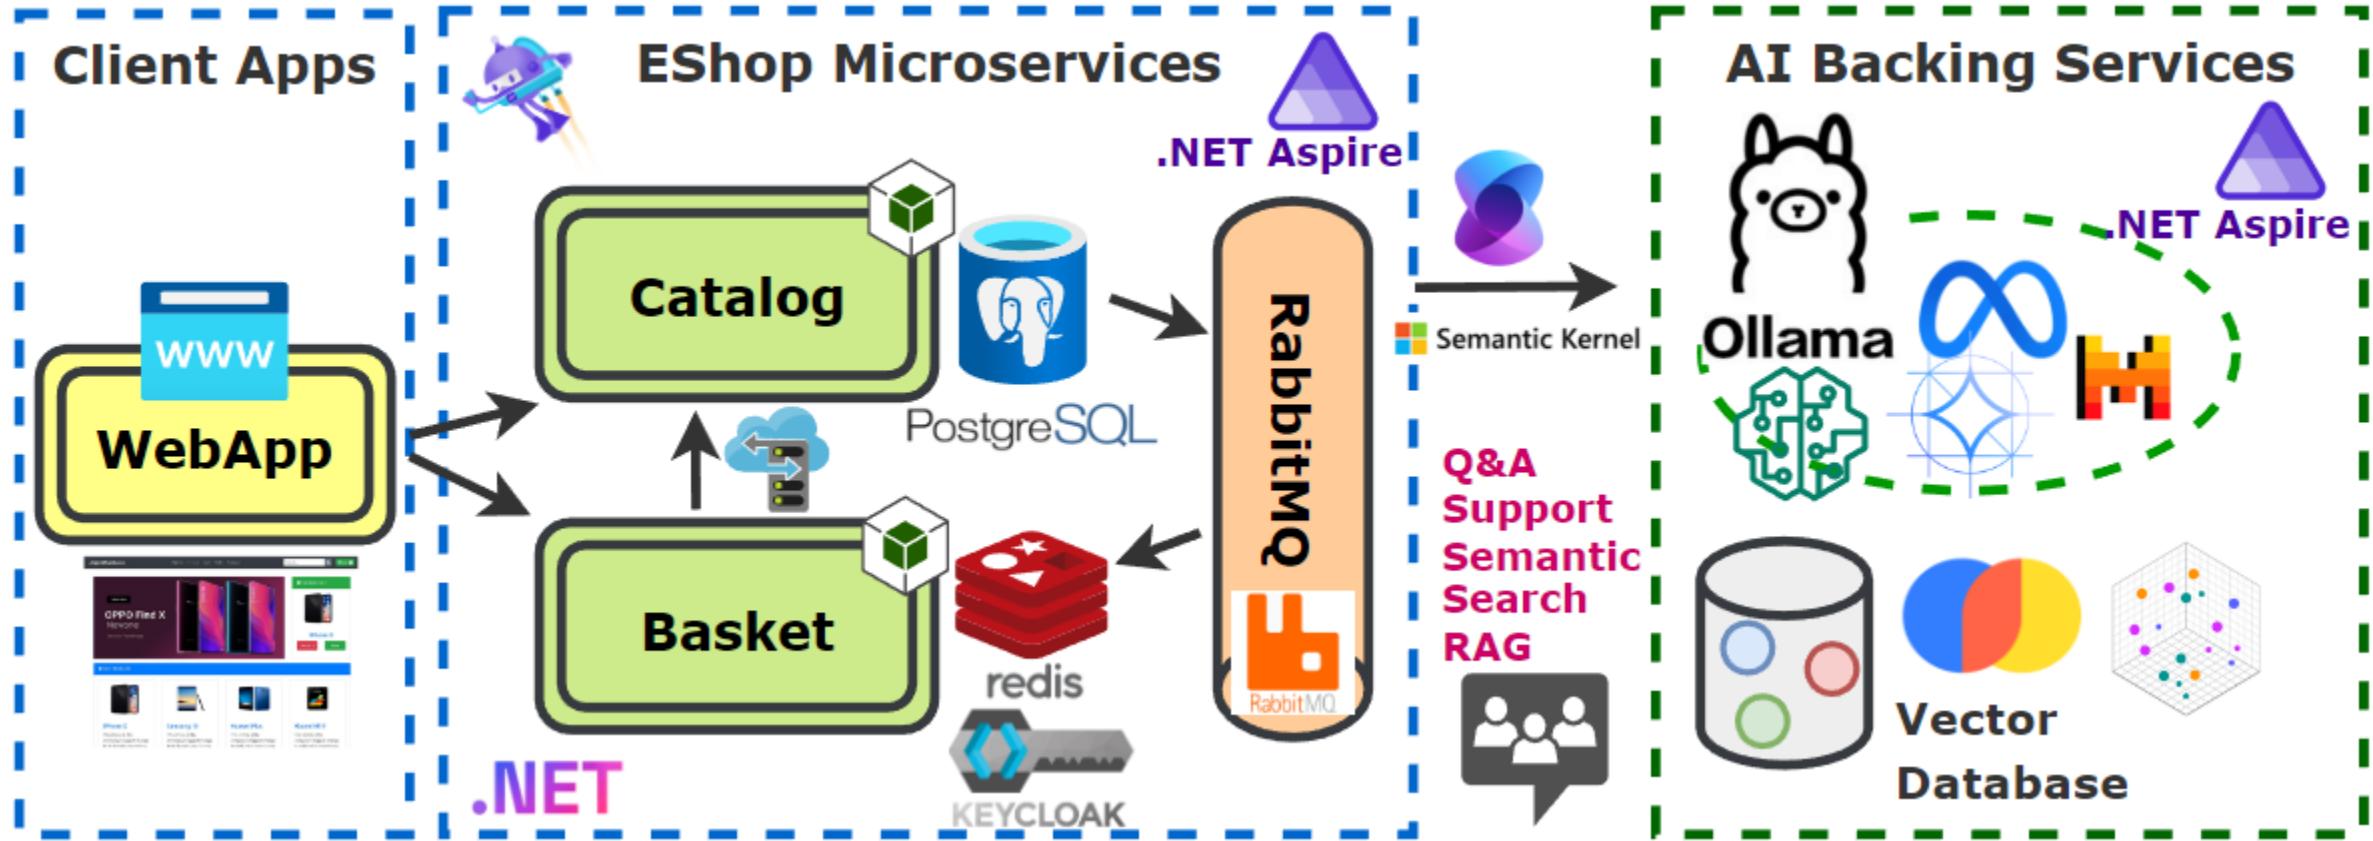
AI is transforming how we interact with software

Combining the power of microservices with AI

The course is very practical, about *95%+* of the lessons will involve you coding



.NET Distributed Architectures w/ Aspire & GenAI



Prerequisites – What You Should Know

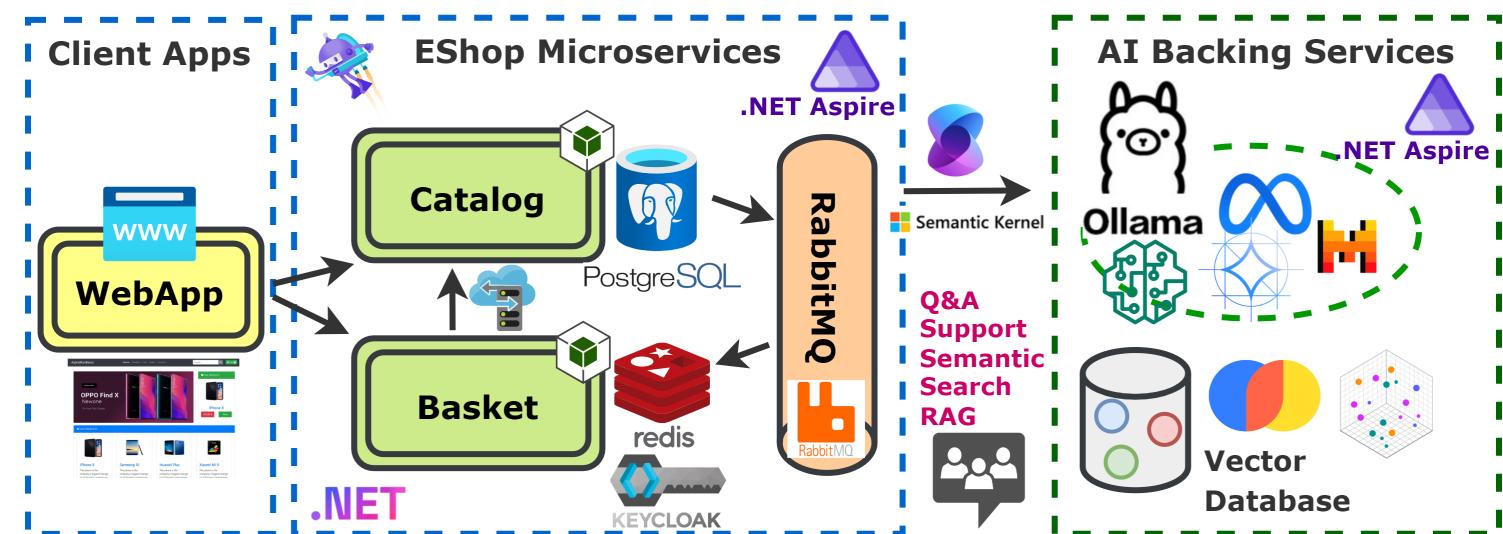
.NET 9 (Includes .NET Aspire by default)



Visual Studio 2022 or Later



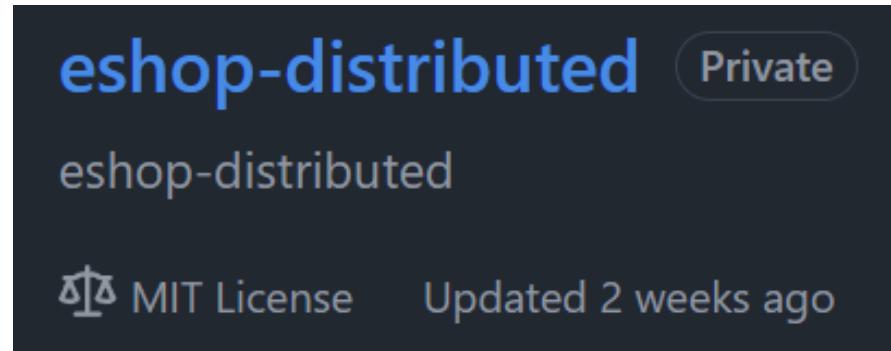
Docker Desktop



Slide Source Code on GitHub

All the code written and demonstrated in this course is available on GitHub

- <https://github.com/mehmetozkaya/eshop-distributed>



Download the source codes section-by-section

- <https://github.com/mehmetozkaya/eshop-distributed-sections>

Course Slides and All Resources

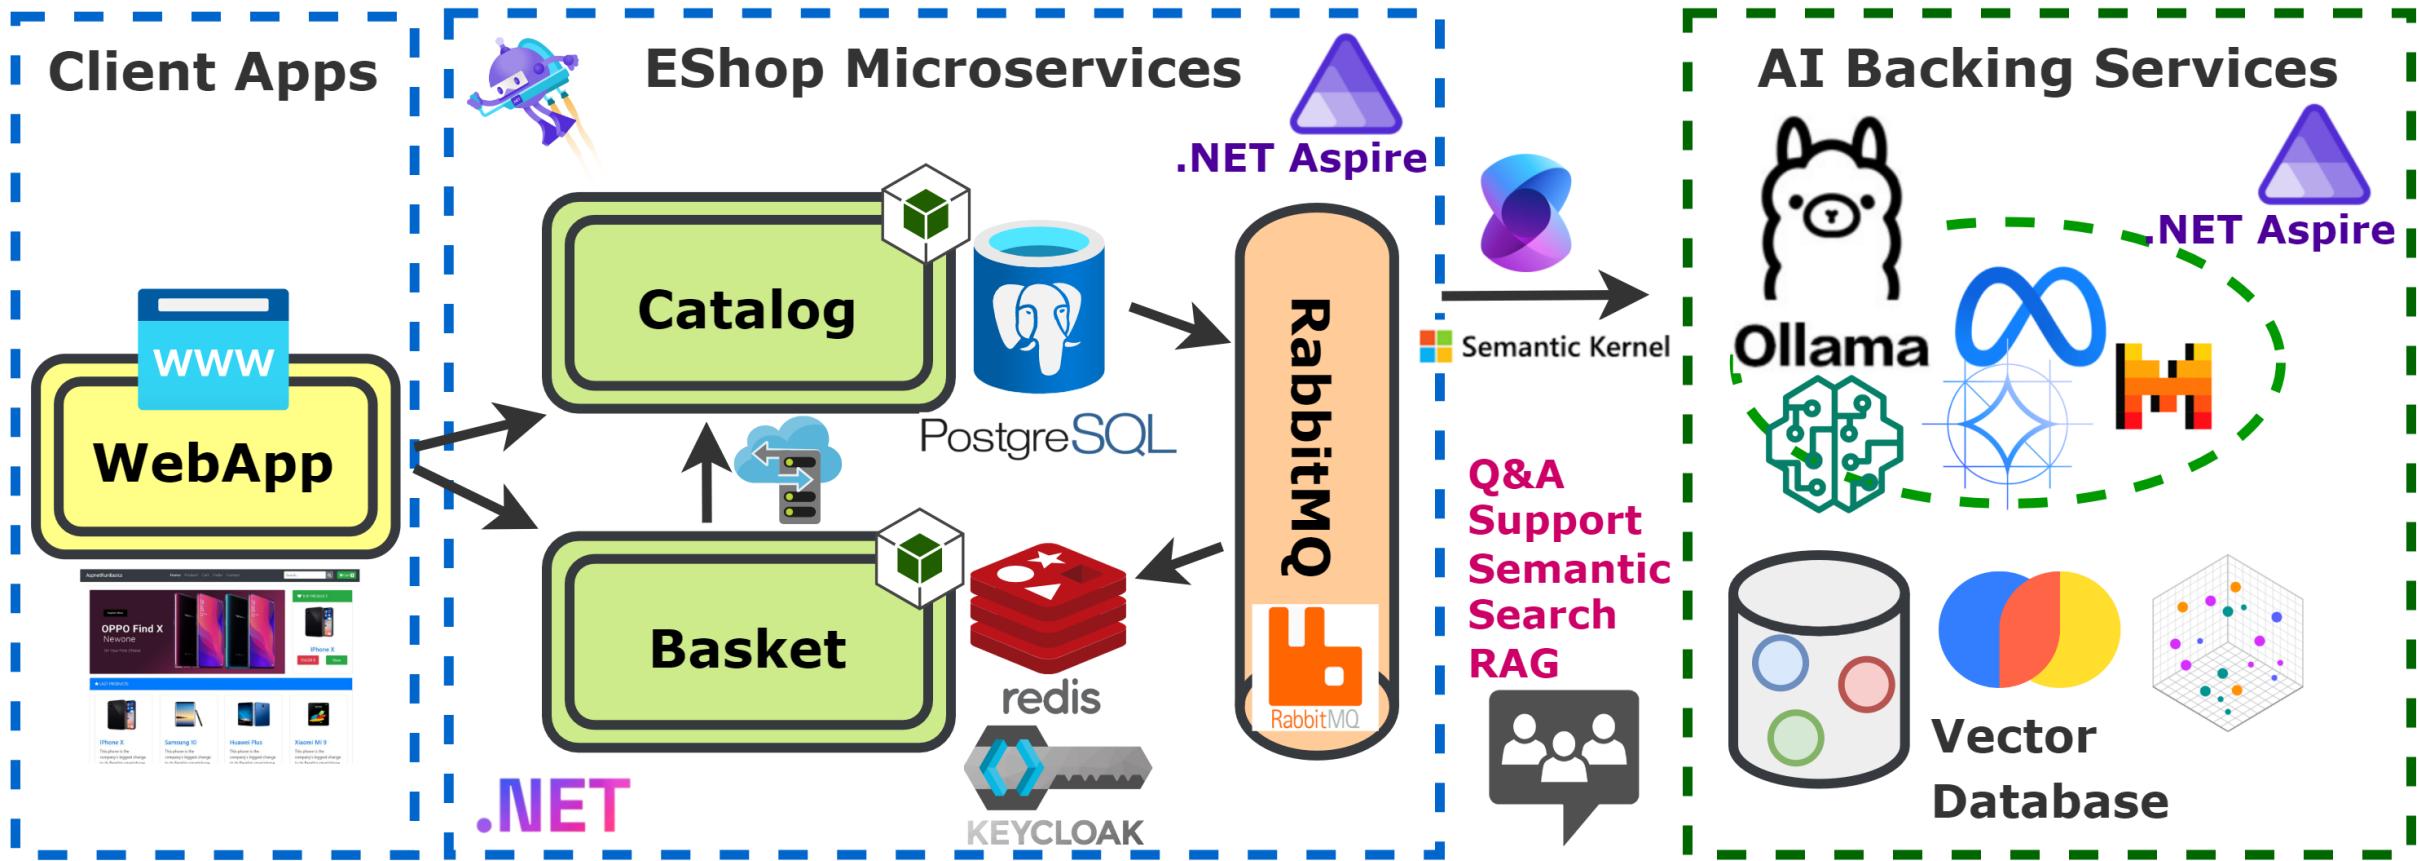
Download Course Slides and All Resources from This Lecture Resources

Course content X

- 23. Variables 9min Resources ▾
- 24. Starting out with Expressions 7min Resources ▾
- 25. Primitive Types 12min Resources ▾
- 26. byte, short, long and width 10min Resources ▾

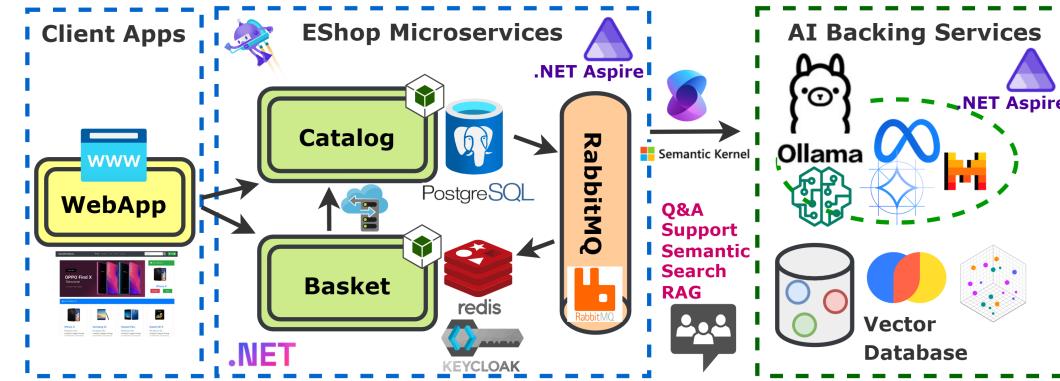


Course Project: EShop Microservices with AI-Powered Capabilities using LLMs



Course Project: EShop Microservices with AI-Powered Capabilities using LLMs

1. Clone the Repository
2. Start Docker Desktop
3. Setting Up in Visual Studio
4. Review .NET Aspire Project and EShop Microservices
5. Review AppHost Project – Cloud-Native Backing Services
6. Run the AppHost Project
7. Review Aspire Dashboard and Resources
8. Blazor Web App Test
9. API Test (Need Keycloak Realm Configurations !!)
10. Messaging and Event-Driven Patterns with RabbitMQ



What is Cloud-Native Distributed Architectures ?

What is Cloud-Native Distributed Architectures ?

Backing Services for Cloud-Native Architectures

Using LLMs and VectorDBs as Cloud-Native Backing Services in
Microservices Architecture

Mehmet Ozkaya



Defining «Cloud-Native»

Apps designed from the ground up to leverage cloud platforms and practices

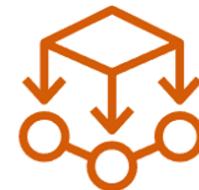
Designed to benefit from **elasticity**, **on-demand resources**, and **automated scaling**

Breaking down large apps into **smaller units** and **packaging them in containers** for consistency

Relying on **DevOps pipelines** for **continuous integration, delivery and deployment**, ensuring **rapid, reliable releases**



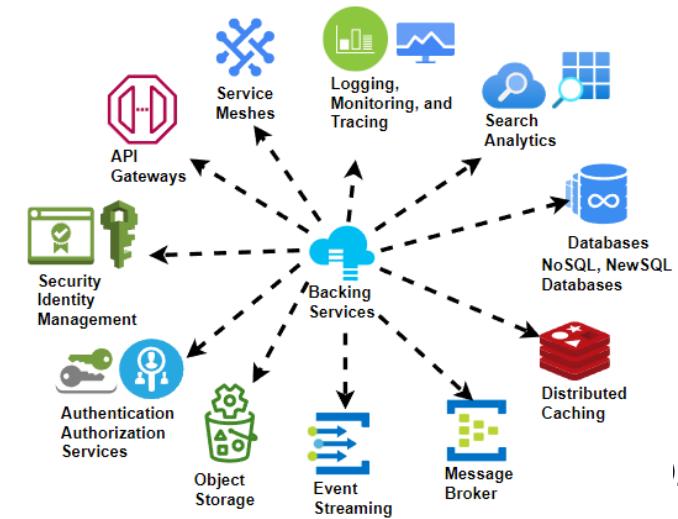
Microservices Containers



Devops



CI/CD



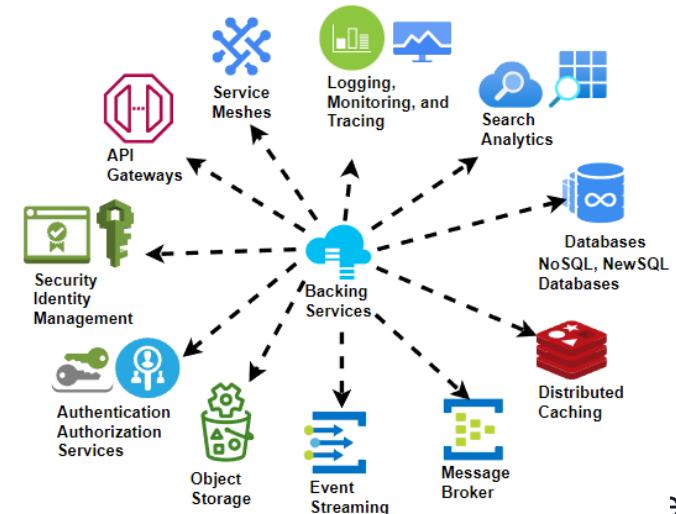
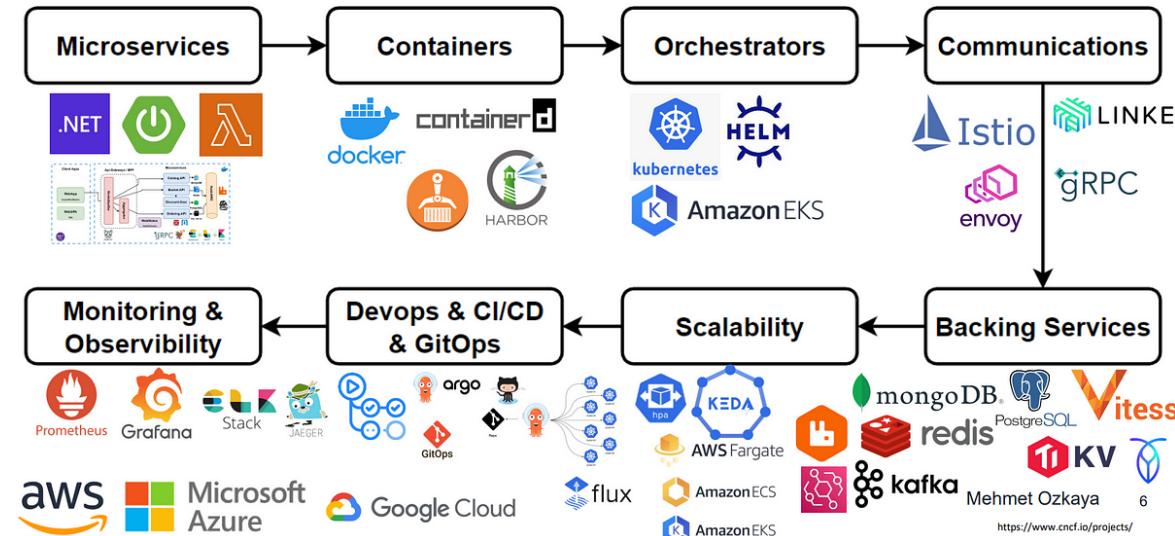
What Makes an Architecture “Distributed”?

Functionality is spread across multiple independent services that communicate over a network

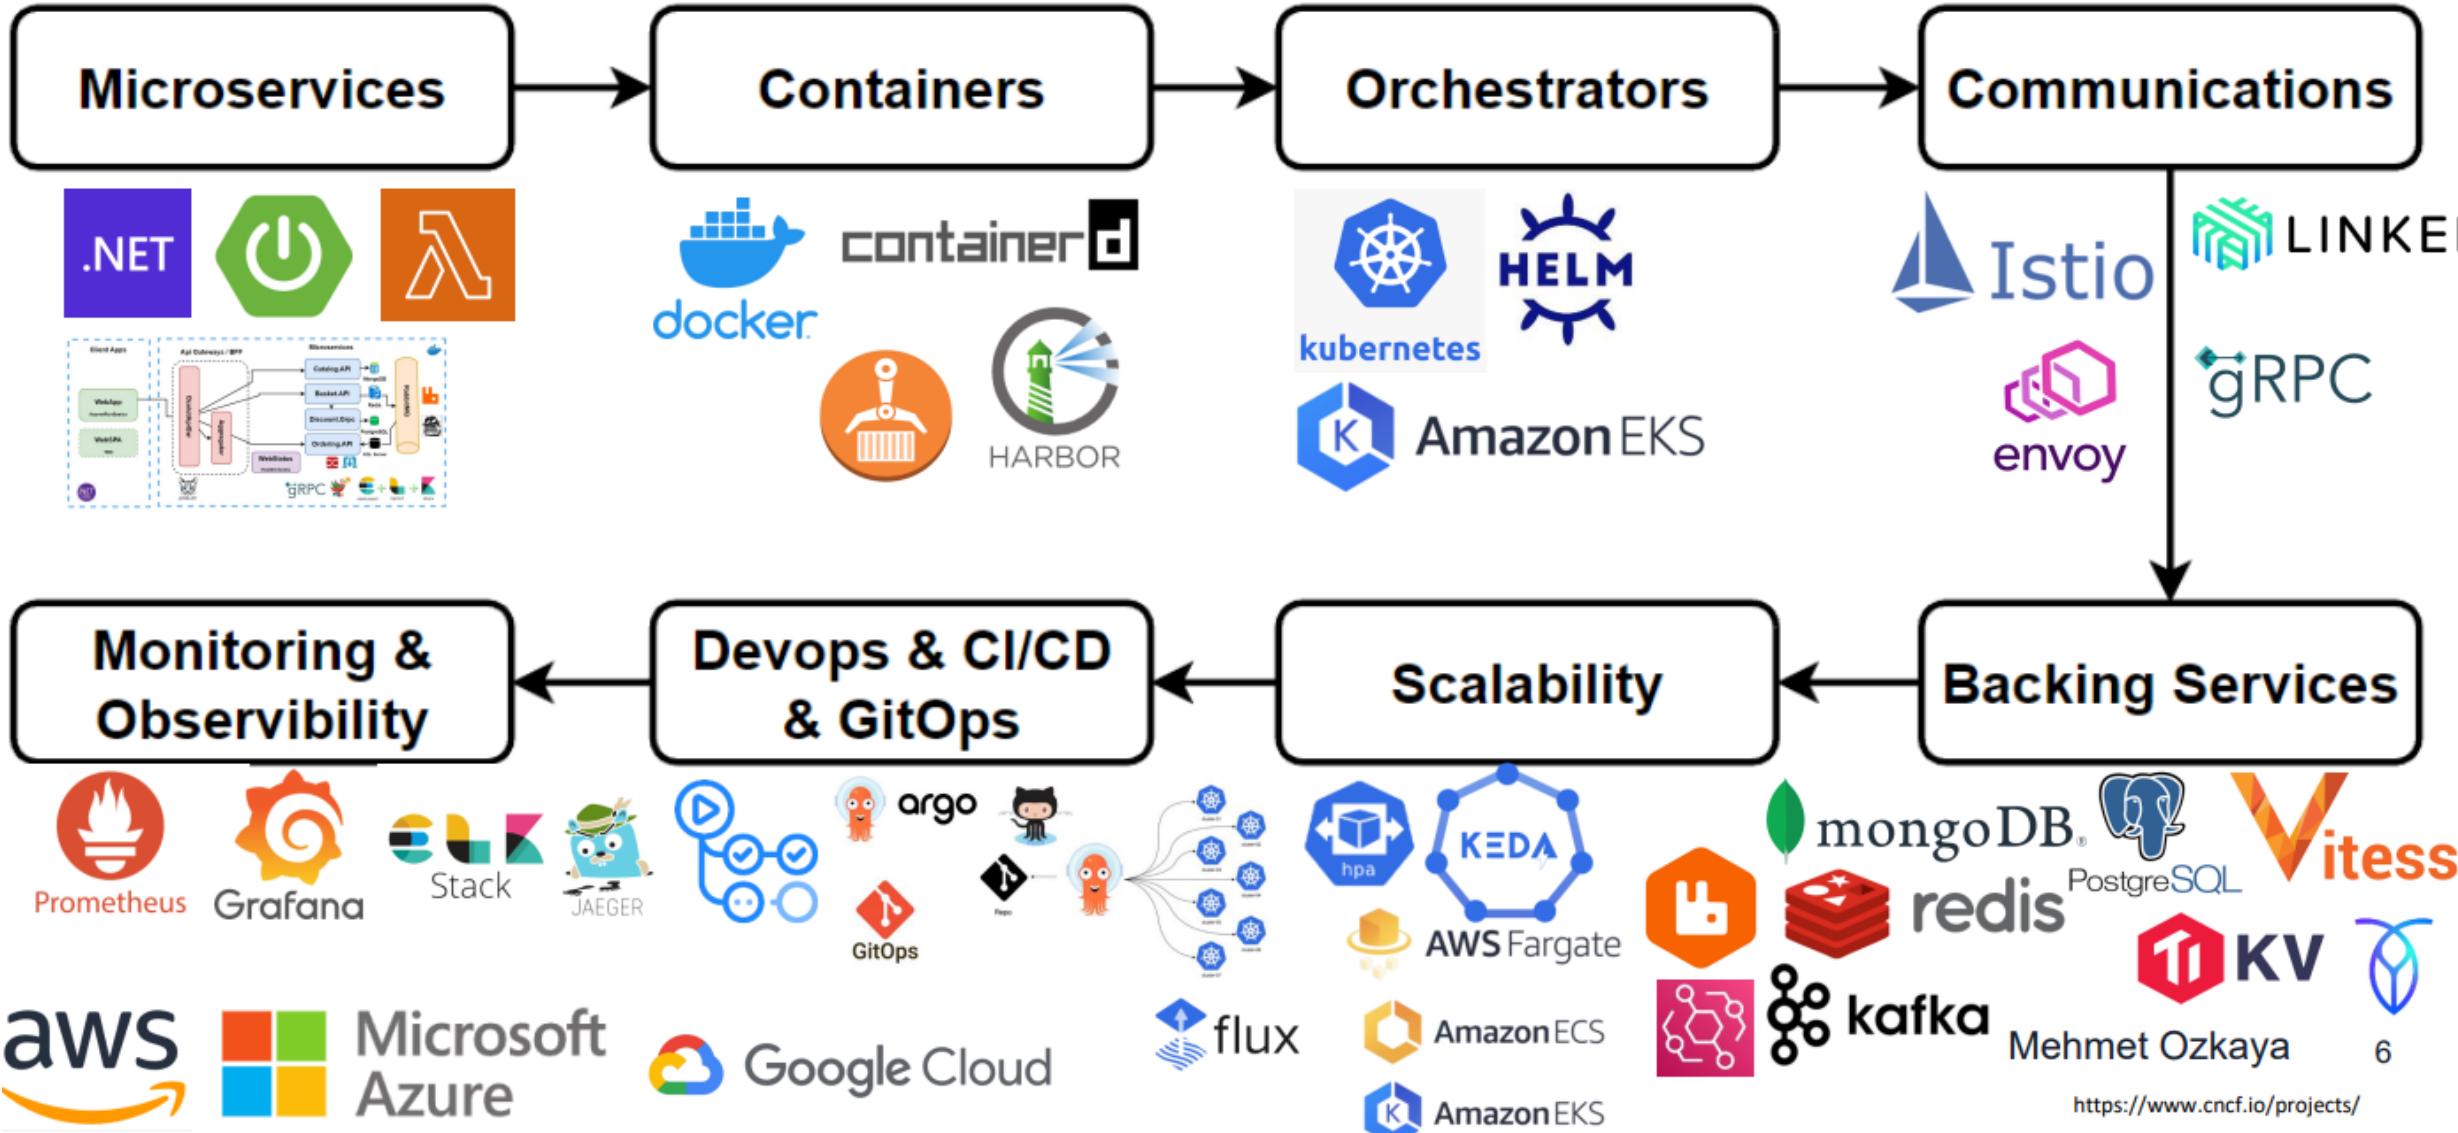
Each service can be developed, deployed and scaled independently

If one service fails, others can continue operating, reducing system-wide downtime

Services handling heavy loads can be scaled up (or down) without affecting others



Key Characteristics of Cloud-Native Distributed Apps



Why Choose Cloud-Native Distributed ?

Scalability & Agility

Scale only the services that need it, add new services when necessary, and quickly adapt to changing demands

Resilience & Fault Tolerance

Failures in one part of the system don't bring down the entire application

Faster Time-to-Market

Independent teams can develop, test, and deploy services in parallel, speeding up innovation

What is Backing Services for Cloud-Native Architectures

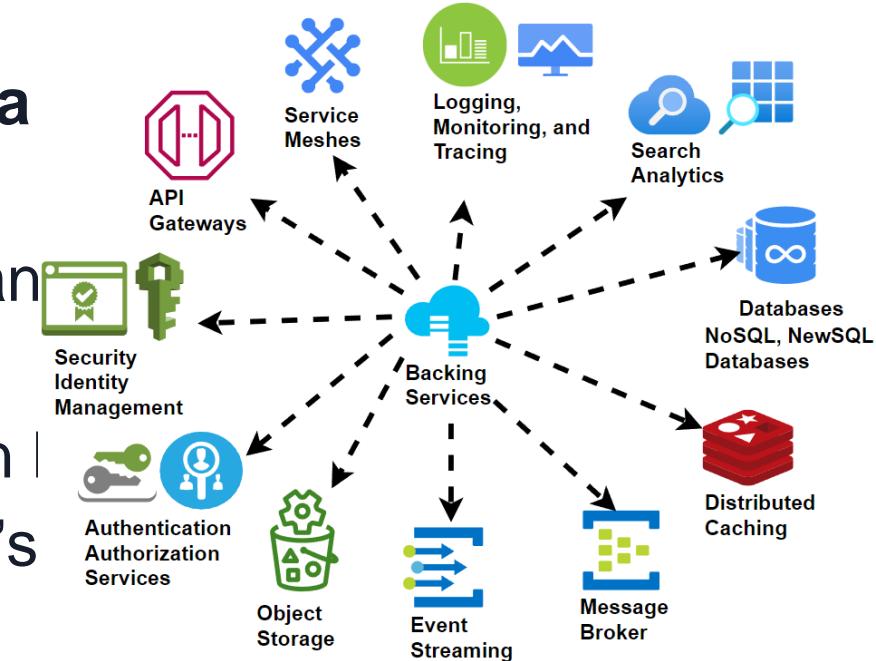
Backing services are the **external components** that applications depend on for their operation

Provide support for **various functionalities**, such as **data storage, messaging, caching and authentication**

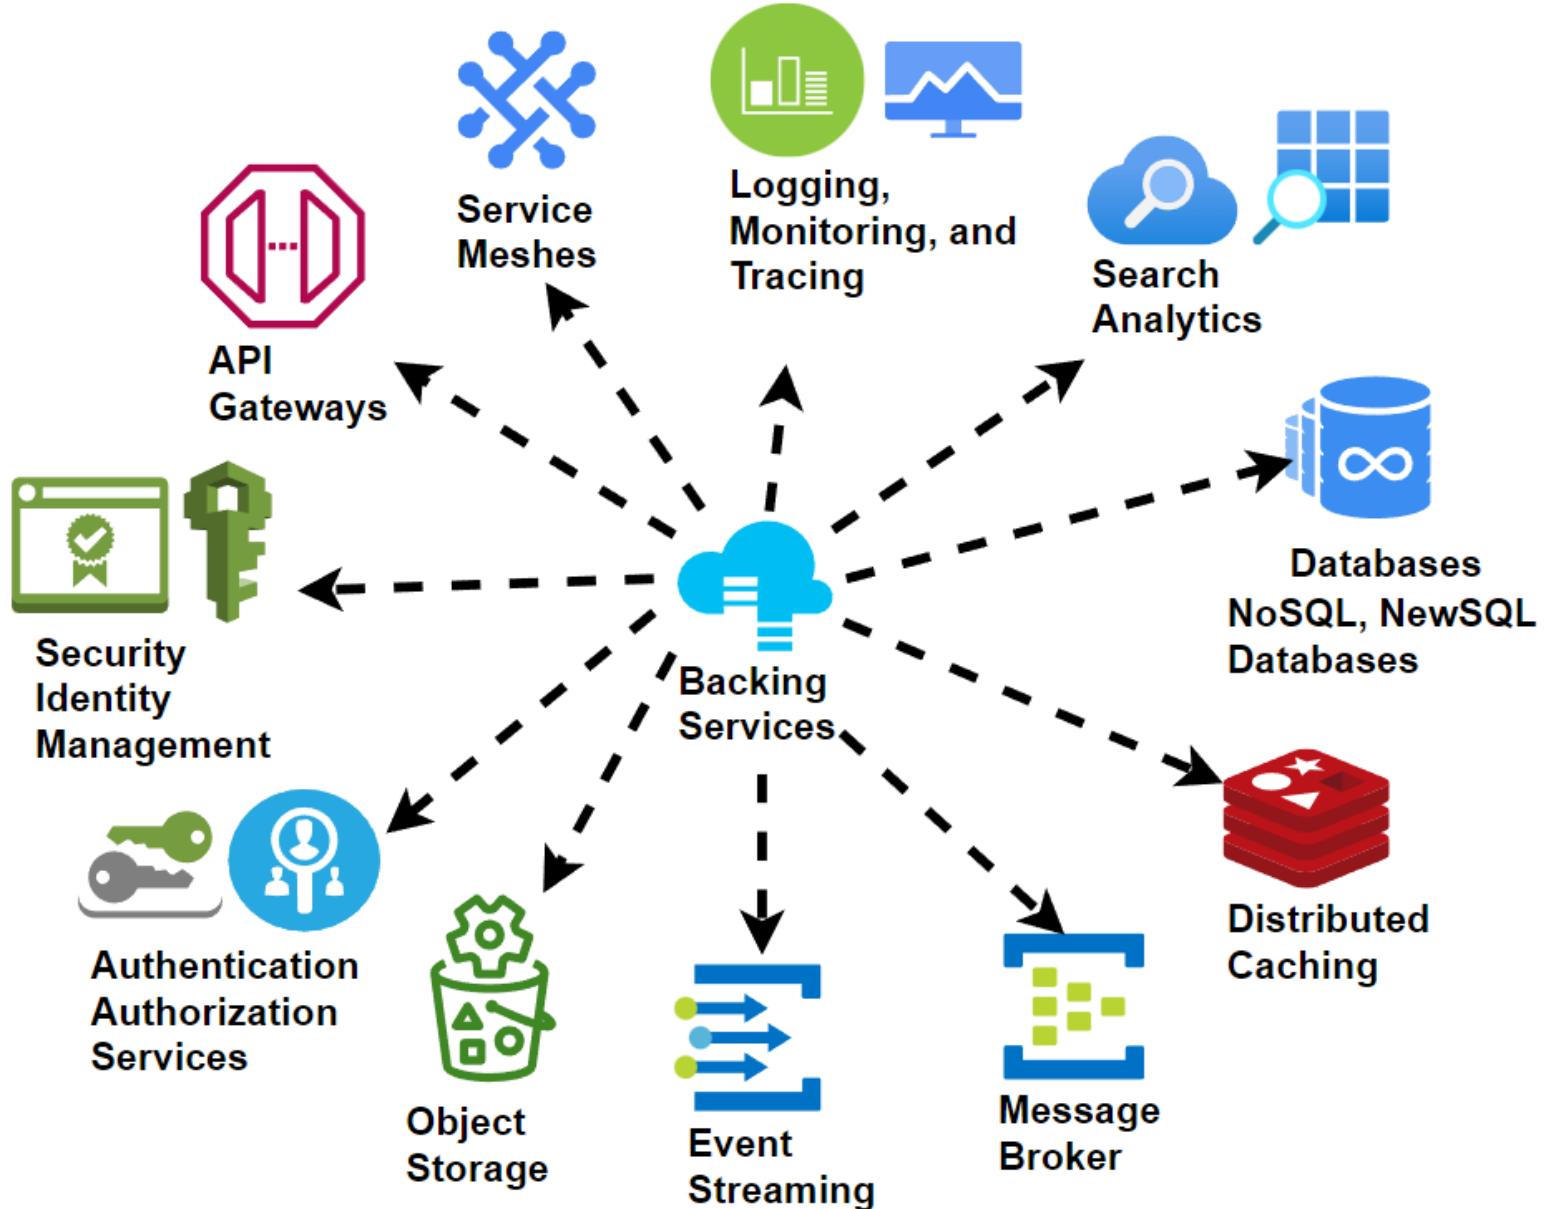
Backing services like databases, messaging systems, and caching services are **treated as attached resources**

Backing services are external to the application and can be swapped or replaced without changing the application's core logic

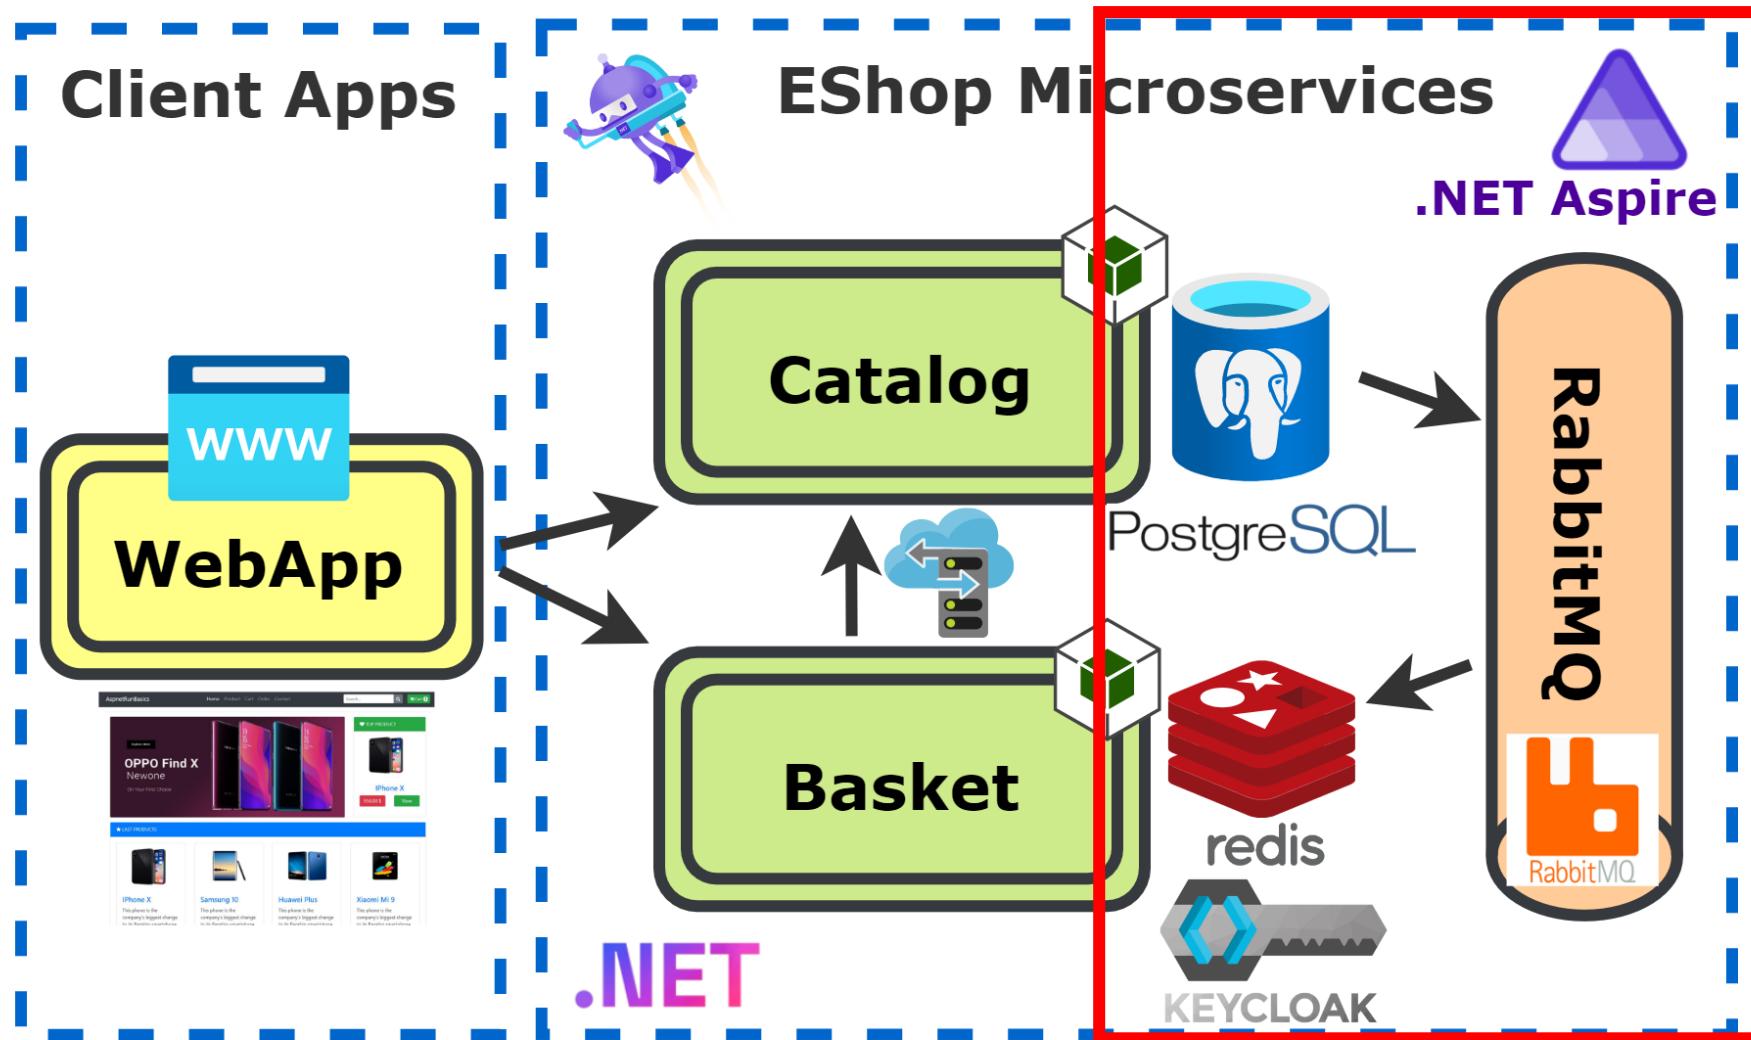
Decoupled from the application themselves, promoting flexibility, scalability, and easier maintenance



Backing Services for Cloud-Native Architectures

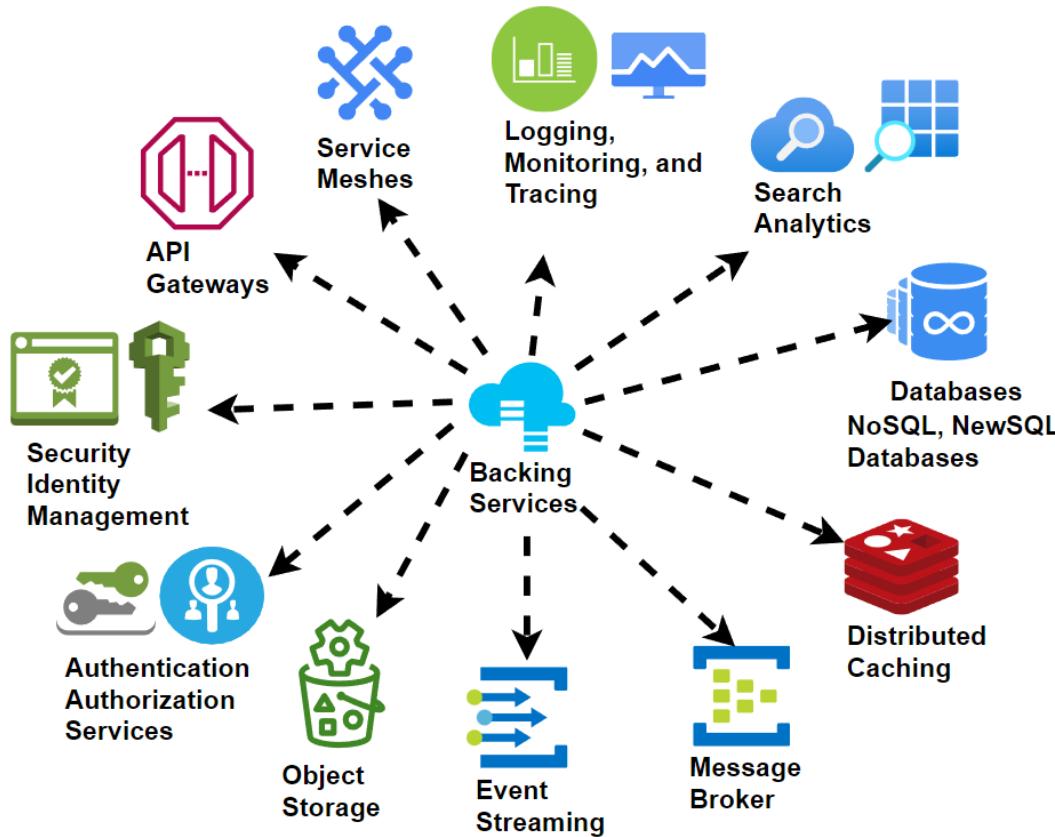


EShop Microservices Backing Services

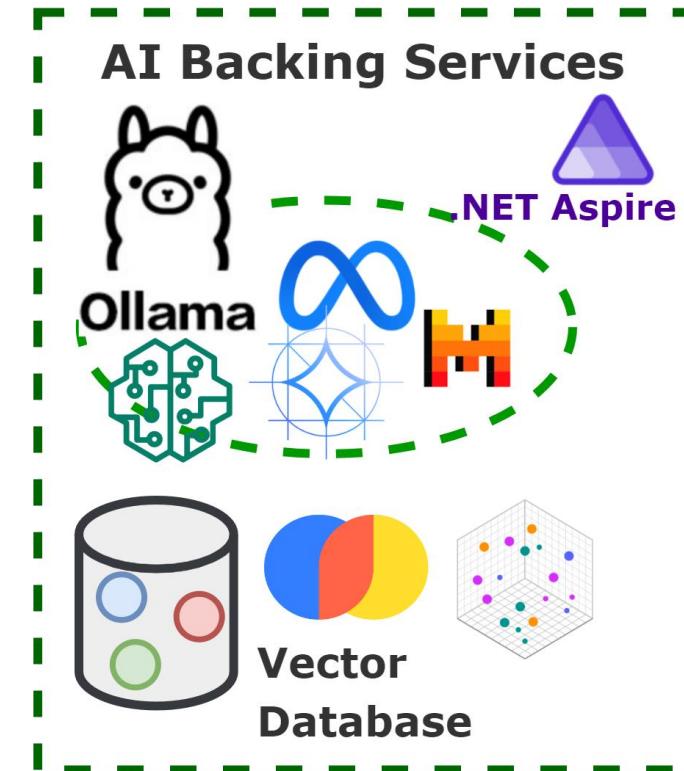


The New Era of Cloud-Native Services

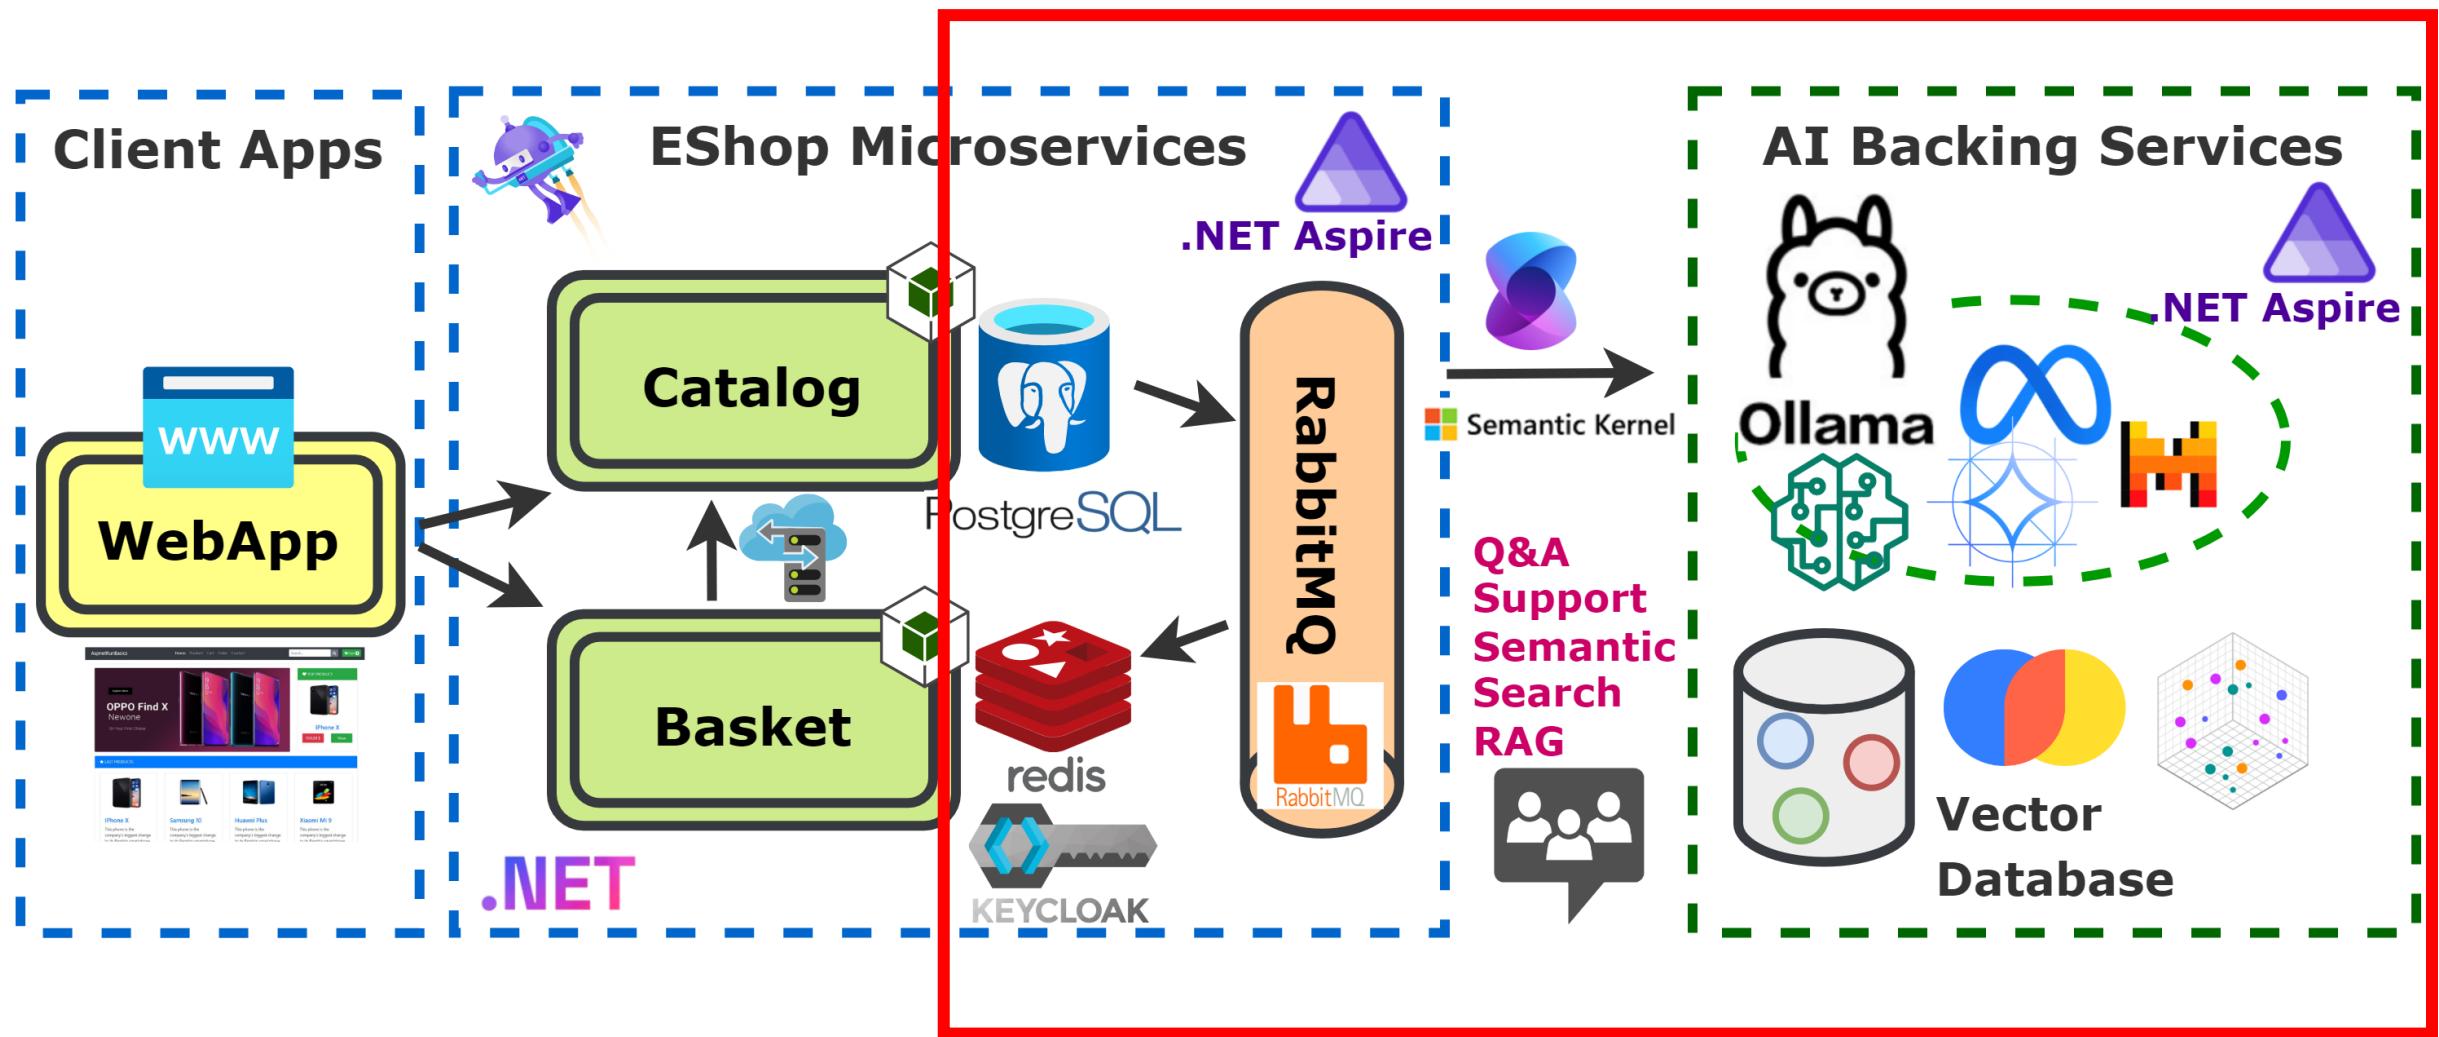
Traditional Backing Services



AI Backing Services



Use LLMs and VectorDBs as Cloud-Native Backing Service



AI as a Backing Service (LLMs and VectorDBs)

LLM Backing Service

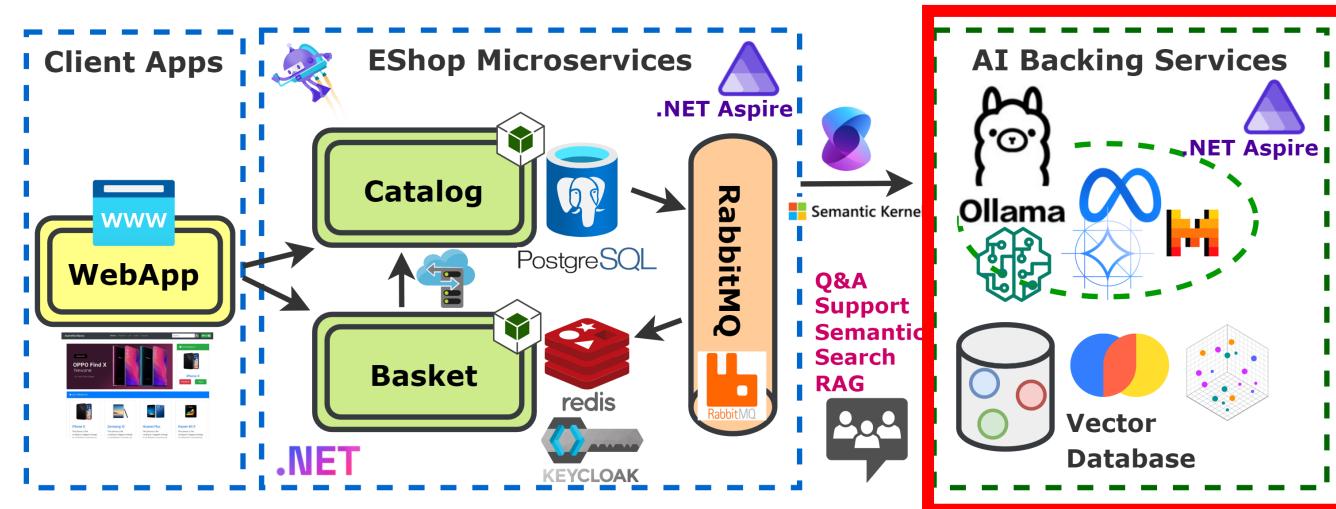
EShop uses Ollama, which enables local inference of Llama or Phi models, Ollama runs entirely on your infrastructure

Vector Databases

Specialized datastore for embeddings used in semantic search and AI workflows

Integration Points

Semantic Kernel and Microsoft.Extensions.AI library provide to connect AI Services from Microservices



.NET Aspire - Distributed App Development Framework

.NET Aspire Core Concepts: Orchestration, Integrations

.NET Aspire Orchestration - Automate Container Management

.NET Aspire Integration - Built-in Connectors

.NET Aspire Service Discovery - Finding Each Other

Mehmet Ozkaya



What is .NET Aspire ?

Cloud-native development framework built on top of the **.NET platform** to **streamline** and **standardize** the **process of creating distributed, microservices-based apps**

Leverages the best practices and features of **.NET**

Adding **scaffolding**, **orchestration** and **preconfigured integrations** for common cloud-backed services

Collection of **templates**, **libraries** that **simplify** the process of **building microservices** and **orchestrating** them in **containerized environments**



Key Features of .NET Aspire

Prebuilt Templates and Scaffolding

Project templates that jump-start your microservice creation



.NET Aspire

Built for Cloud-Native Architectures

Integrates with Docker and container orchestration platforms like ACA or K8s

Infrastructure as Code & Environment Management

Tooling that handle spinning up backing services for local development or test env

Integration with Common Backing Services

Supports popular databases (PostgreSQL, MongoDB), caches (Redis), and message brokers (RabbitMQ, Azure Service Bus)

Modular & Extensible Architecture

Modular structure to integrate AI services, logging platforms, or monitoring solutions

Why .NET Aspire ?

Standard templates and tooling help teams quickly stand up new microservices

Built-in patterns for logging, configuration and container orchestration

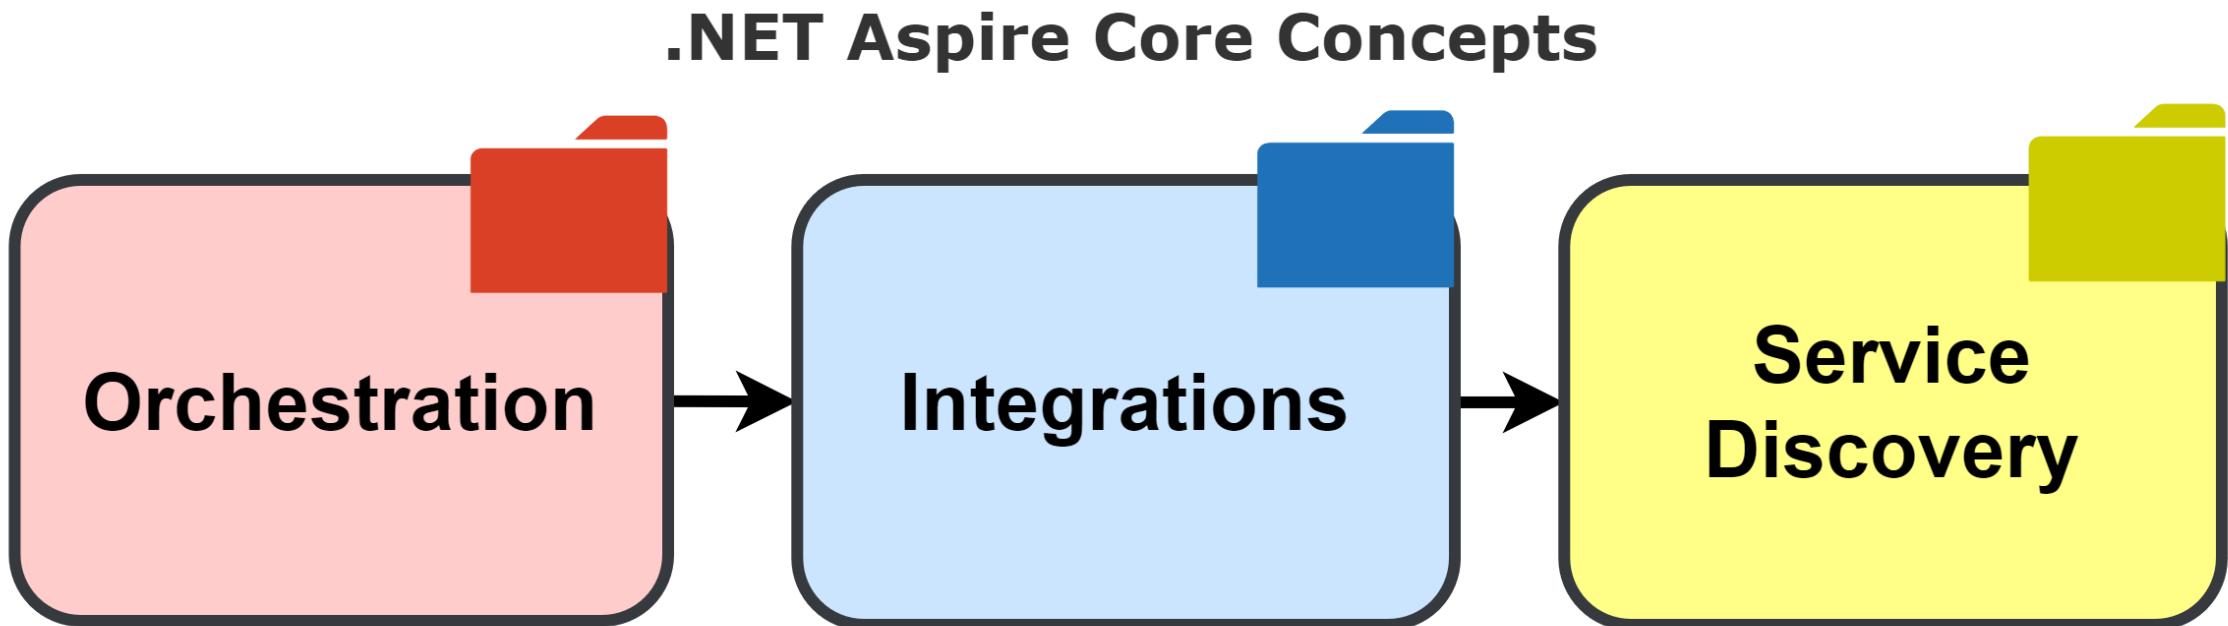
Event-driven communication, containerization, and horizontal scalability from day one

Connects readily to external services (databases, message brokers, identity providers) with **minimal extra plumbing**

Faster development of cloud-native microservices, combining **the power of the .NET platform with proven patterns** and **out-of-the-box integrations**



.NET Aspire Concepts: Orchestration, Integrations and Service Discovery



.NET Aspire



Automating Container Management

Automate container-based deployments without requiring manual Dockerfiles or complex docker-compose setups for each microservice

Environment Awareness

Uses environment configuration (local, test, production) to decide what services to start and how to configure them

Service Dependencies

If a microservice needs to wait for a database to spin up, Aspire can coordinate that

Scalable Infrastructure

Straightforward to scale individual microservices, whether you're using Docker, Kubernetes or ACA, Aspire manages the foundational scripts and tooling

Integrations in .NET Aspire



.NET Aspire

Integrations

Built-in Connectors

Prebuilt integrations with common cloud and on-premises services, such as databases (PostgreSQL, MongoDB), caches (Redis), message brokers (RabbitMQ, Azure Service Bus), and identity providers (Keycloak)

Cross-Cutting Features

Provides default setups for logging, monitoring, and telemetry, saving time and promoting consistent patterns across different microservices

Modular & Extensible

Adding a new service—like AI-based features or a vector database—requires minimal adjustments thanks to Aspire's pluggable design



Dynamic Routing

Service endpoints can change dynamically—for instance, when scaling up or down, keep track of microservice locations, avoiding the need for hard-coded URLs or IP addresses

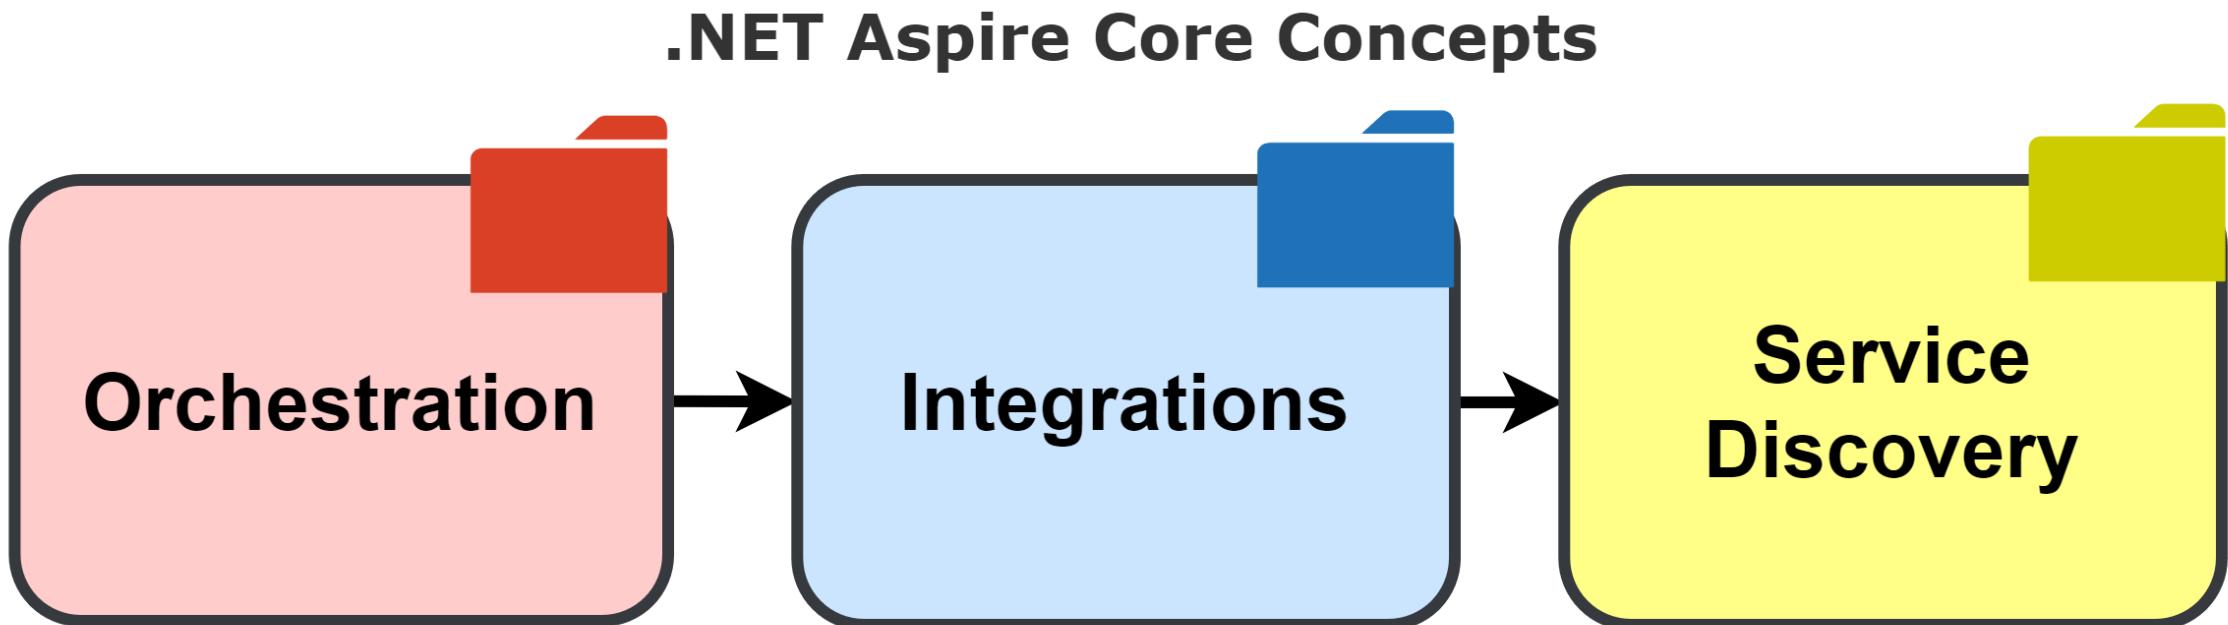
Automatic Registration

When each microservice starts, it registers itself in Aspire's service registry (or a backing service registry), making it discoverable to other services

Resilience & Scaling

If one instance of a service becomes unavailable, calls are routed to another instance automatically, increasing resiliency and ensuring uninterrupted service

.NET Aspire Concepts: Orchestration, Integrations and Service Discovery



.NET Aspire

Terms of Orchestration in .NET Aspire

App Model

The heart of your distributed application

App Host / Orchestrator

The project responsible for running all resources

Resource

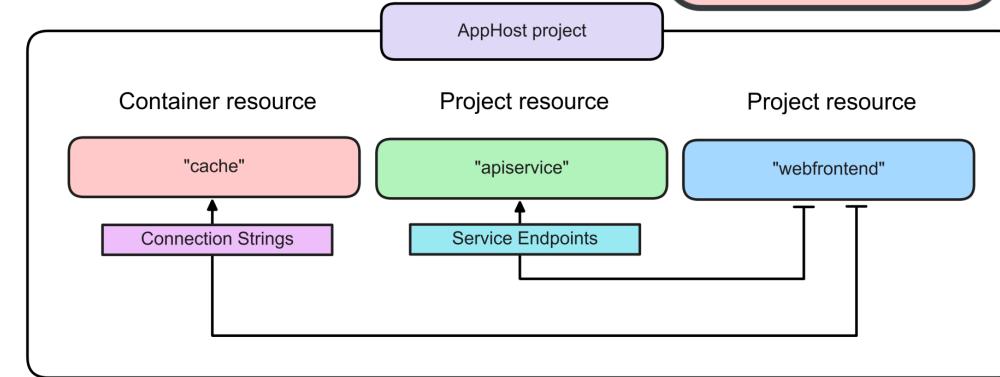
Any piece of your app (databases, microservices, containers)

Integration

NuGet packages that model or configure services

Reference

Declares dependencies between resources



Defining the App Model of Orchestration

App Model

The heart of your distributed application

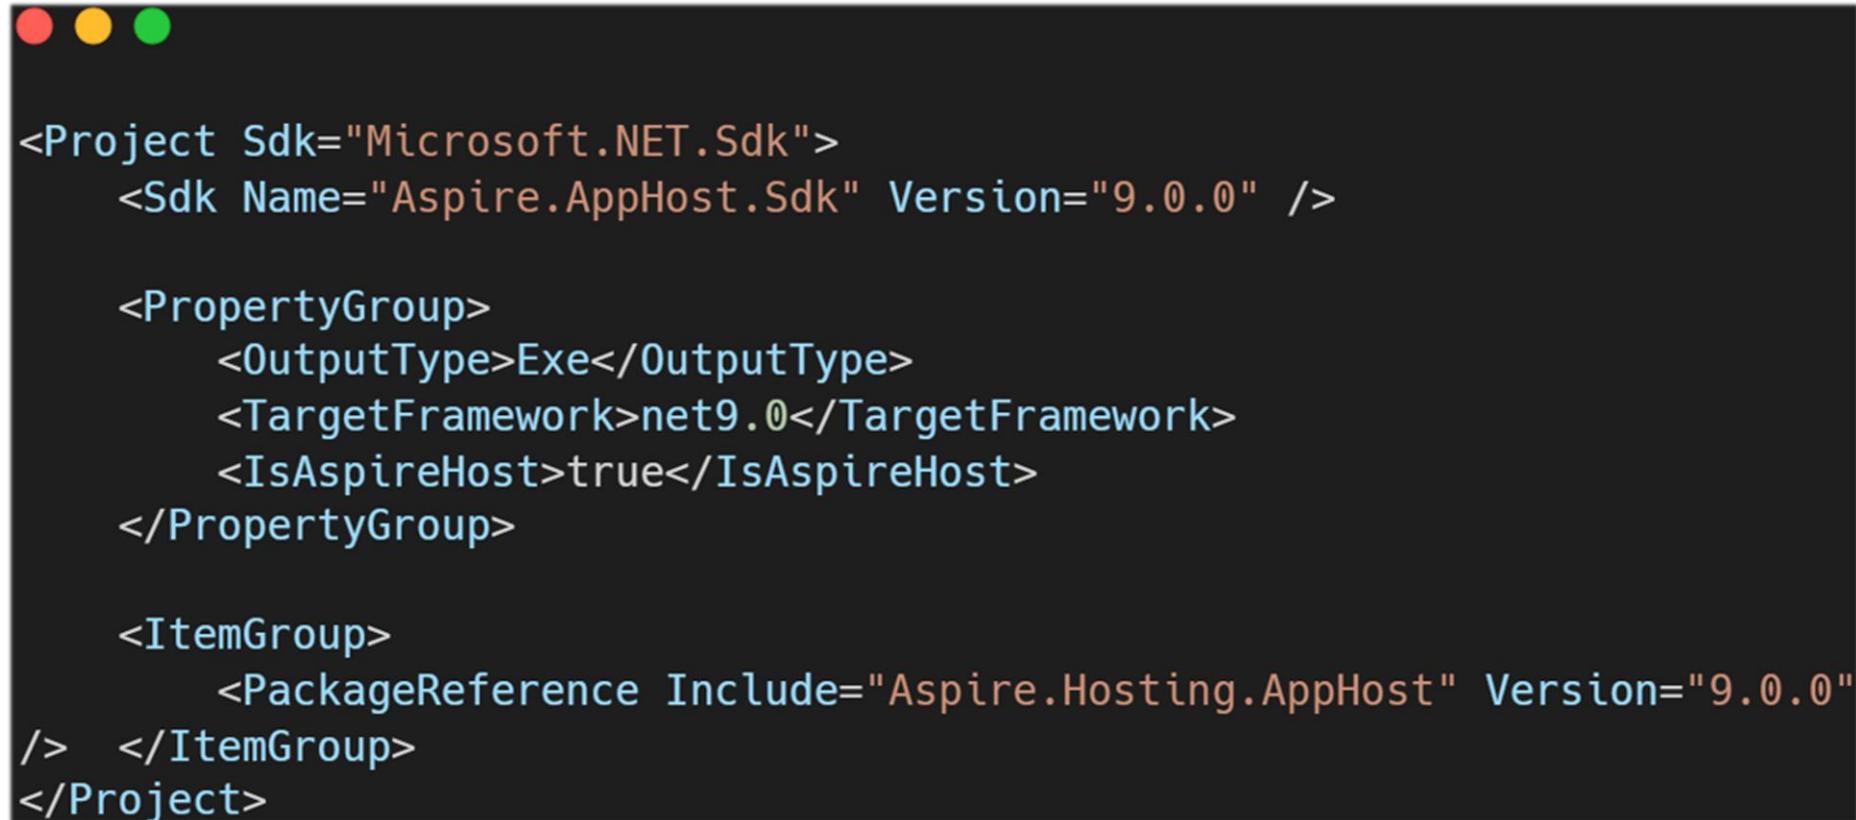


```
var builder =  
    DistributedApplication.CreateBuilder(args);  
    // Add resources  
    // Express dependencies  
  
    builder.Build().Run();
```

App Host / Orchestrator Project

App Host / Orchestrator

The project responsible for running all resources



```
<Project Sdk="Microsoft.NET.Sdk">
  <Sdk Name="Aspire.AppHost.Sdk" Version="9.0.0" />

  <PropertyGroup>
    <OutputType>Exe</OutputType>
    <TargetFramework>net9.0</TargetFramework>
    <IsAspireHost>true</IsAspireHost>
  </PropertyGroup>

  <ItemGroup>
    <PackageReference Include="Aspire.Hosting.AppHost" Version="9.0.0"
  /> </ItemGroup>
</Project>
```

Coding Example – Redis + Two Projects

```
var builder = DistributedApplication.CreateBuilder(args);

var cache = builder.AddRedis("cache");

var apiService = builder.AddProject<Projects.AspireApp_ApiService>("apiservice");
builder.AddProject<Projects.AspireApp_Web>("webfrontend")
    .WithExternalHttpEndpoints()
    .WithReference(cache)
    .WaitFor(cache)
    .WithReference(apiService)
    .WaitFor(apiService);

builder.Build().Run();
```

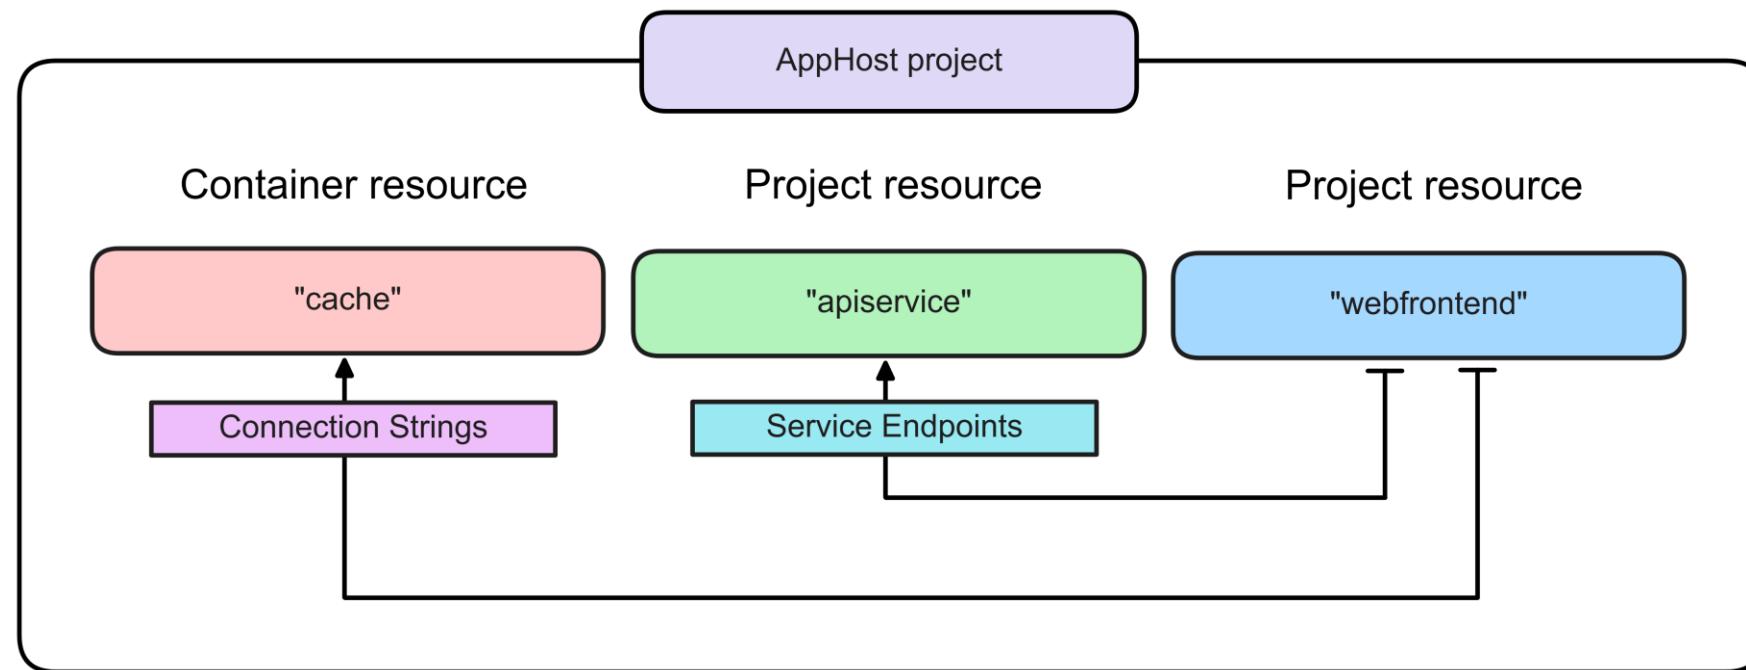
Built-In Resource Types of Orchestration

AddProject – ProjectResource

AddContainer – ContainerResource

AddExecutable – ExecutableResource

AddParameter – ParameterResource



References & Dependencies

```
var builder =  
    DistributedApplication.CreateBuilder(args);  
var cache = builder.AddRedis("cache");  
  
builder.AddProject<Projects.AspireApp_Web>  
    ("webfrontend")  
    .WithReference(cache);
```

Waiting for Resources



```
var postgres = builder.AddPostgres("postgres");
var postgresdb =
postgres.AddDatabase("postgresdb");
builder.AddProject<Projects.AspireApp_ApiService>
("apiservice")
.WithReference(postgresdb)
.WaitFor(postgresdb);
```

Putting It All Together - Orchestration in .NET Aspire

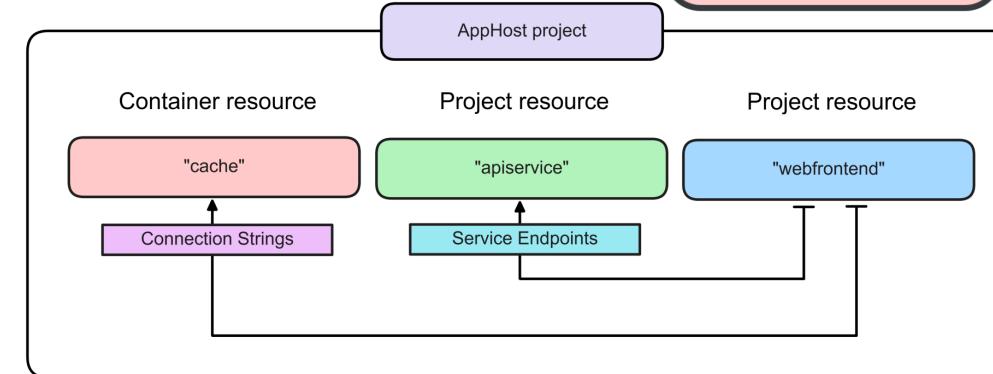


Define your App Model in an App Host project (.AppHost)

Add & Configure Resources—like microservice projects, caches, and databases—using built-in integrations

Reference & WaitFor to express dependencies and correct startup order

Build & Run: .NET Aspire automatically launches the containers, orchestrates networking, and injects environment variables



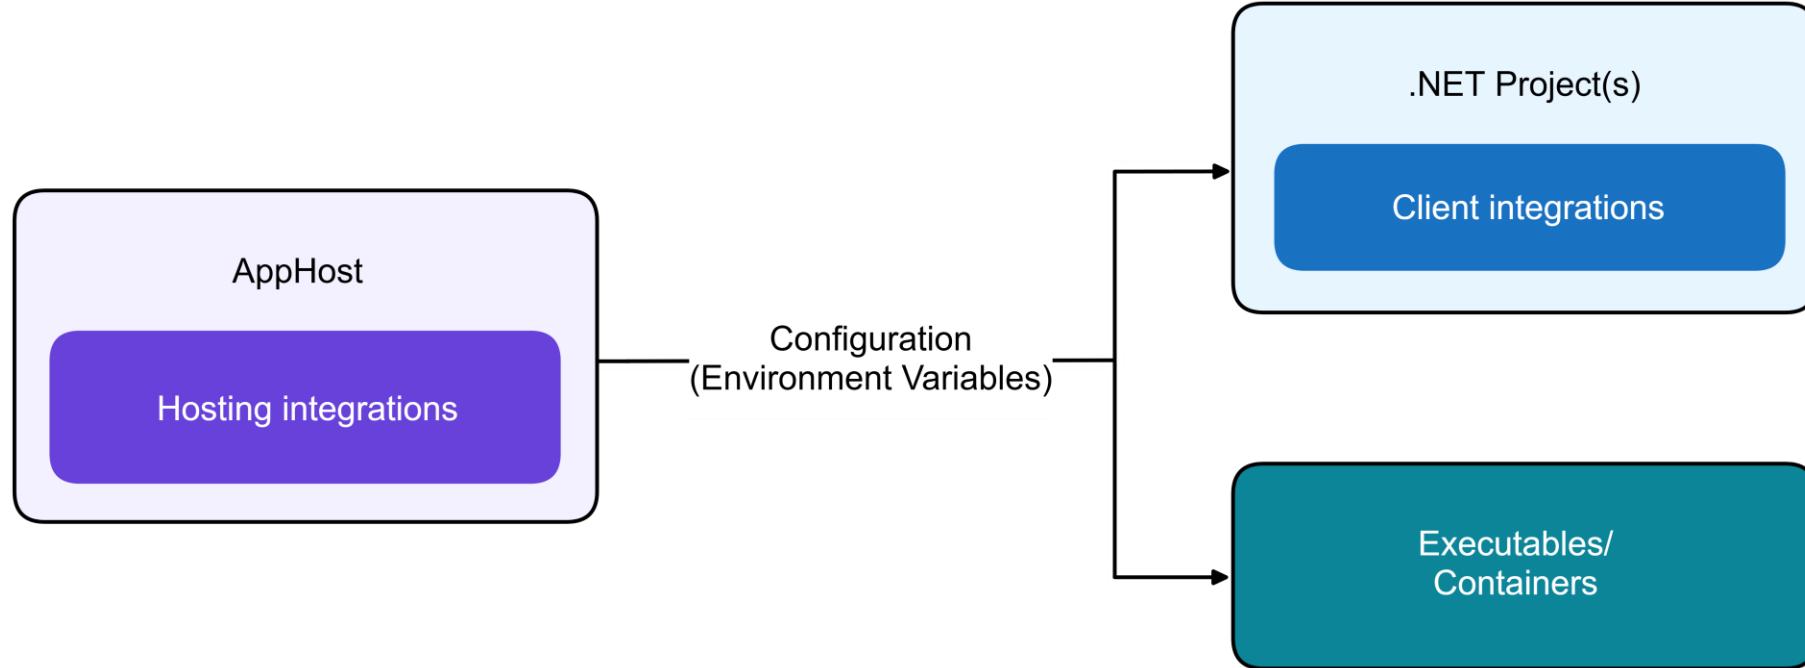
.NET Aspire Integrations



.NET Aspire

Integrations

Curated NuGet packages designed to make it **simple** to **integrate** your **microservices** with **external resources**—such as **caches**, **databases**, or **message brokers**

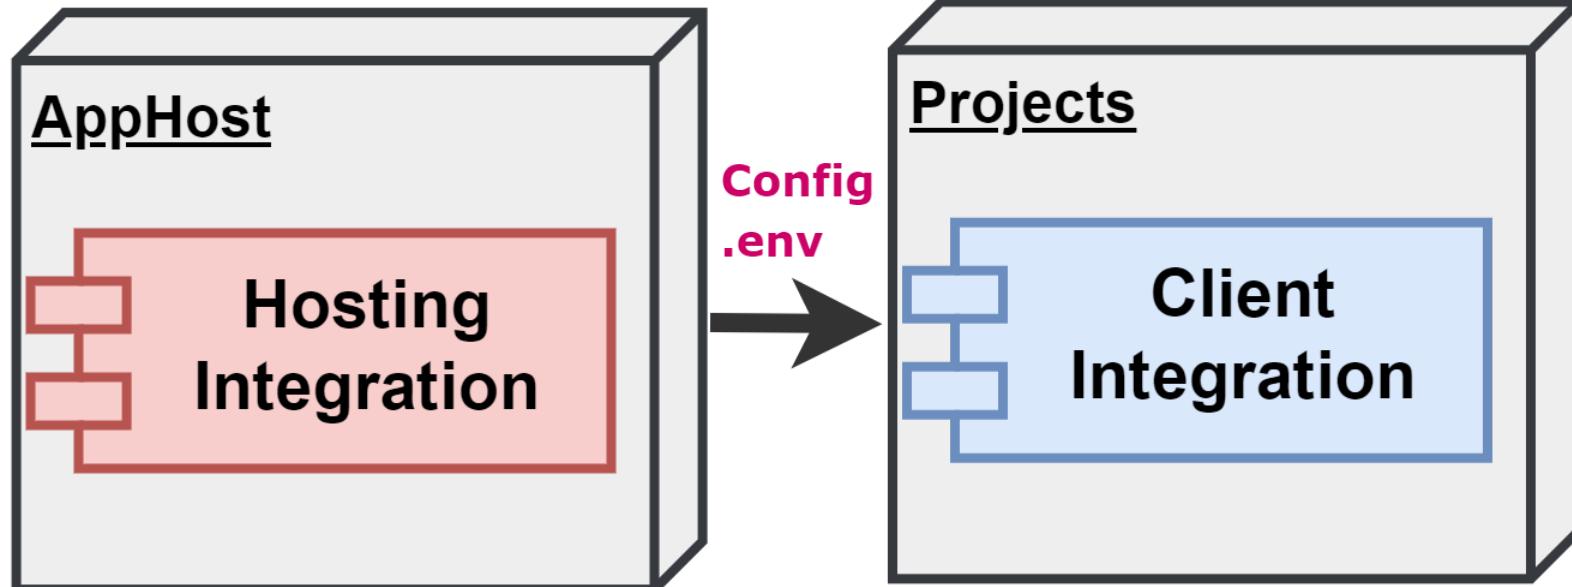


Hosting vs. Client Integrations

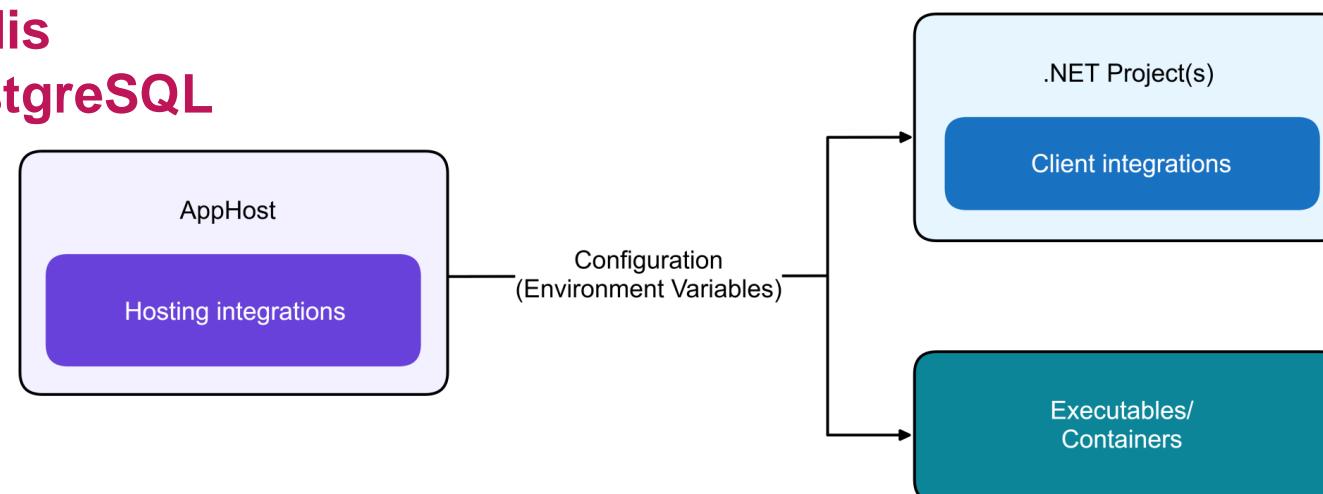


.NET Aspire

Integrations



Aspire.Hosting.Redis
Aspire.Hosting.PostgreSQL



Aspire.StackExchange.Redis
Aspire.Npgsql.EntityFrameworkCore.PostgreSQL

Hosting Integrations – Provisioning Resources

Extend the **IDistributedApplicationBuilder** interface, adding methods like **AddRedis("name")** or **AddPostgres("db")**

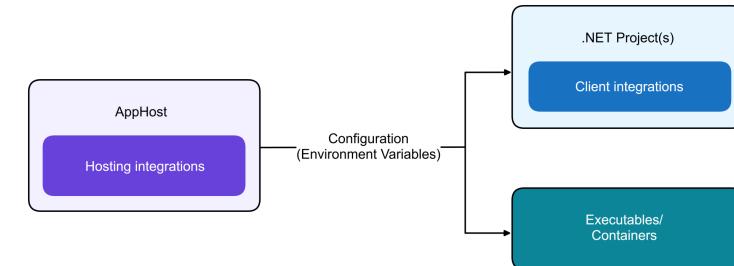
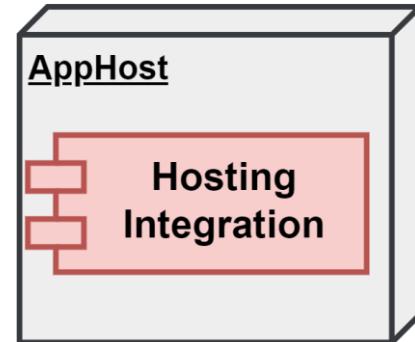


Provision container-based resources automatically

Use **environment variables** to pass **connection strings** to other projects that depend on them



```
var builder =  
    DistributedApplication.CreateBuilder(args);  
    // Hosting integration call:  
    var cache = builder.AddRedis("cache");  
  
    // This sets up a Redis container named "cache"
```



Client Integrations – Connecting to Resources

Extend **IHostApplicationBuilder** in Consuming Projects;
builder.AddRedisDistributedCache(`cache`)

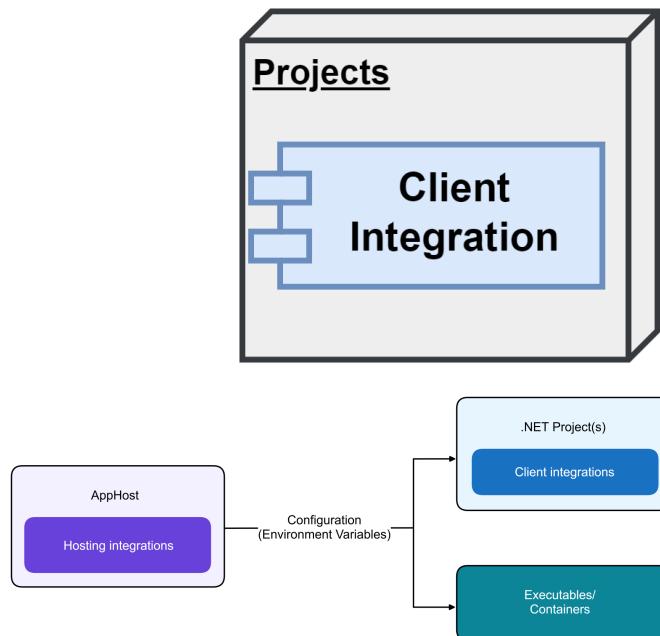
Add Health Checks, Logging, Telemetry

Example: Aspire.StackExchange.Redis

```
● ○ ●  
var builder = WebApplication.CreateBuilder(args);  
  
// AddServiceDefaults is often called automatically in Aspire-based  
// projects; you can call it explicitly, too.  
builder.AddServiceDefaults();  
  
// This might be from Aspire.StackExchange.Redis  
builder.AddRedisDistributedCache(connectionName: "cache");
```

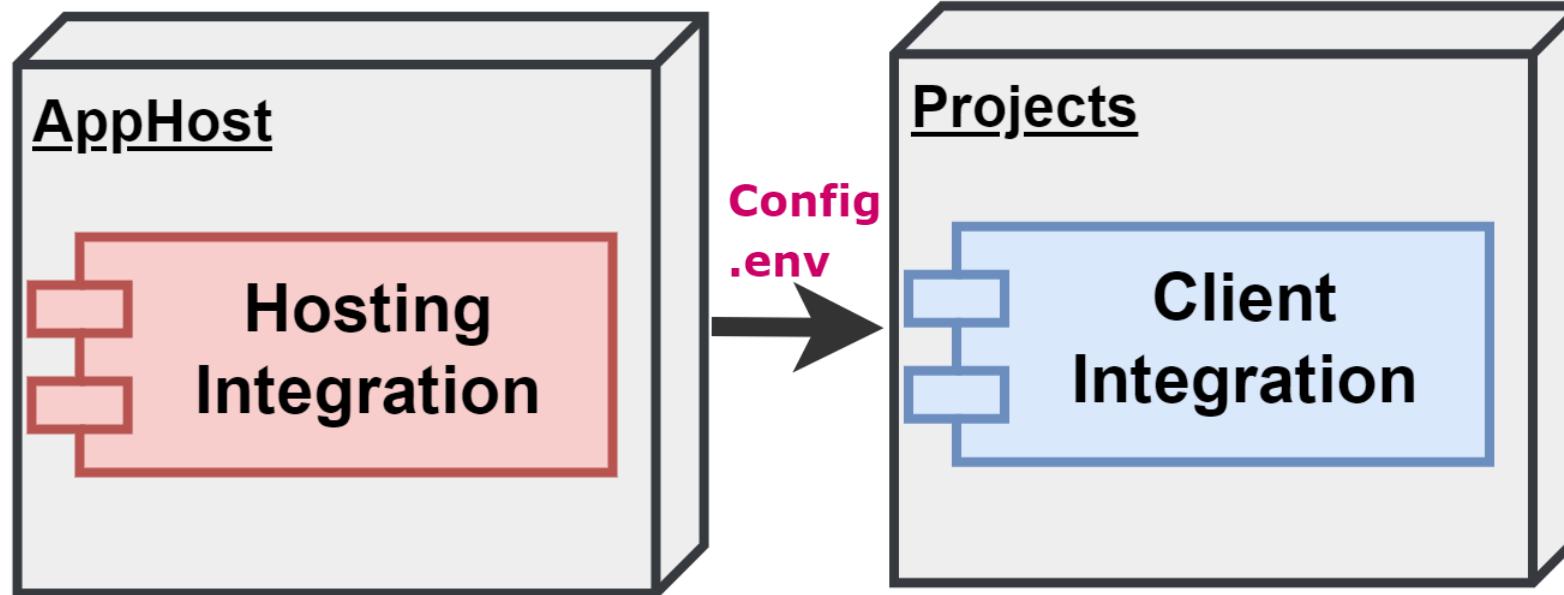


.NET Aspire

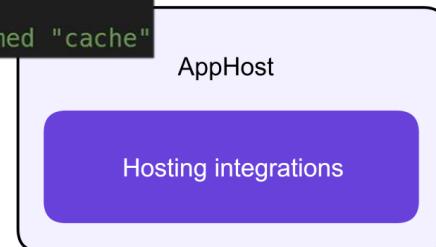


Relationship Hosting vs. Client Integrations

Integrations

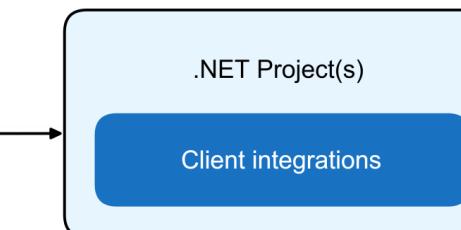


```
var builder =  
DistributedApplication.CreateBuilder(args);  
// Hosting integration call:  
var cache = builder.AddRedis("cache");  
  
// This sets up a Redis container named "cache"
```



Observability
Logging, tracing
and metrics
Health Checks
Resiliency

Configuration
(Environment Variables)



```
var builder = WebApplication.CreateBuilder(args);  
// AddServiceDefaults is often called automatically  
builder.AddServiceDefaults();  
  
// This might be from Aspire.StackExchange.Redis  
builder.AddRedisDistributedCache(connectionName:  
"cache");
```



Official Integrations

Redis, PostgreSQL, RabbitMQ, Kafka, Keycloak

Azure & AWS Packages

Community Toolkit

[Official Integrations Table](#)

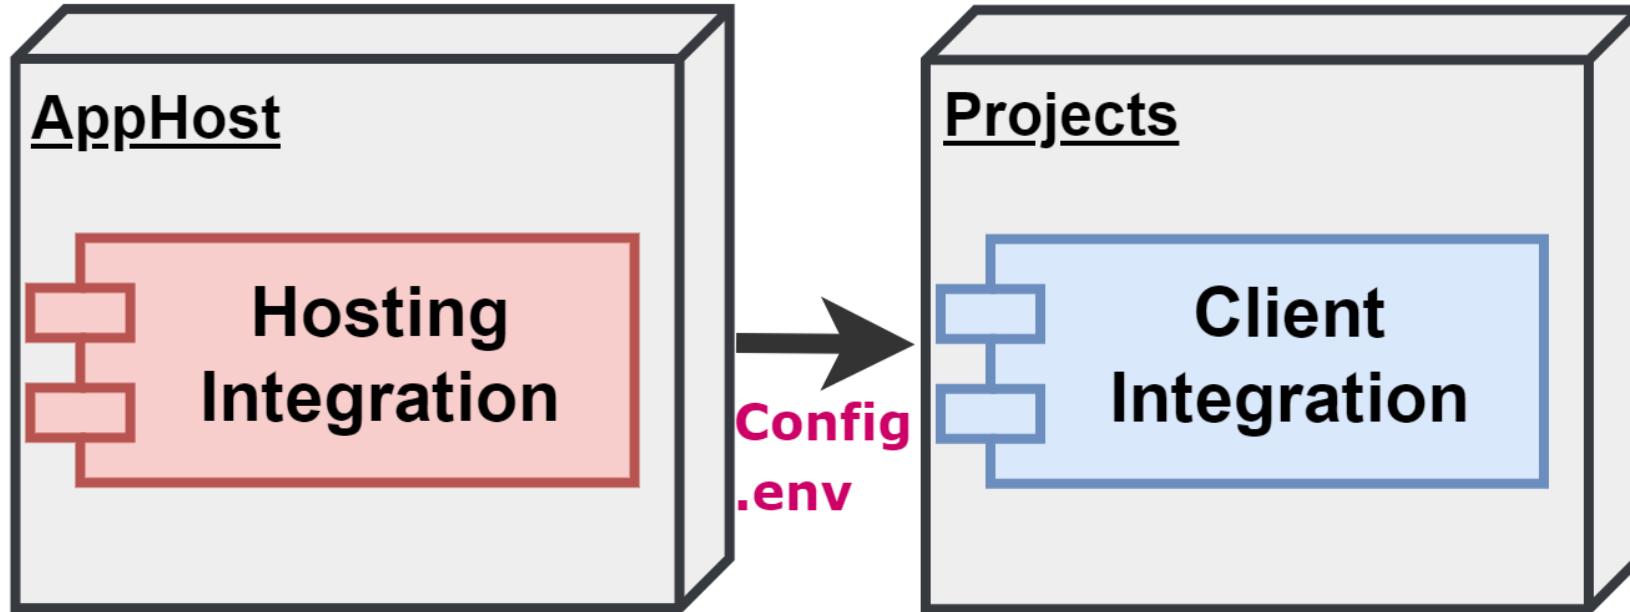
Integration docs and NuGet packages	Description
<ul style="list-style-type: none"> - Learn more: Apache Kafka - Hosting: Aspire.Hosting.Kafka - Client: Aspire.Confluent.Kafka 	A library for producing and consuming messages from an Apache Kafka broker.
<ul style="list-style-type: none"> - Learn more: Dapr - Hosting: Aspire.Hosting.Dapr - Client: N/A 	A library for modeling Dapr as a .NET Aspire resource.
<ul style="list-style-type: none"> - Learn more: Elasticsearch - Hosting: Aspire.Hosting.Elasticsearch - Client: Aspire.Elastic.Clients.Elasticsearch 	A library for accessing Elasticsearch databases.
<ul style="list-style-type: none"> - Learn more: Keycloak - Hosting: Aspire.Hosting.Keycloak - Client: Aspire.Keycloak.Authentication 	A library for accessing Keycloak authentication.
<ul style="list-style-type: none"> - Learn more: Milvus - Hosting: Aspire.Hosting.Milvus - Client: Aspire.Milvus.Client 	A library for accessing Milvus databases.
<ul style="list-style-type: none"> - Learn more: MongoDB Driver - Hosting: Aspire.Hosting.MongoDB - Client: Aspire.MongoDB.Driver 	A library for accessing MongoDB databases.



.NET Aspire

Conclusion: Hosting vs. Client Integrations

Integrations



```
● ● ●  
var builder =  
DistributedApplication.CreateBuilder(args);  
// Hosting integration call:  
var cache = builder.AddRedis("cache");  
  
// This sets up a Redis container named "cache"
```

```
● ● ●  
var builder = WebApplication.CreateBuilder(args);  
  
// AddServiceDefaults is often called automatically  
builder.AddServiceDefaults();  
  
// This might be from Aspire.StackExchange.Redis  
builder.AddRedisDistributedCache(connectionName:  
"cache");
```

.NET Aspire Service Discovery

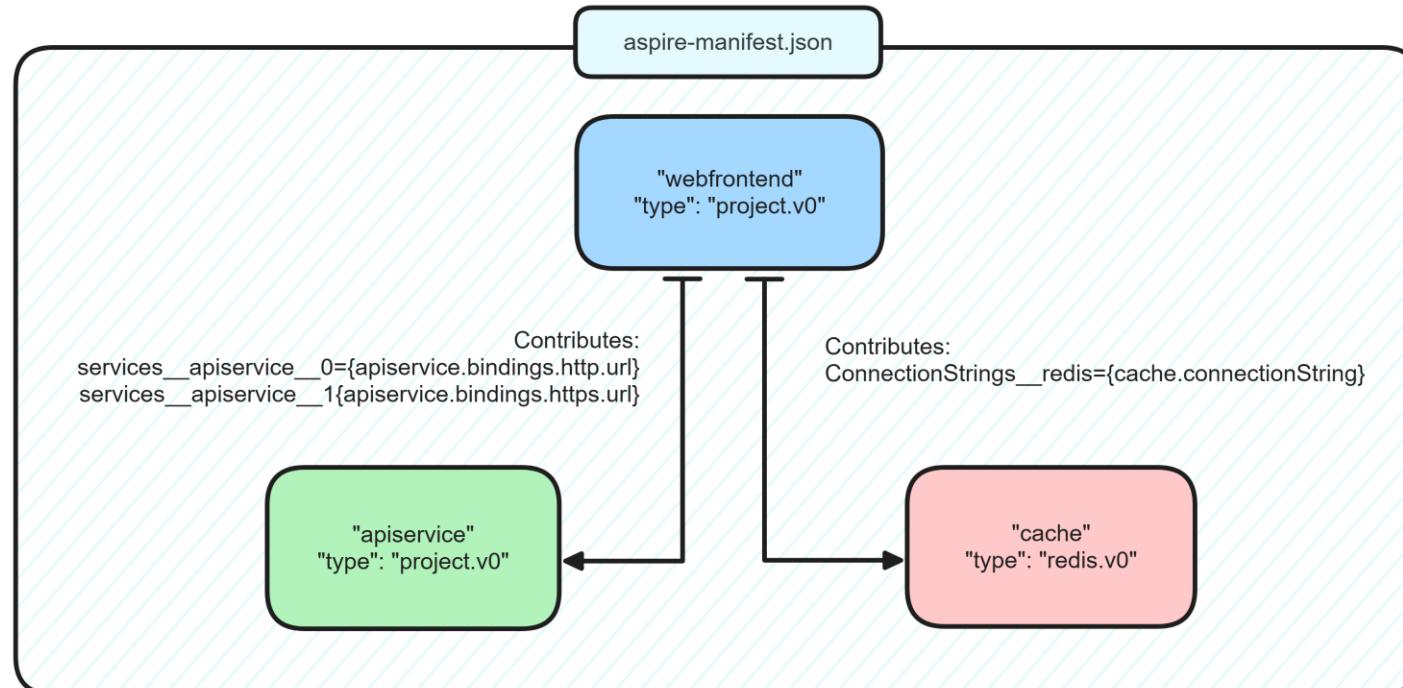


Service
Discovery

Microservices frequently **scale** or move around, **Service discovery** is how those **microservices** locate each other

Service A needs to **call Service B's REST endpoint**—but can't rely on a **hard-coded host/port**

Rely on **configuration** that's **injected automatically**, making each service **discoverable by name**



Implicit Service Discovery by Reference



.NET Aspire

Service
Discovery



```
var builder = DistributedApplication.CreateBuilder(args);

var cache = builder.AddRedis("cache");

var catalog = builder.AddProject<Projects.CatalogService>
("catalog");
var basket = builder.AddProject<Projects.BasketService>("basket");

var frontend = builder.AddProject<Projects.MyFrontend>("frontend")
    .WithReference(cache)
    .WithReference(catalog)
    .WithReference(basket);
```

Built-in resource types

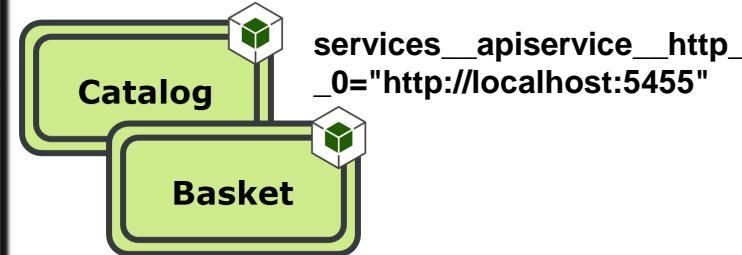


.NET Aspire

Service
Discovery

Built-in resource types

Method	Resource type	Description
AddProject	ProjectResource	A .NET project, for example, an ASP.NET Core web app.
AddContainer	ContainerResource	A container image, such as a Docker image.
AddExecutable	ExecutableResource	An executable file, such as a Node.js app.
AddParameter	ParameterResource	A parameter resource that can be used to express external parameters.



Implicit Service Discovery by Reference



.NET Aspire

Service
Discovery

Containers → injects → connection string

Projects → injects → microservice endpoints

```
var builder = DistributedApplication.CreateBuilder(args);

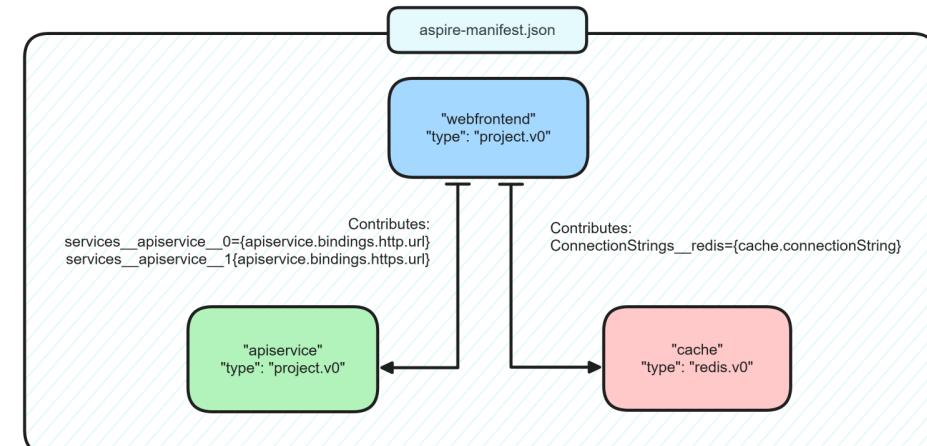
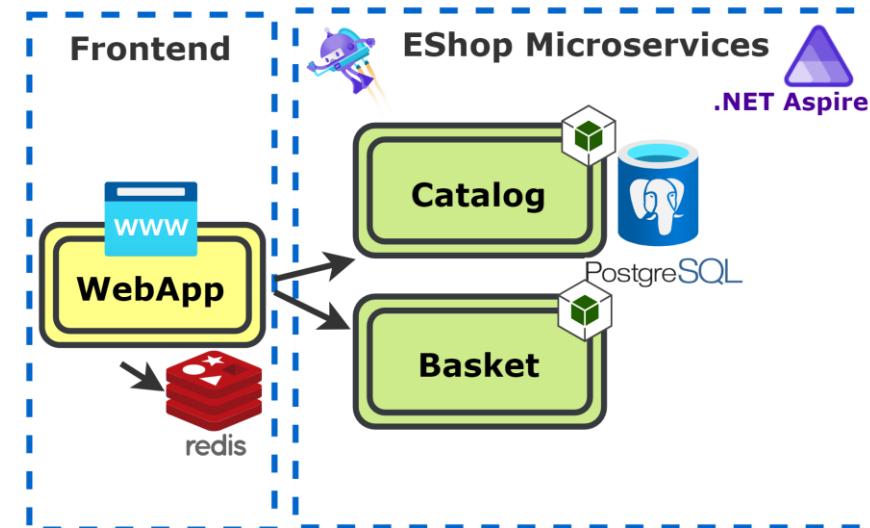
var cache = builder.AddRedis("cache");

var catalog = builder.AddProject<Projects.CatalogService>("catalog");
var basket = builder.AddProject<Projects.BasketService>("basket");

var frontend = builder.AddProject<Projects.MyFrontend>("frontend")
    .WithReference(cache)
    .WithReference(catalog)
    .WithReference(basket);
```

```
WithReference(cache)
    ConnectionStrings__cache="localhost:62354"
```

```
WithReference(apiservice)
    services__apiservice__http__0="http://localhost:5455"
    services__apiservice__https__0="https://localhost:7356"
```



:kaya

50

Configuring Container Endpoints



.NET Aspire

Service
Discovery



```
var builder = DistributedApplication.CreateBuilder(args);

var customContainer = builder.AddContainer("myapp", "mycustomcontainer")
    .WithHttpEndpoint(port: 9043, name: "endpoint");

var endpoint = customContainer.GetEndpoint("endpoint");

var apiservice = builder.AddProject<Projects.AspireApp_ApiService>
    ("apiservice") .WithReference(endpoint);
```

Implicit Service Discovery by Reference



.NET Aspire

Service
Discovery

Containers → injects → connection string

Projects → injects → microservice endpoints

```
var builder = DistributedApplication.CreateBuilder(args);

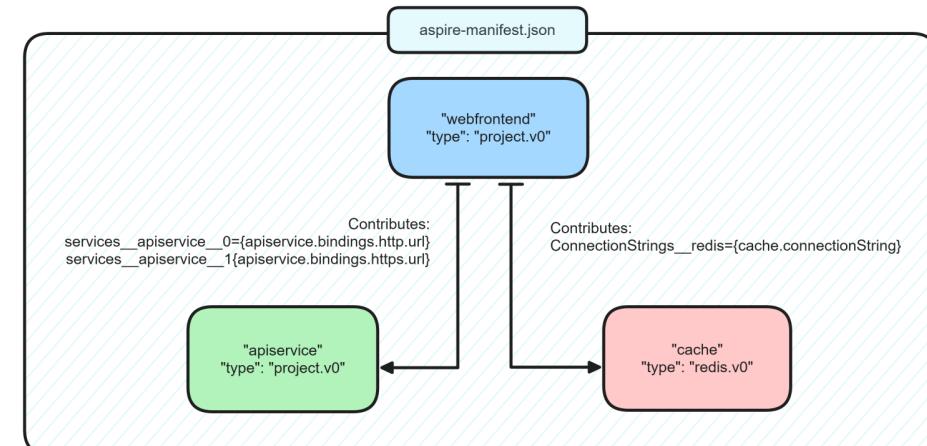
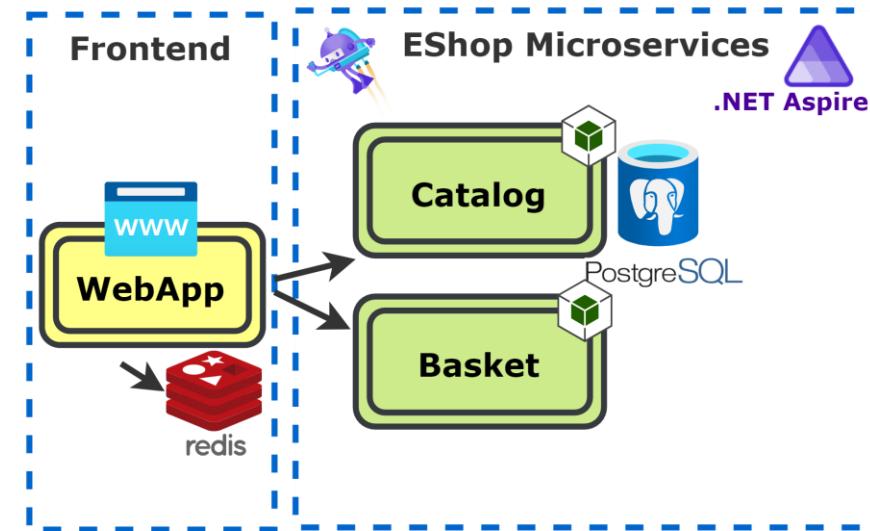
var cache = builder.AddRedis("cache");

var catalog = builder.AddProject<Projects.CatalogService>("catalog");
var basket = builder.AddProject<Projects.BasketService>("basket");

var frontend = builder.AddProject<Projects.MyFrontend>("frontend")
    .WithReference(cache)
    .WithReference(catalog)
    .WithReference(basket);
```

```
WithReference(cache)
    ConnectionStrings__cache="localhost:62354"
```

```
WithReference(apiservice)
    services__apiservice__http__0="http://localhost:5455"
    services__apiservice__https__0="https://localhost:7356"
```



:kaya

52

Building Your First .NET Aspire Application

Start the .NET Aspire Project for Redis-Api-Frontend Architecture

Run Aspire App and Explore the .NET Aspire dashboard

Understanding .NET Aspire Project Organization

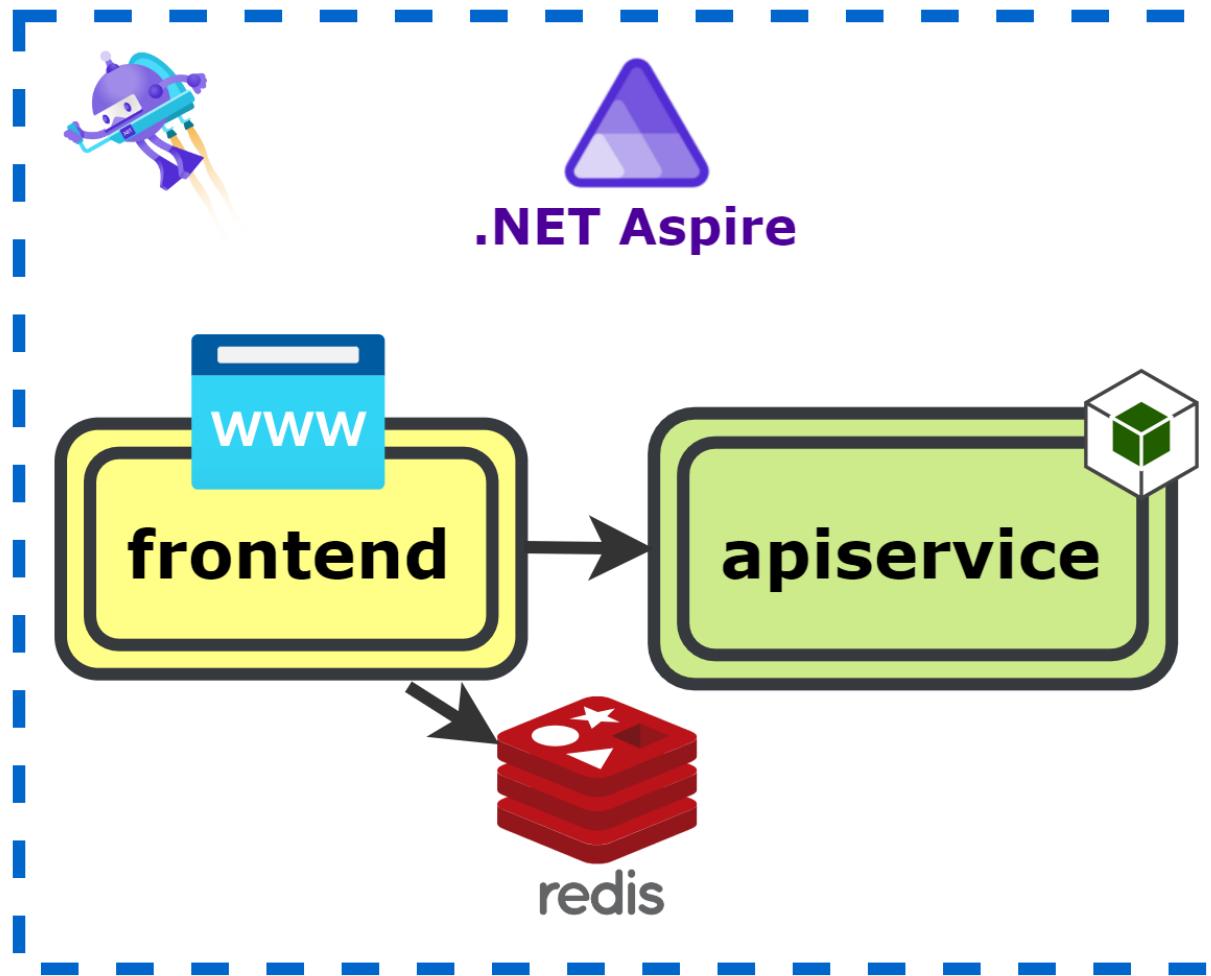
Examine .NET Aspire Host Project

Examine .NET Aspire Service Defaults Project

Mehmet Ozkaya

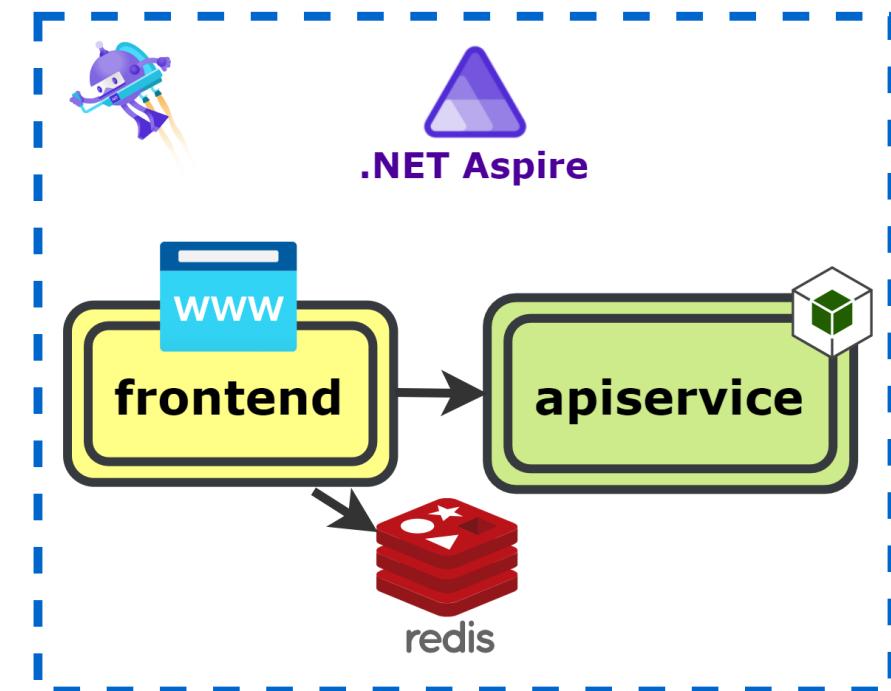


Start the .NET Aspire Project for Redis-Api-Frontend



Start the .NET Aspire Project for Redis-Api-Frontend

1. Create a brand-new .NET Aspire Starter App
2. Check Redis as our caching solution
3. Build an API project within the solution and see how Aspire handles service discovery
4. Connect everything—API, Frontend, and Redis—under one solution, all orchestrated by Aspire



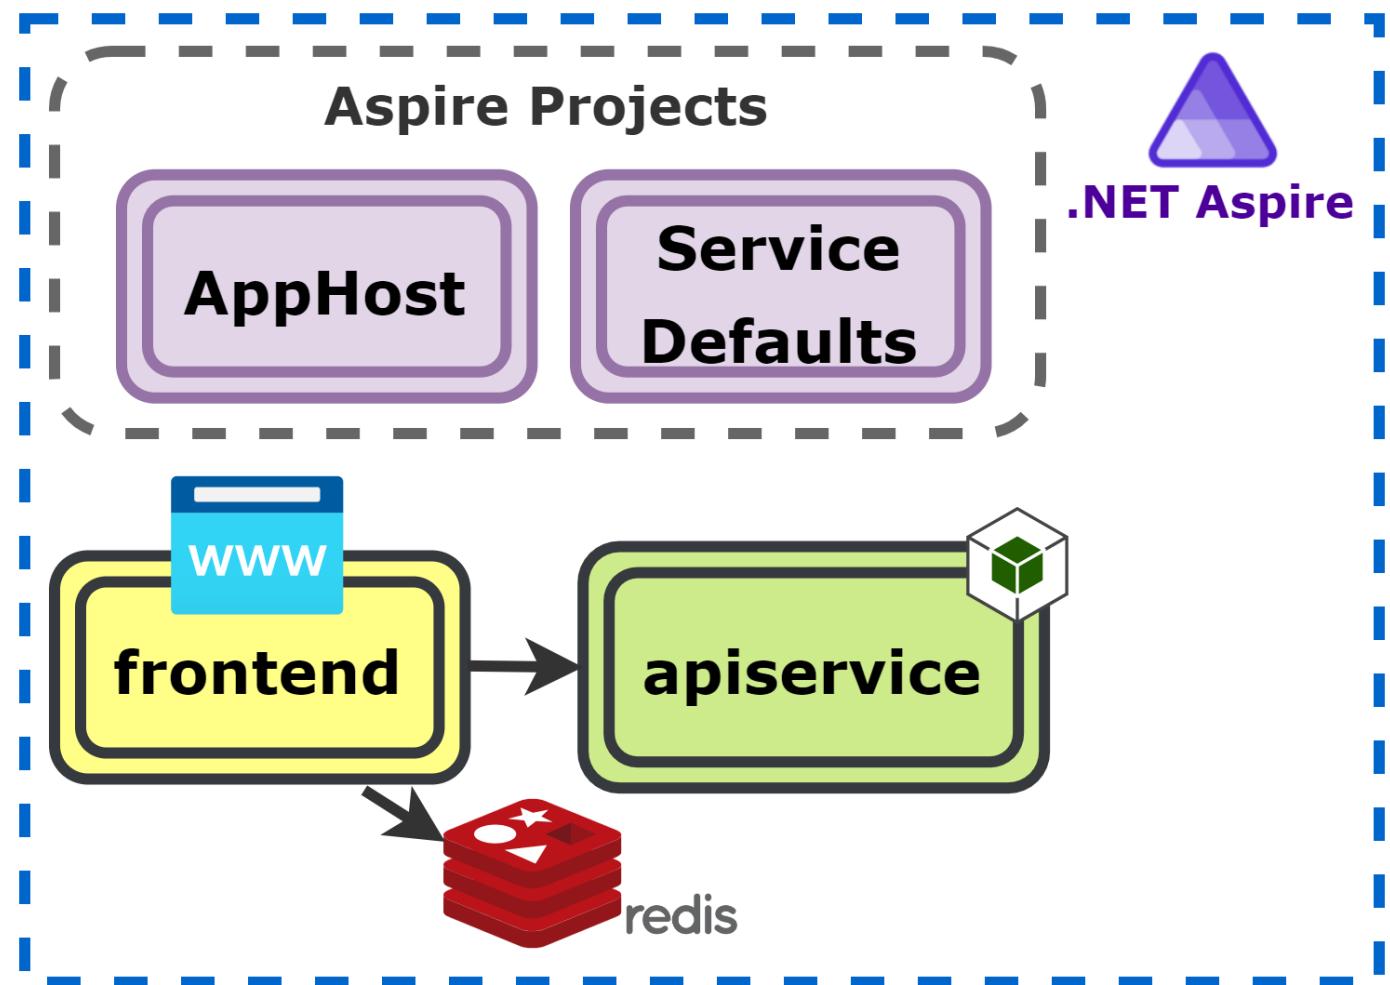
Start the .NET Aspire Project for Redis-Api-Frontend

AspireSample.AppHost

AspireSample.ServiceDefaults

AspireSample.ApiService

AspireSample.Web

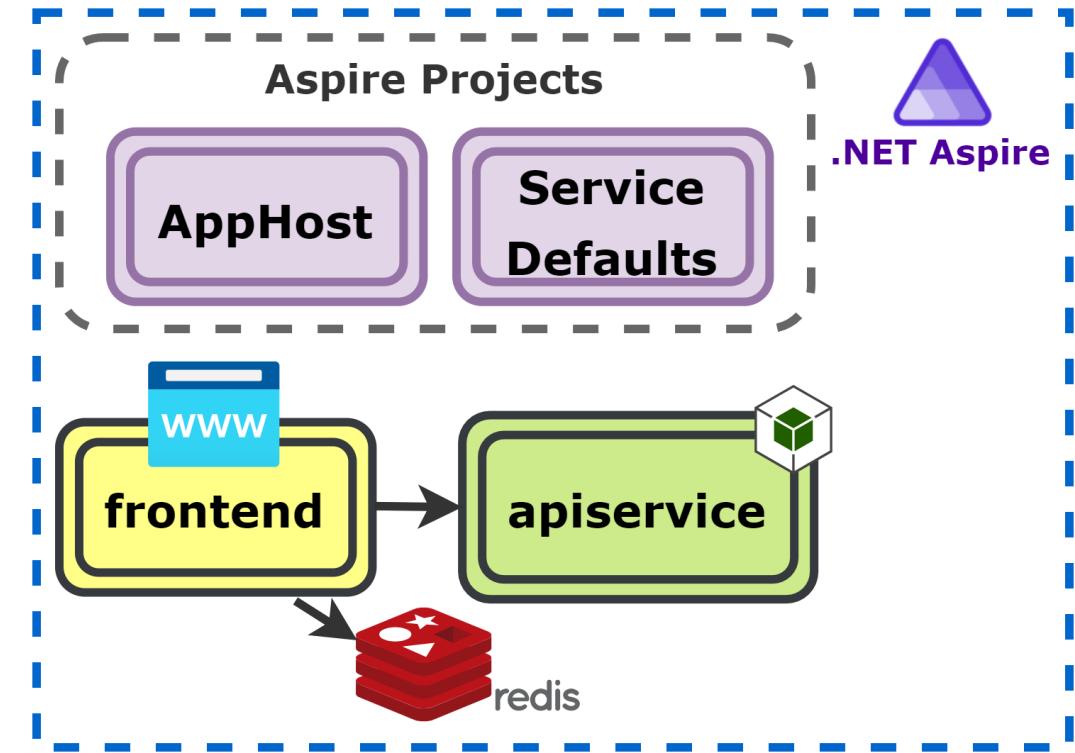


AspireSample.AppHost

Central Hub for All Projects & Services

Program.cs Defines Orchestration

Must Be Set as the Startup Project



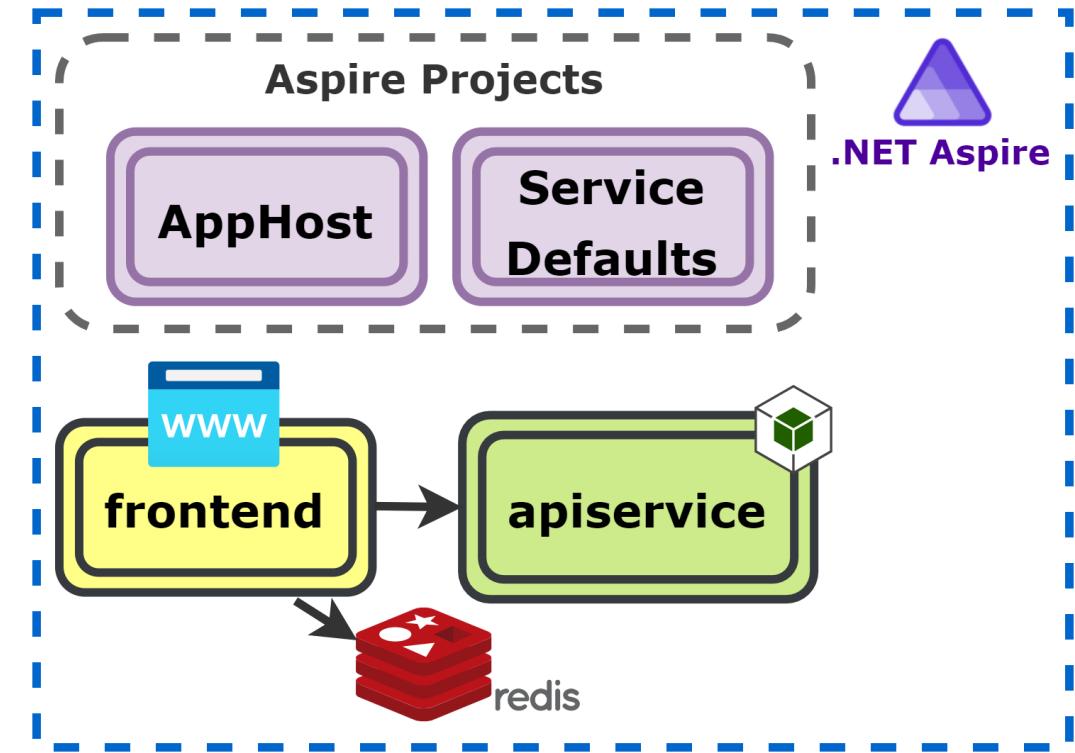
AspireSample.ServiceDefaults

Centralized Cross-Cutting Concerns

Logging, Tracing, Health Checks

Service Discovery defaults

Referenced by ApiService & Web

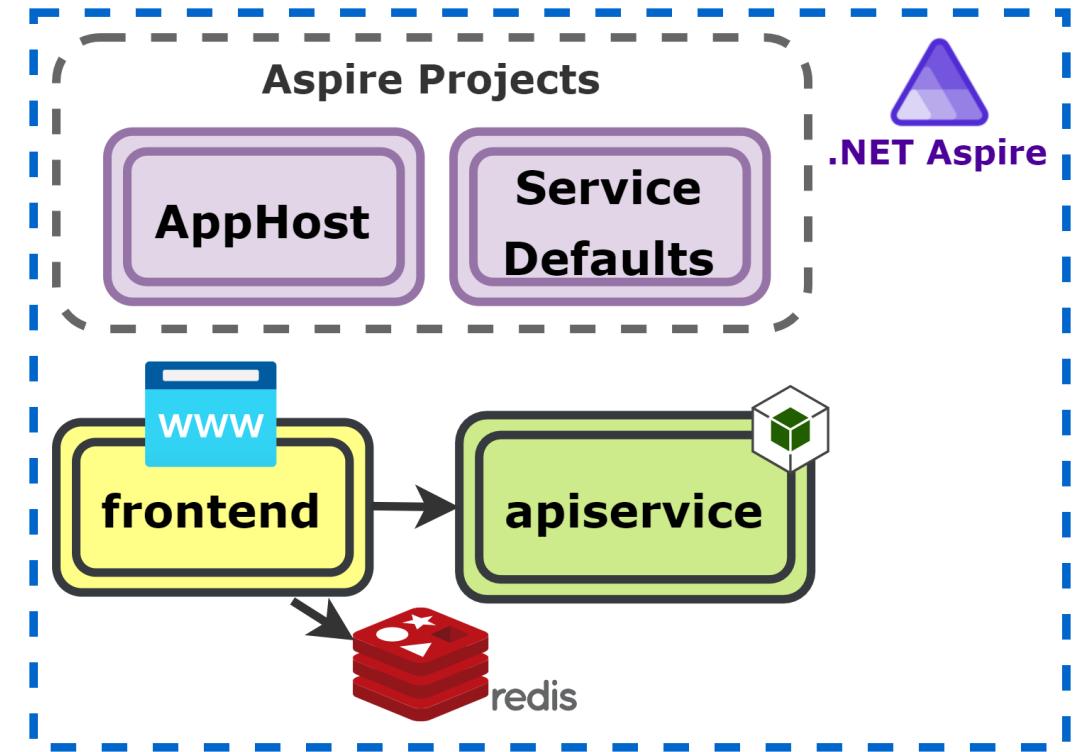


AspireSample.ApiService

Exposes endpoints that the Web project can call

Depends on the shared **ServiceDefaults** for cross-cutting concerns

Typical ASP.NET Core app, minimal set of Routes for data retrieval



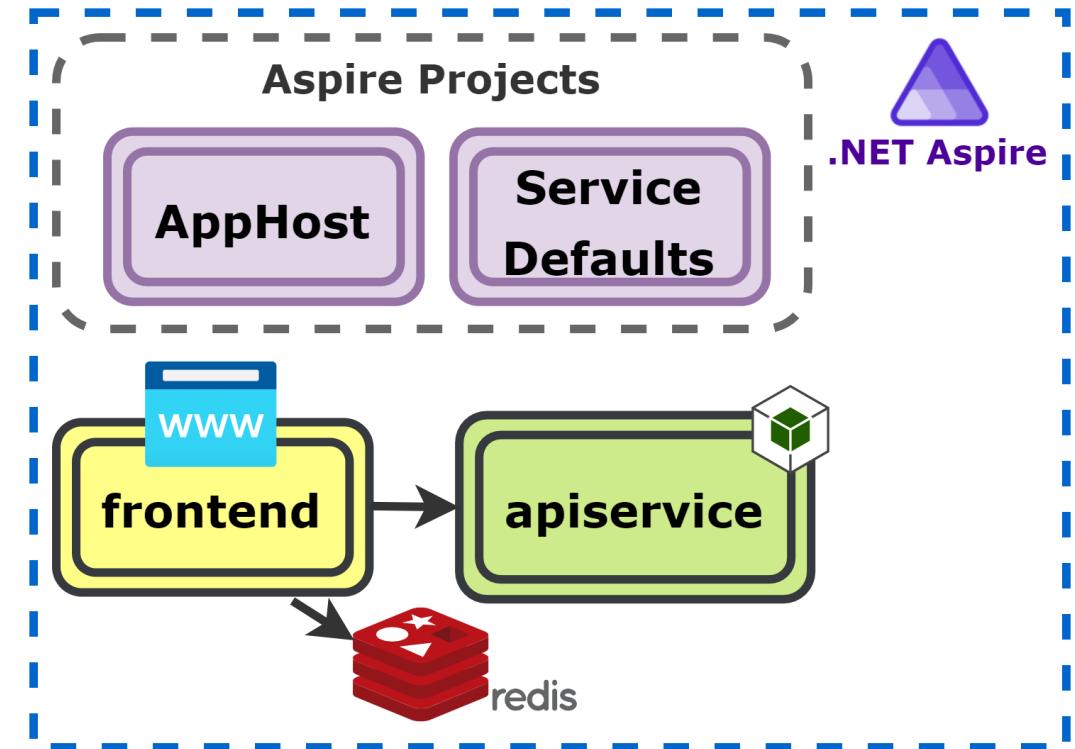
AspireSample.Web

Houses **UI components, layouts and pages**

Depends on the shared **ServiceDefaults** for cross-cutting concerns

Communicates with the **ApiService**—strongly typed **HTTP client** to fetch data

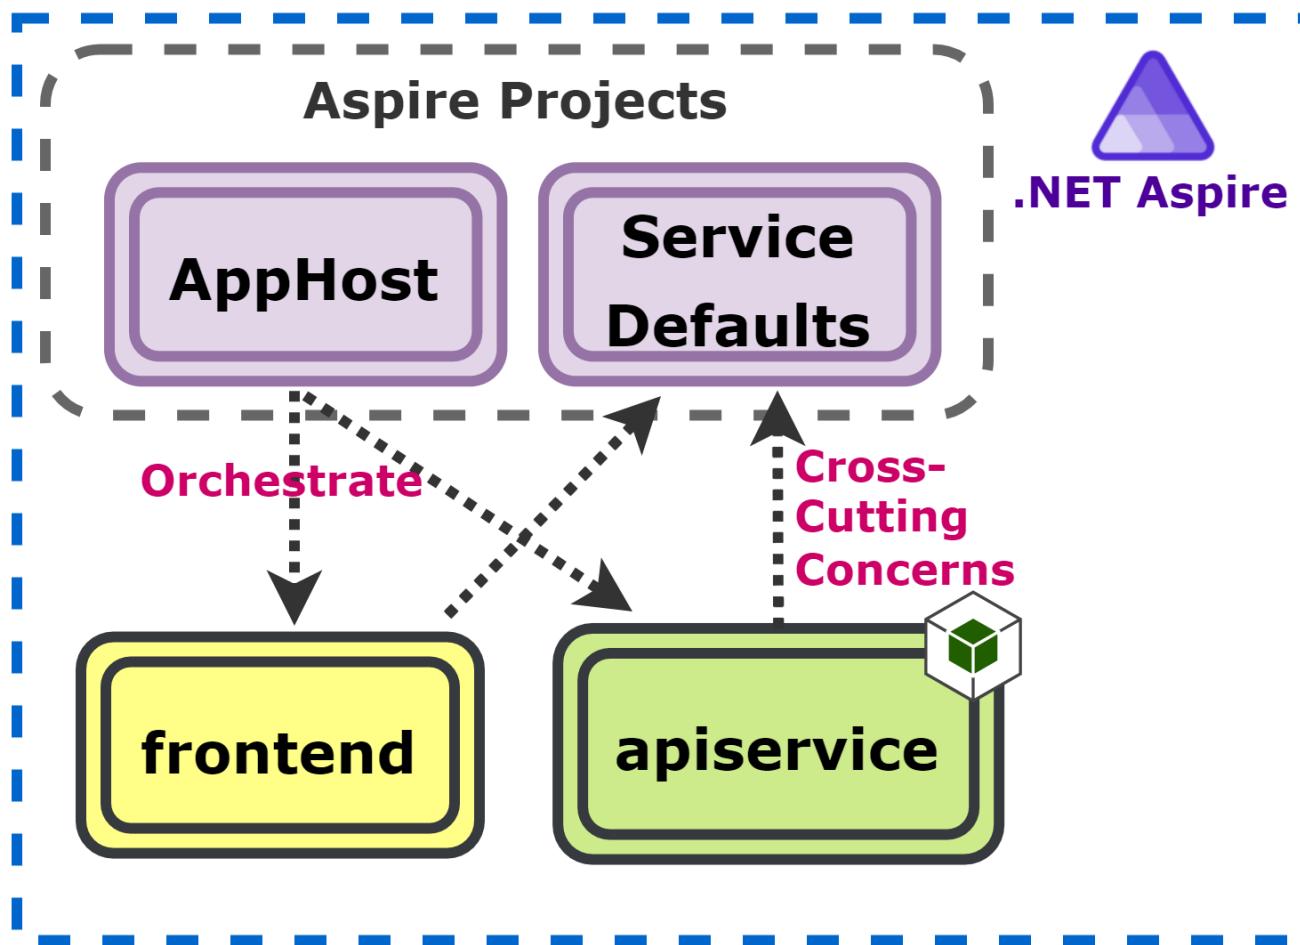
Discovered via .NET Aspire's **built-in service discovery**



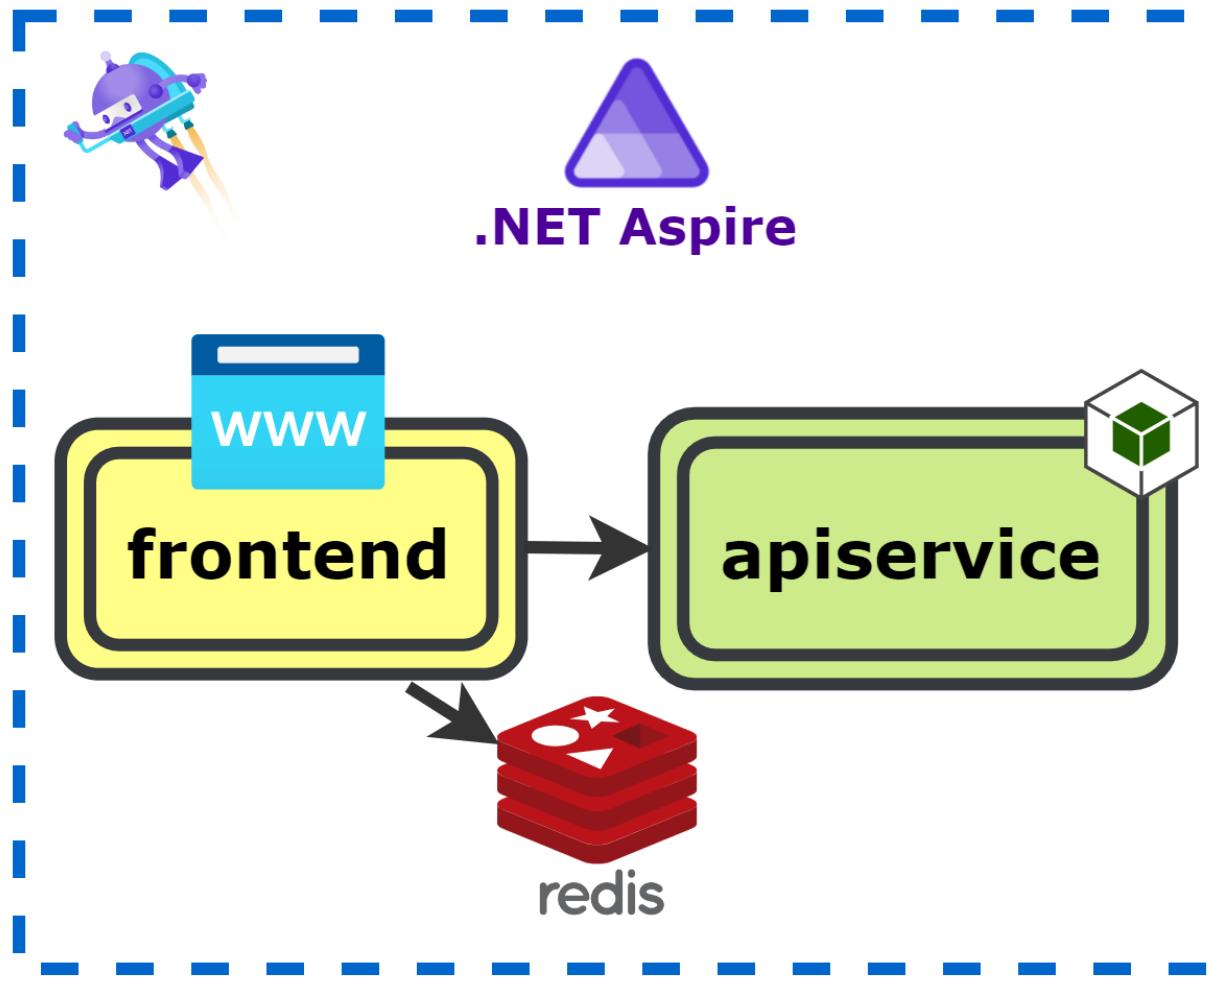
Project Dependencies

AppHost → references **ApiService & Web** for orchestration

ApiService/Web → reference **ServiceDefaults** for cross-cutting concerns



Start the .NET Aspire Project for Redis-Api-Frontend



Remember: Service Discovery



Service
Discovery

Containers → injects → connection string

Projects → injects → microservice endpoints

```
var builder = DistributedApplication.CreateBuilder(args);

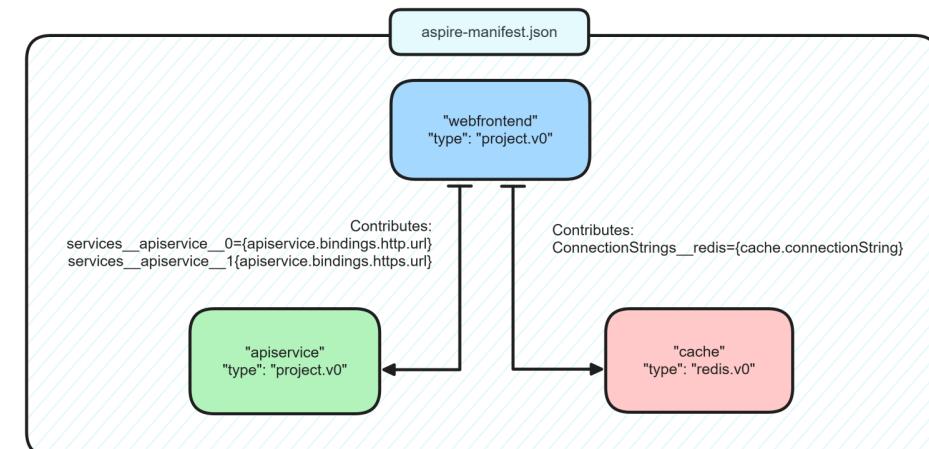
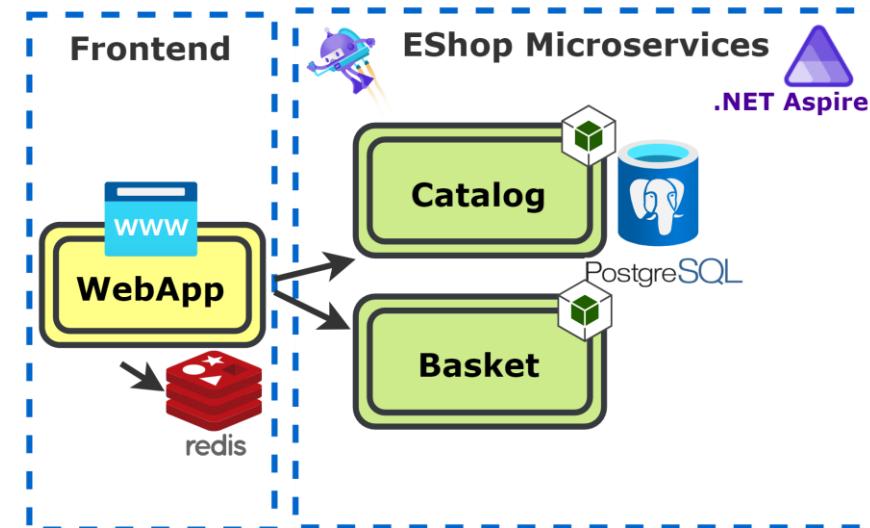
var cache = builder.AddRedis("cache");

var catalog = builder.AddProject<Projects.CatalogService>("catalog");
var basket = builder.AddProject<Projects.BasketService>("basket");

var frontend = builder.AddProject<Projects.MyFrontend>("frontend")
    .WithReference(cache)
    .WithReference(catalog)
    .WithReference(basket);
```

```
WithReference(cache)
    ConnectionStrings__cache="localhost:62354"
```

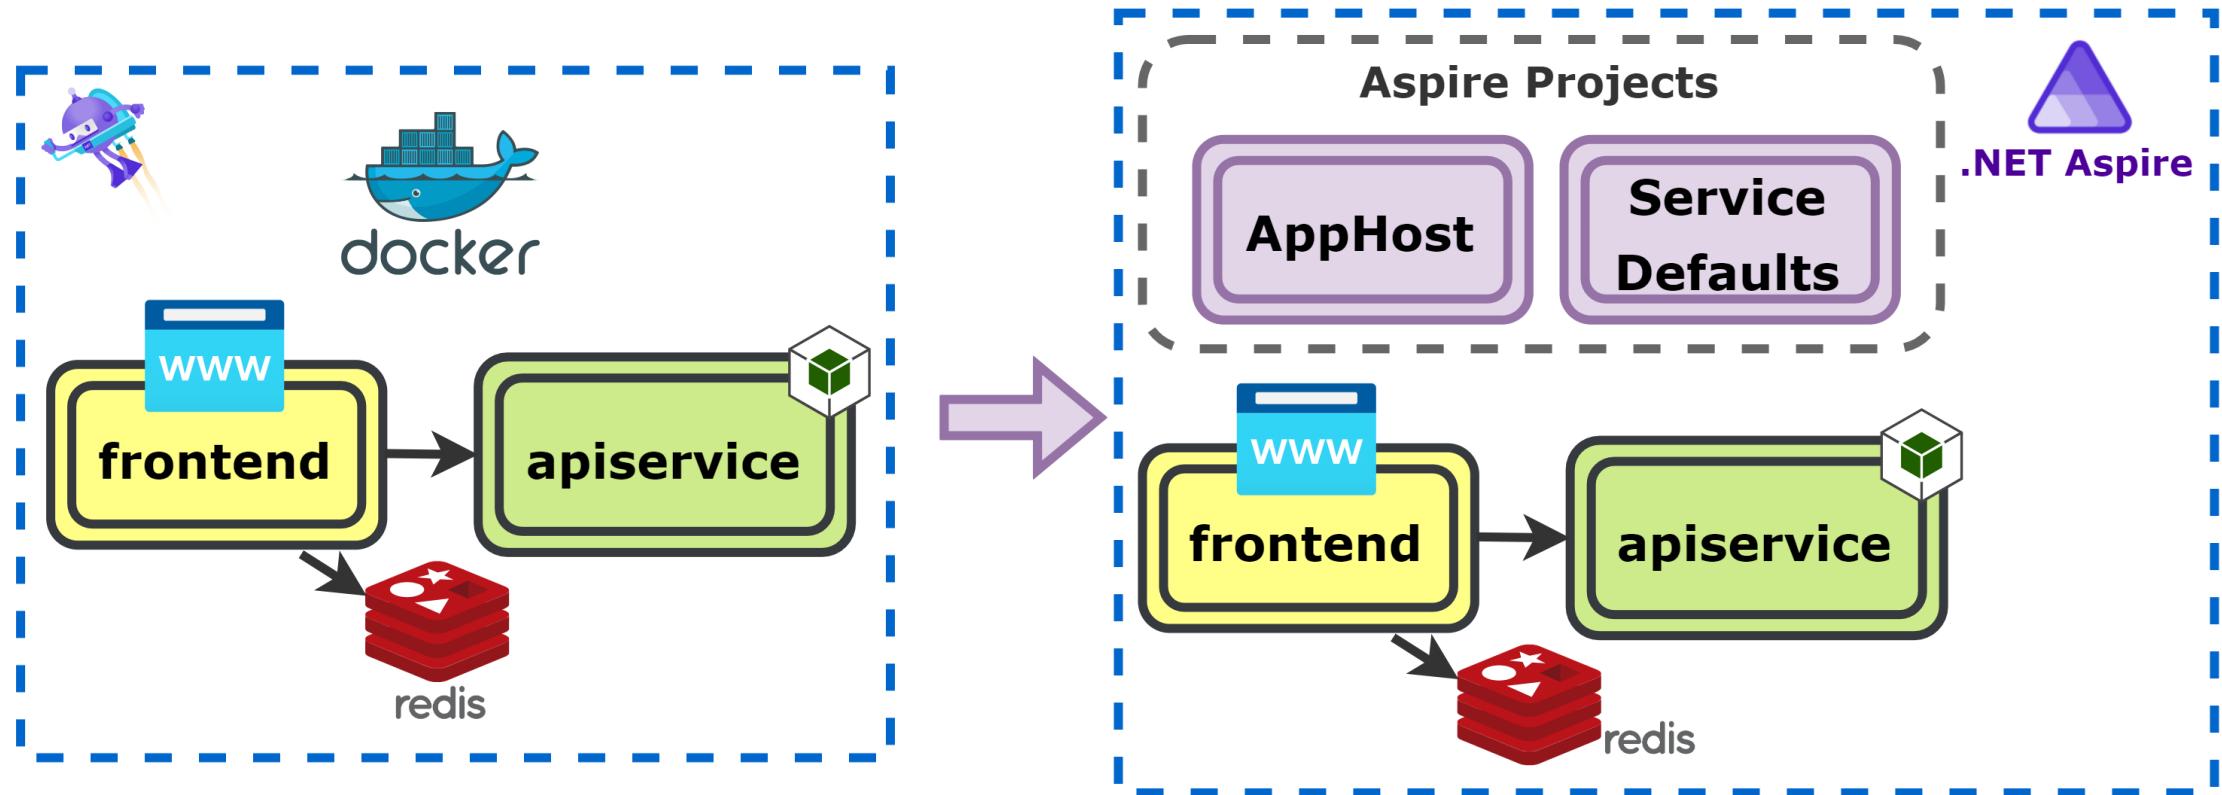
```
WithReference(apiservice)
    services__apiservice__http__0="http://localhost:5455"
    services__apiservice__https__0="https://localhost:7356"
```



:kaya

63

Adding .NET Aspire to Existing .NET Applications



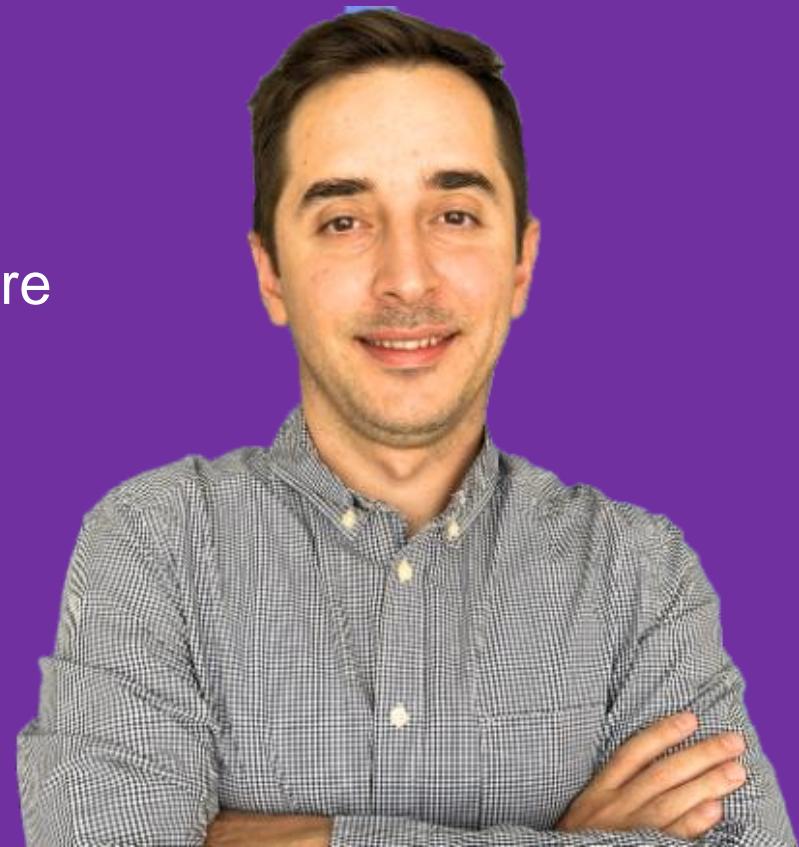
Building EShop Distributed Microservices Architecture with .NET Aspire

Create New Blank Solution and Organize Solution Folder Structure

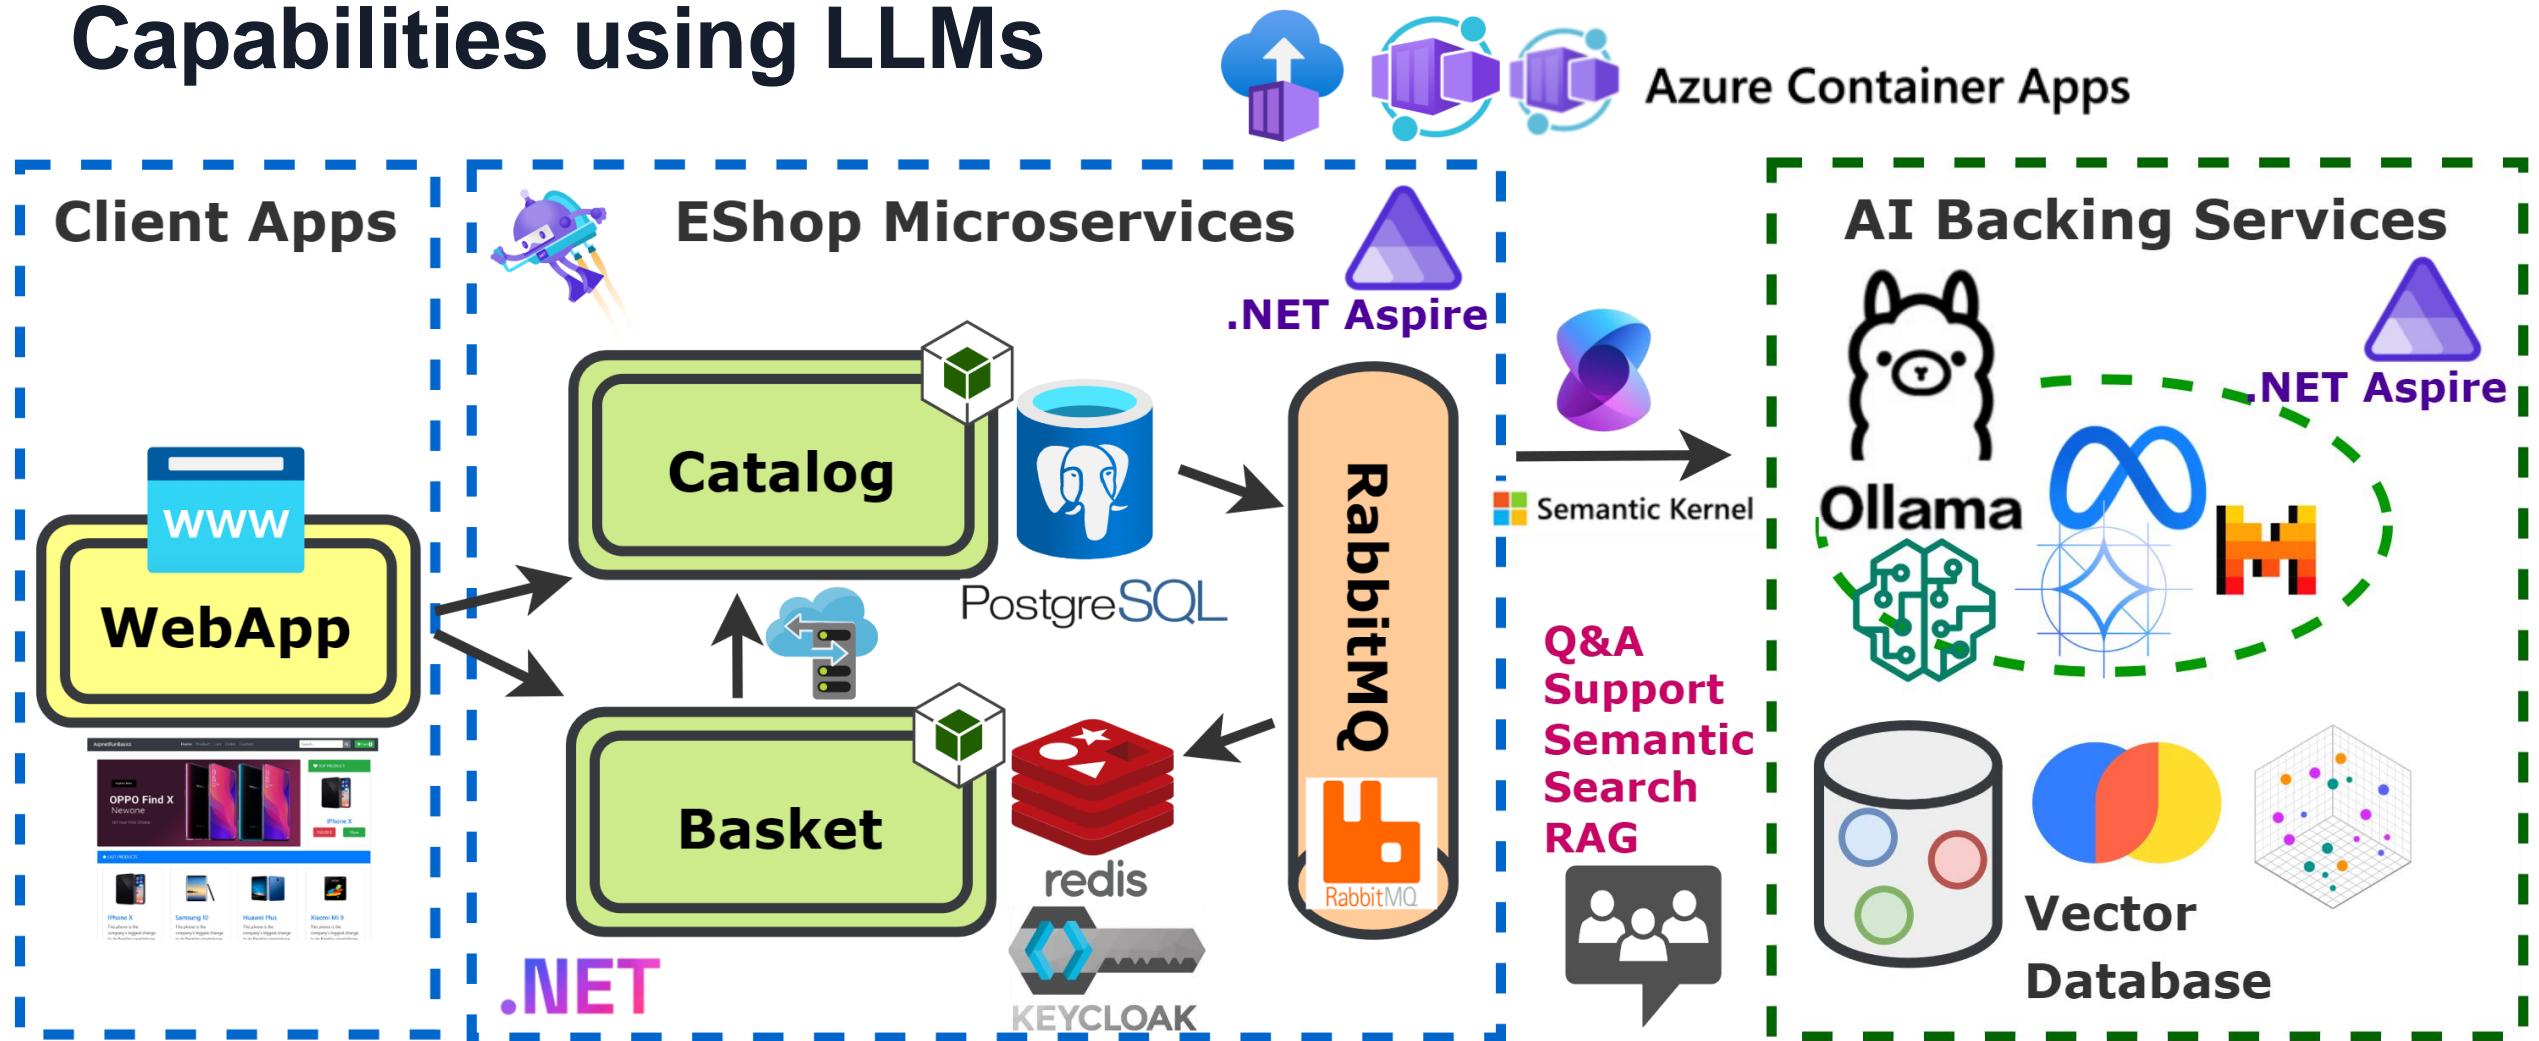
Create .NET Aspire Projects: AppHost (Orchestrator),
ServiceDefaults(Shared config)

Step by Step Development Plan for Building EShop Distributed
Architecture with .NET Aspire

Mehmet Ozkaya



Course Project: EShop Microservices with AI-Powered Capabilities using LLMs



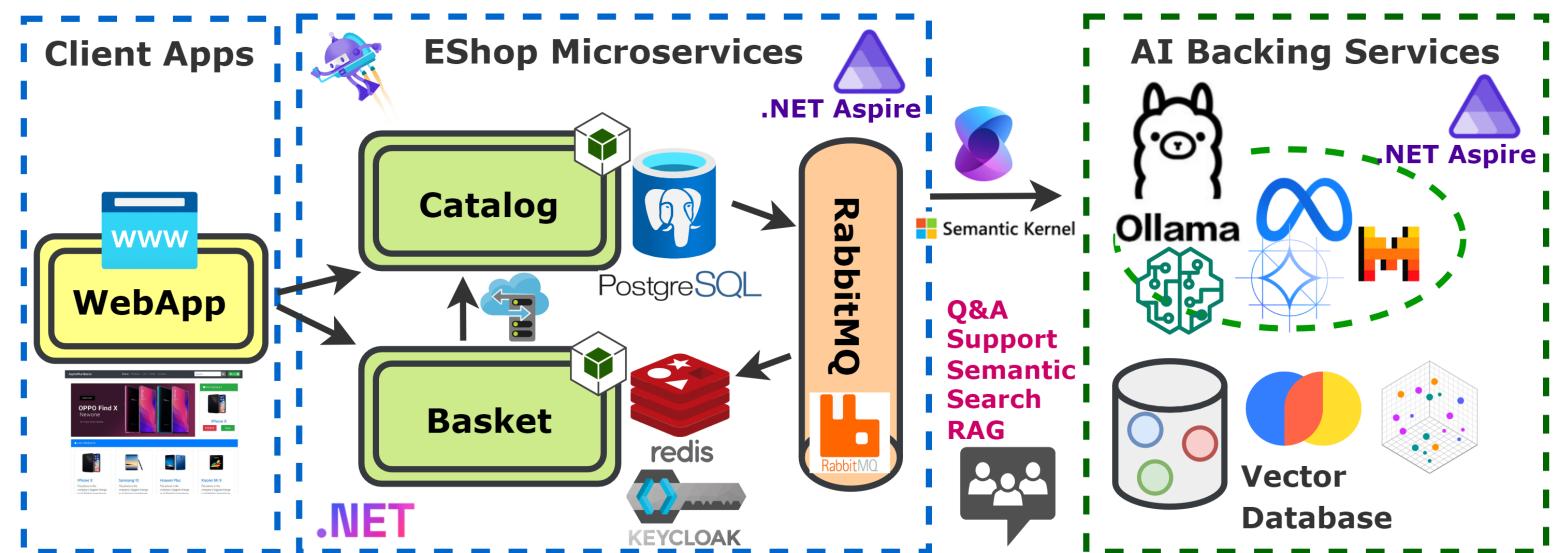
During this section

Create New Blank Solution

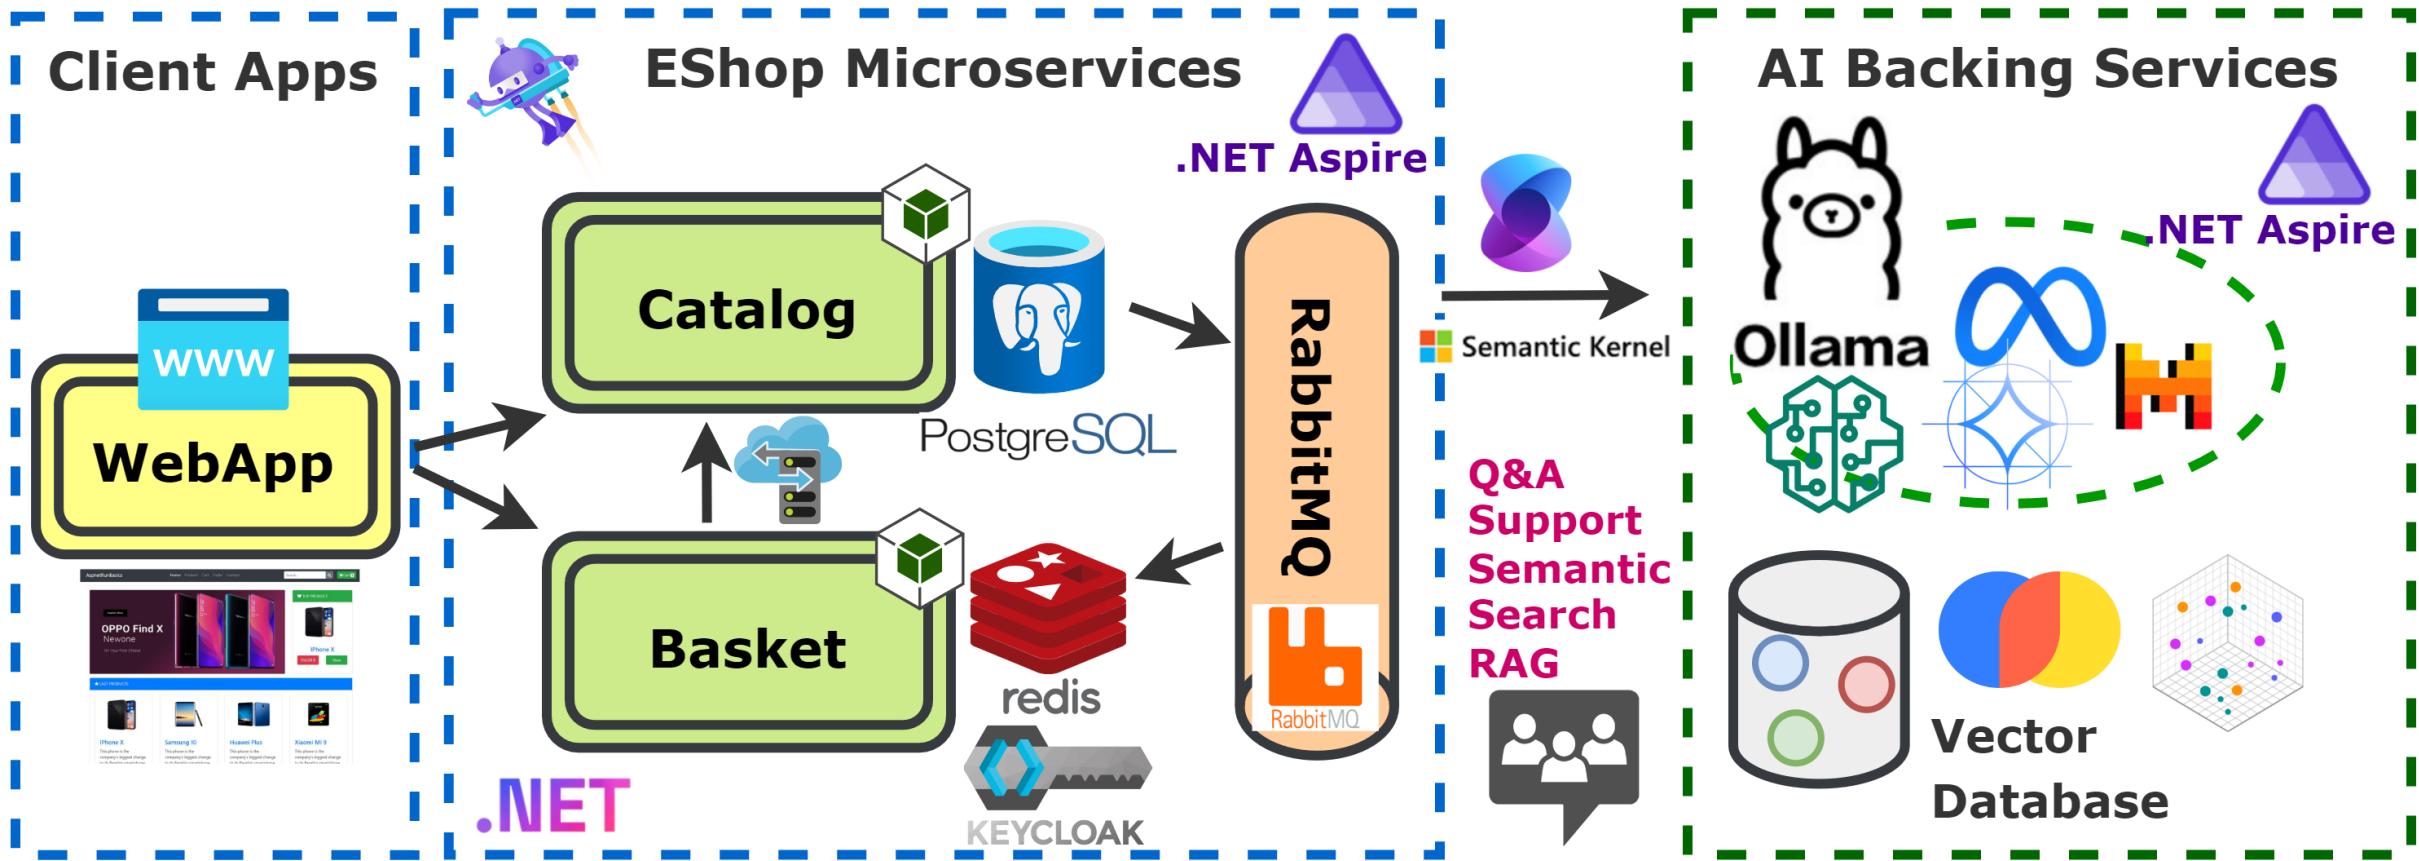
Create Visual Studio Solution Structure: aspire, services, frontend

Create .NET Aspire Projects

Make Step by Step Development Plan



Course Project: EShop Microservices with AI-Powered Capabilities using LLMs



Step by Step Development Plan

Setting Up Core Microservices

Catalog and Basket Microservices Development w/ Backing Services

Microservices Communications with .NET Aspire

Synchronous and Asynchronous Messaging between Catalog & Basket using RabbitMQ Message Broker

Authentication & Security

Secure Basket Endpoints w/ Keycloak using OpenID Connect Jwt Tokens

Frontend Development

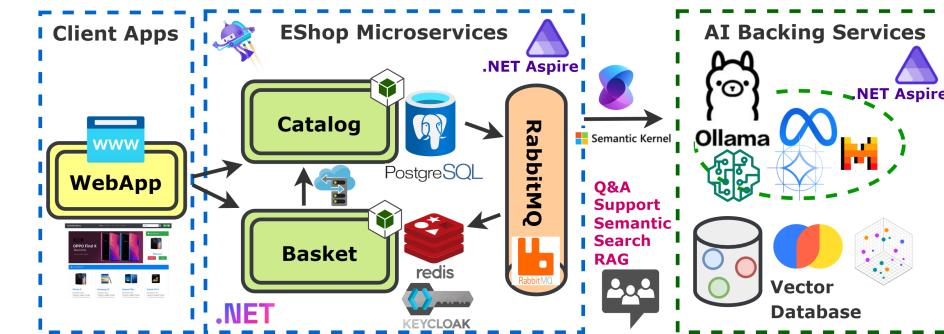
Blazor WebApp Products page development

Deployment to Azure Container Apps

.NET Aspire project to deploy ACA using azd up/down

Adding GenAI Features

Integrate Ollama LLMs and Vector Store for Semantic Search



Develop Catalog Microservice with PostgreSQL orchestrate in .NET Aspire

Domain and Technical Analysis of Catalog Microservice

Create Catalog Microservice Web API and Add .NET Aspire

Catalog Hosting Integration: Add PostgreSQL Resource in AppHost

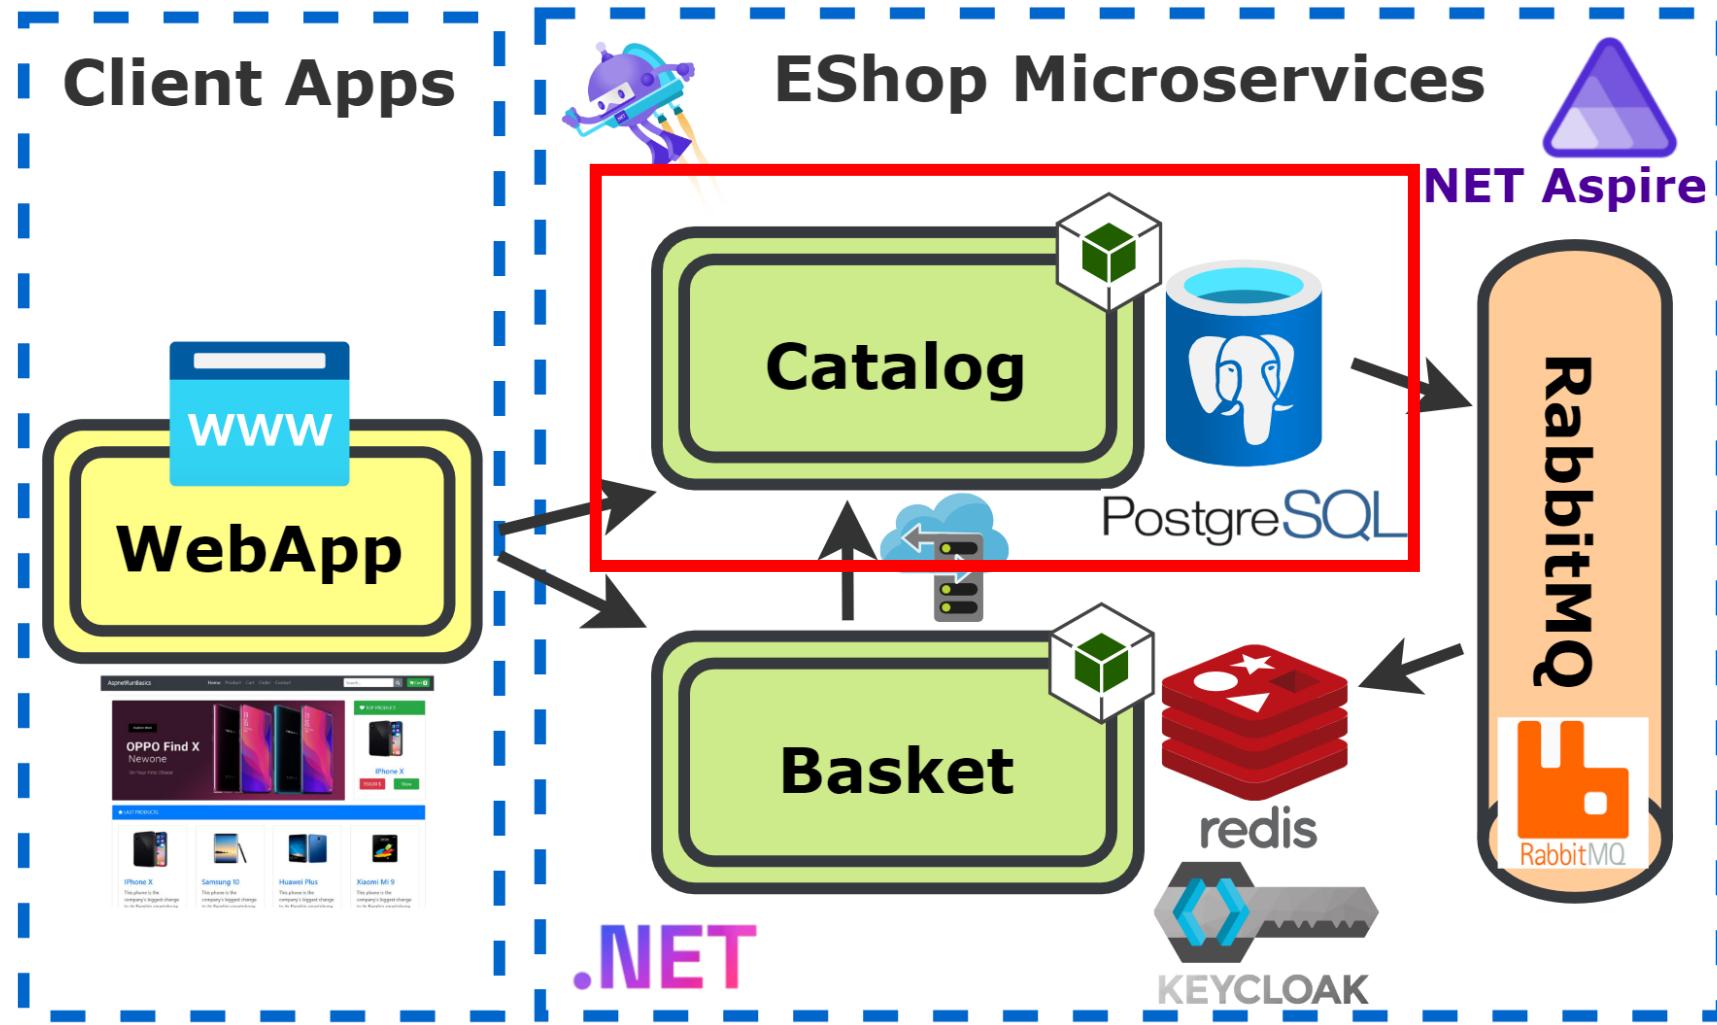
Catalog Client Integration: Connect PostgreSQL DB in Microservice

Develop and Inject EF Core ProductDbContext.cs in Data Layer

Mehmet Ozkaya



Catalog Microservice with PostgreSQL



Catalog Microservice with PostgreSQL

Domain Analysis of Catalog Microservices:
Models, UCs, Rest APIs, Databases

Technical Analysis of Catalog Microservices:
Architectures, Patterns, Libraries, Folders

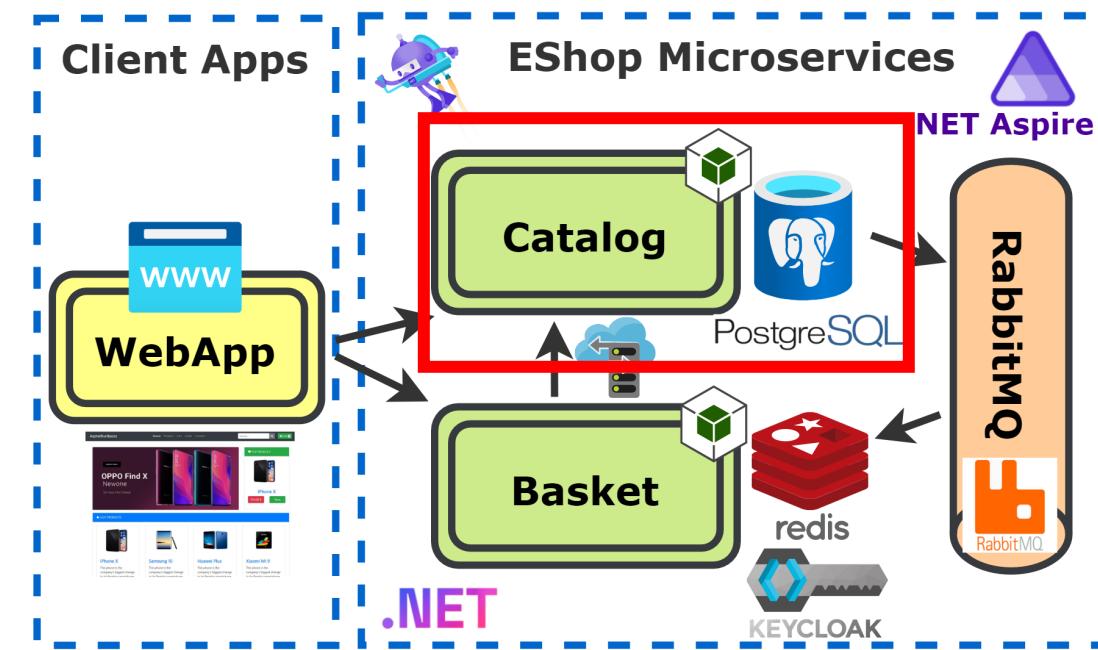
Develop Product Domain Entity Models

Hosting Integration(Backing Services)

Client Integration(EF Core Connection)

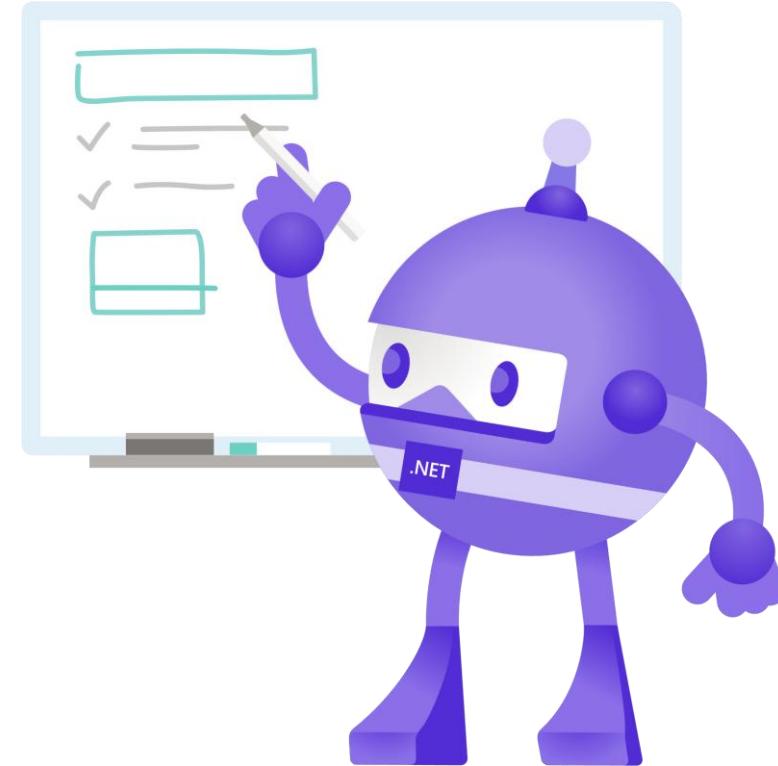
Catalog Application Use Case Development

**Develop Catalog API Endpoints Exposing
Minimal API Endpoints**



Domain Analysis of Catalog Microservice

1. Domain Models
2. Application Use Cases
3. Rest API Endpoints
4. Underlying Data Structures



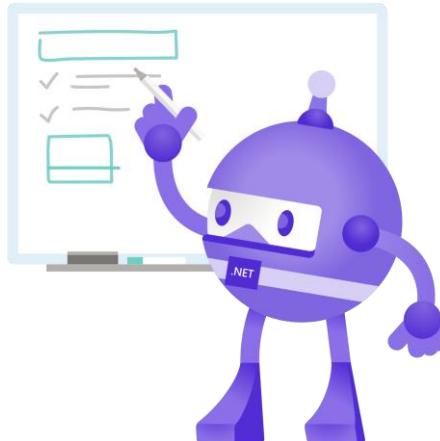
1. Domain Models of Catalog Microservice

- Primary domain model is 'Product'



```
public class Product
{
    public int Id { get; set; }
    public string Name { get; set; } = default!;
    public string Description { get; set; } = default!;
    public decimal Price { get; set; }
    public string ImageUrl { get; set; } = default!;
}
```

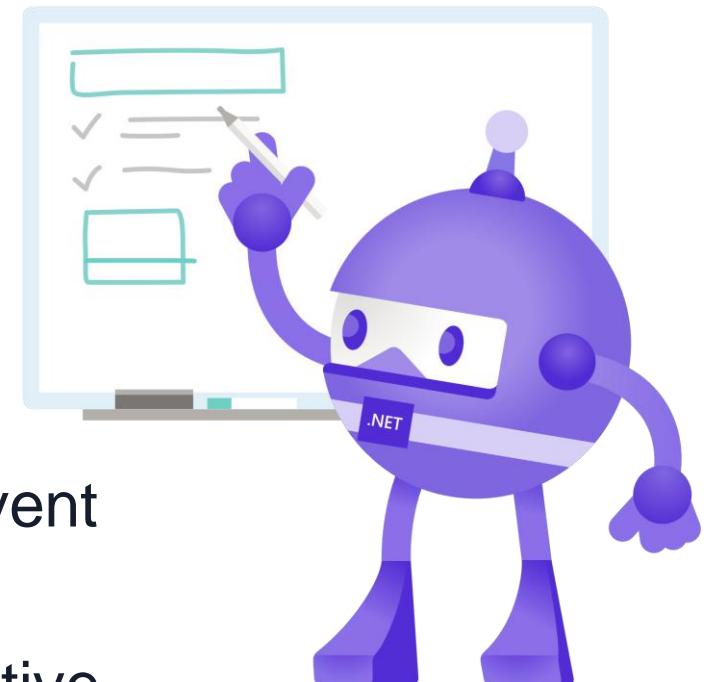
- Consider Domain event for **ProductPriceChanges**, leading to integration events.



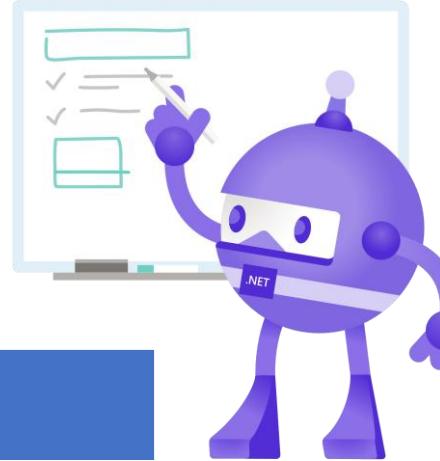
2. Application Use Cases of Catalog Microservice

CRUD Operations:

- Listing Products
- Get Product with Product Id
- Create new Product
- Update Product
- Delete Product
- If the price has changed, raise a **ProductPriceChanged** event that will lead to an integration event.
- Along with this design our APIs according to REST perspective.



3. Rest API Endpoints of Catalog Microservices



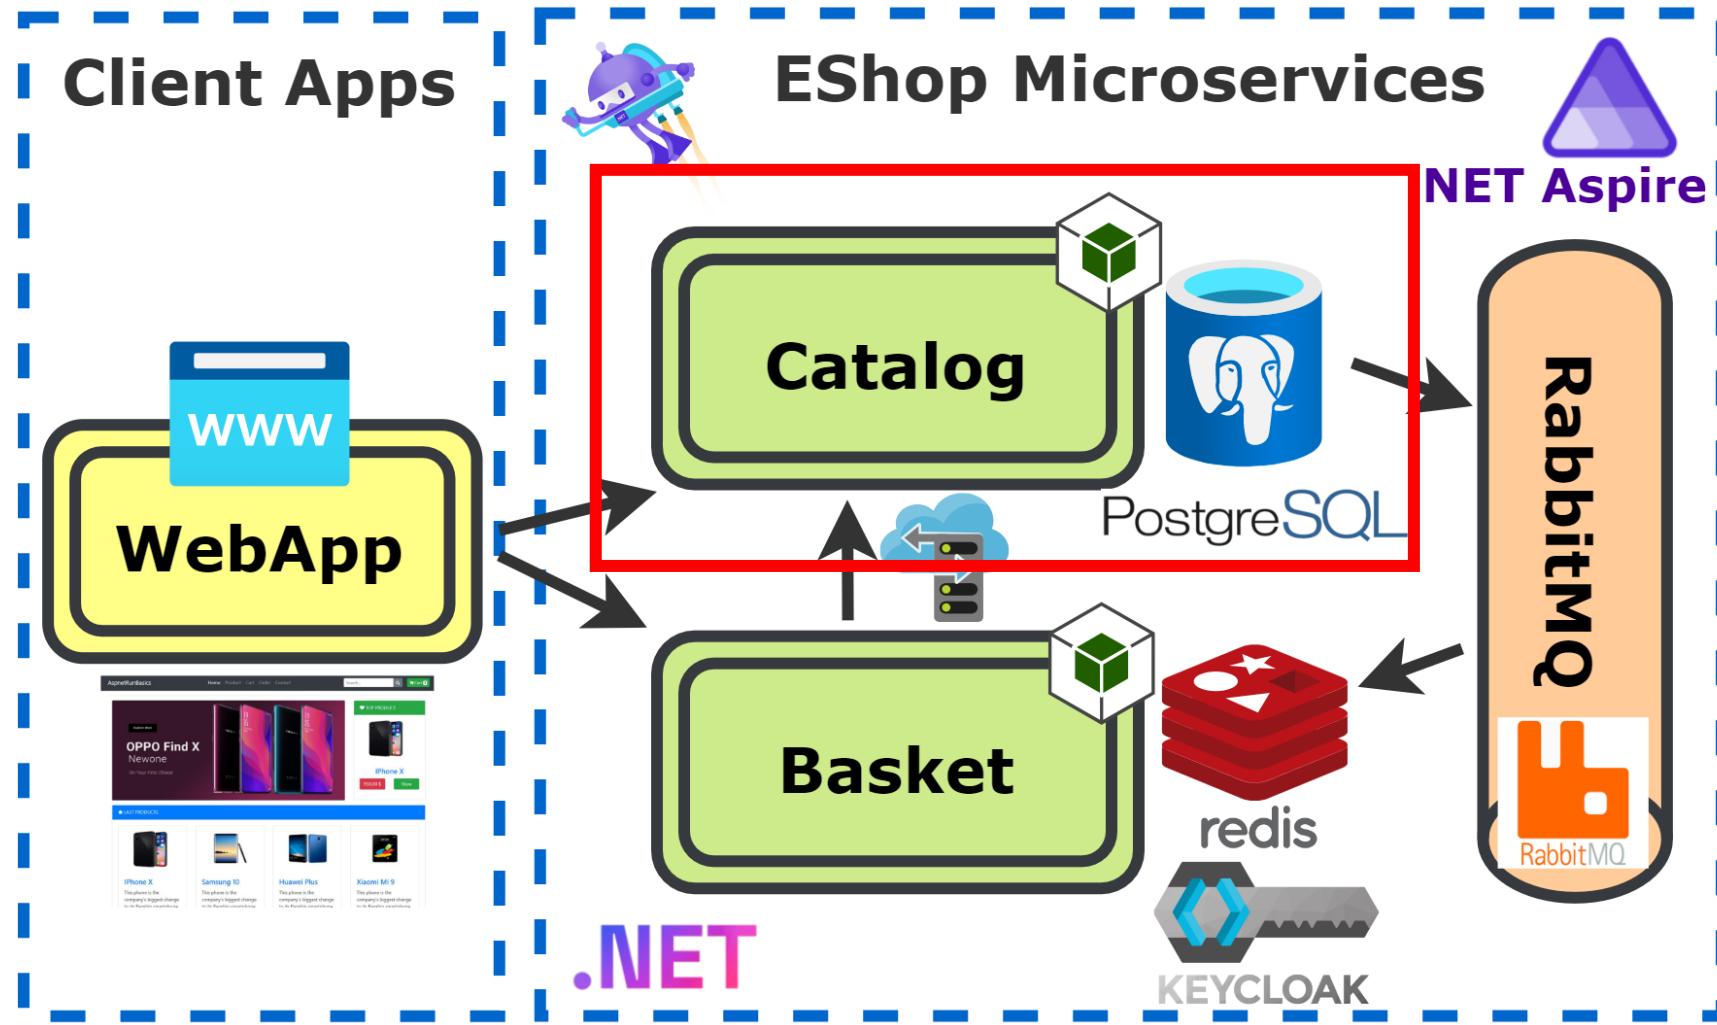
Method	Request URI	Use Cases
GET	/products	List all products
GET	/products/{id}	Fetch a specific product
POST	/products	Create a new product
PUT	/products/{id}	Update a product
DELETE	/products/{id}	Remove a product

4. Underlying Data Structures of Catalog Microservices

- **Catalog Microservice** will be stored in a **PostgreSQL** relational database.
- Use **Entity Framework Core's Code-First approach** for **database migrations** and **entity configurations**.
- Configure our **entities** and their **relationships** using **EF Core**.

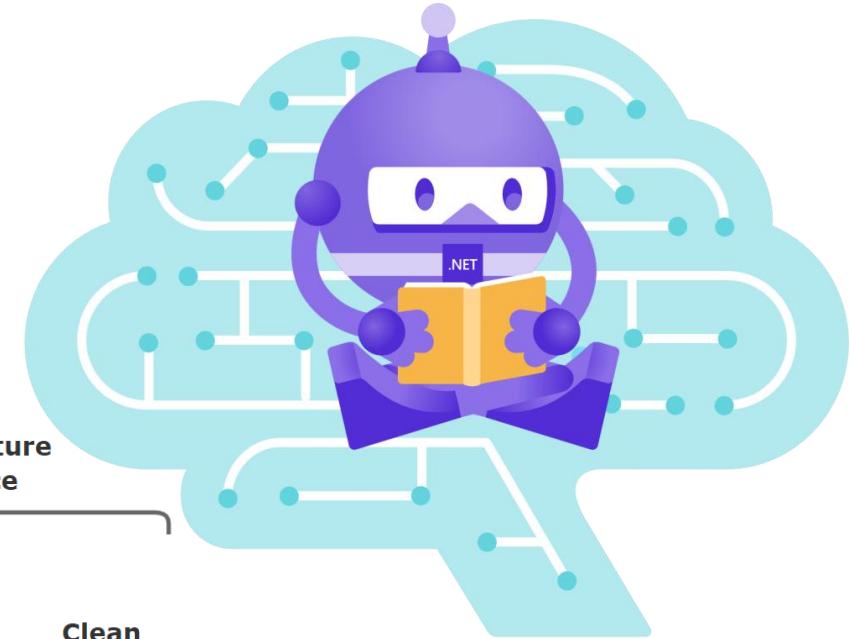
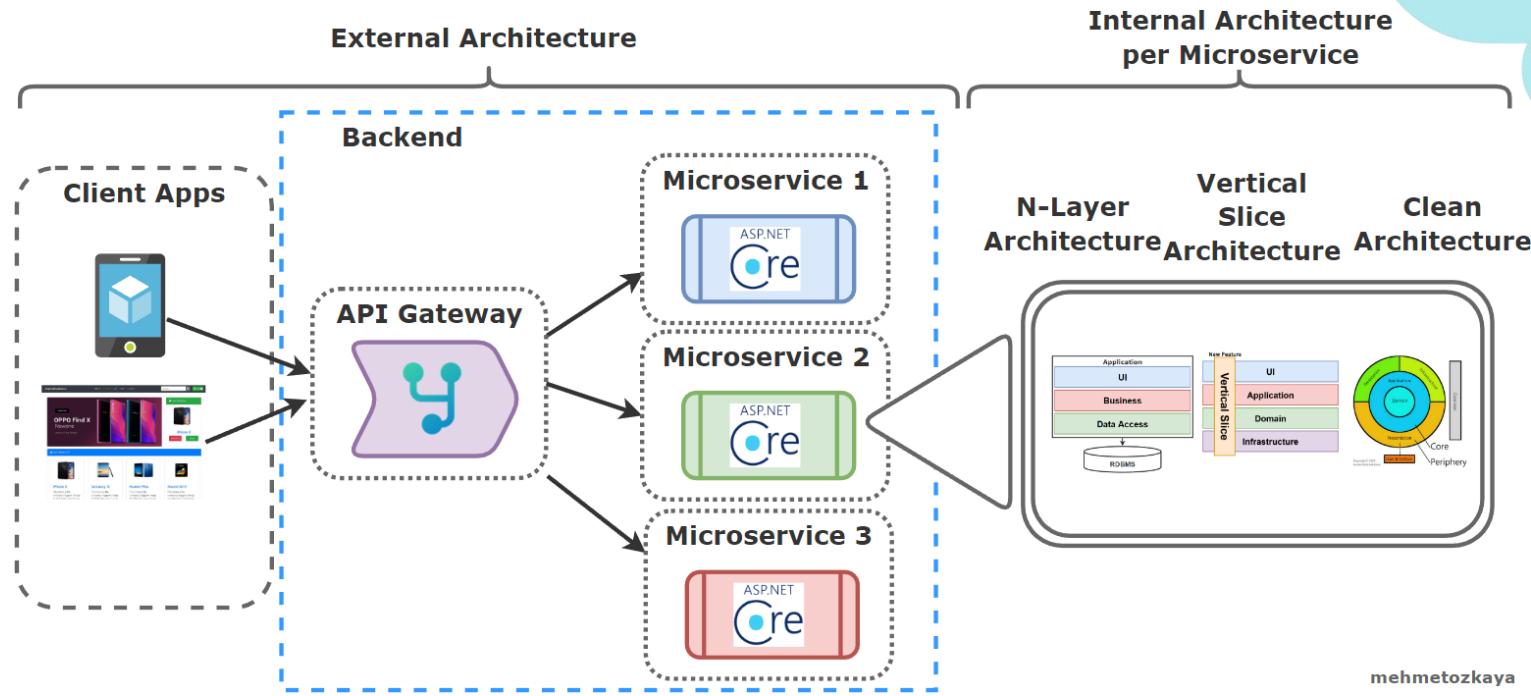


Catalog Microservice with PostgreSQL



Technical Analysis of Catalog Microservice

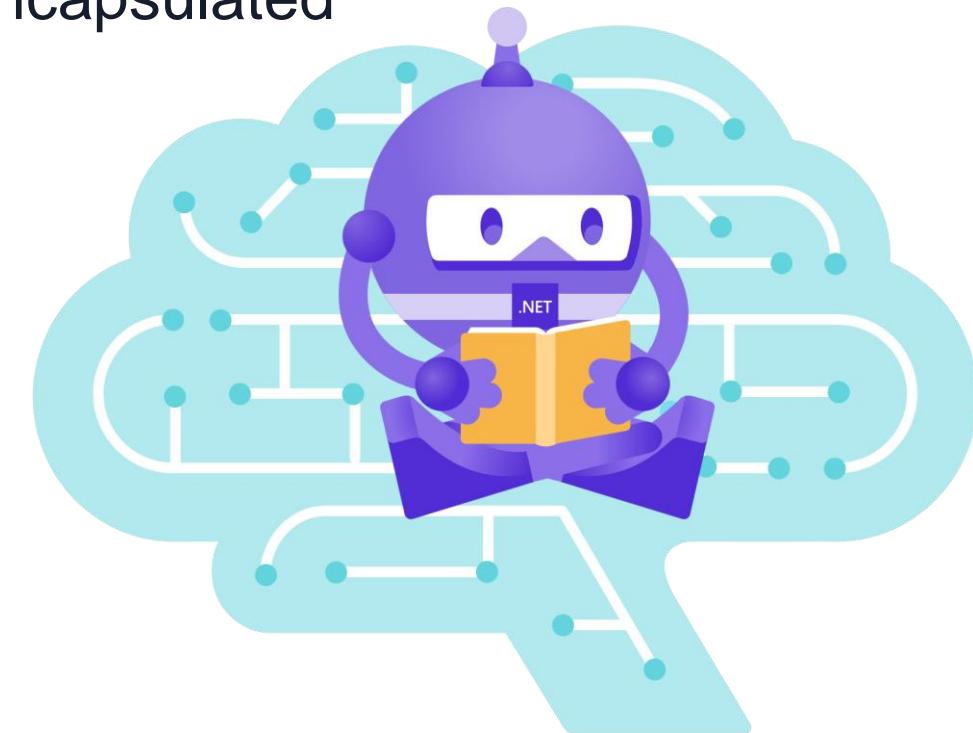
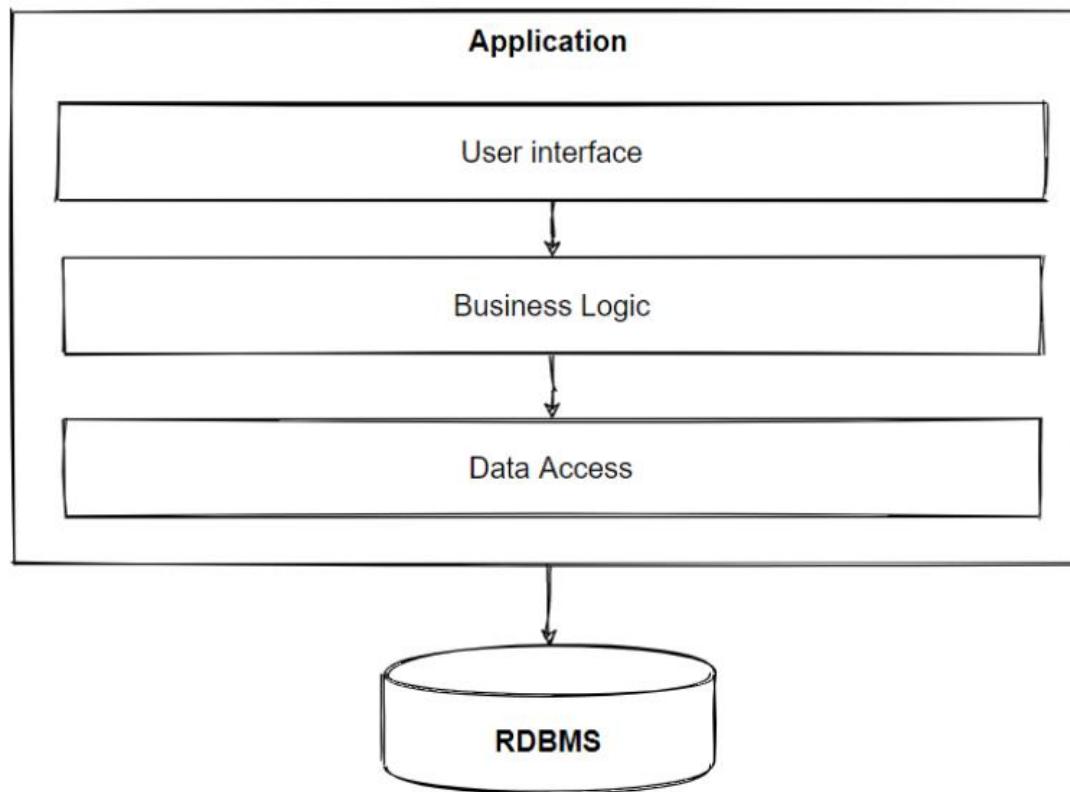
1. Application Architecture Style
2. Patterns & Principles
3. Libraries Nuget Packages
4. Project Folder Structure



1. Application Architecture Style of Catalog Microservice

Classical N-Layer Architecture

Organizes our code into layer folders, each feature encapsulated in a folder



N-Layer(Layered) Architecture

Domain Layer

Stores and Handles domain entities and domain events

Data Access Layer

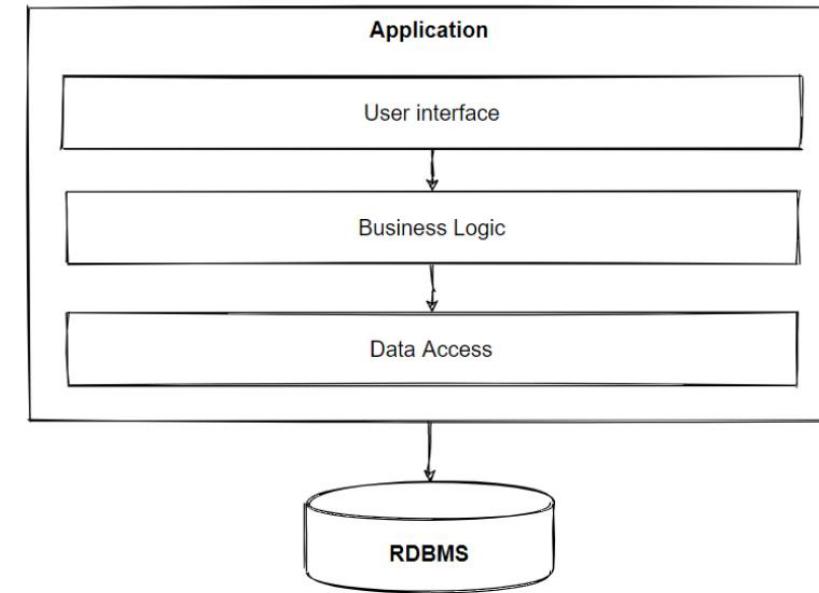
Manages interactions with the database or other storage systems

Business Logic Layer

Processes the application's core business logic and rules

Presentation Layer

Displaying data to users and capturing user input



2. Patterns & Principles of Catalog Microservice

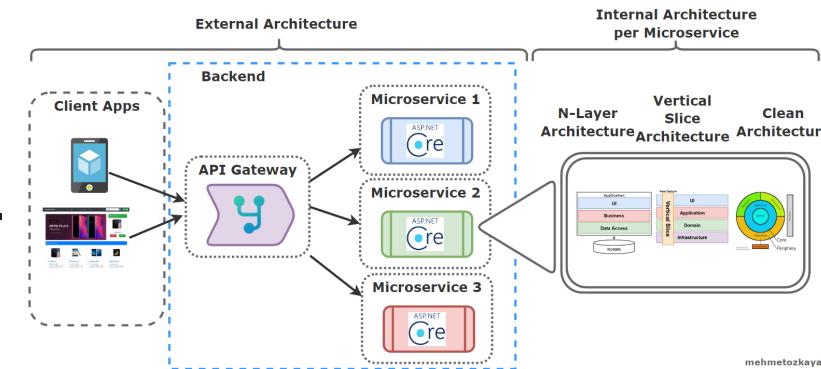
KISS principles and **use less patterns** in our **microservices**

Focus on integrations and orchestrations with **.NET Aspire**

DI in ASP.NET Core: Dependency Injection is a core feature, allowing us to inject dependencies.

Minimal APIs and Routing in ASP.NET 8: ASP.NET Minimal APIs simplify endpoint definitions, providing lightweight syntax for routing and handling HTTP requests.

ORM Pattern: Object-Relational Mapping abstracts database interactions, work with database objects using high-level codes.



3. Libraries Nuget Packages of Catalog Microservice



- **Aspire Hosting Packages for Orchestration:**
Aspire.Hosting.PostgreSQL
- **Aspire EF Core for PostgreSQL Interaction: (Client Integration)** : Aspire.Npgsql.EntityFrameworkCore.PostgreSQL
- EF Core provides a robust ORM for interacting with PostgreSQL, enabling us to manage database operations with ease.



4. Project Folder Structure of Catalog Microservice

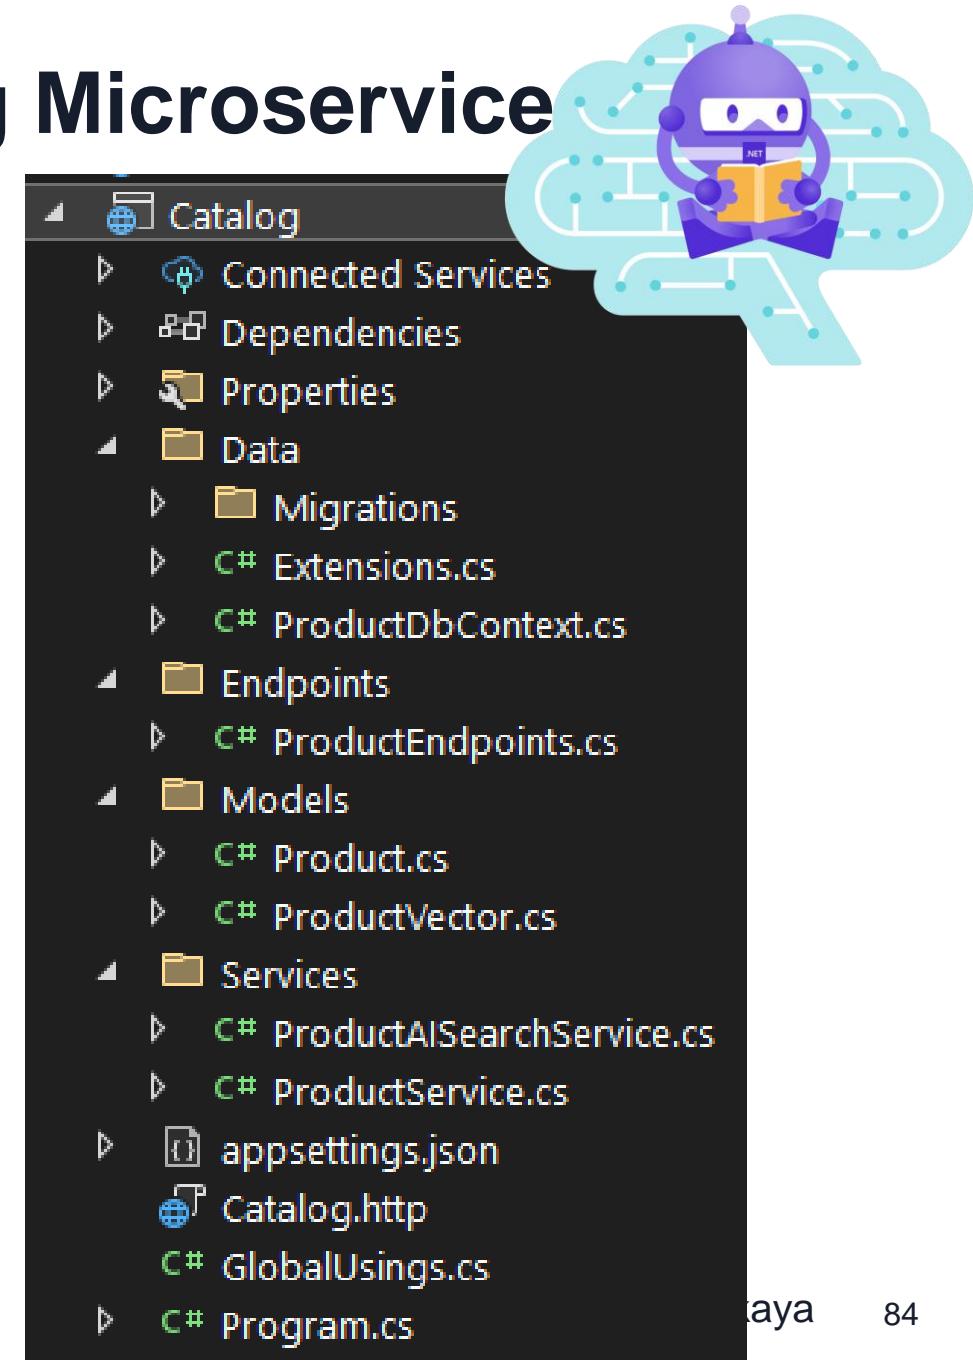
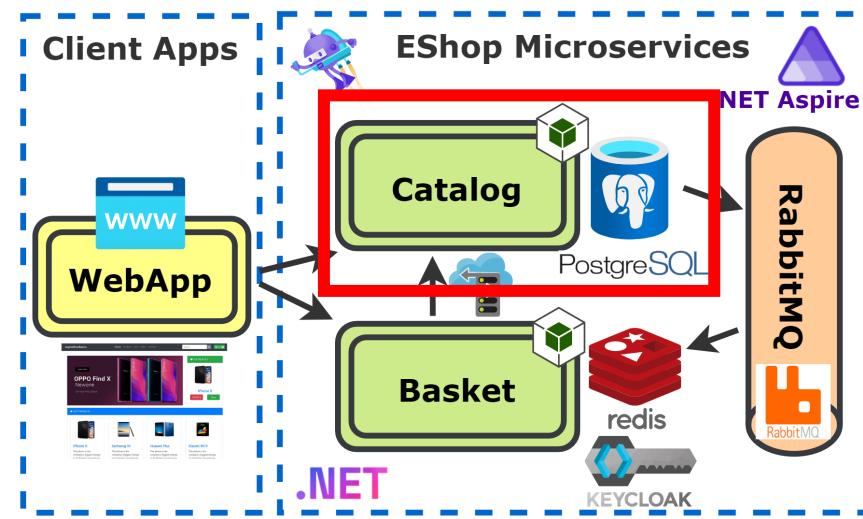
Organized into **Models**, **Data**, **Services** and **Endpoints**

Models holds our **domain entities** (e.g., Product)

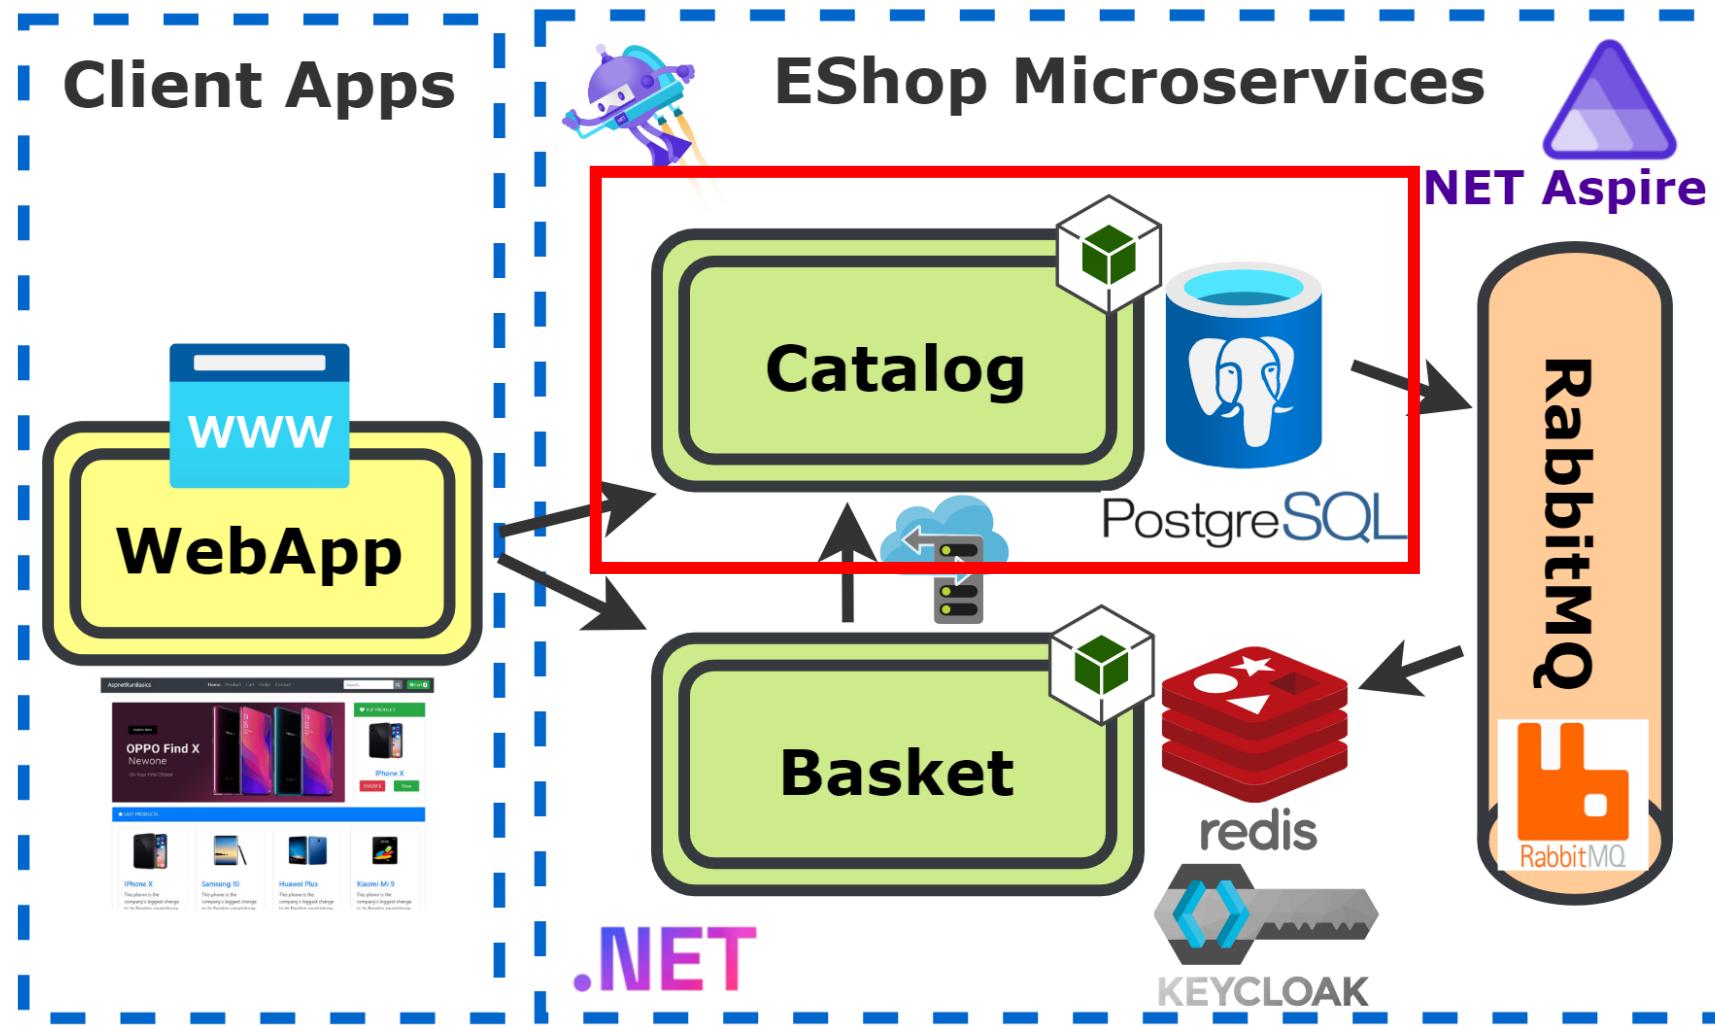
Data includes **DbContext** classes for data access logic

Services contains **business logic**

Endpoints minimal API endpoints that call into the services layer



Catalog Microservice with PostgreSQL

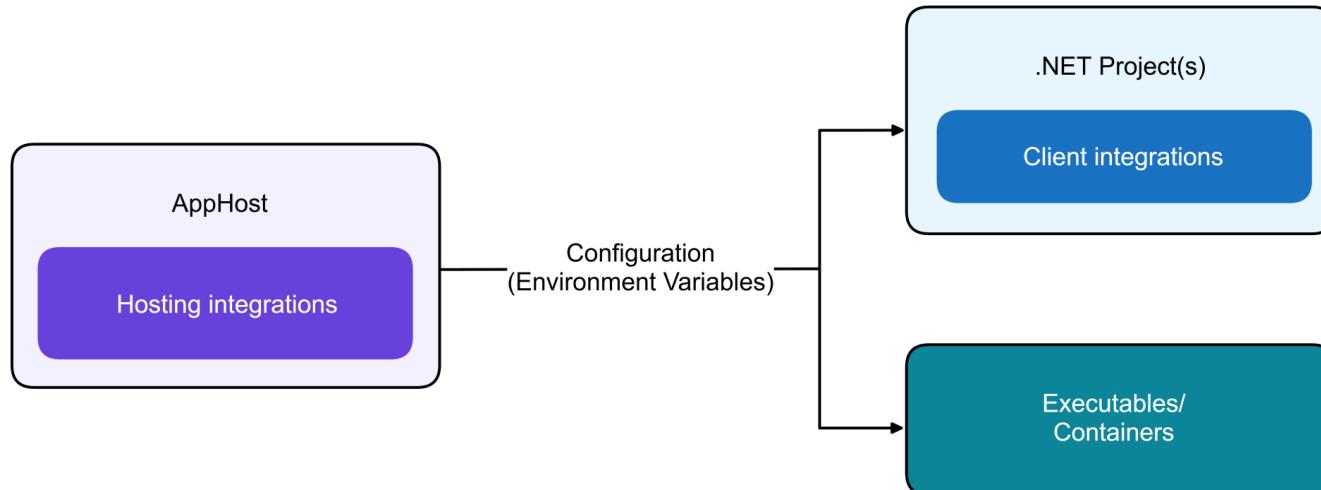
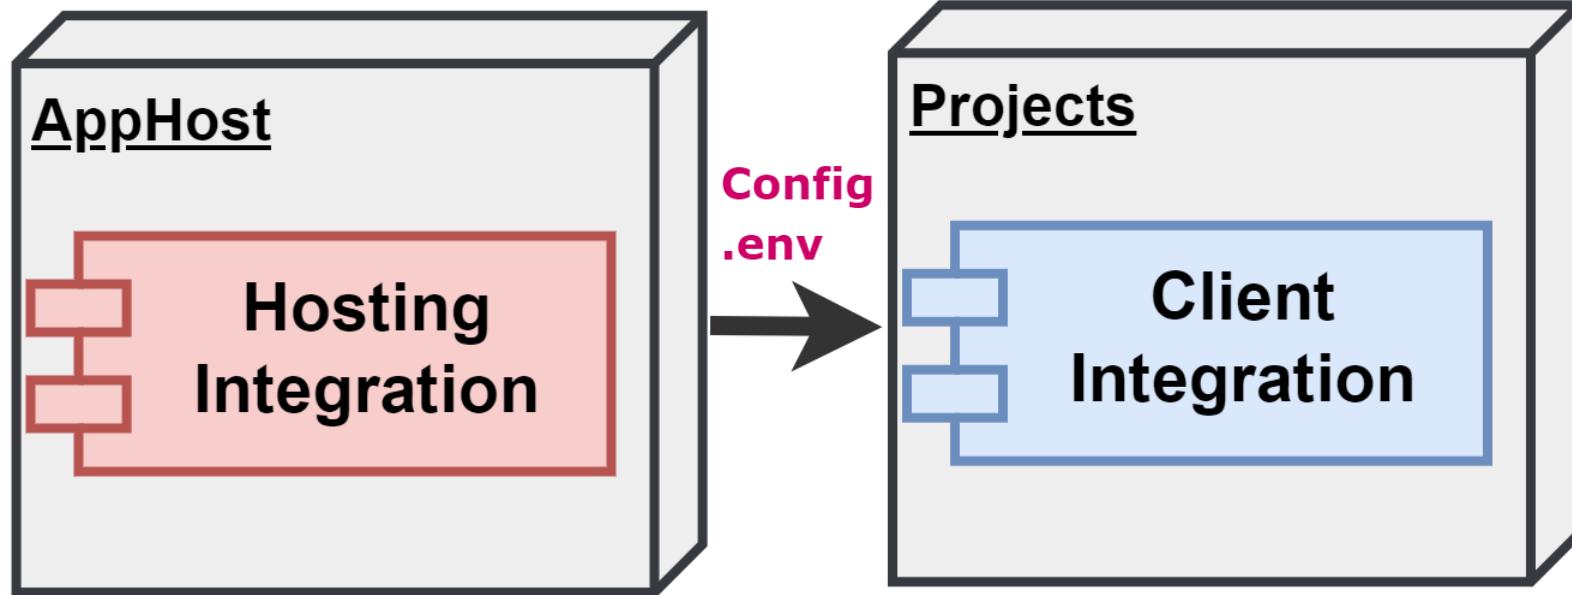


Hosting vs. Client Integrations



.NET Aspire

Integrations



Hosting Integrations – Provisioning Resources

Extend the **IDistributedApplicationBuilder** interface, adding methods like **AddRedis("name")** or **AddPostgres("db")**

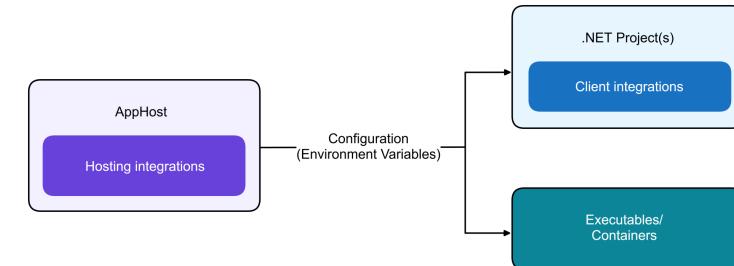
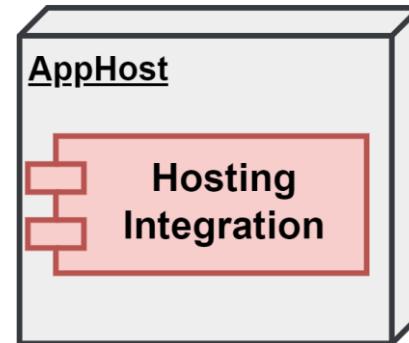


Provision container-based resources automatically

Use **environment variables** to pass **connection strings** to other projects that depend on them



```
var builder =  
    DistributedApplication.CreateBuilder(args);  
    // Hosting integration call:  
    var cache = builder.AddRedis("cache");  
  
    // This sets up a Redis container named "cache"
```



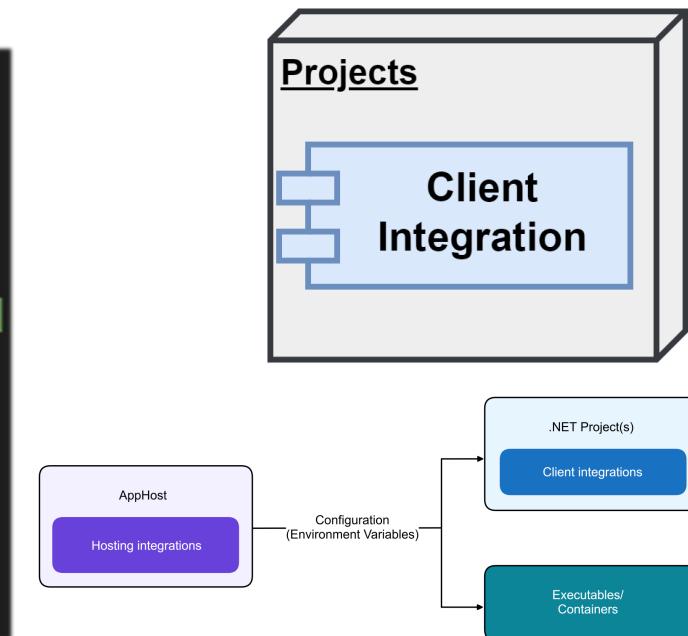
Client Integrations – Connecting to Resources

Extend **IHostApplicationBuilder** in Consuming Projects;
builder.Services.AddStackExchangeRedis(...)

Add Health Checks, Logging, Telemetry

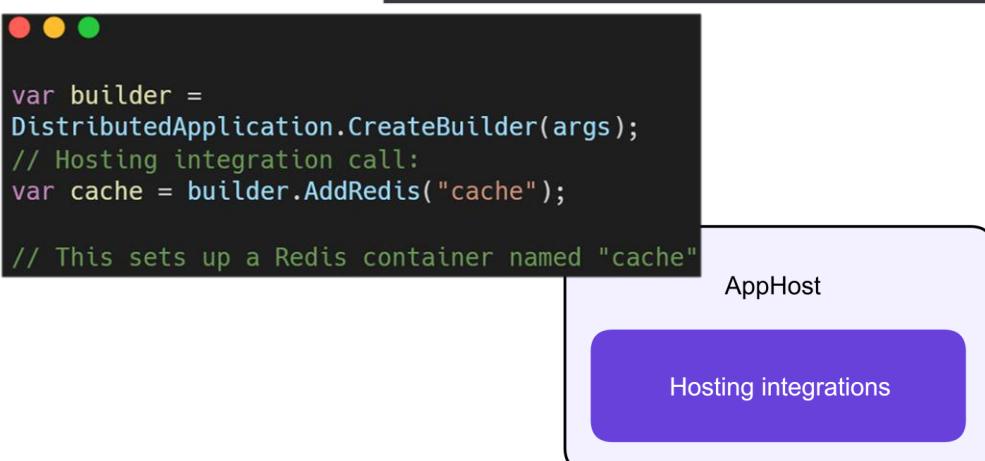
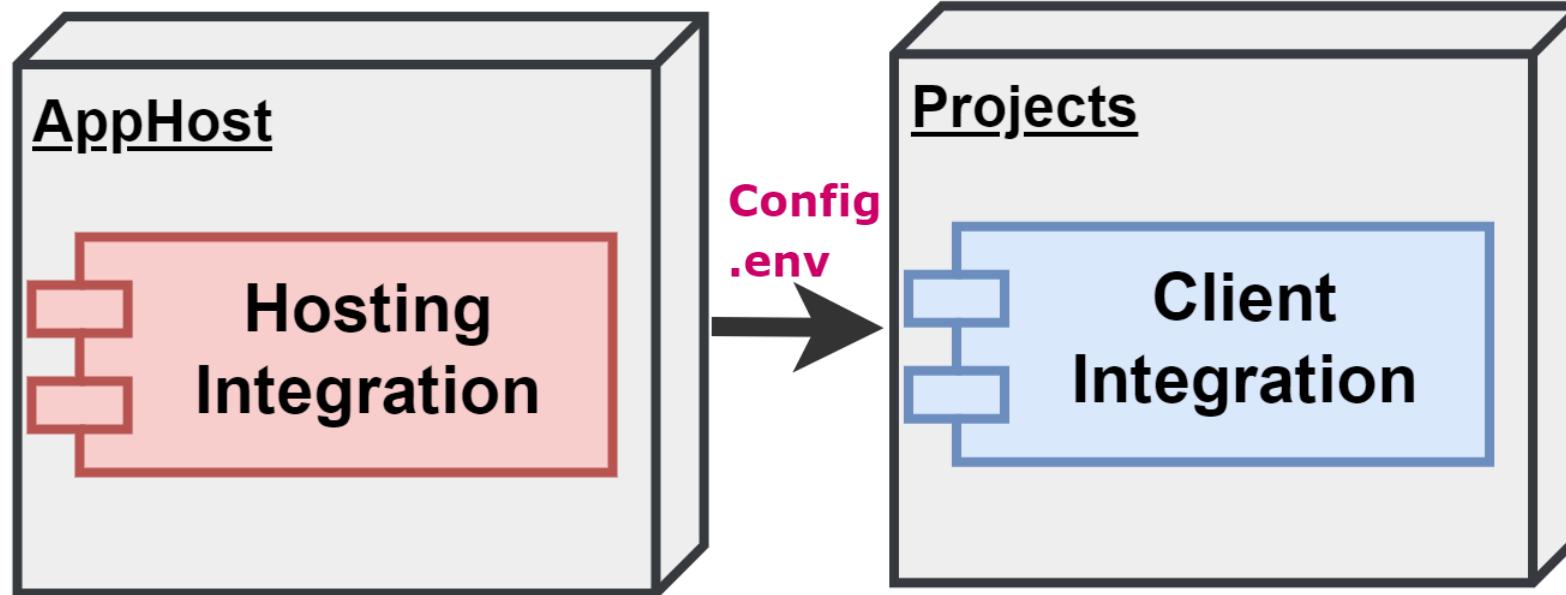
Example: Aspire.StackExchange.Redis

```
● ○ ●  
var builder = WebApplication.CreateBuilder(args);  
  
// AddServiceDefaults is often called automatically in Aspire-based  
// projects; you can call it explicitly, too.  
builder.AddServiceDefaults();  
  
// This might be from Aspire.StackExchange.Redis  
builder.AddRedisDistributedCache(connectionName: "cache");
```

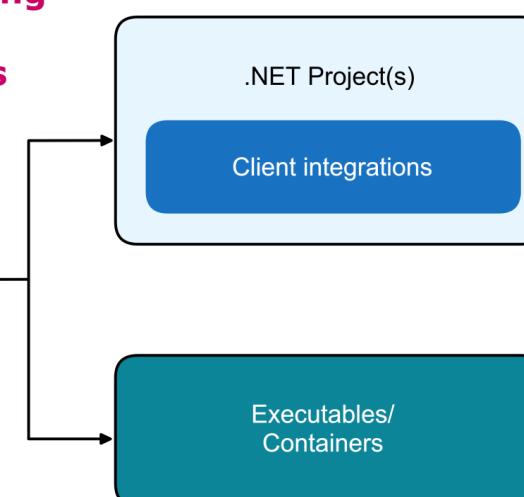


Relationship Hosting vs. Client Integrations

Integrations

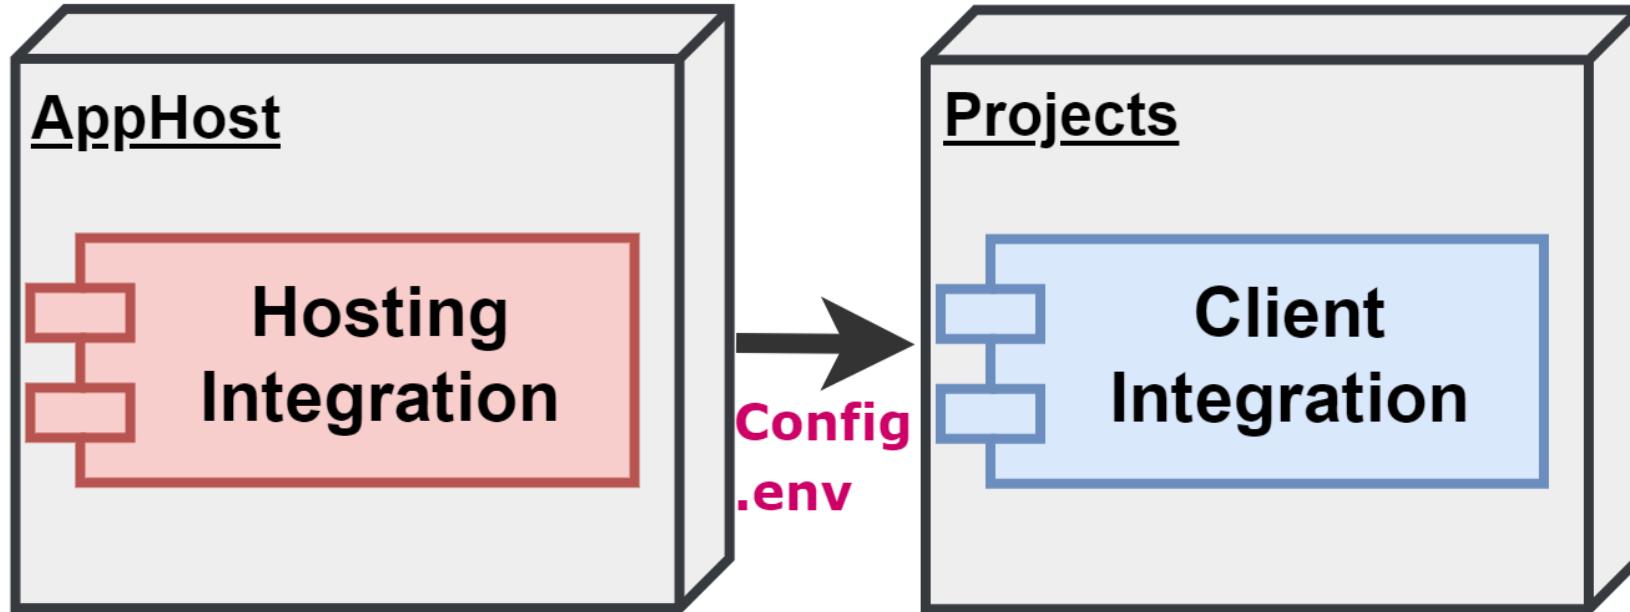


Observability
Logging, tracing
and metrics
Health Checks
Resiliency



Conclusion: Hosting vs. Client Integrations

Integrations



```
var builder =  
DistributedApplication.CreateBuilder(args);  
// Hosting integration call:  
var cache = builder.AddRedis("cache");  
  
// This sets up a Redis container named "cache"
```



```
var builder = WebApplication.CreateBuilder(args);  
  
// AddServiceDefaults is often called automatically  
builder.AddServiceDefaults();  
  
// This might be from Aspire.StackExchange.Redis  
builder.AddRedisDistributedCache(connectionName:  
"cache");
```

Official Integrations

Redis, PostgreSQL, RabbitMQ, Kafka, Keycloak

Azure & AWS Packages

Community Toolkit

[Official Integrations Table](#)

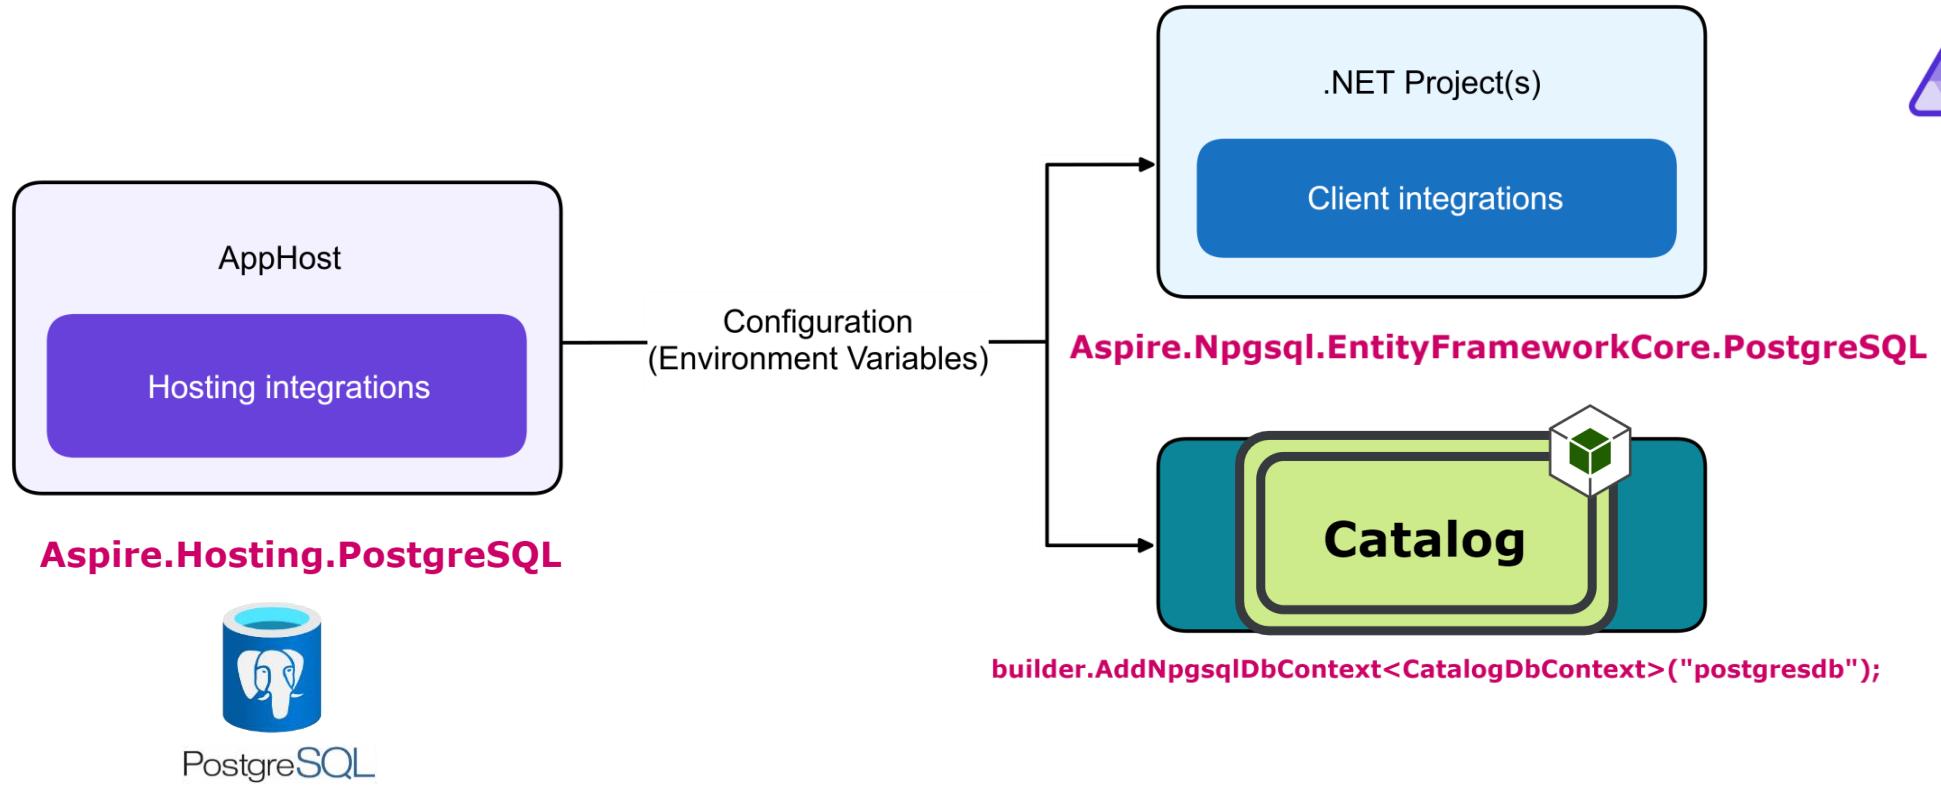
Integration docs and NuGet packages	Description
<ul style="list-style-type: none"> - Learn more: Apache Kafka - Hosting: Aspire.Hosting.Kafka - Client: Aspire.Confluent.Kafka 	A library for producing and consuming messages from an Apache Kafka broker.
<ul style="list-style-type: none"> - Learn more: Dapr - Hosting: Aspire.Hosting.Dapr - Client: N/A 	A library for modeling Dapr as a .NET Aspire resource.
<ul style="list-style-type: none"> - Learn more: Elasticsearch - Hosting: Aspire.Hosting.Elasticsearch - Client: Aspire.Elastic.Clients.Elasticsearch 	A library for accessing Elasticsearch databases.
<ul style="list-style-type: none"> - Learn more: Keycloak - Hosting: Aspire.Hosting.Keycloak - Client: Aspire.Keycloak.Authentication 	A library for accessing Keycloak authentication.
<ul style="list-style-type: none"> - Learn more: Milvus - Hosting: Aspire.Hosting.Milvus - Client: Aspire.Milvus.Client 	A library for accessing Milvus databases.
<ul style="list-style-type: none"> - Learn more: MongoDB Driver - Hosting: Aspire.Hosting.MongoDB - Client: Aspire.MongoDB.Driver 	A library for accessing MongoDB databases.



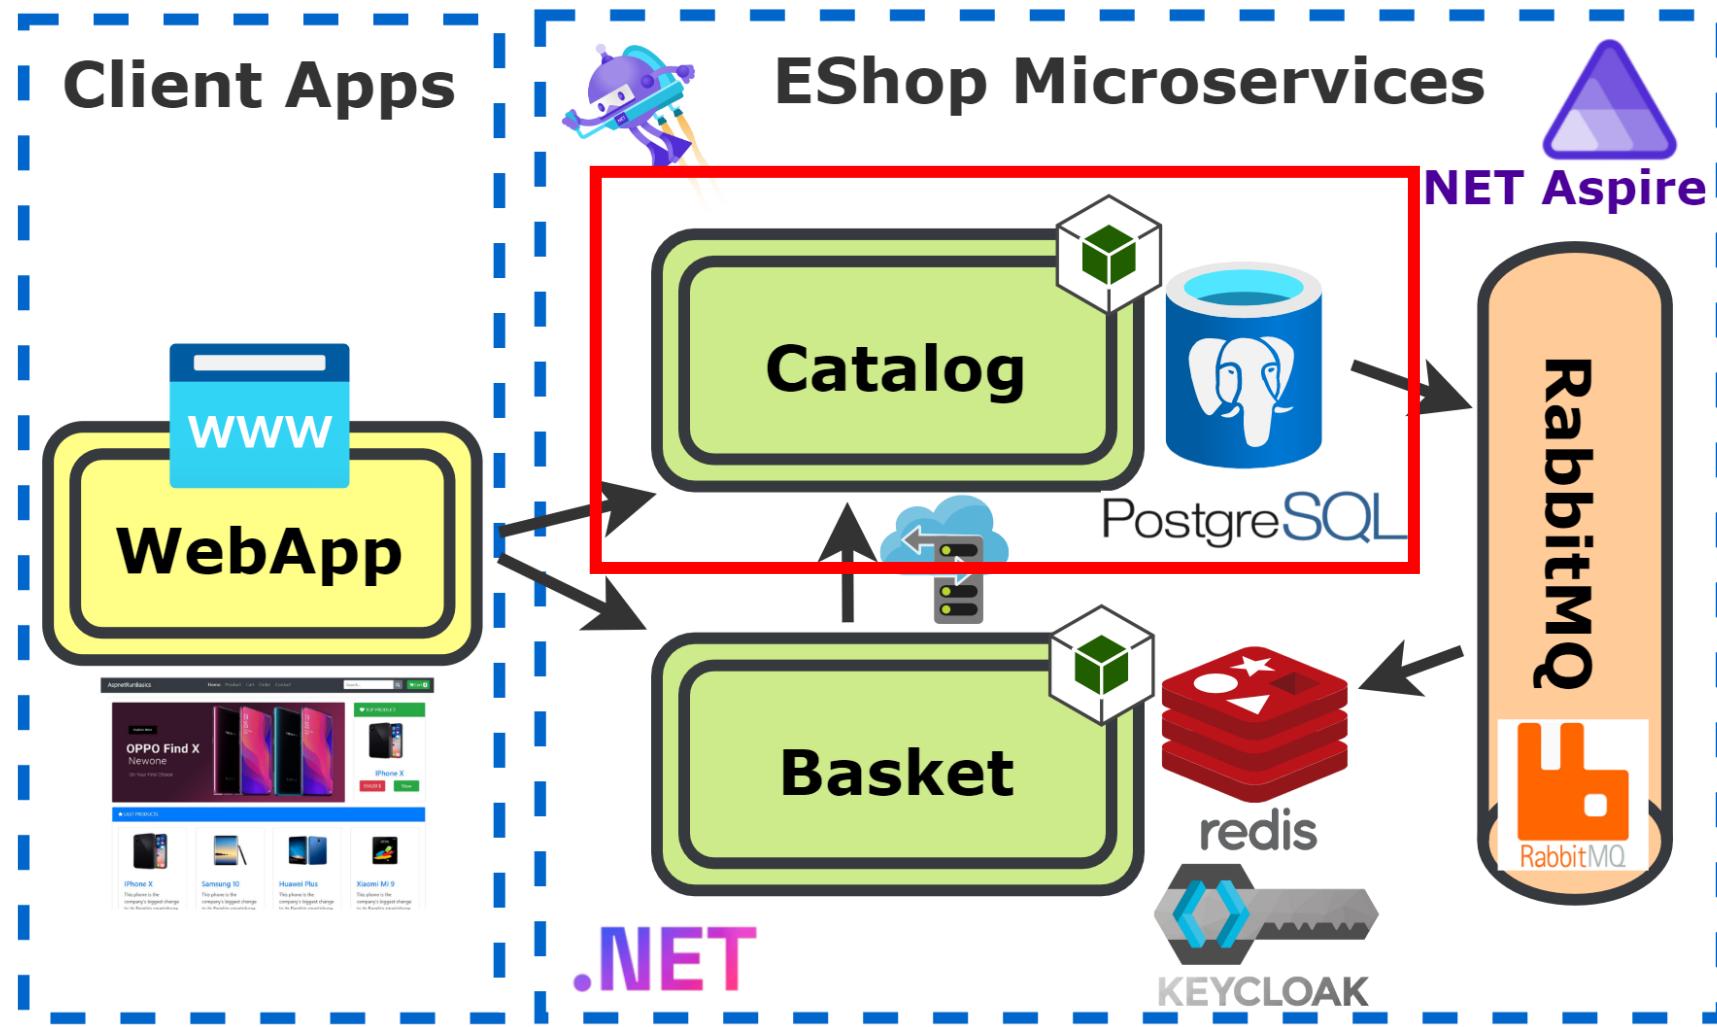
.NET Aspire

Catalog Microservice with PostgreSQL

Integrations



Catalog Microservice with PostgreSQL



Develop Basket Microservice with Redis orchestrate in .NET Aspire

Domain and Technical Analysis of Basket Microservice

Basket Hosting Integration: Add Redis Resource in AppHost

Basket Client Integration: Connect Redis Cache in Microservice

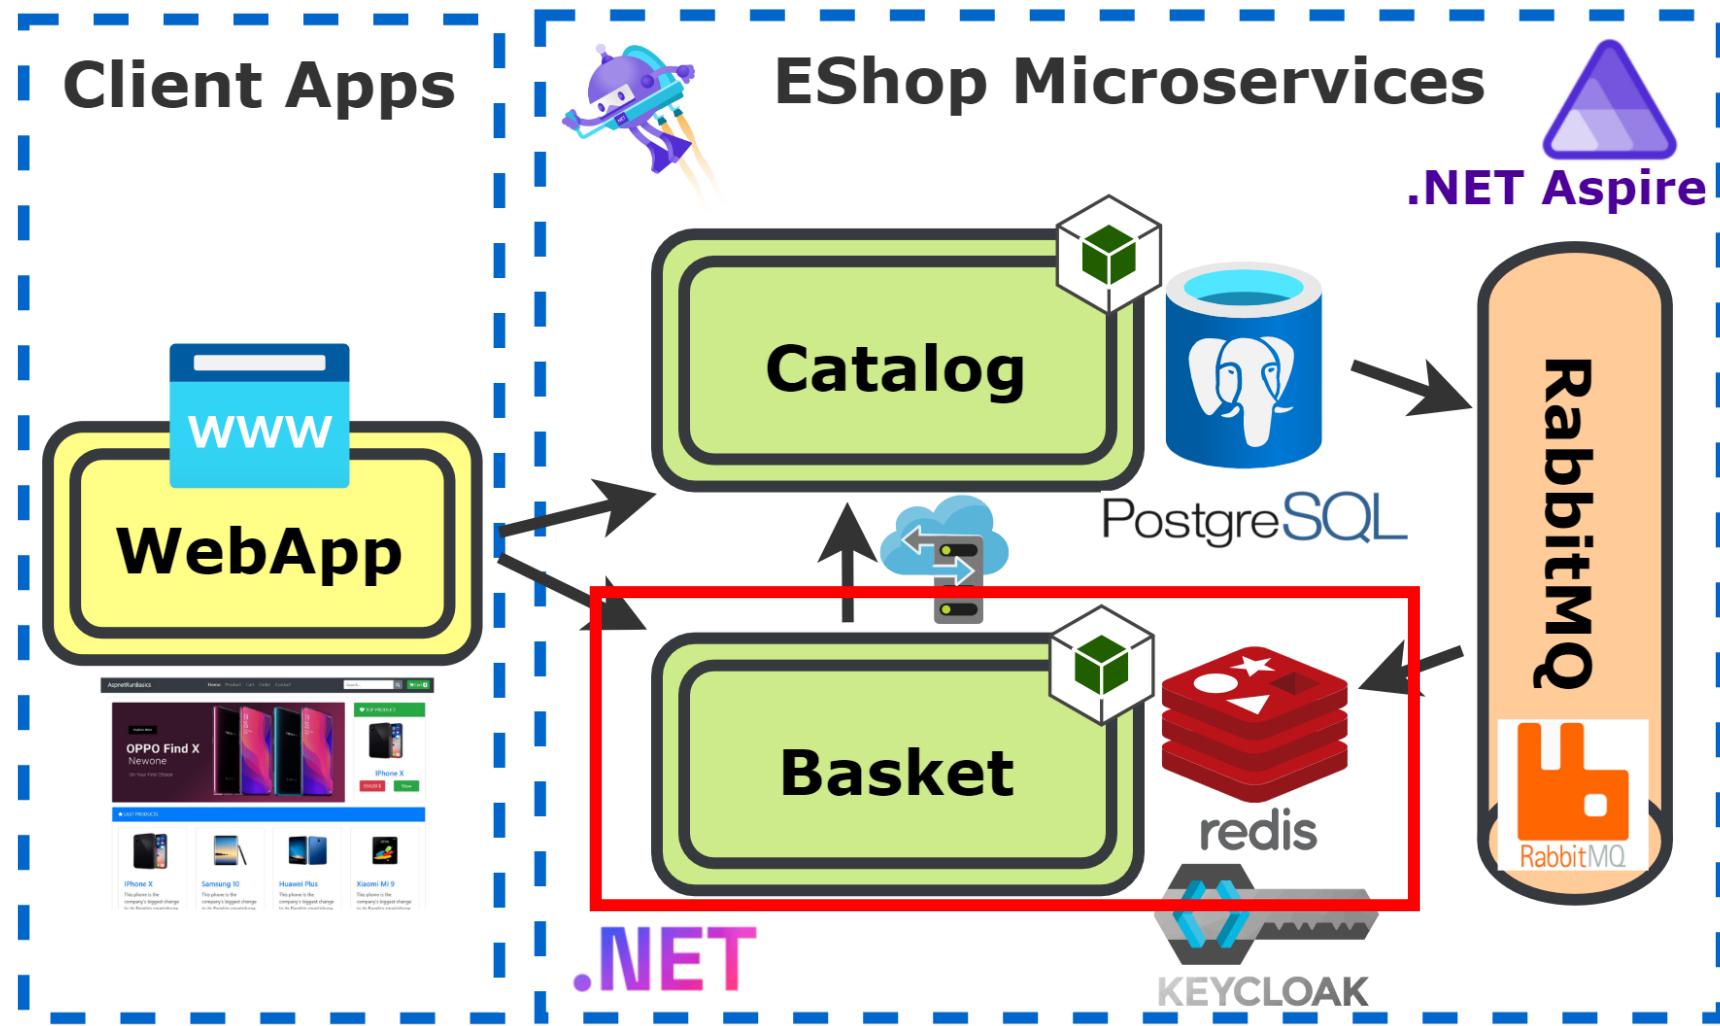
Develop BasketService class with injecting IDistributedCache

Develop Basket API Endpoints with ASP.NET Minimal APIs

Mehmet Ozkaya



Basket Microservice with Redis



Basket Microservice with Redis

Domain Analysis of Basket Microservices:
Models, UCs, Rest APIs, Databases

Technical Analysis of Basket Microservices:
Architectures, Patterns, Libraries, Folders

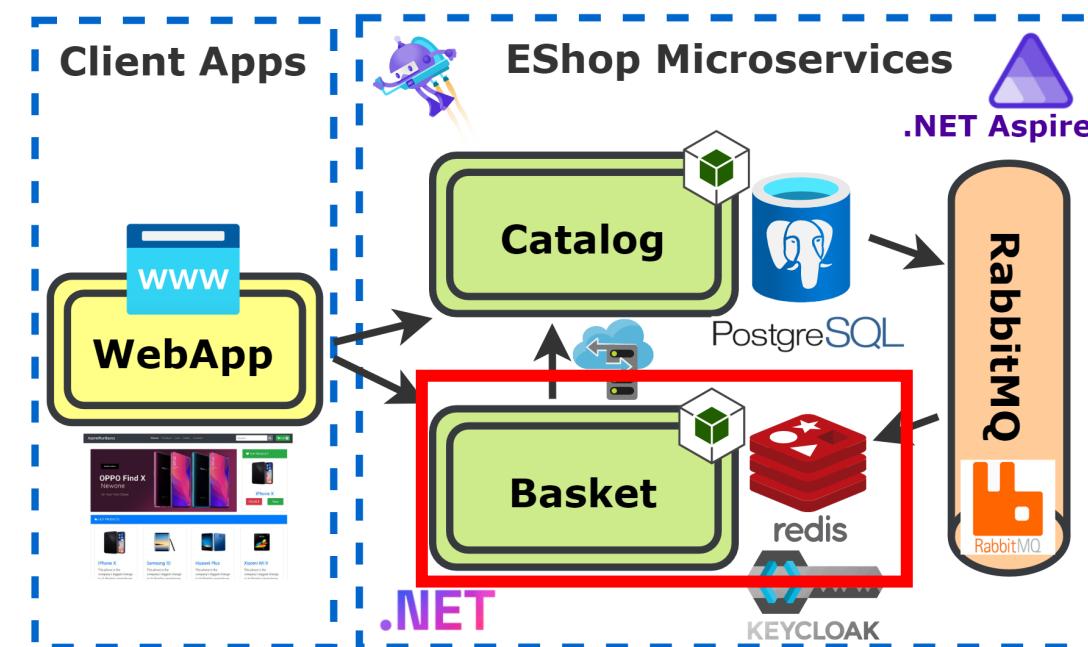
Develop **Basket Domain Entity Models**

Hosting Integration(Backing Services)

Client Integration(StackExchange Redis)

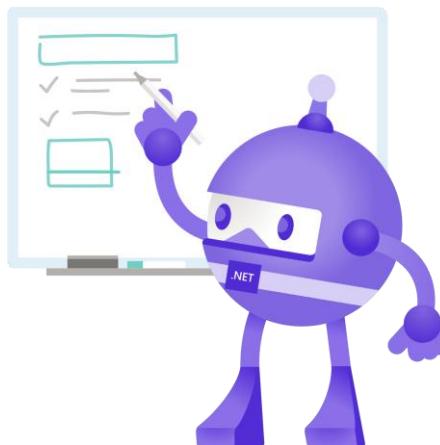
Basket Application Use Case Development

Develop **Basket API Endpoints** Exposing Minimal
API Endpoints



1. Domain Models of Basket Microservice

- Primary domain model is '**ShoppingCart**' and '**ShoppingCartItem**'



```
public class ShoppingCart
{
    public string UserName { get; set; } = default!;
    public List<ShoppingCartItem> Items { get; set; } = new();
    public decimal TotalPrice => Items.Sum(x => x.Price * x.Quantity);
}
```

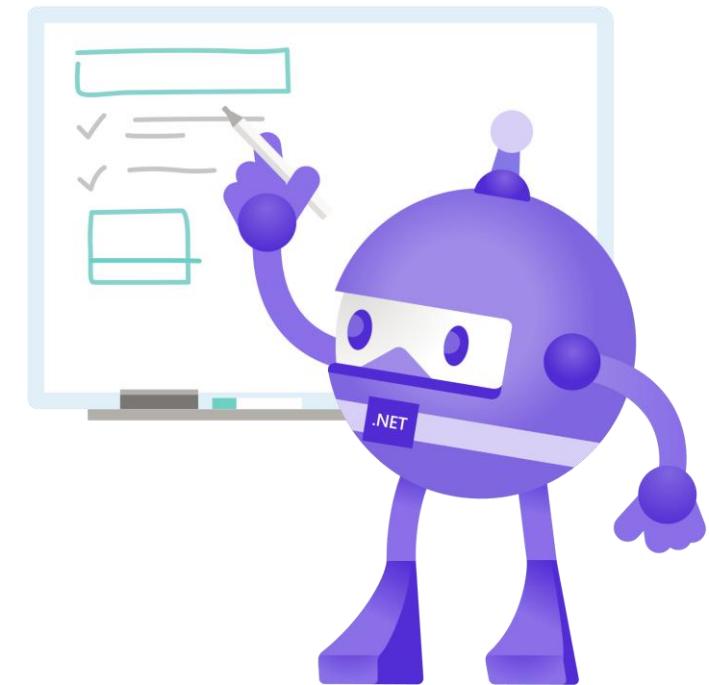
```
public class ShoppingCartItem
{
    public int Quantity { get; set; } = default!;
    public string Color { get; set; } = default!;
    public int ProductId { get; set; } = default!;

    // will comes from Catalog module
    public decimal Price { get; set; } = default!;
    public string ProductName { get; set; } = default!;
}
```

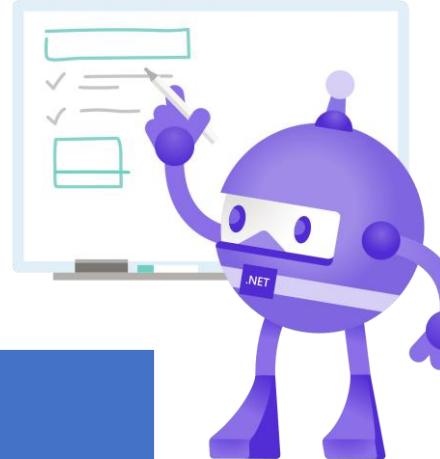
2. Application Use Cases of Basket Microservice

CRUD Operations:

- GetBasket
 - UpdateBasket - whole basket
 - DeleteBasket
 - UpdateBasketItemProductPrices
-
- Before update SC, we should **call Catalog Microservice GetProductById endpoint sync**
 - **Get latest product information and set Price and ProductName** when adding



3. Rest API Endpoints of Basket Microservice



Method	Request URI	Use Cases
GET	/basket/{userName}	Fetch a specific basket
POST	/basket	Create or update basket
DELETE	/basket/{id}	Remove a basket

4. Underlying Data Structures of Basket Microservices

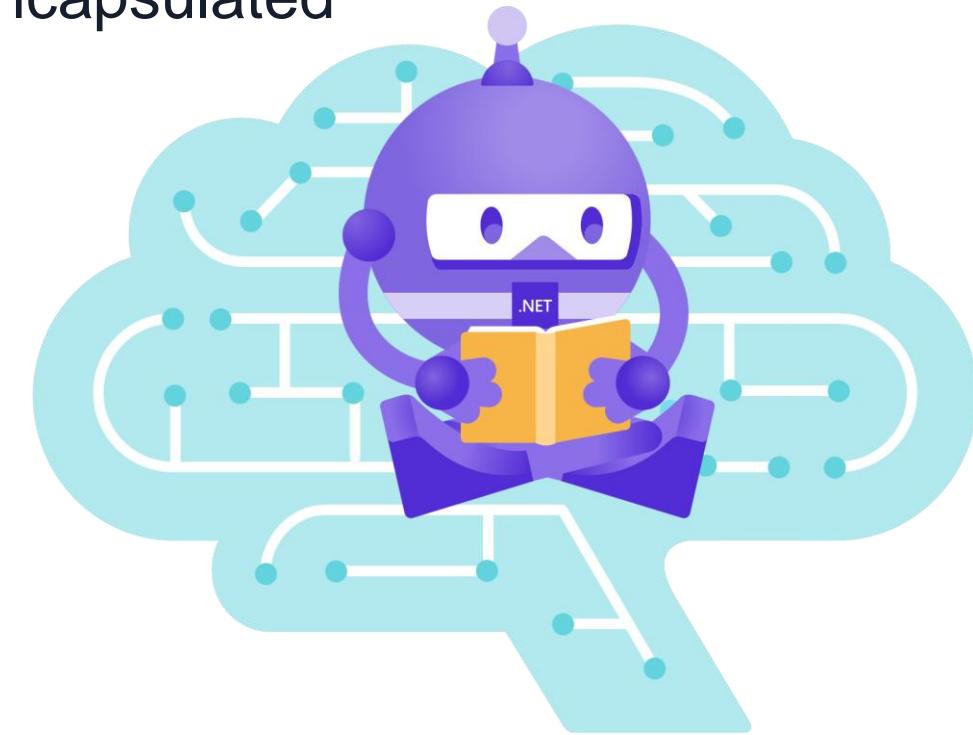
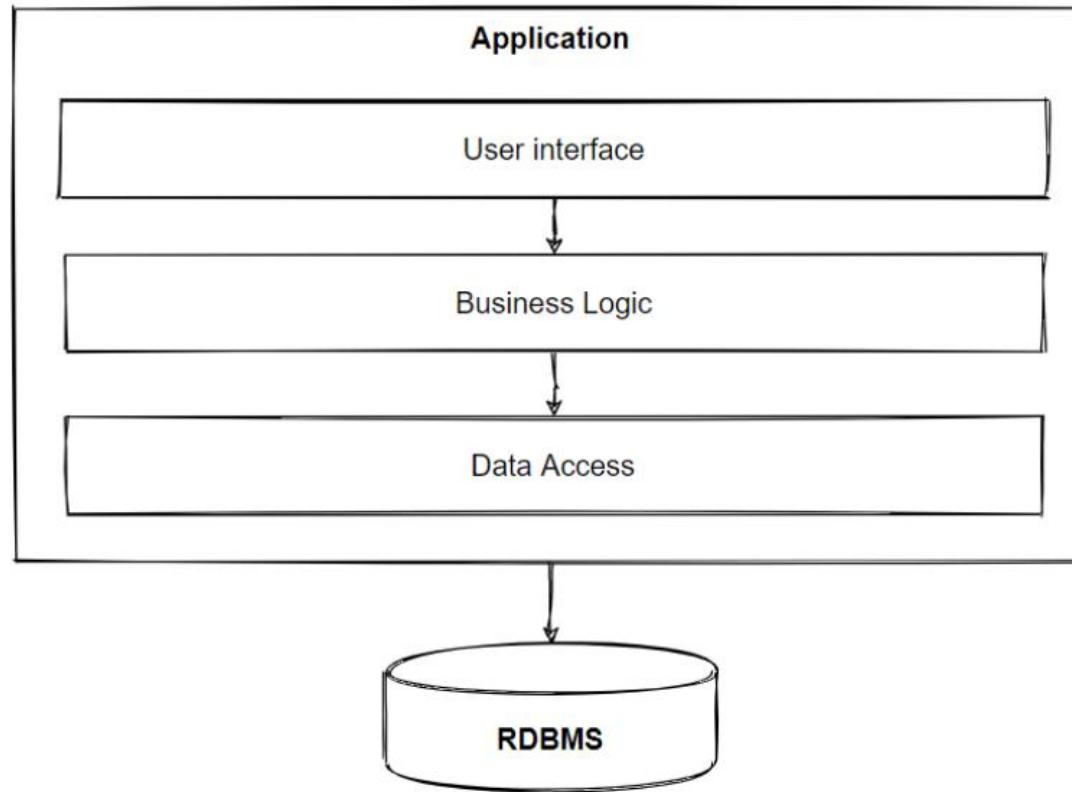
- **Basket Microservice** will be stored in a **Redis Distributed Cache**
- Use **Aspire.StackExchange.Redis.DistributedCaching**
- **IDistributedCache** cache



1. Application Architecture Style of Basket Microservice

Classical N-Layer Architecture

Organizes our code into layer folders, each feature encapsulated in a folder



N-Layer(Layered) Architecture

Domain Layer

Stores and Handles domain entities and domain events

Data Access Layer

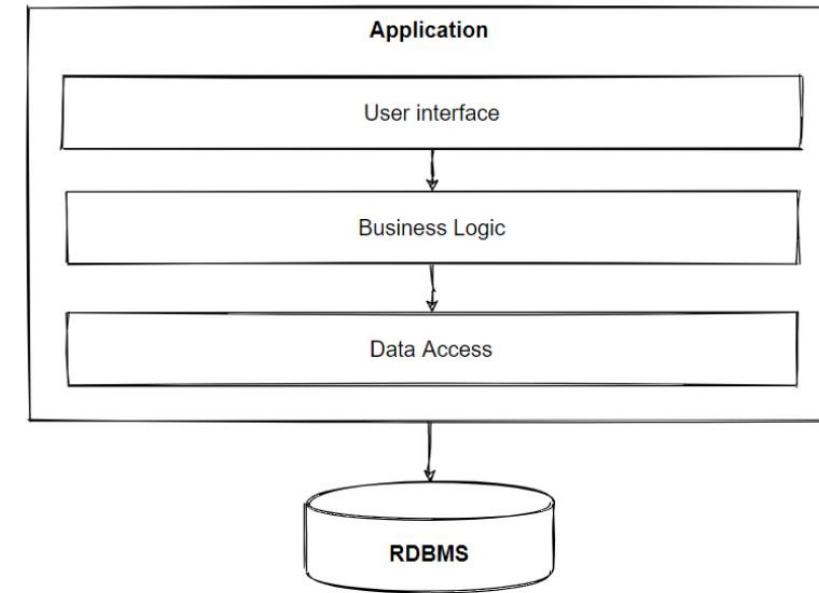
Manages interactions with the database or other storage systems

Business Logic Layer

Processes the application's core business logic and rules

Presentation Layer

Displaying data to users and capturing user input



2. Patterns & Principles of Basket Microservice



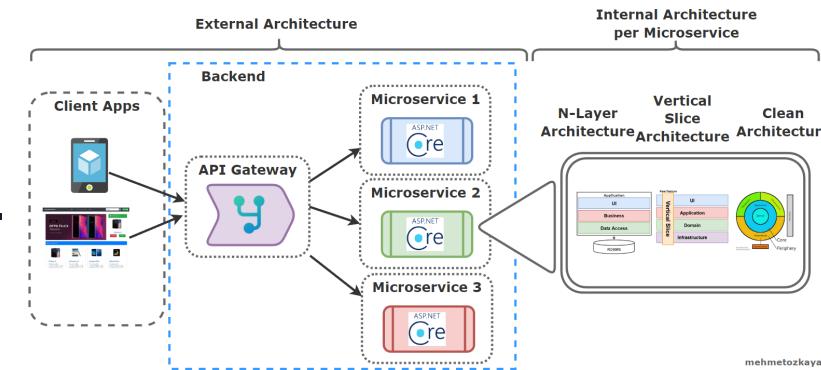
KISS principles and **use less patterns** in our **microservices**

Focus on integrations and orchestrations with **.NET Aspire**

DI in ASP.NET Core: Dependency Injection is a core feature, allowing us to inject dependencies.

Minimal APIs and Routing in ASP.NET 8: ASP.NET 8's Minimal APIs simplify endpoint definitions, providing lightweight syntax for routing and handling HTTP requests.

ORM Pattern: Object-Relational Mapping abstracts database interactions, work with database objects using high-level codes.



3. Libraries Nuget Packages of Basket Microservice



- **Aspire Hosting Packages for Orchestration:**

Aspire.Hosting.Redis

- **Aspire for Redis Interaction: (Client Integration) :**

Aspire.StackExchange.Redis.DistributedCaching



- `builder.AddRedisDistributedCache(connectionName: "cache");`

- `public class ExampleService(IDistributedCache cache)`

4. Project Folder Structure of Catalog Microservice



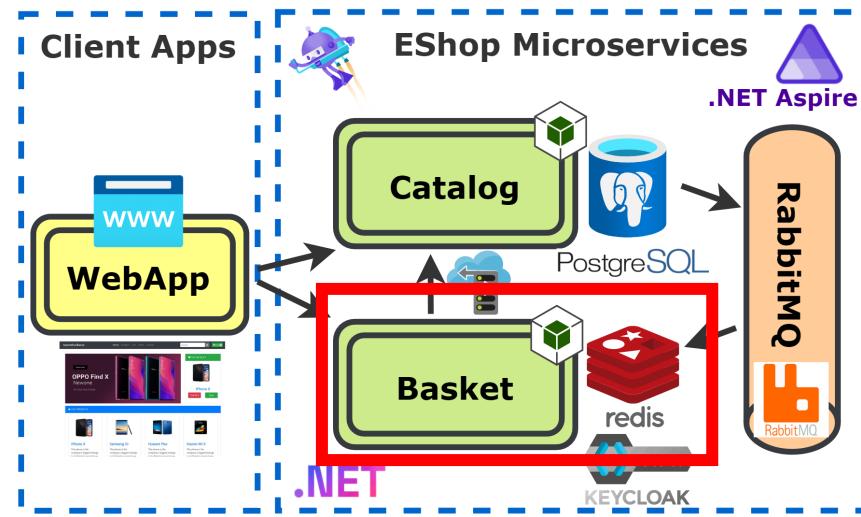
Organized into **Models**, **Services**, **Endpoints** and **ApiClient**s

Models holds our **domain entities** (e.g. SC)

Services contains **business logic**

Endpoints minimal API endpoints that call into the services layer

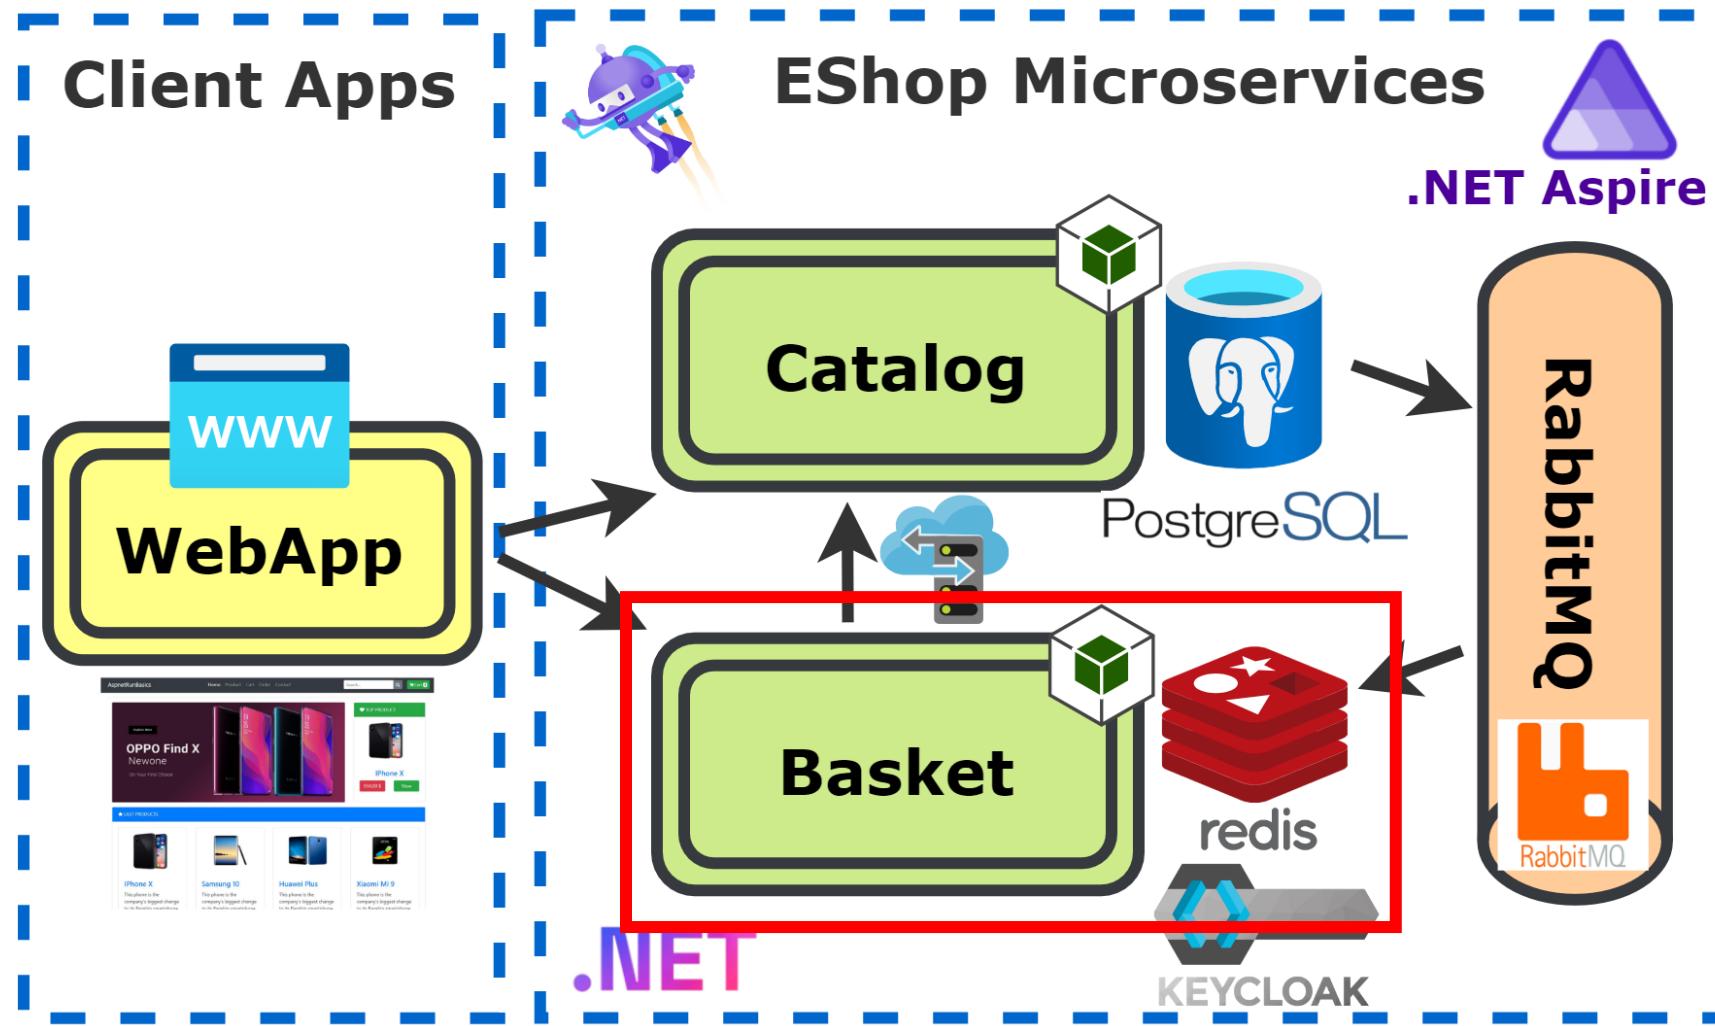
ApiClients includes **HTTP Client** for calling Catalog



A screenshot of the Visual Studio file explorer showing the project structure for the Basket microservice. The structure includes:

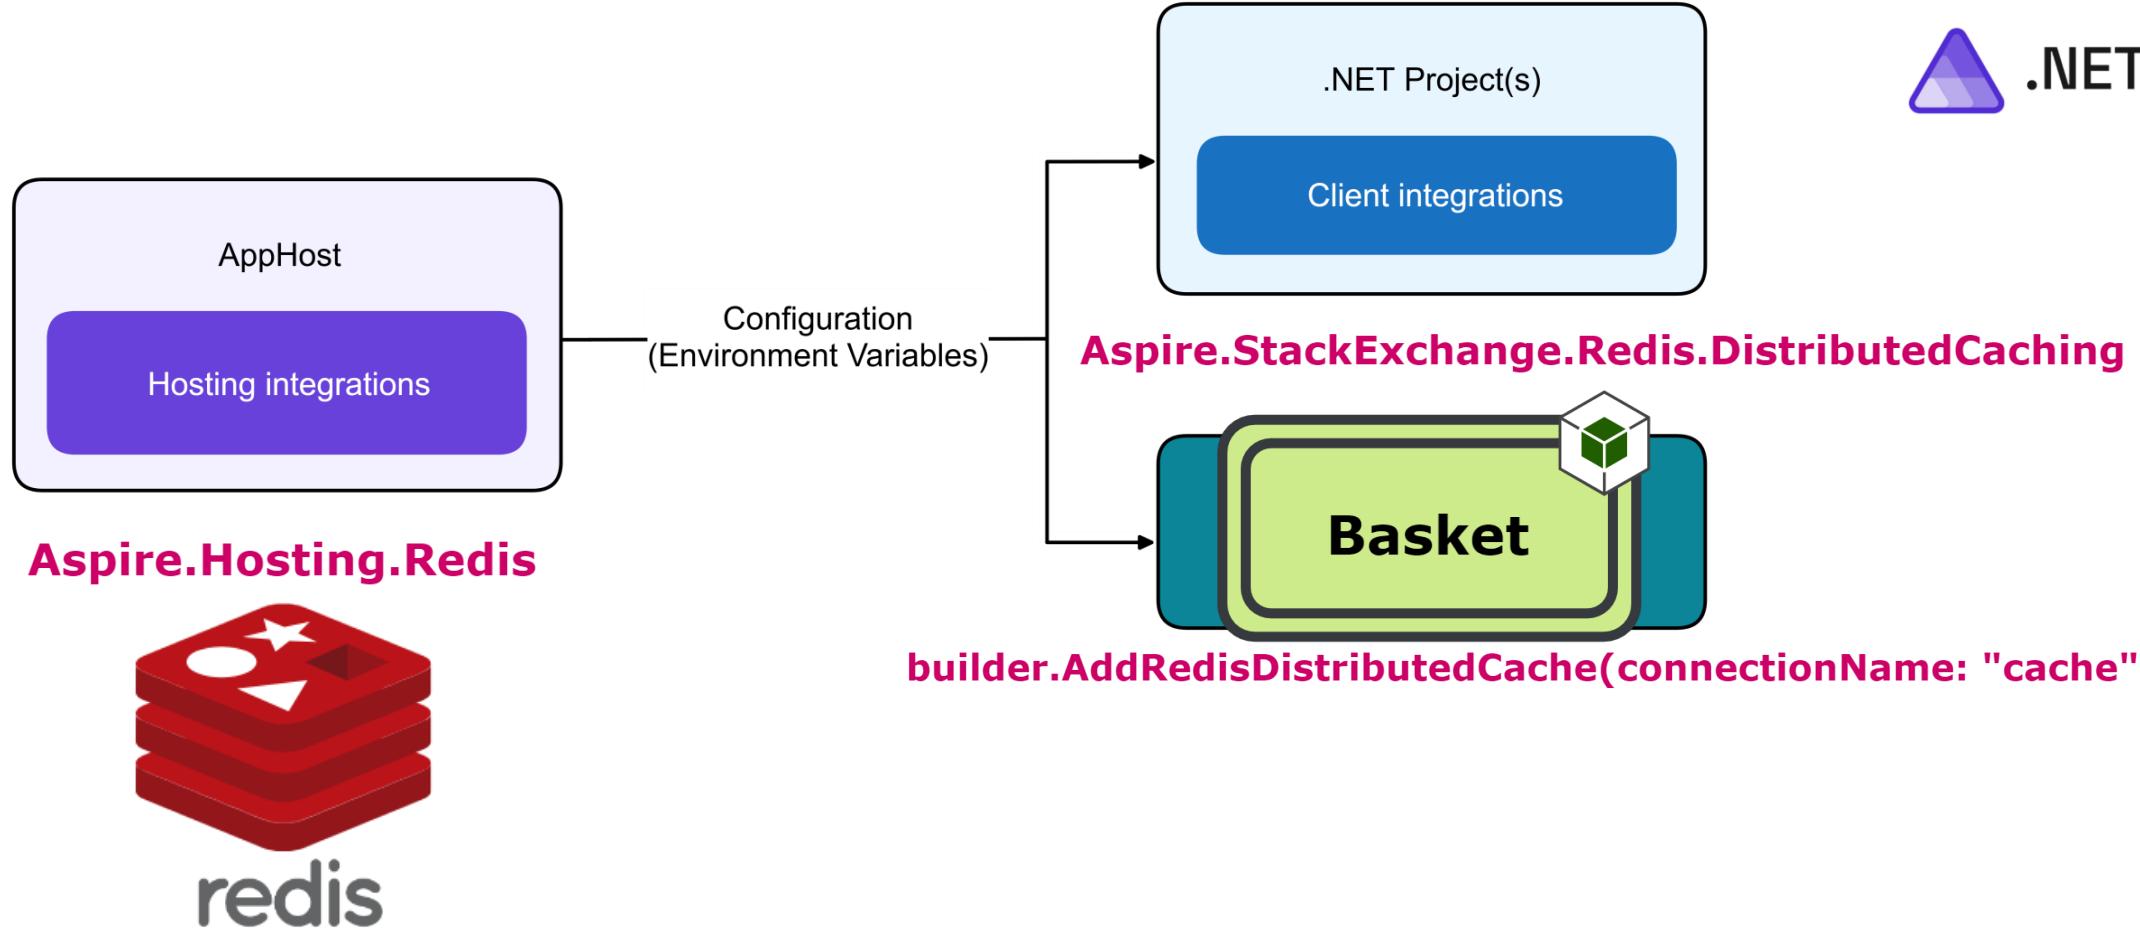
- Connected Services
- Dependencies
- Properties
- ApiClient (containing CatalogApiClient.cs)
- Endpoints (containing BasketEndpoints.cs)
- EventHandlers (containing ProductPriceChangedIntegrationEventHandler.cs)
- Models (containing Product.cs, ShoppingCart.cs, ShoppingCartItem.cs)
- Services (containing BasketService.cs)
- appsettings.json
- Basket.http
- GlobalUsings.cs
- Program.cs

Catalog Microservice with PostgreSQL



Basket Microservice with Redis

Integrations



Sync Communications between Catalog-Basket w/ .NET Aspire Service Discovery

Sync Communications Use Case: When Update Basket, Get Latest Product Prices

Add WithReference(catalog) into Basket for Service Discovery

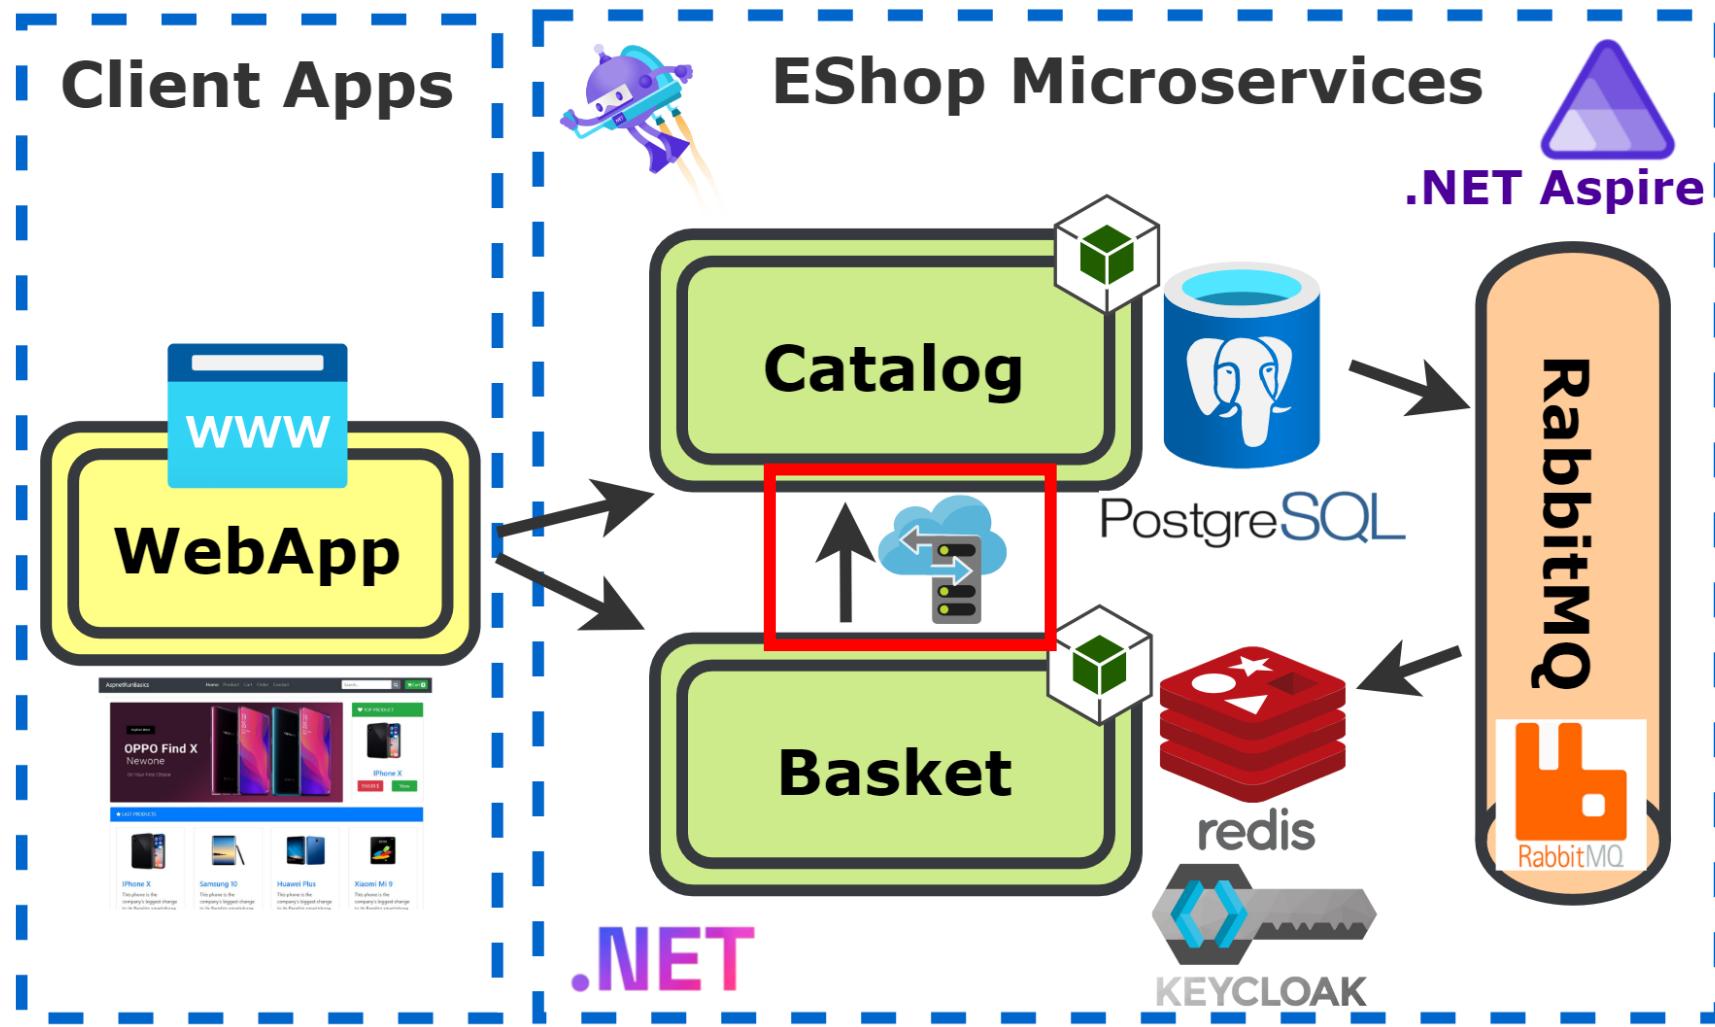
Develop CatalogApiClient fro Sync Http Communication

Integrate w/ Catalog into BasketService.cs using CatalogApiClient

Mehmet Ozkaya



Sync Communications between Catalog-Basket



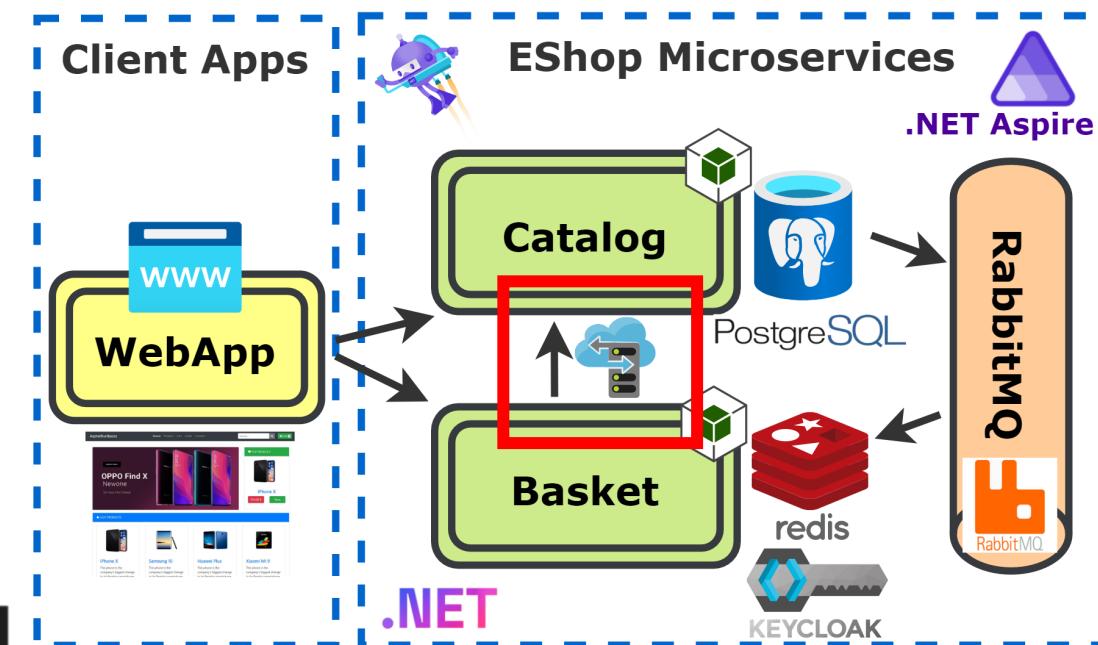
Sync Communications between Catalog-Basket

Adding .WithReference(catalog)

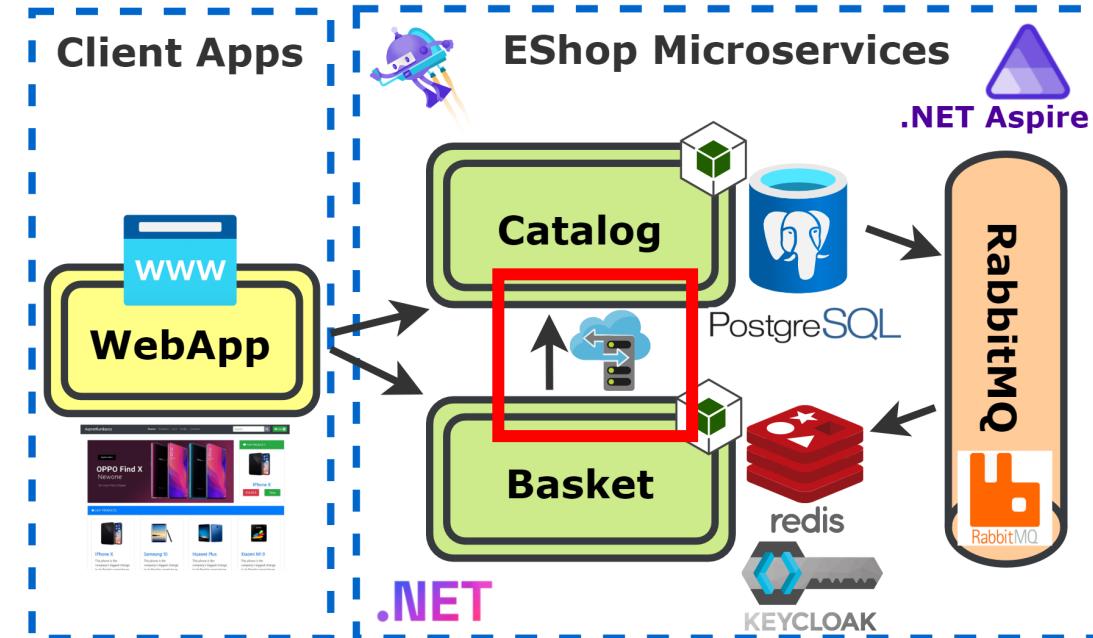
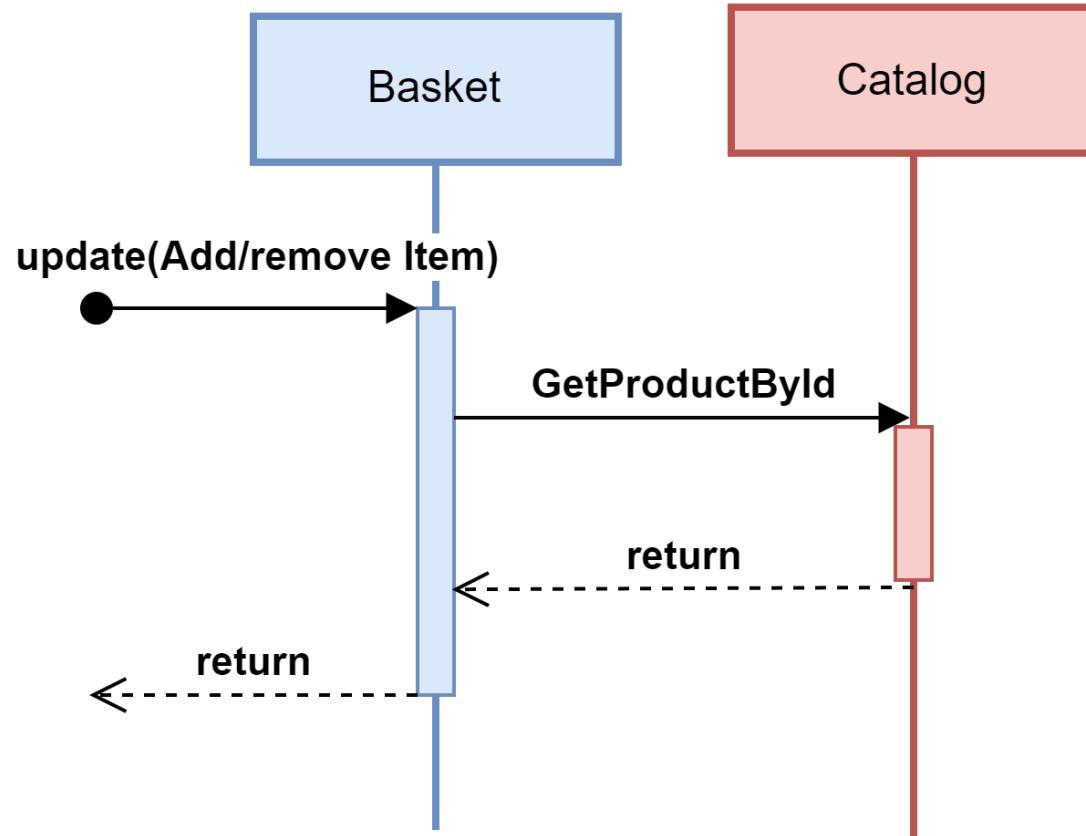
```
var basket = builder
    .AddProject<Projects.Basket>("basket")
    .WithReference(catalog)
    .WaitFor(catalog);
```

Creating a Typed HttpClient consume Catalog from Basket

```
builder.Services.AddHttpClient<CatalogApiClient>(client
=> {
    client.BaseAddress = new("https+http://catalog");
});
```



When Update Basket, Get Latest Product Prices



```
public async Task UpdateBasket(ShoppingCart basket)
{
    //TODO: Before update/Add/remove Item into SC, we should call Catalog Module GetProductById method
    // Get latest product information and set Price and ProductName when adding item into SC
```

Remember: Service Discovery



Service
Discovery

Containers → injects → connection string

Projects → injects → microservice endpoints

```
var builder = DistributedApplication.CreateBuilder(args);

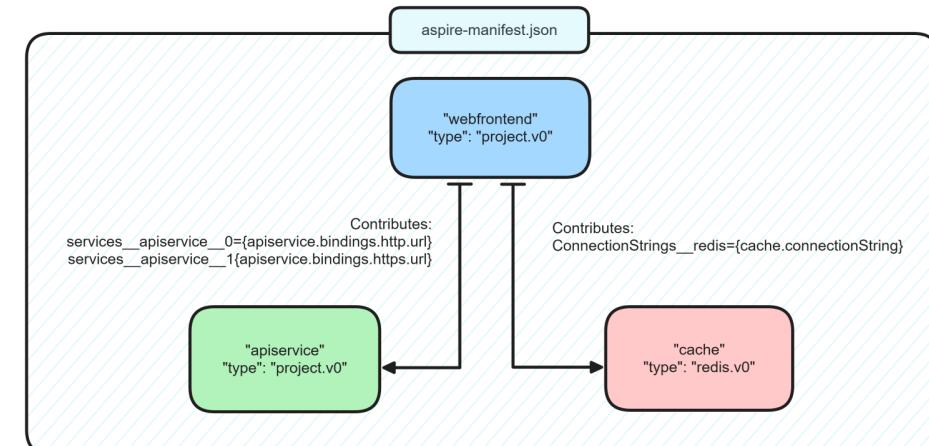
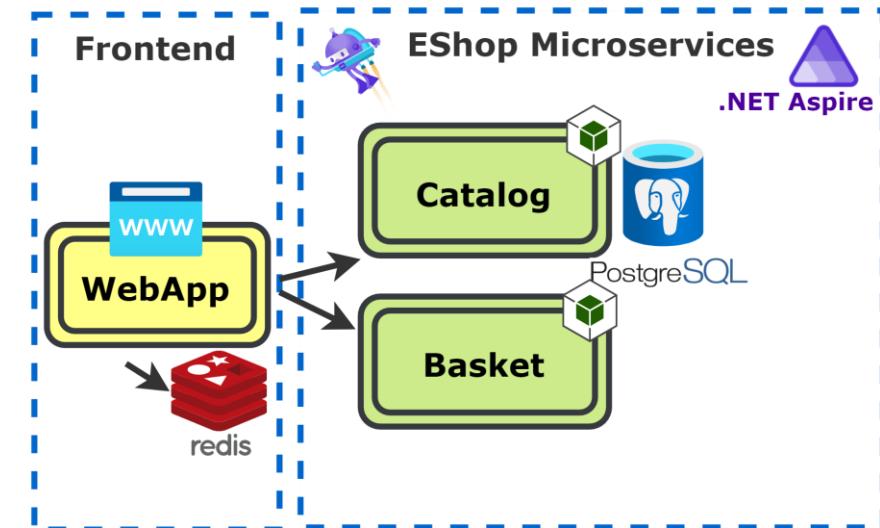
var cache = builder.AddRedis("cache");

var catalog = builder.AddProject<Projects.CatalogService>("catalog");
var basket = builder.AddProject<Projects.BasketService>("basket");

var frontend = builder.AddProject<Projects.MyFrontend>("frontend")
    .WithReference(cache)
    .WithReference(catalog)
    .WithReference(basket);
```

```
WithReference(cache)
    ConnectionStrings__cache="localhost:62354"
```

```
WithReference(apiservice)
    services__apiservice__http__0="http://localhost:5455"
    services__apiservice__https__0="https://localhost:7356"
```



When Update Basket, Get Latest Product Prices

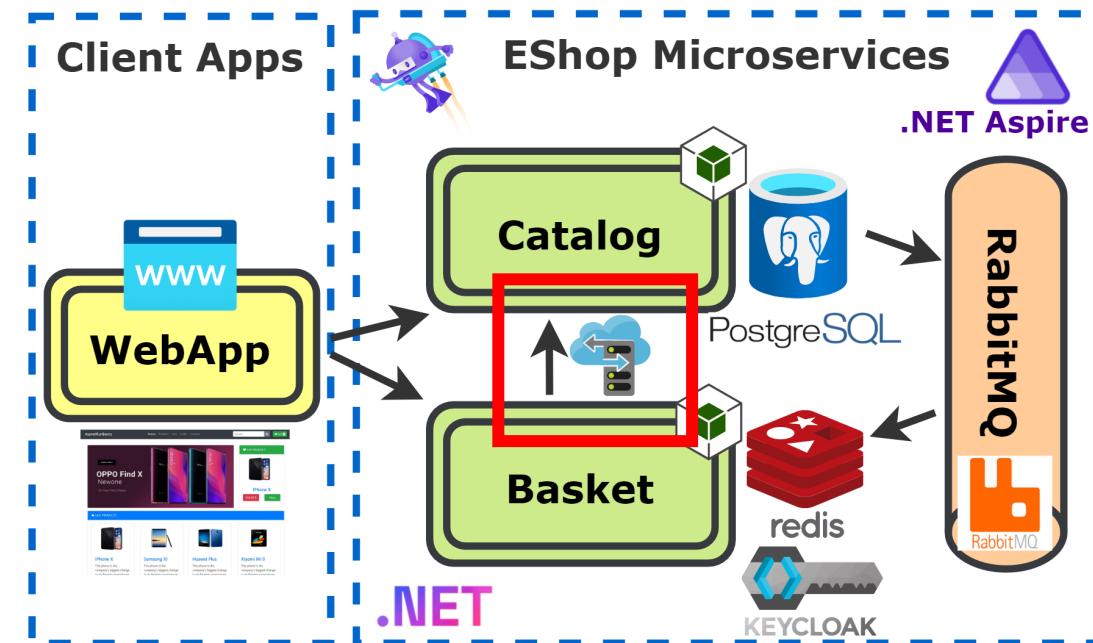
.WithReference(catalog) for Basket → Catalog

AddHttpClient<CatalogClient> referencing
"catalog"

Synchronous calls via strongly typed methods:
await
catalogClient.GetProductByIdAsync(productId);

Goto BasketService – UpdateBasket Method

```
public async Task UpdateBasket(ShoppingCart basket)
{
    //TODO: Before update/Add/remove Item) into SC, we should call Catalog Module GetProductById method
    // Get latest product information and set Price and ProductName when adding item into SC
```



Async Communications w/ RabbitMQ & MassTransit orchestrate .NET Aspire

RabbitMQ Hosting Integration .NET Aspire

Create Shared Messaging Folders and Classes for
ProductPriceChangedIntegrationEvent

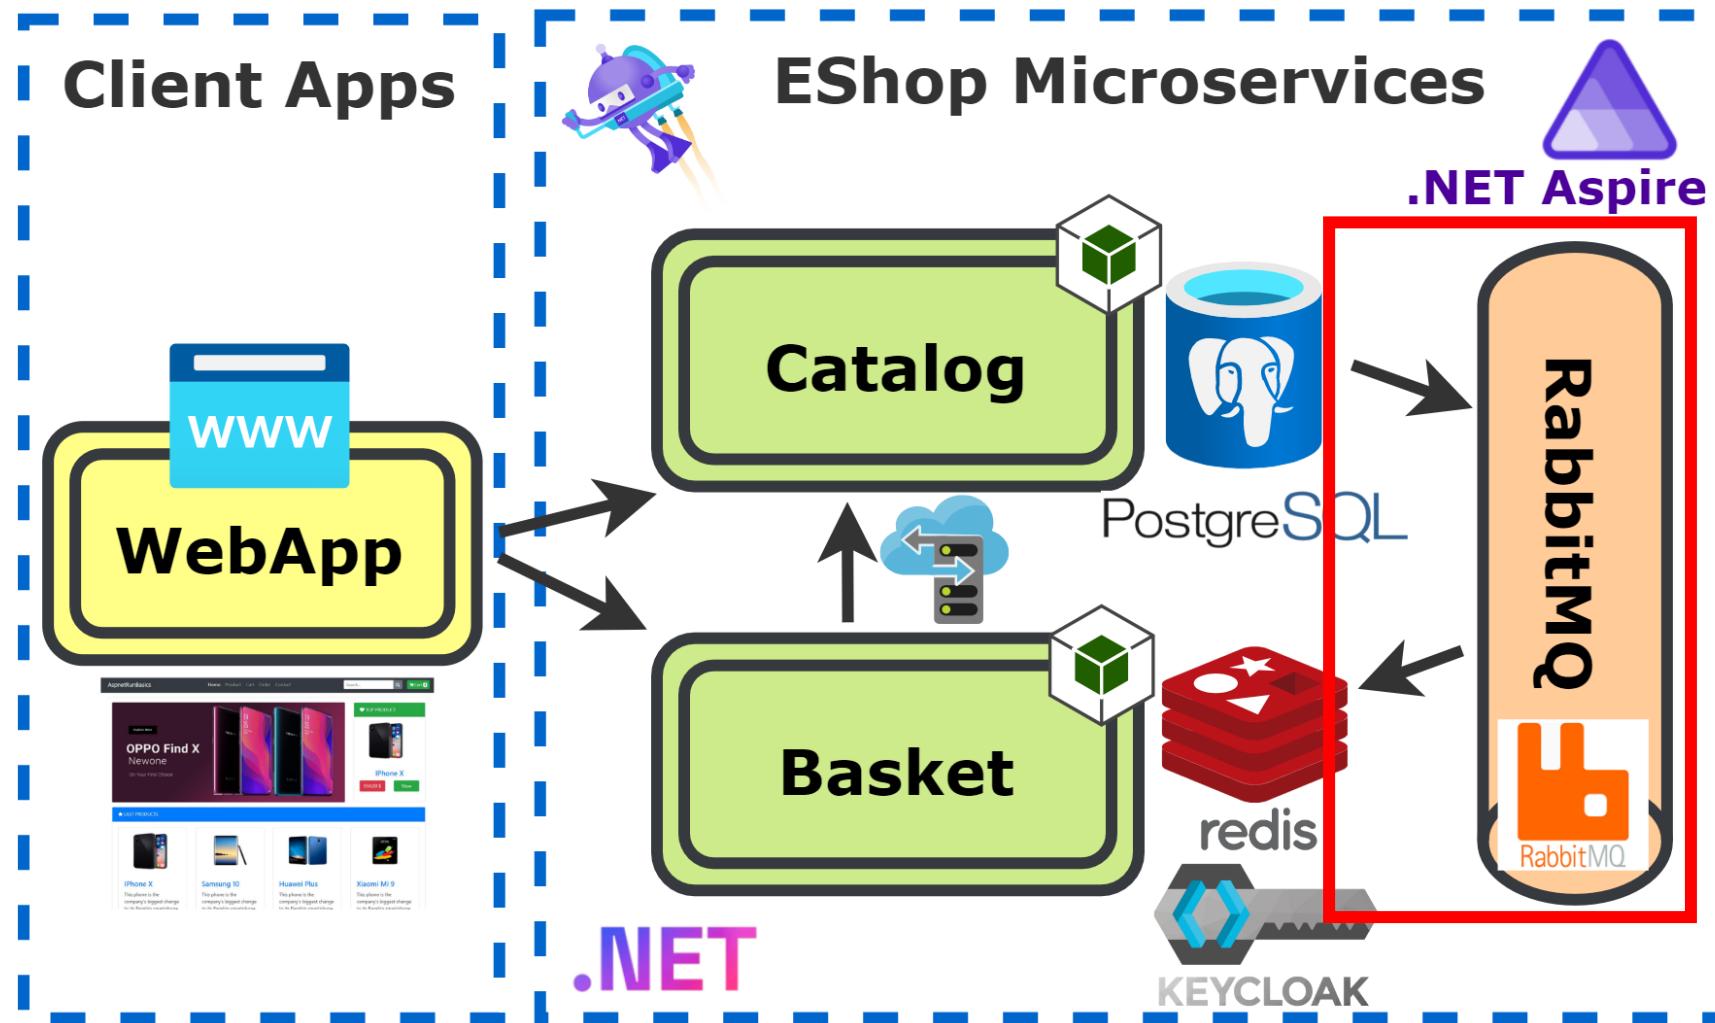
Catalog ms Publish ProductPriceChanged Integration Event

Basket Subscribe and consume ProductPriceChanged Integration Event

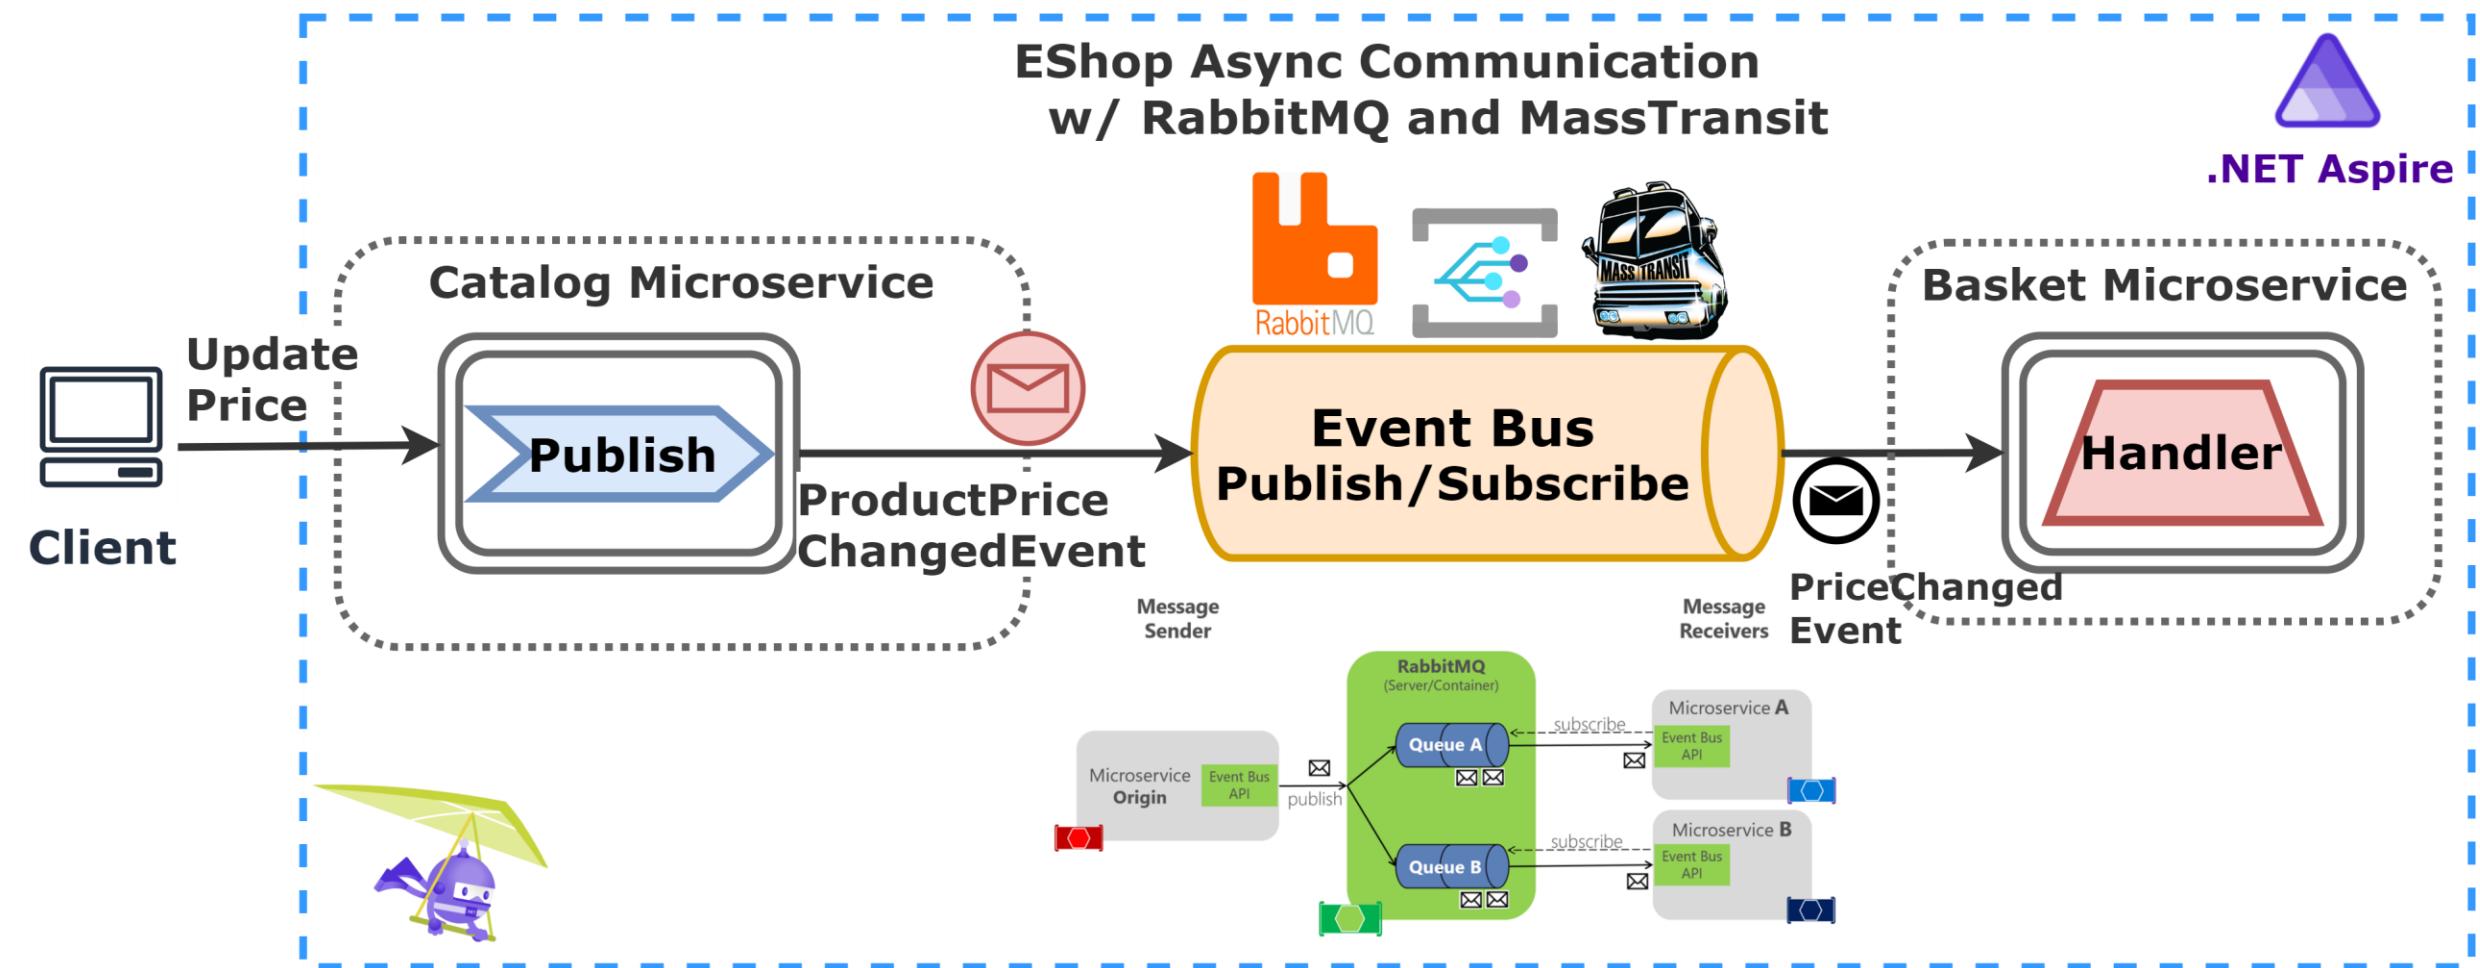
Mehmet Ozkaya



Async Communications w/ RabbitMQ & MassTransit



Async Communications w/ RabbitMQ & MassTransit



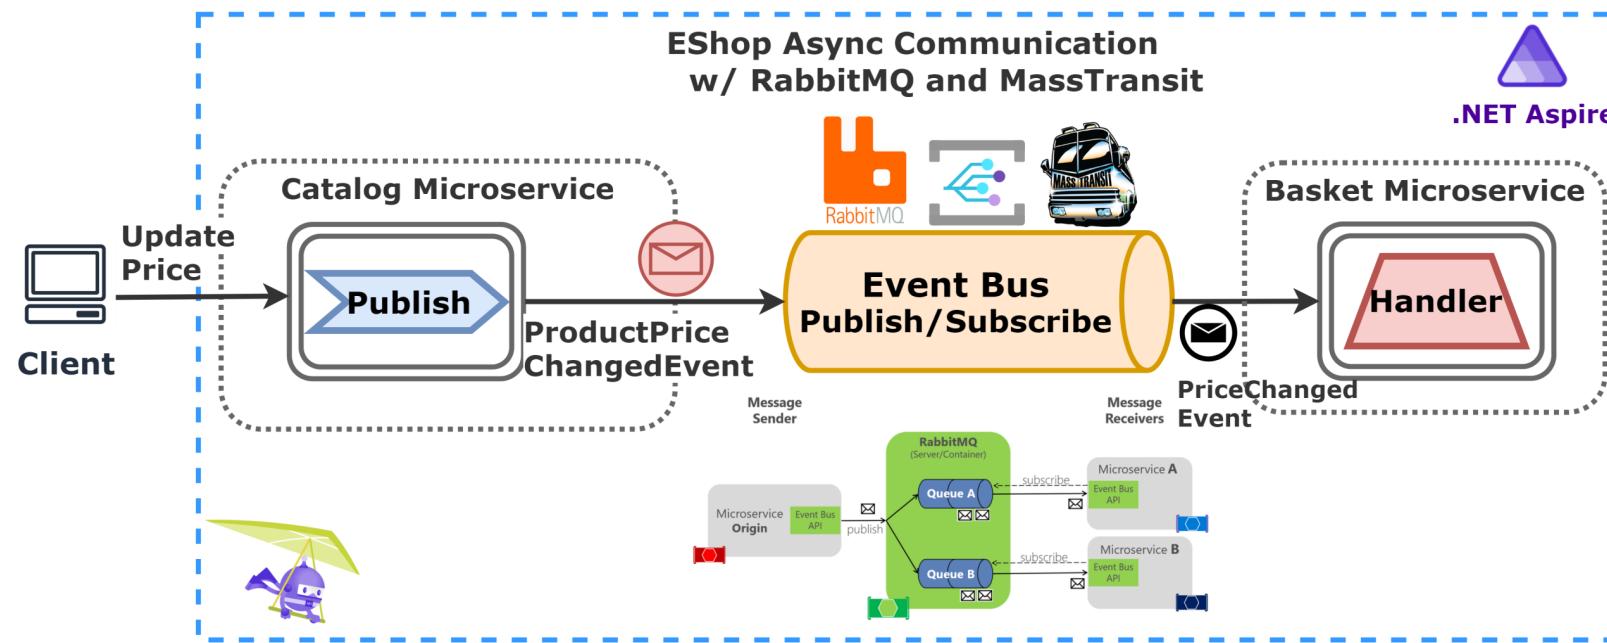
Steps for Async Communication

Hosting Integration (.AddRabbitMq("messagebroker"))

Client Integration

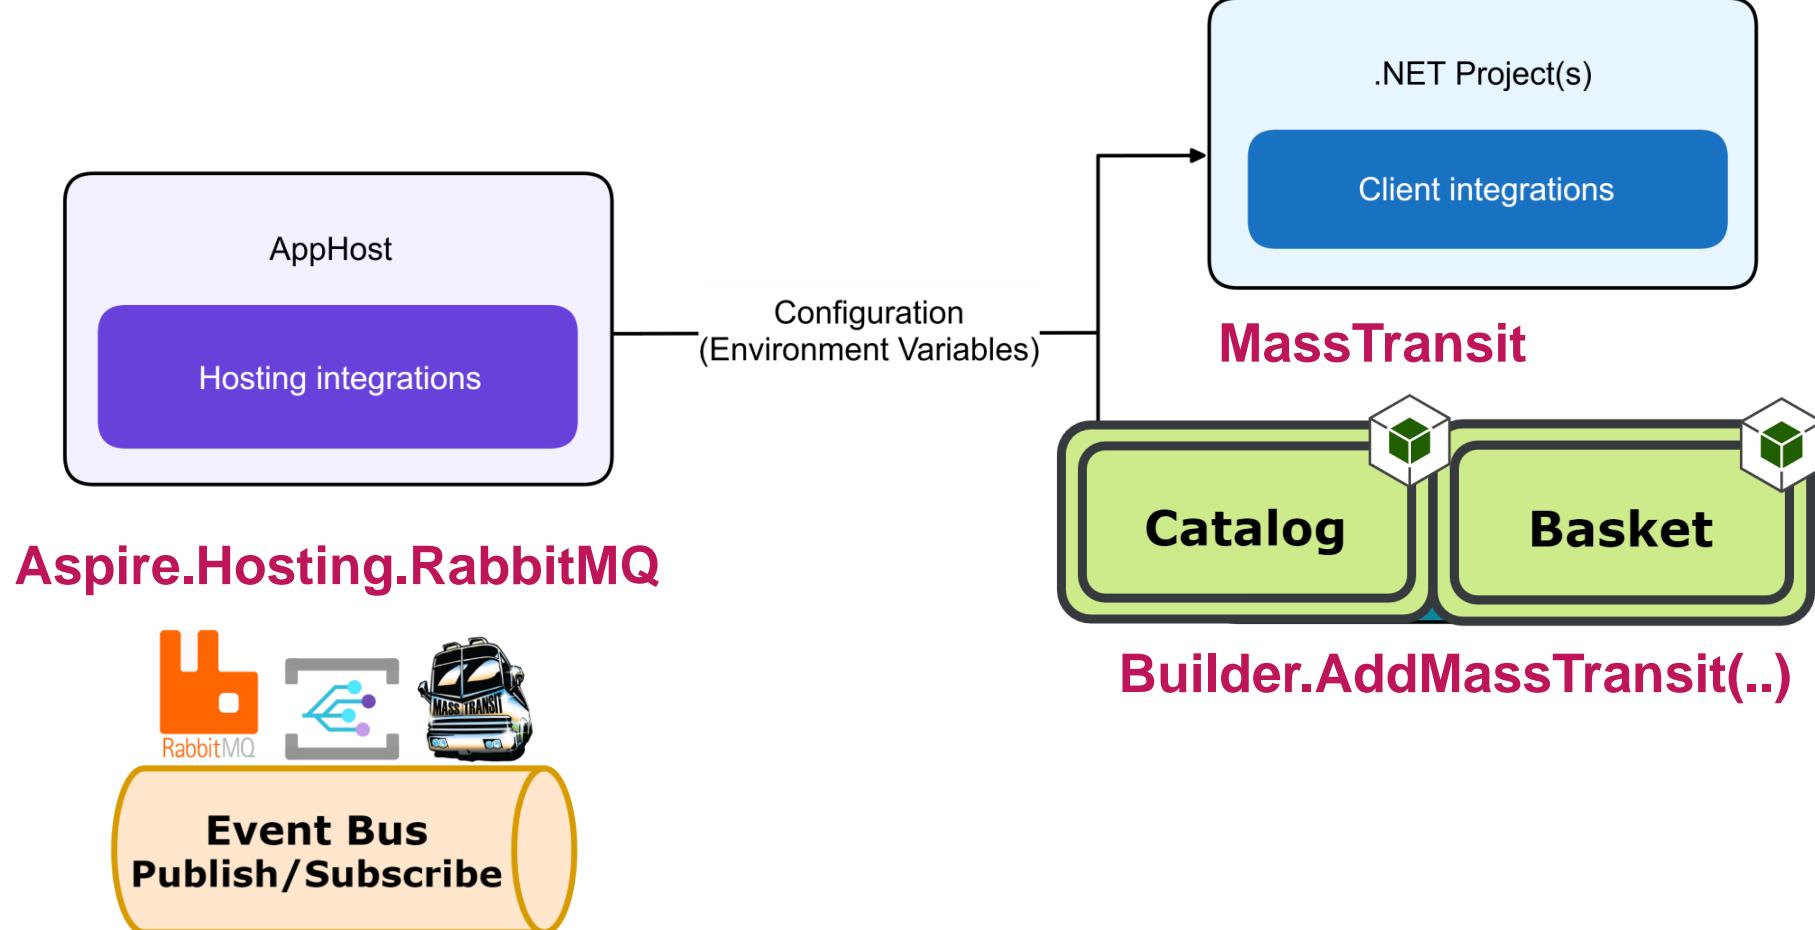
Shared Messaging Events

MassTransit



Async Communications w/ RabbitMQ & MassTransit

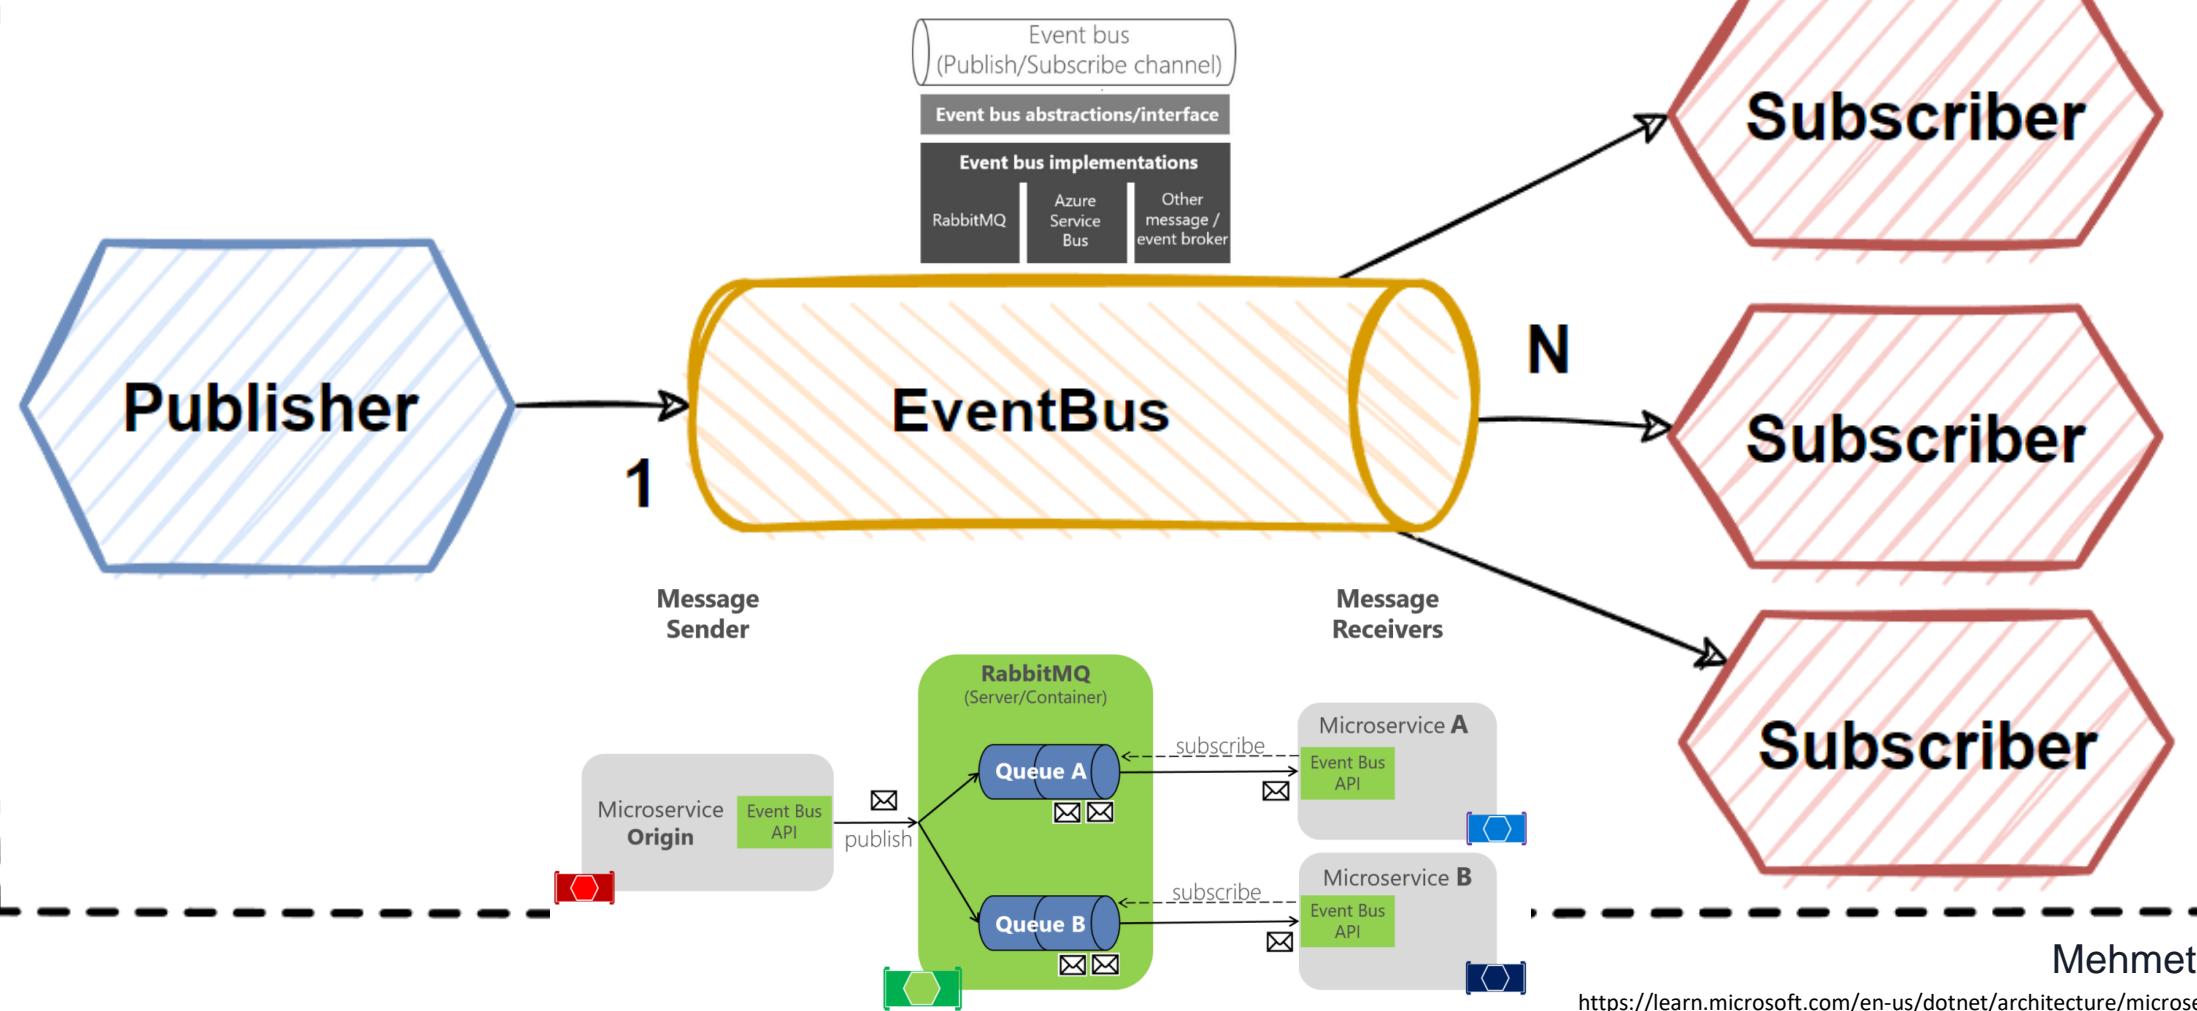
Integrations



Pub/Sub RabbitMQ Architecture



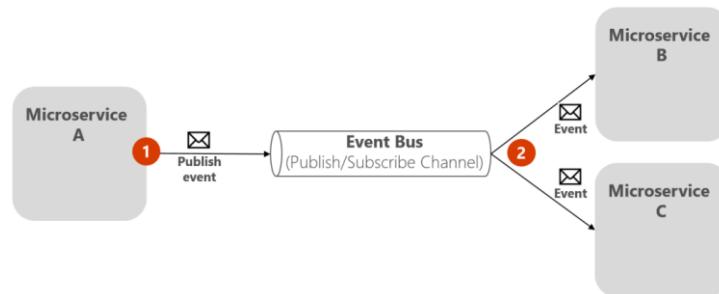
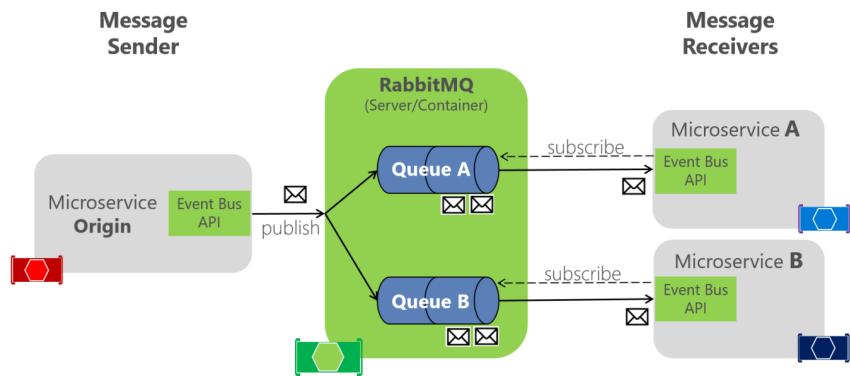
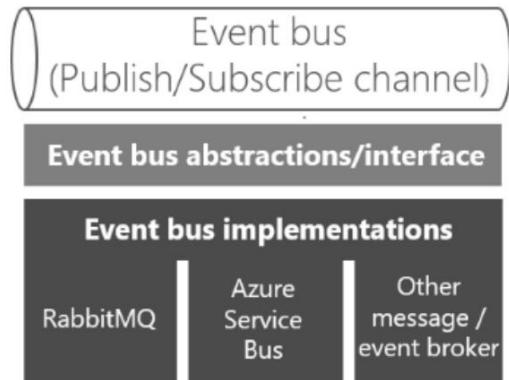
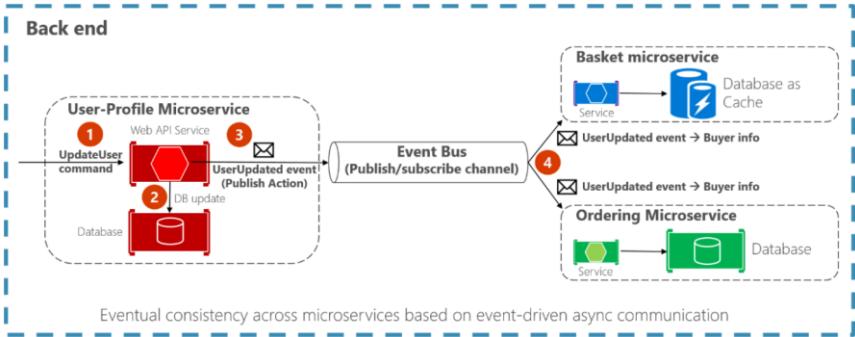
Publish–Subscribe Design Pattern



Mehmet Ozkaya

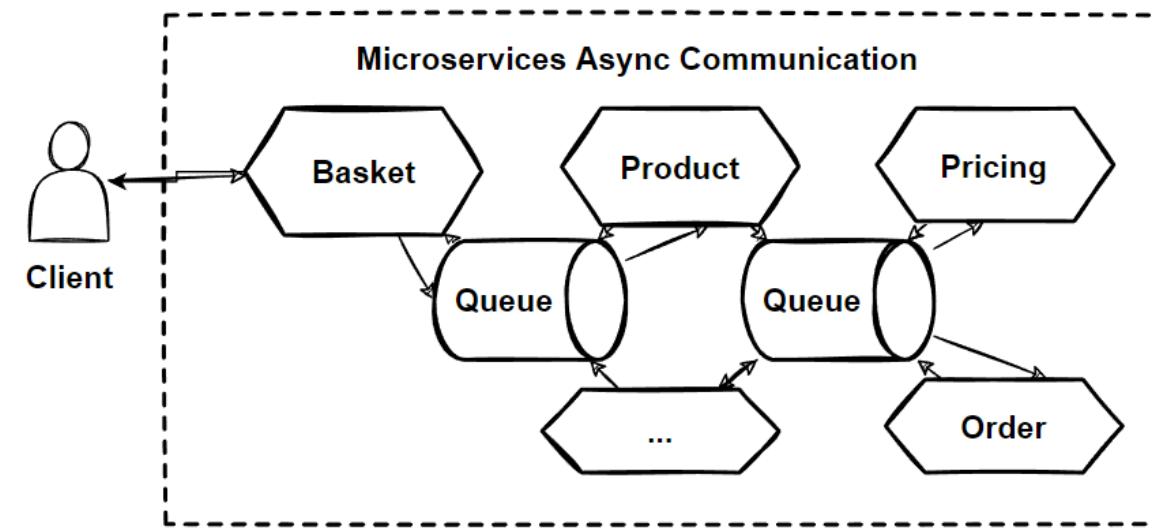
119

Microservices Async Communication w/ RabbitMQ & MassTransit



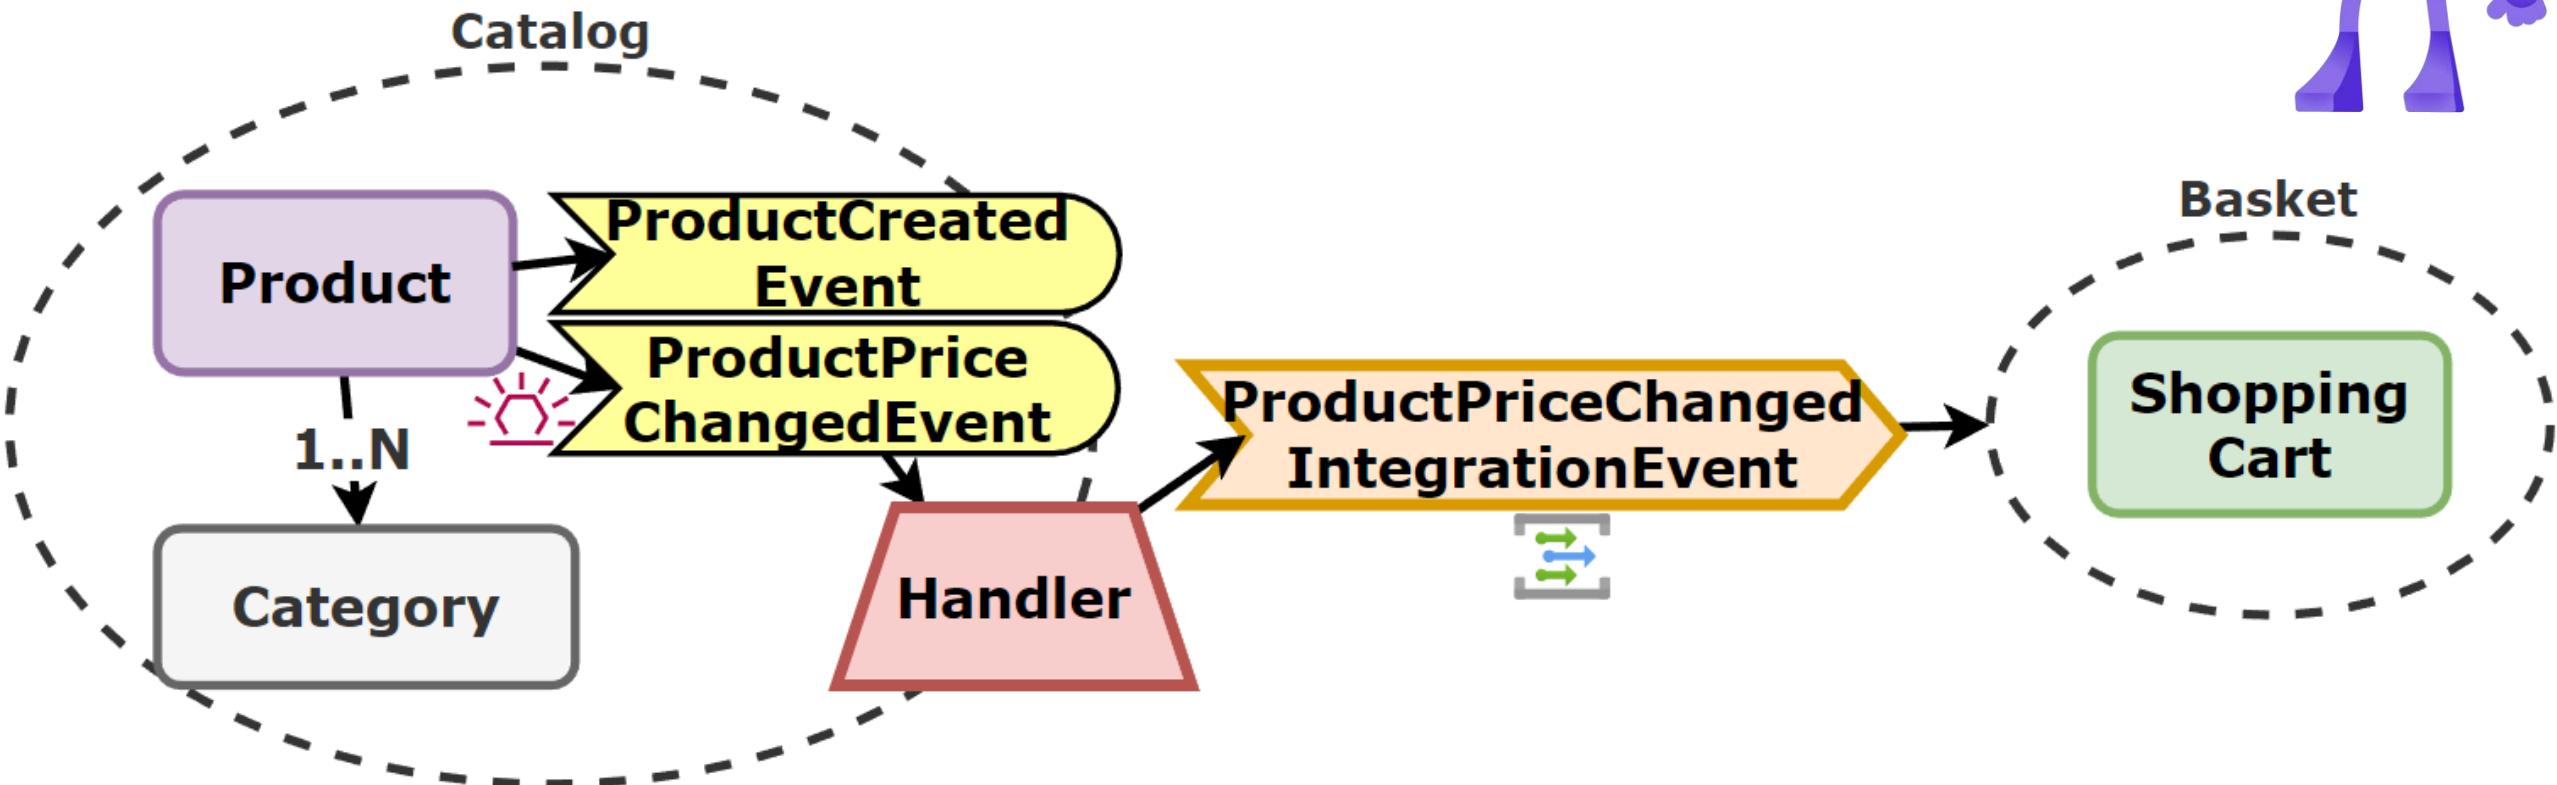
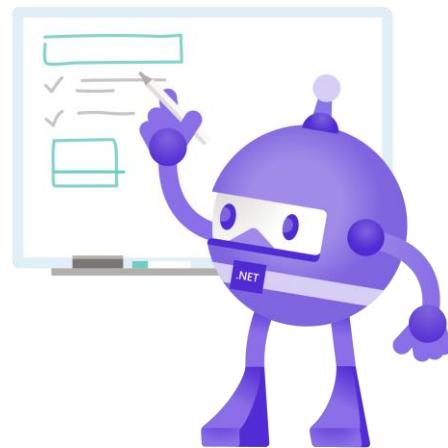
Why Use Asynchronous Communication?

- **Decoupling:** Asynchronous communication decouples the sender and receiver, allowing them to operate independently.
- **Scalability:** By not waiting for responses, the system can handle more requests and scale effectively.
- **Resilience:** If a module fails, messages can be retried or processed later, improving fault tolerance.
- **Responsiveness:** Users receive immediate feedback, even if the operation takes time to complete.



1. Event Models of ProductPriceChangedEvent

- Primary event model is 'ProductPriceChangedEvent'



- Consider Domain event for ProductPriceChanged, leading to integration events.

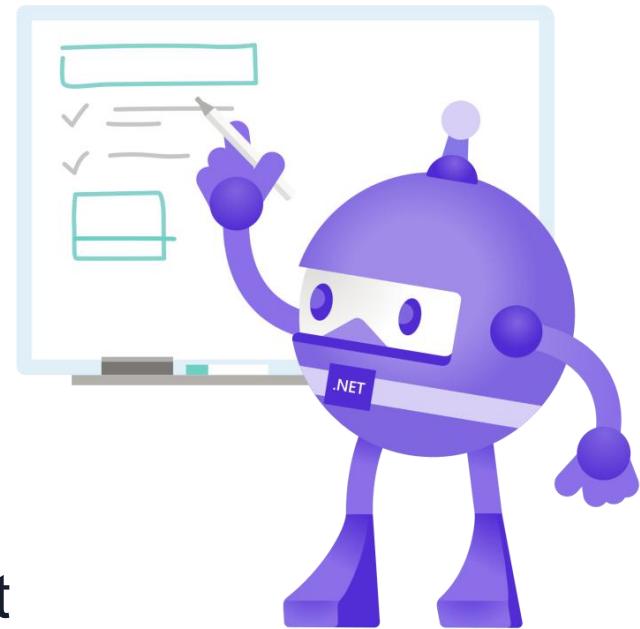
2. Application Use Cases between Catalog & Basket

Async Cases for Catalog:

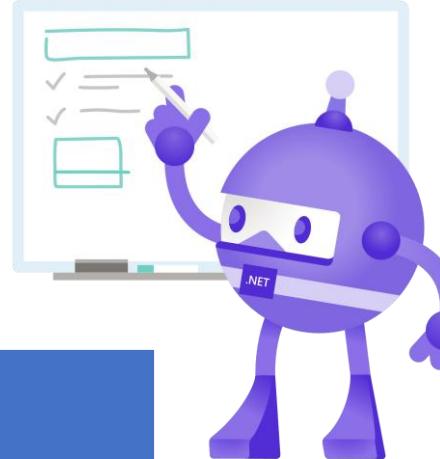
- **Update Product Price:** publish domain event that lead to an integration event that contains the updated product information.

Async Cases for Basket:

- **Basket Update Price:** Consume ProductPriceChangedEvent integration event from RabbitMQ using MassTransit.
 - Get existing basket with changed productId
 - Set ProductName and Price from incoming integration event
 - Update all basket shopping carts according to changed product information



3. Rest API Endpoints of Catalog Module



Method	Request URI	Use Cases
POST	/products	Update Product Price

RabbitMQ AMQP Protocol

- AMQP: Consume ProductPriceChangedEvent Event from RabbitMQ using MassTransit

4. Underlying Data Structures of Async Communications

For Async Communicatin Between Modules has 1 Backing Services:

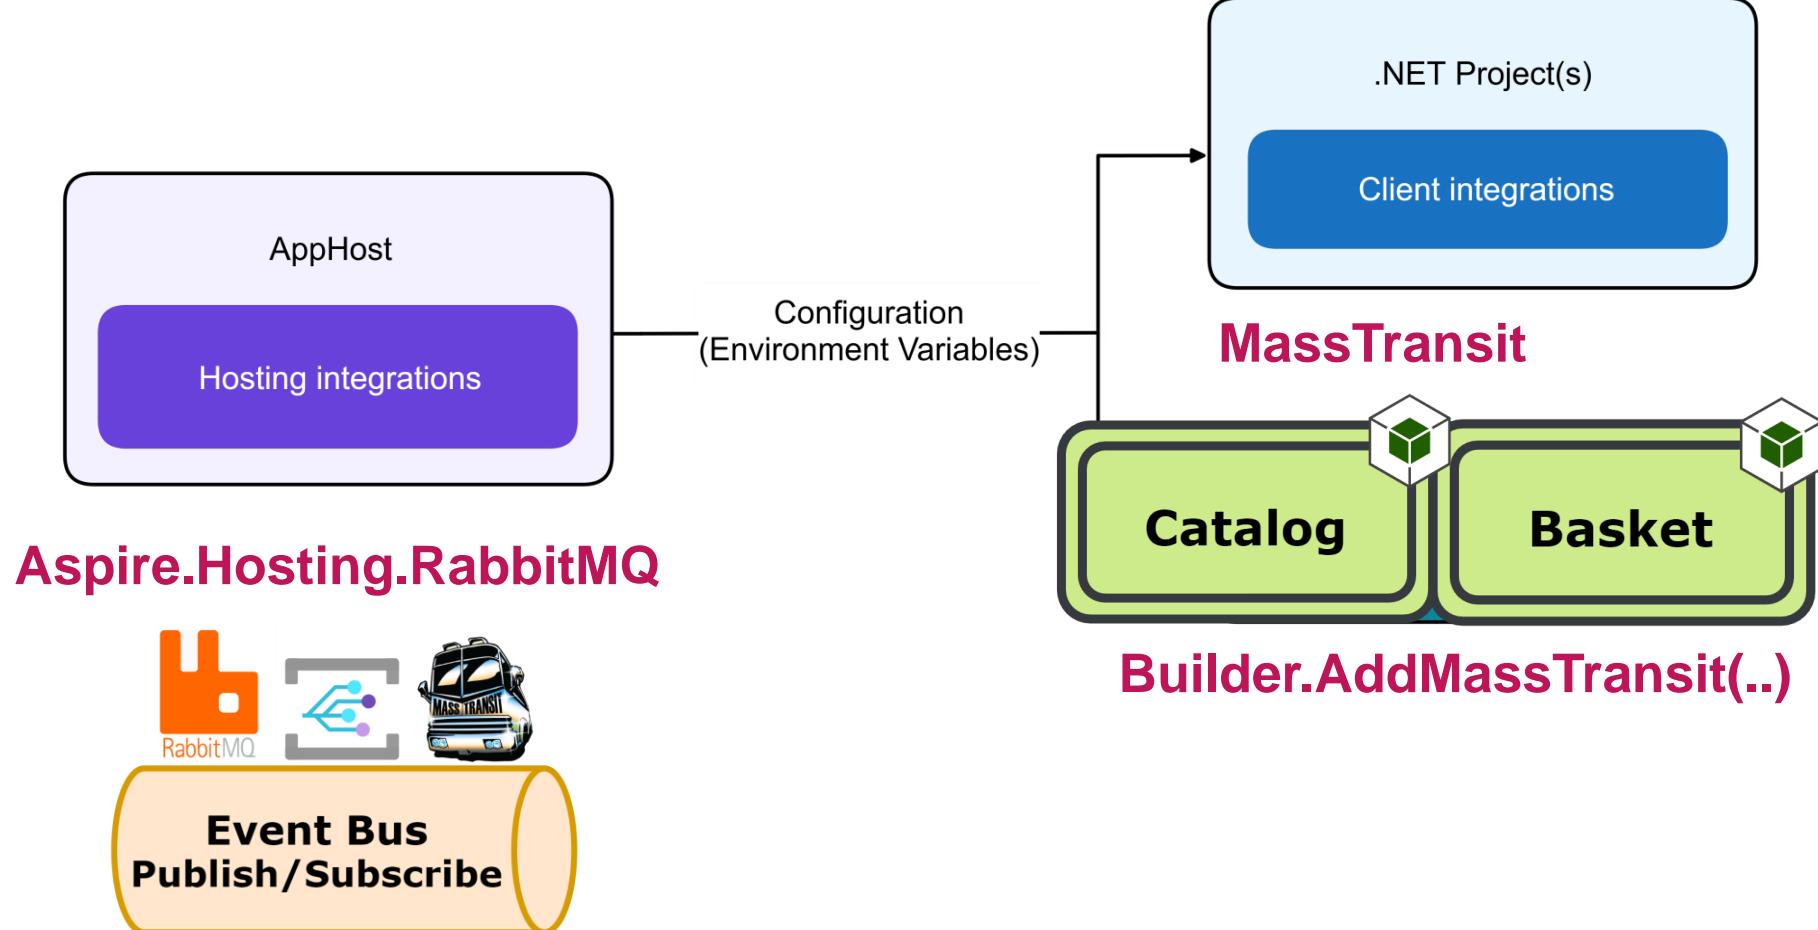
1. RabbitMQ Distributed Message Broker
2. Client Integrations w/ MassTransit



.NET Aspire

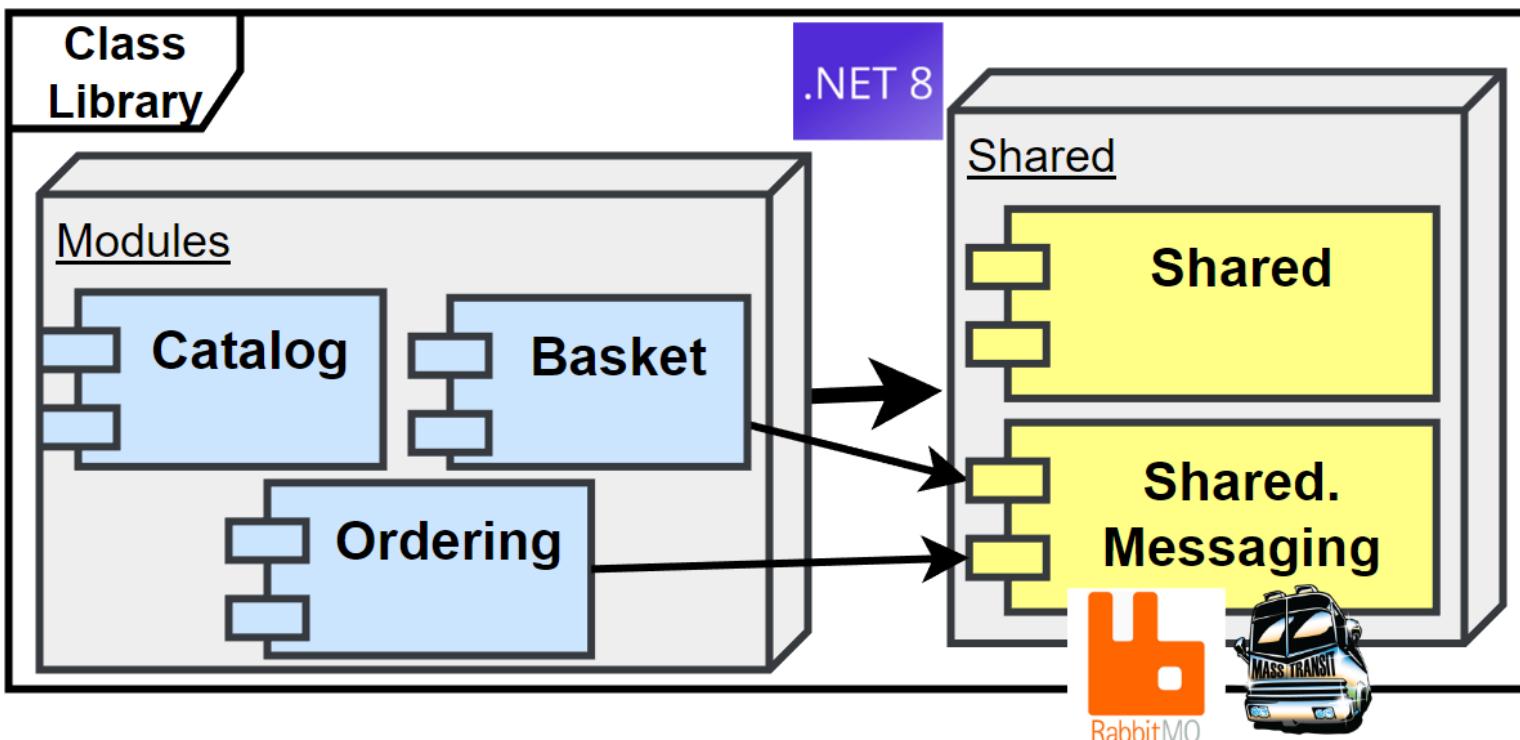
Async Communications w/ RabbitMQ & MassTransit

Integrations



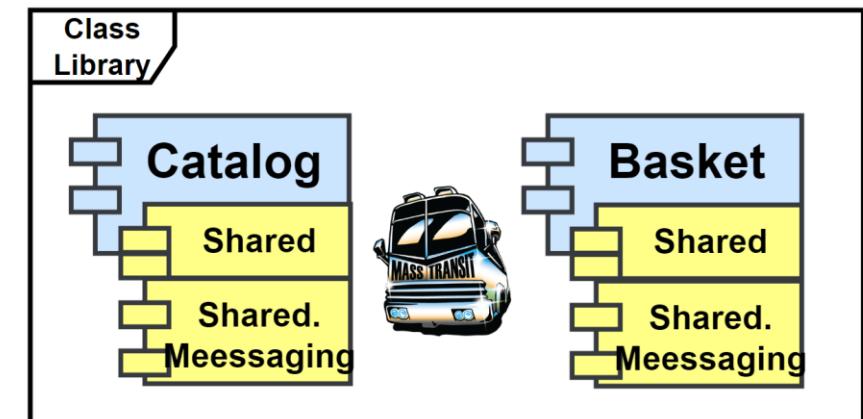
Class Library Architecture of Async Communication

- Catalog -- referenced Shared and Shared.Messaging
- Basket -- referenced Shared and Shared.Messaging
- Identity
- Ordering-- referenced Shared and Shared.Messaging



```
Shared.Messaging
  Dependencies
  Events
    BasketCheckoutIntegrationEvent.cs
    IntegrationEvent.cs
    ProductPriceChangedIntegrationEvent.cs
  Extentions
    MassTransitExtentions.cs
```

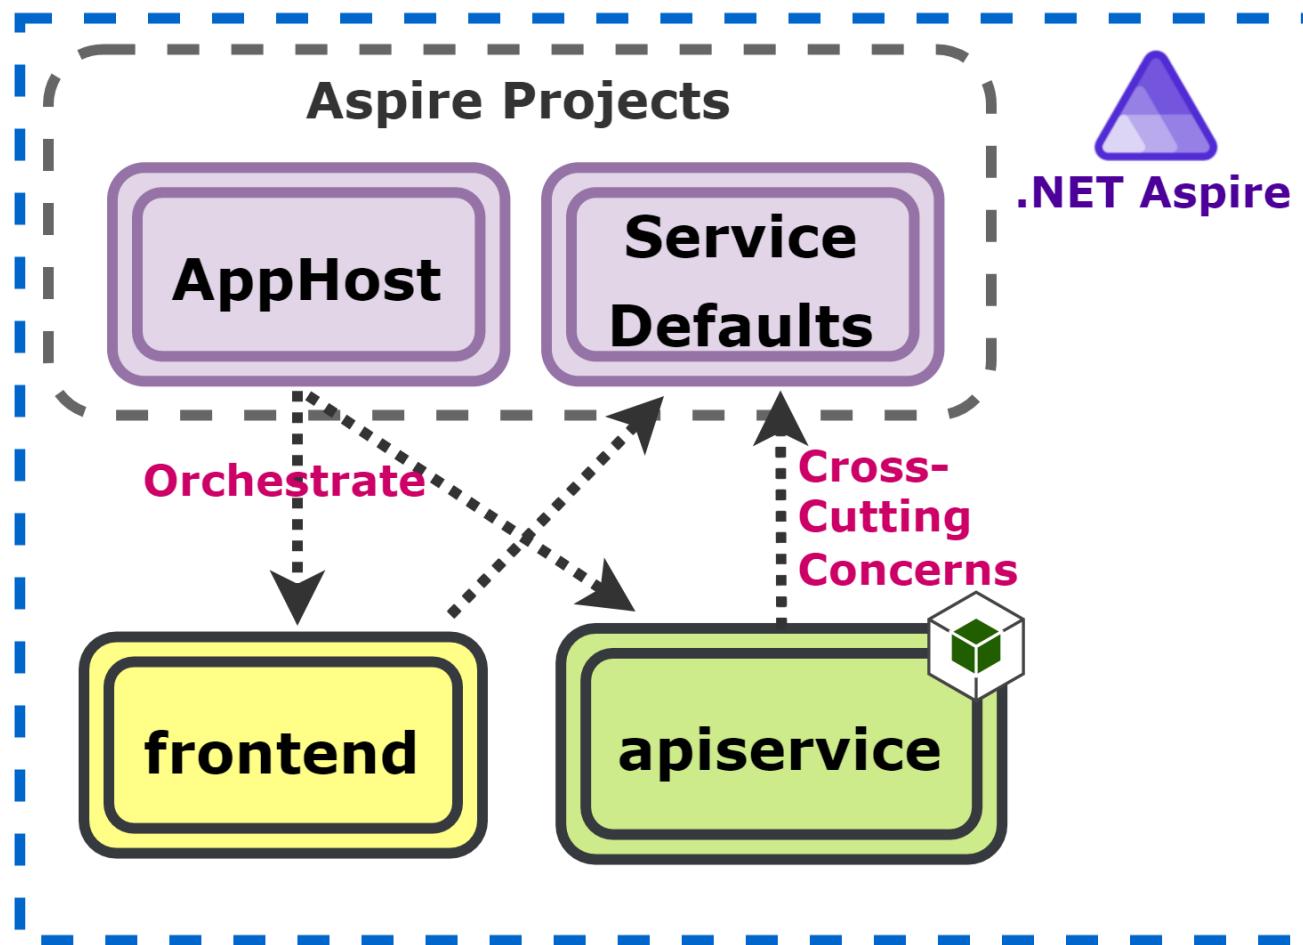
Module Project References for Async Call



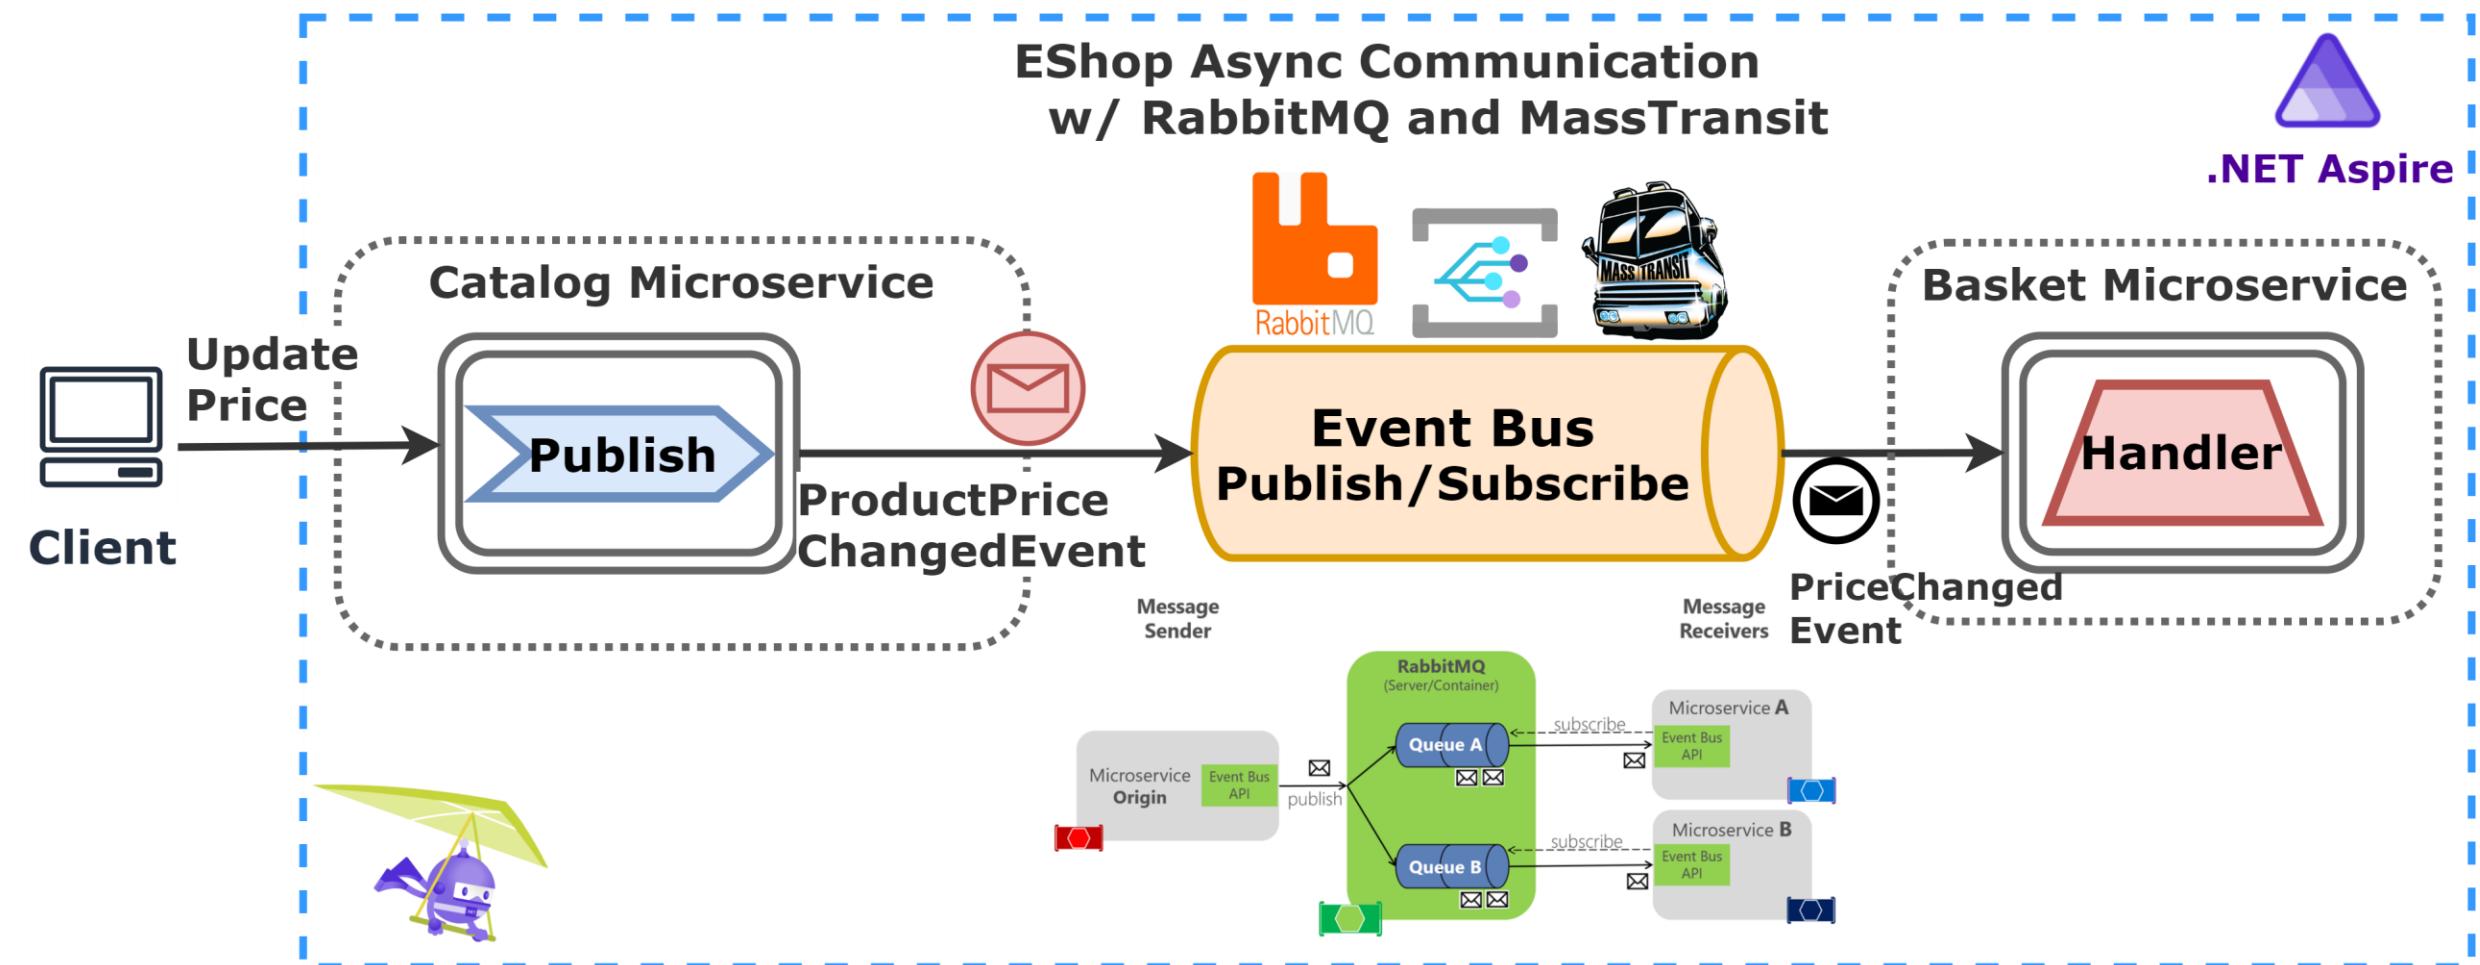
Project Dependencies

AppHost → references **ApiService & Web** for orchestration

ApiService/Web → reference **ServiceDefaults** for cross-cutting concerns



Async Communications w/ RabbitMQ & MassTransit



Secure Basket with Keycloak Authentication orchestrate .NET Aspire

Keycloak Identity and Access Management

Keycloak Hosting Integration .NET Aspire

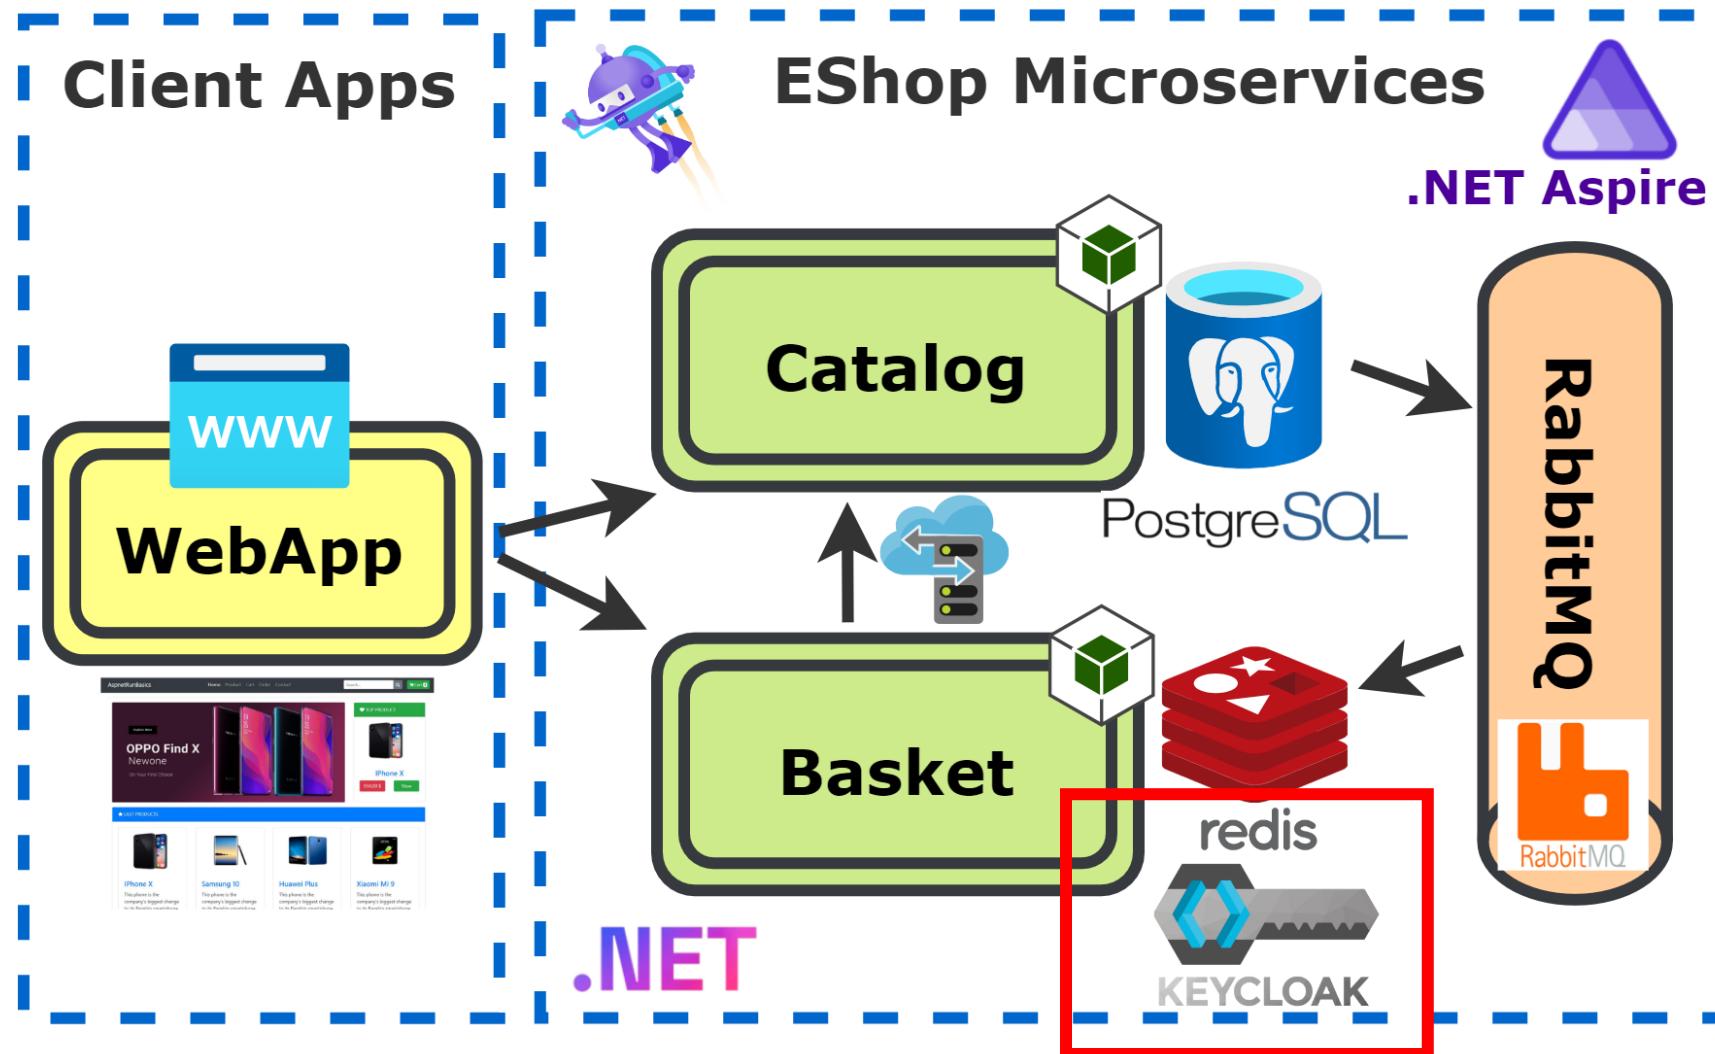
Create Realm, User and Client for OpenID Connect with Keycloak Identity Provider

Keycloak Client Integration - Secure Basket ms endpoints w/ .Net Aspire

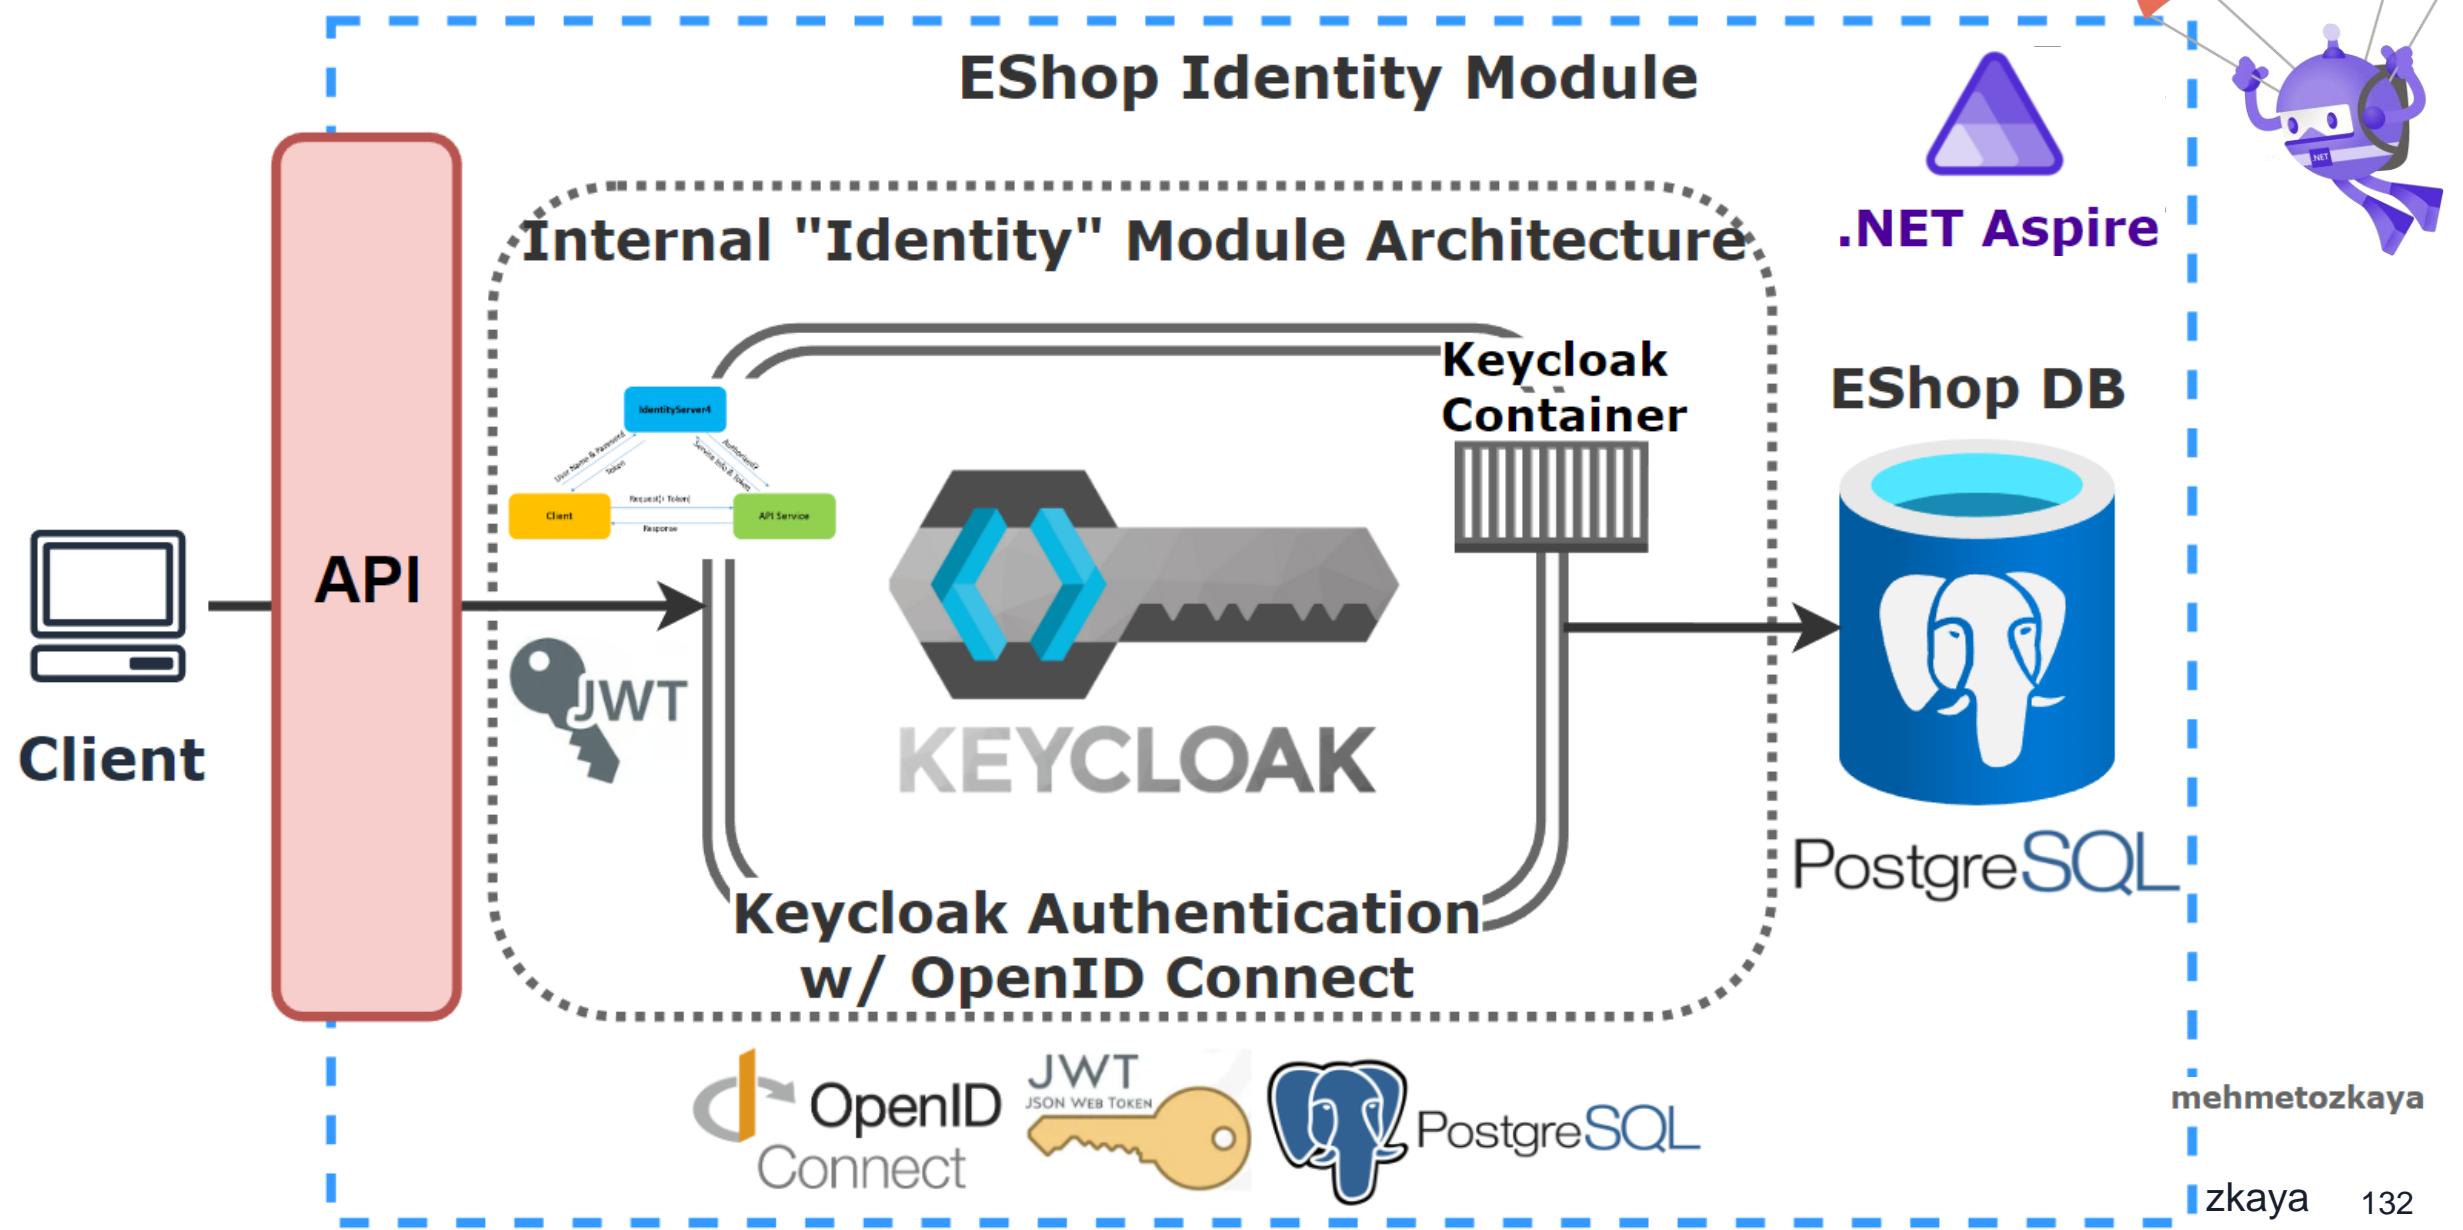
Mehmet Ozkaya



Secure Basket with Keycloak Authentication orchestrate .NET Aspire



Identity Microservice Internal Architecture



mehmetozkaya

zkaya 132

132

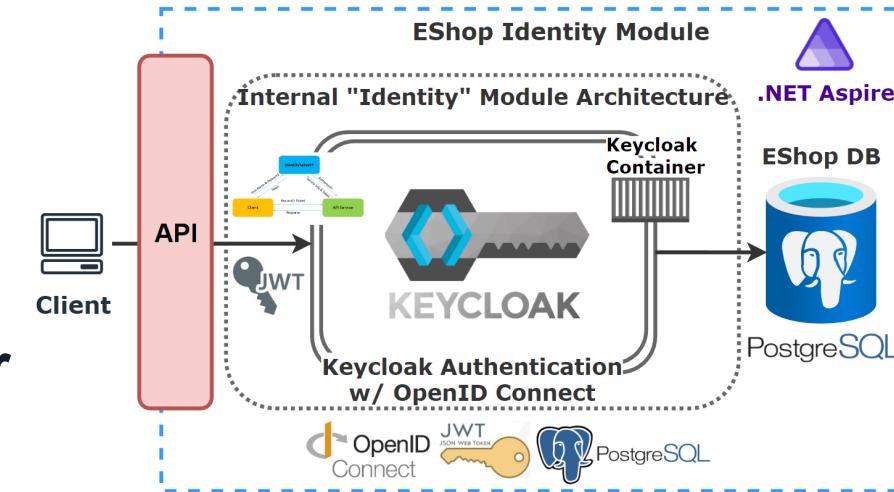
Secure Basket with Keycloak Authentication

Keycloak Identity and Access Management

OAuth2 + OpenID Connect Flows with Keycloak

Keycloak as a Backing Services into .NET Aspire

Setup Keycloak into .NET Aspire for Identity Provider



Definitions:

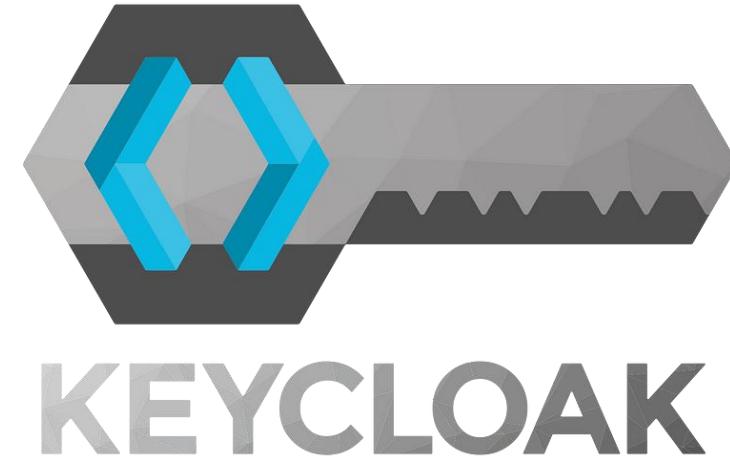
Create Realm, User and Client for OpenID Connect with Keycloak Identity

JwtBearer token for OpenID Connect with Keycloak Identity

Get Current User from Token with ClaimsPrincipal in Aspnet Authentication

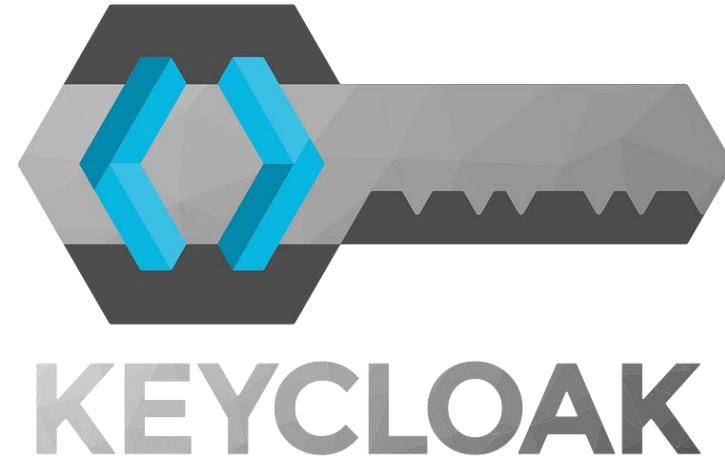
Keycloak Identity and Access Management

- **Open-source Identity and Access Management** solution aimed at modern applications and services.
- Provides **comprehensive support** for **authentication** and **authorization**, offering features such as **Single Sign-On (SSO)**, **Identity Brokering**, and **Social Login**.
- **Simplifies the process of securing apps and services with minimal coding effort.**
- **Single Sign-On (SSO)**: Users can log in once and access multiple app.
- **Identity Brokering and Social Login**: Integrate with external identity providers like Google, Facebook, etc.
- **User Federation**: Connect to existing user databases.
- **Centralized Management**: Manage all user identities and permissions in one place.
- **Standard Protocols**: Supports OAuth 2.0, OpenID Connect, SAML 2.0.



Why Use Keycloak ?

- Significantly **reduce the complexity of managing user identities** and **access controls** in your apps.
- **Security:** Robust security features that adhere to industry standards.
- **Scalability:** Capable of handling millions of users and requests.
- **Flexibility:** Easy to integrate with various applications and services.
- **User Experience:** Enhances user experience with SSO and Social Login.



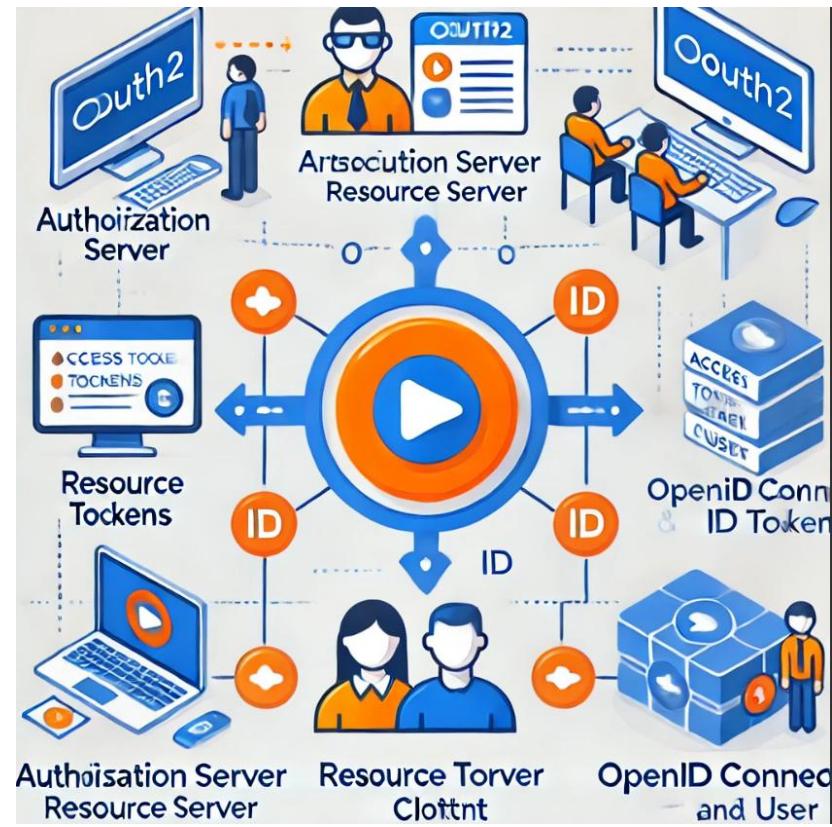
OAuth2 + OpenID Connect

OAuth2

- Authorization framework that allows apps to obtain limited access to user accounts on an HTTP service.
- Delegating user authentication to the service that hosts the user account and authorizing third-party apps to access the user account.
- **Key Concepts of OAuth2:** Resource Owner, Client, Authorization Server, Resource Server.

OpenID Connect (OIDC)

- Identity layer built on top of OAuth2.
- Clients to verify the identity of the end-user based on the authentication performed by an authorization server and to obtain basic profile information about the end-user.
- **ID Token:** A JSON Web Token (JWT) that contains user profile info.
- **UserInfo Endpoint:** An endpoint that returns additional profile information about the user.



OAuth2 + OpenID Connect Flows

Authorization Code Flow

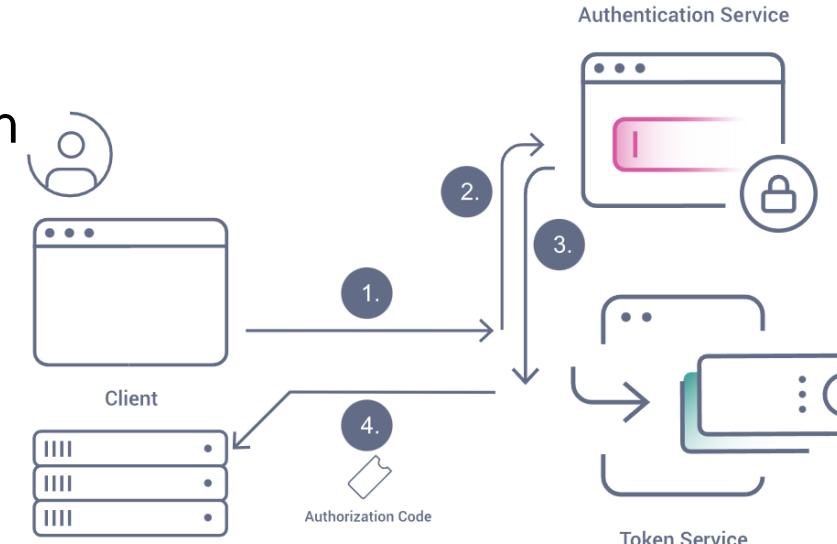
- 1. The client directs the user to the Keycloak login page.
- 2. The user authenticates and consents.
- 3. Keycloak redirects the user back to the client with an authorization code.
- 4. The client exchanges the authorization code for an access token and an ID token.

Implicit Flow

- SPAs where tokens are returned directly to the client. However, it's less secure and generally not recommended for new applications.

Client Credentials Flow

- Used for machine-to-machine (M2M) communication. The client authenticates directly with Keycloak using its client ID and secret to obtain an access token.



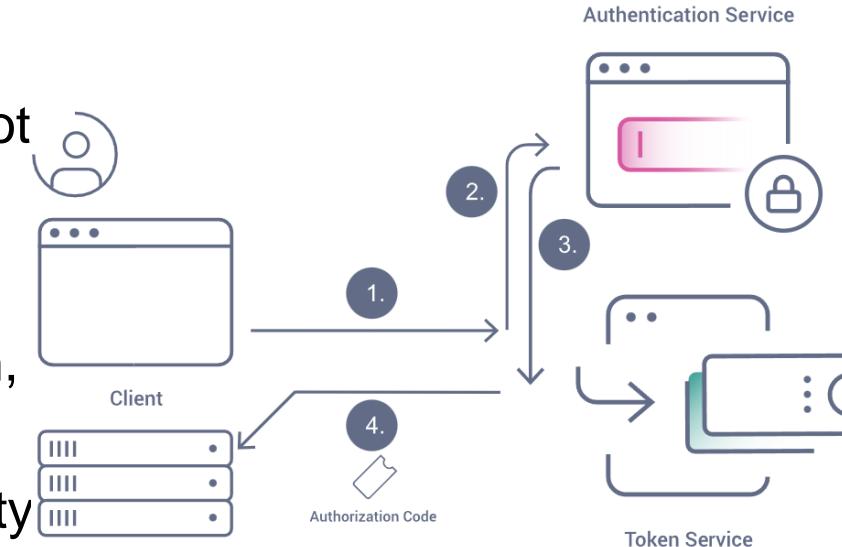
OAuth2 + OpenID Connect Flows - 2

Resource Owner Password Credentials Flow

- The client collects the user's credentials and exchanges them for an access token.
- This flow is suitable only for highly trusted clients and is generally not recommended.

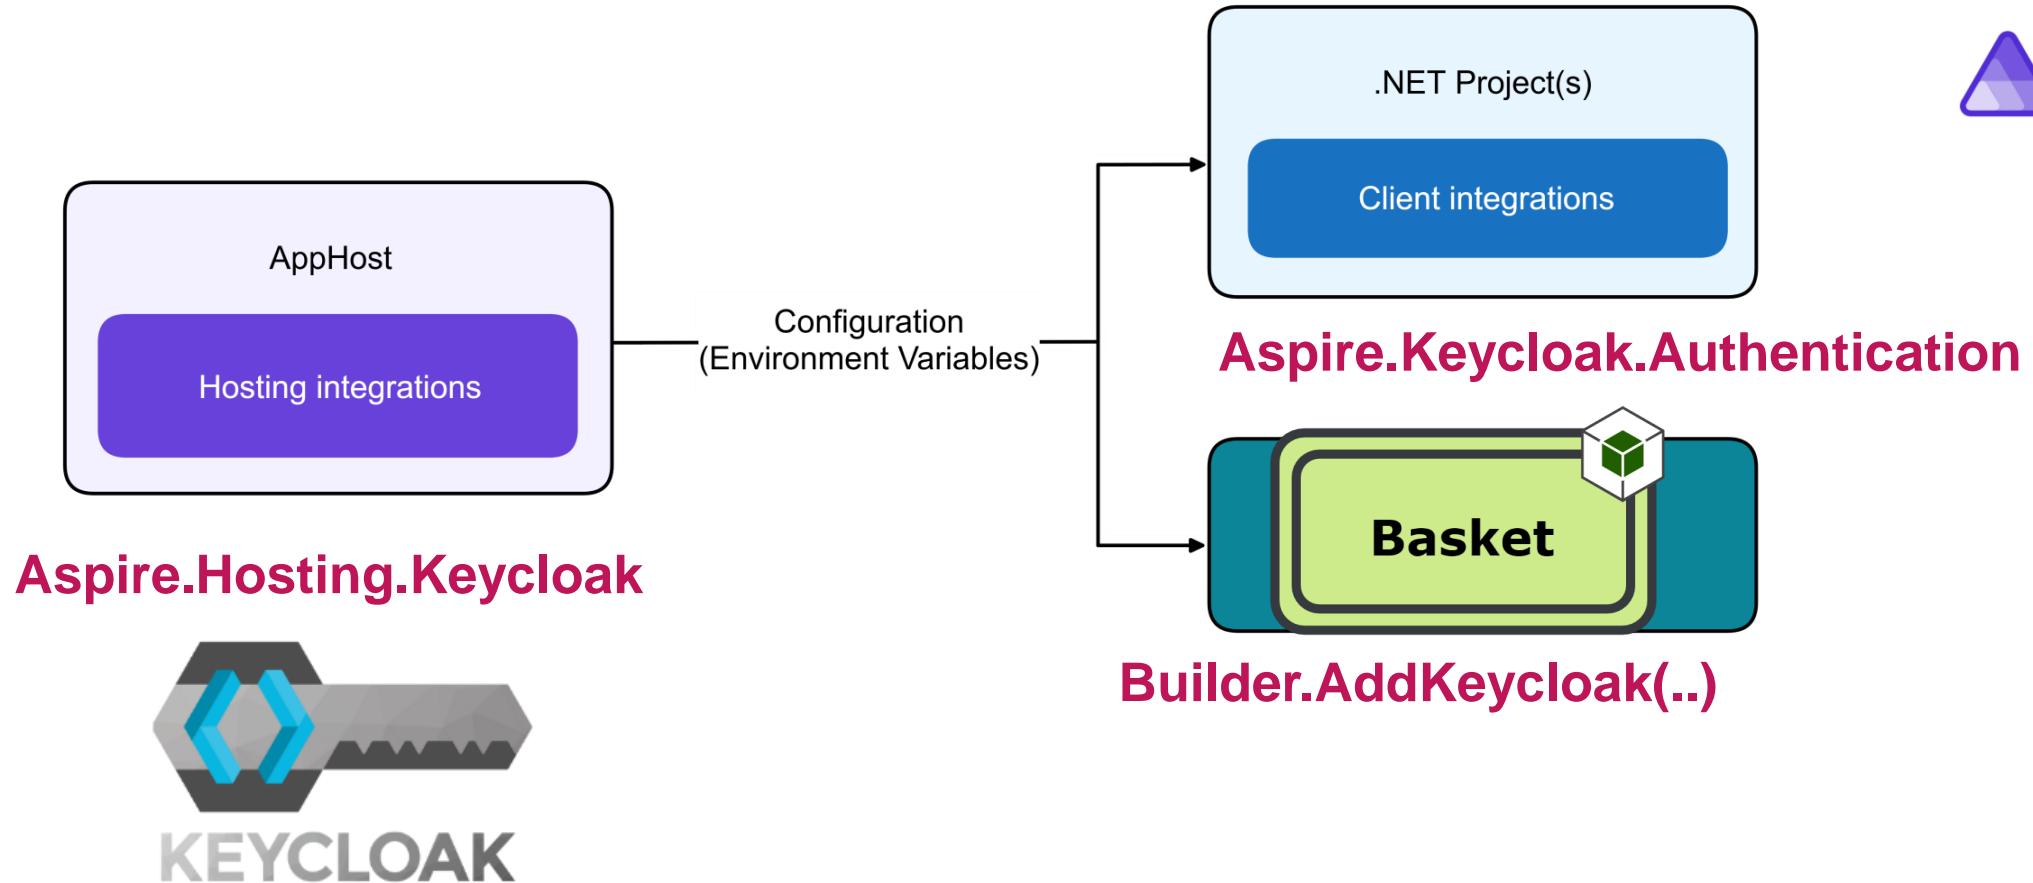
OpenID Connect Flow

- Extends the OAuth2 Authorization Code Flow to include an ID token, which contains user profile information.
- This flow is ideal for applications that need to verify the user's identity and obtain basic profile information.



Secure Basket with Keycloak Authentication

Integrations



Create Realm, User and Client for OpenID Connect w/ Keycloak Identity Provider

The image displays three screenshots of the Keycloak administration interface:

- Screenshot 1 (Top Left):** Shows the "Master" realm dashboard with a "Add realm" button highlighted in blue.
- Screenshot 2 (Top Right):** Shows the "General" view of the "Master" realm, featuring a "Login" button.
- Screenshot 3 (Bottom):** Shows the "Create client" wizard. The "General Settings" step is active, with the following configuration:
 - Client type:** OpenID Connect
 - Client ID:** test-client
 - Name:** (empty)
 - Description:** (empty)
 - Always display in console:** Off

Add user

ID:

Created At:

Username*: myuser

Email:

First Name: Foo

Last Name: Bar

User Enabled: ON

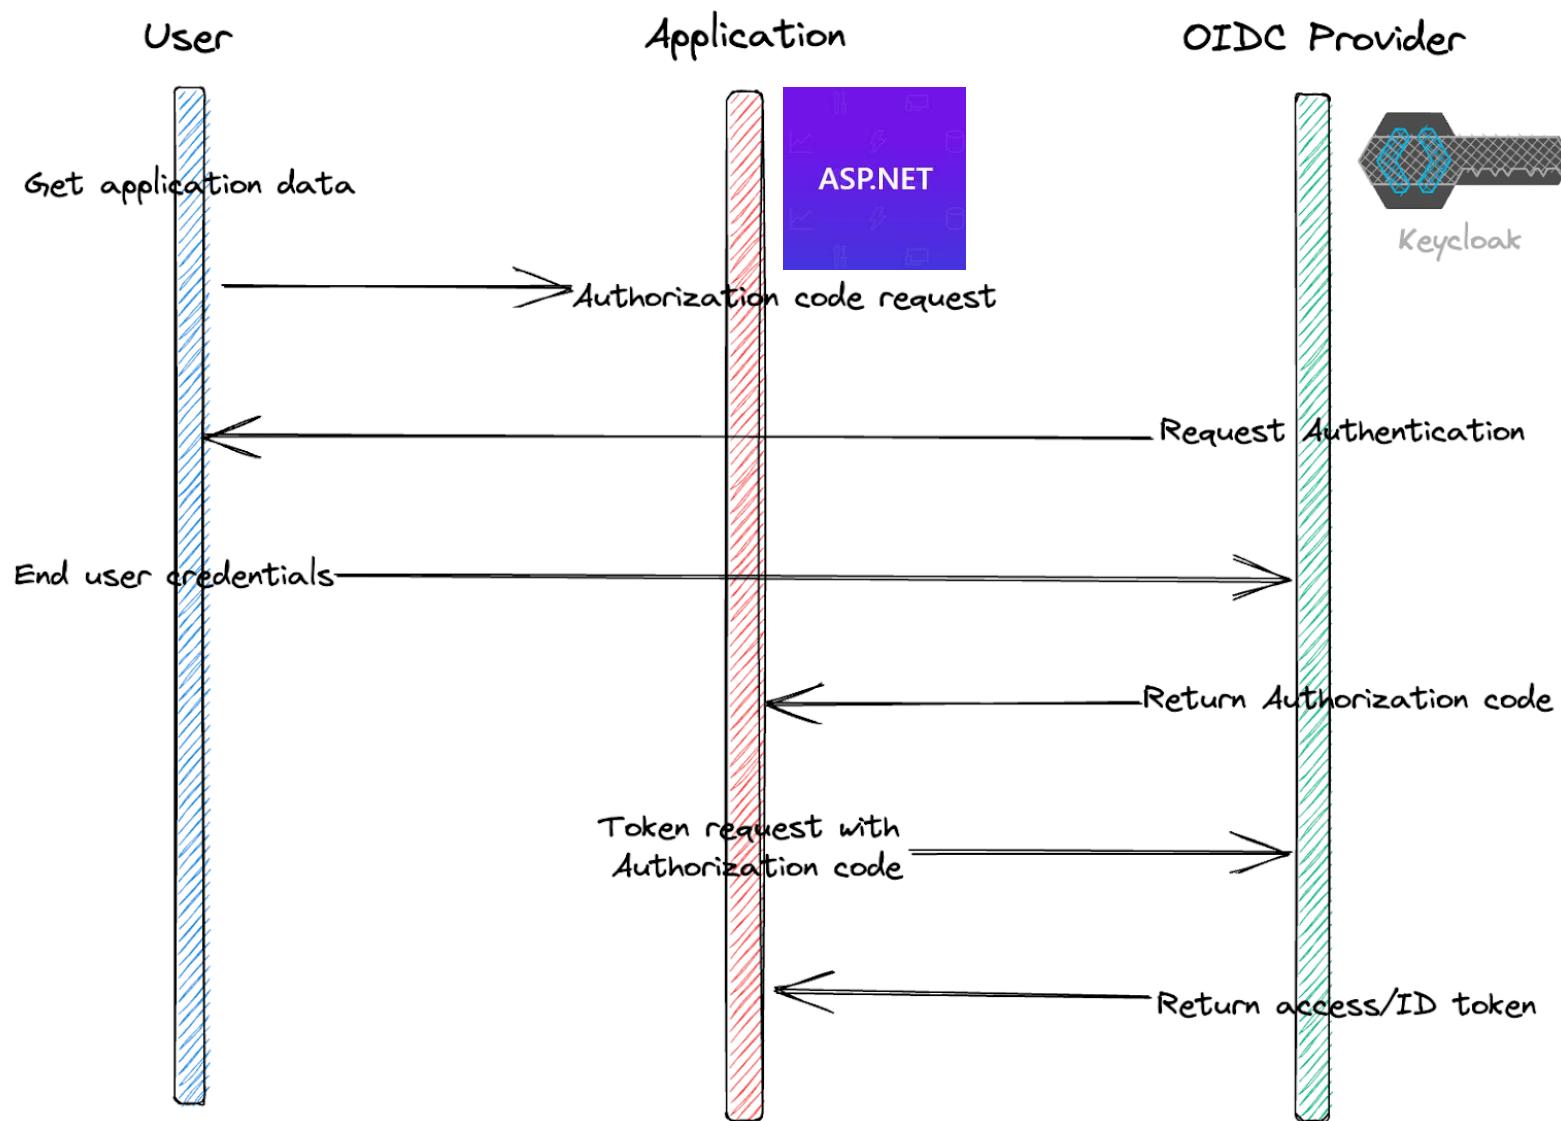
Email Verified: OFF

Required User Actions: Select an action...

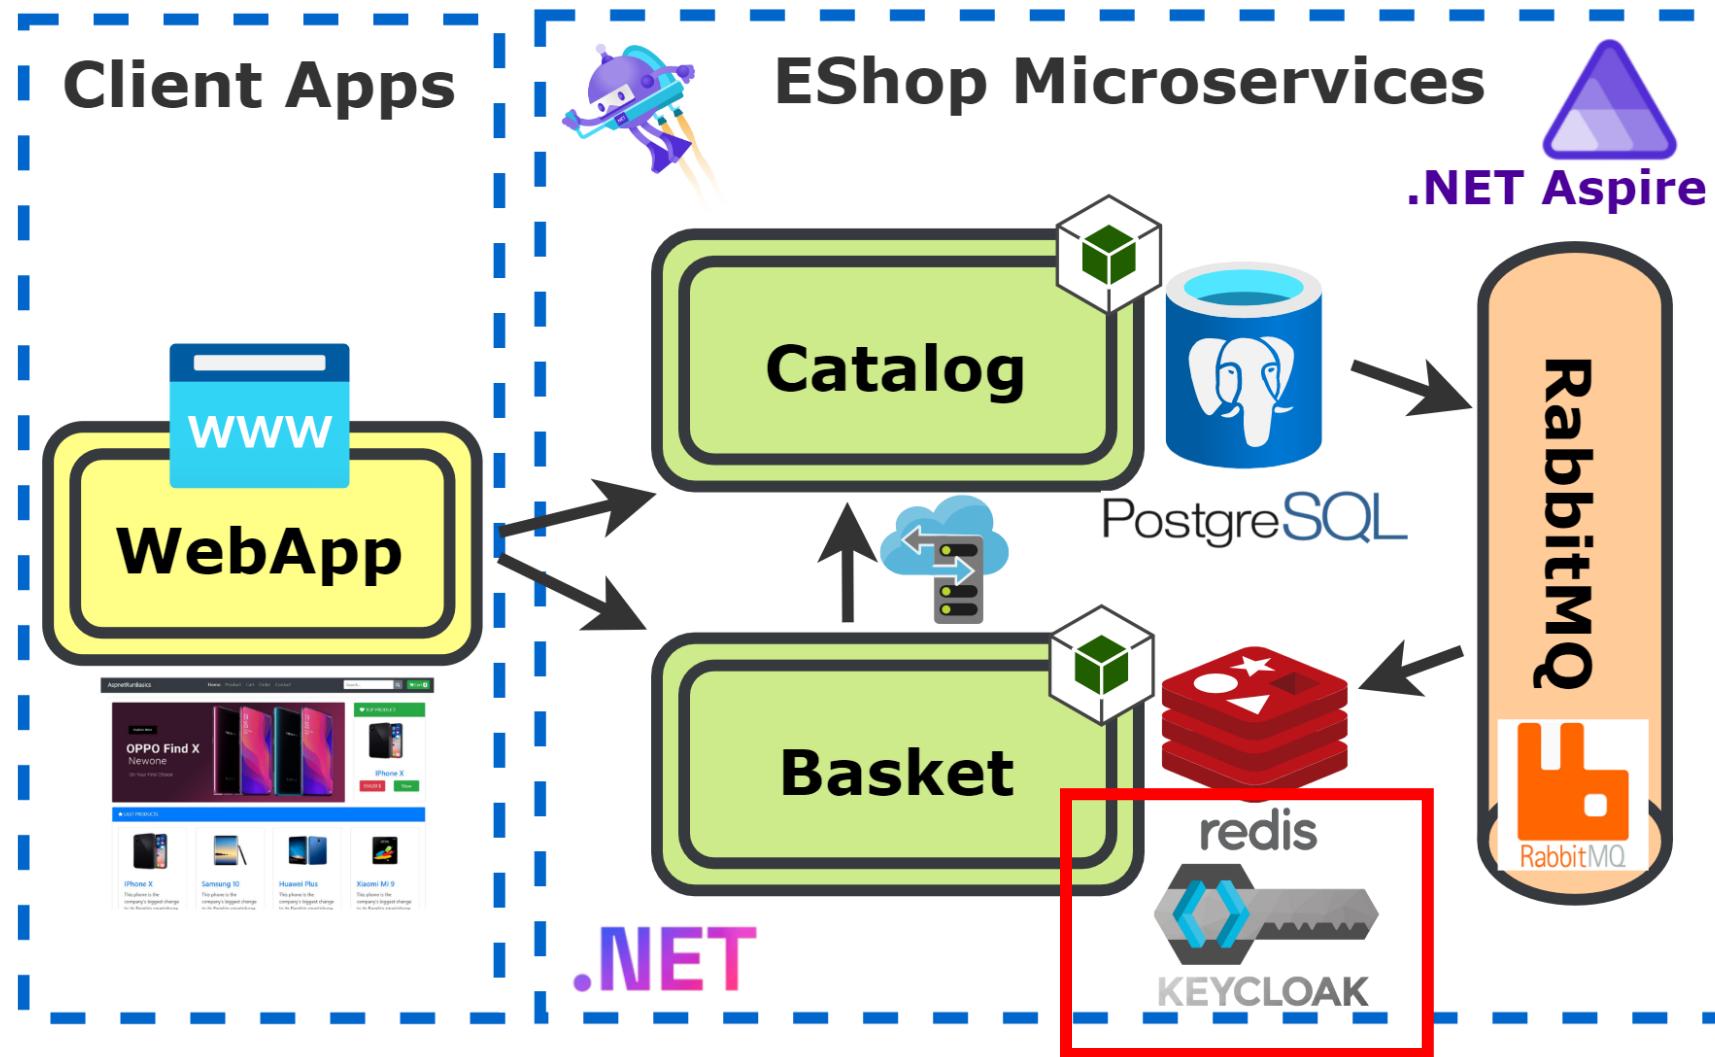
Save **Cancel**



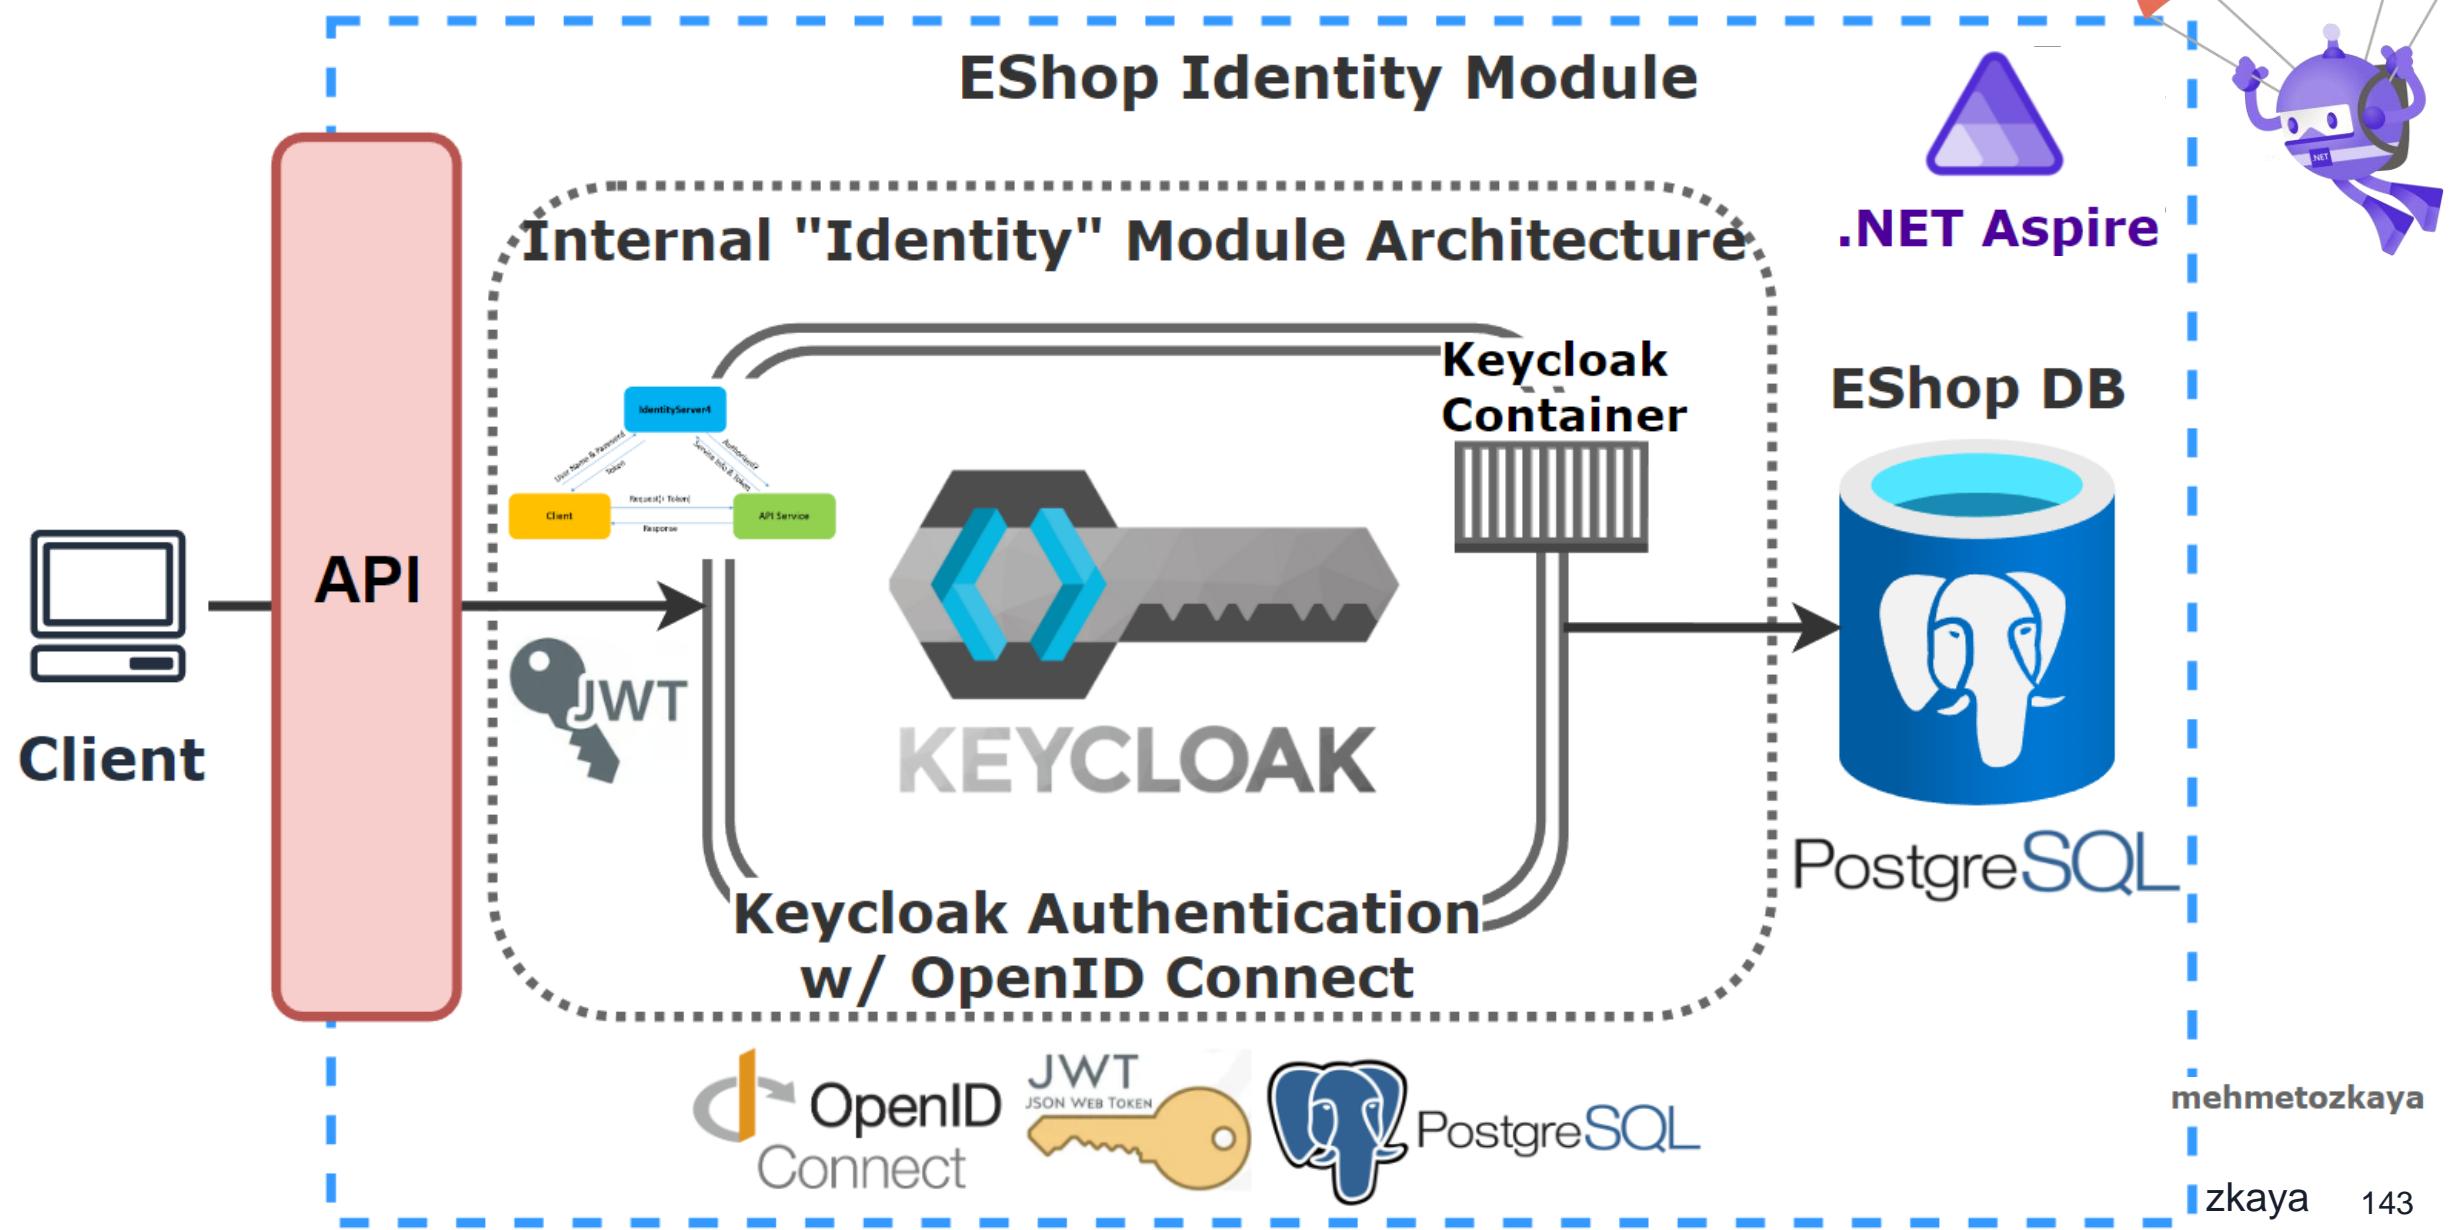
Secure EShop Modules with Keycloak OpenId Connect



Secure Basket with Keycloak Authentication orchestrate .NET Aspire



Identity Microservice Internal Architecture



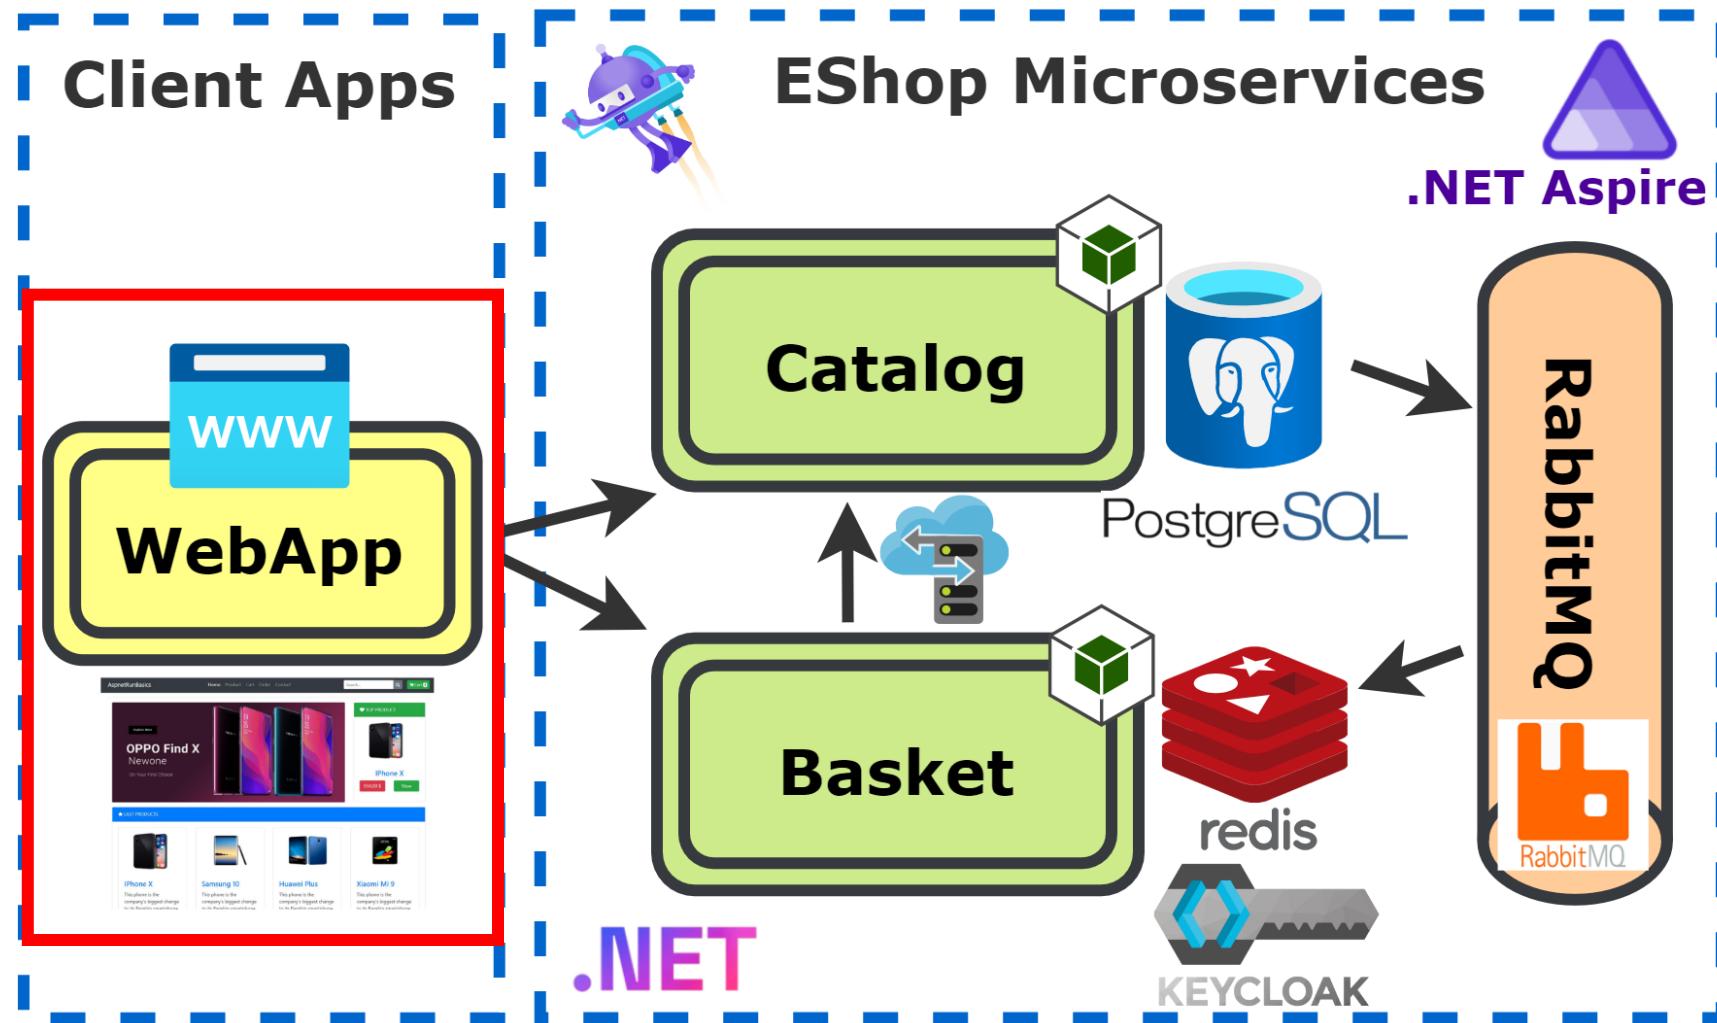
Develop Client Blazor Web Application

Create Blazor Web Application orchestrate .NET Aspire
Client WebApp integrate to Catalog w/ CatalogApiClient.cs
Register CatalogApiClient with Aspire integrations for Service Discovery
WebApp Product Page Development
Output Caching for Products Page

Mehmet Ozkaya



Develop Client Blazor Web Application



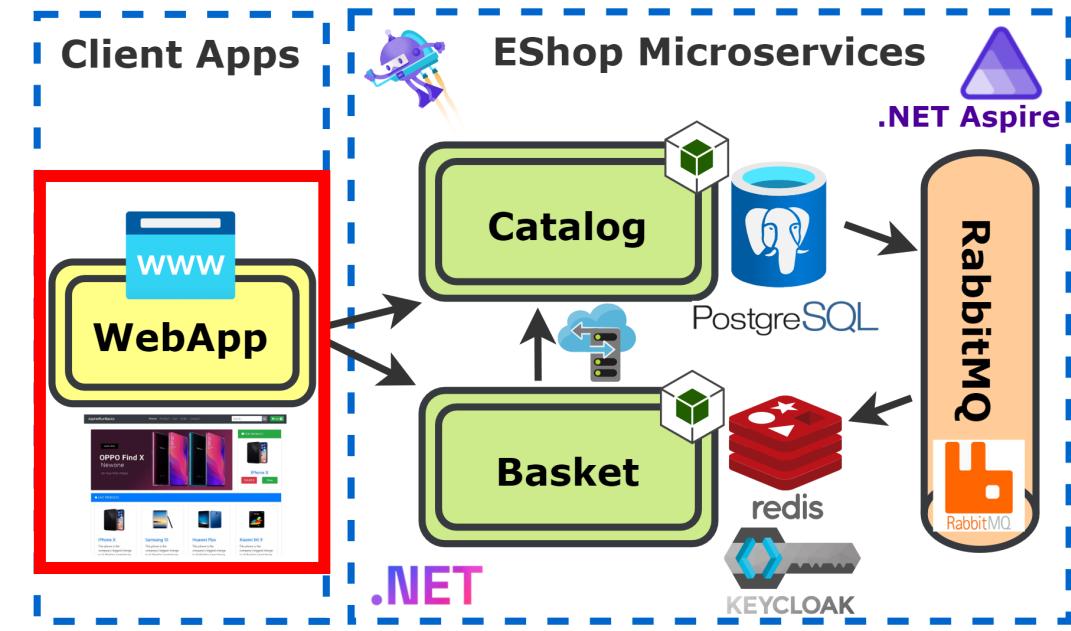
Develop Client Blazor Web Application

Blazor Web App and include .NET Aspire
Products razor page and populate Products from Catalog microservice

Output Caching

Later:

Deploy .NET Aspire distributed apps into ACA
Develop Semantic Search and AI-related pages



Develop Client Blazor Web Application

Products

Here are some of our amazing outdoor products that you can purchase.

Image	Name	Description
	Hiking Poles	Ideal for camping and hiking trips
	Outdoor Rain Jacket	This product will keep you warm and dry in all weathers
	Survival Kit	A must-have for any outdoor adventurer
	Outdoor Backpack	This backpack is perfect for carrying all your outdoor essentials

Remember: Service Discovery



Service
Discovery

Containers → injects → connection string

Projects → injects → microservice endpoints

```
var builder = DistributedApplication.CreateBuilder(args);

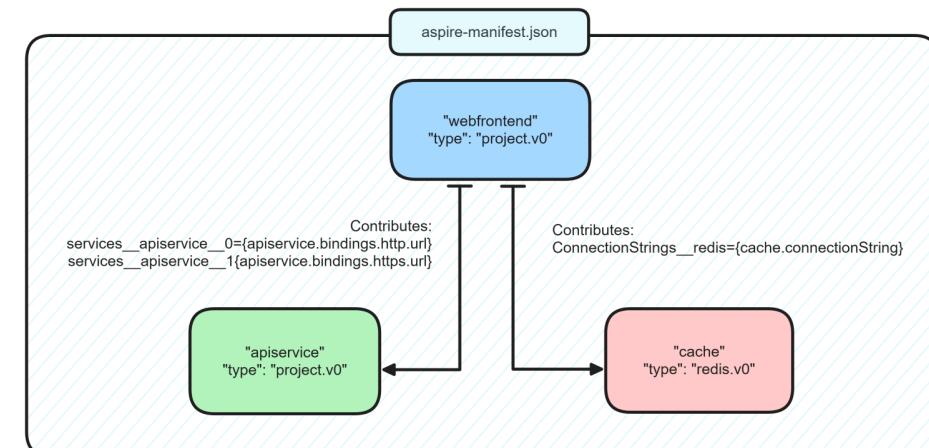
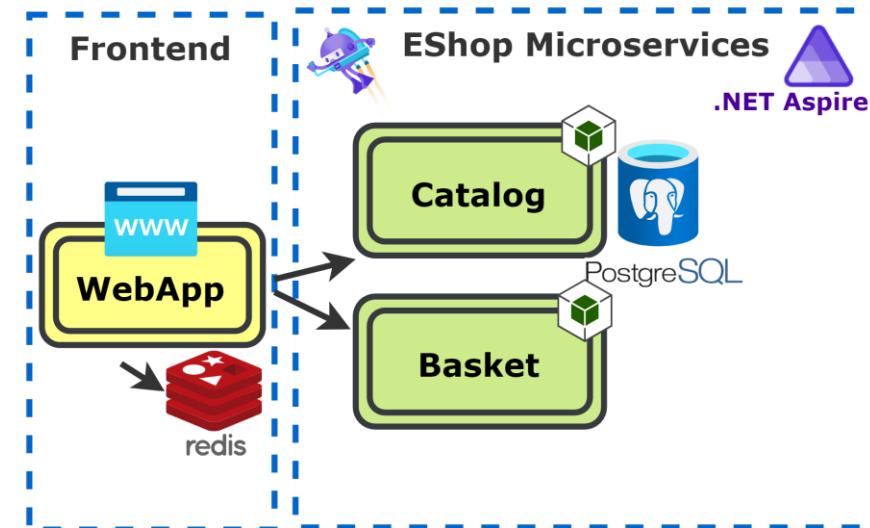
var cache = builder.AddRedis("cache");

var catalog = builder.AddProject<Projects.CatalogService>("catalog");
var basket = builder.AddProject<Projects.BasketService>("basket");

var frontend = builder.AddProject<Projects.MyFrontend>("frontend")
    .WithReference(cache)
    .WithReference(catalog)
    .WithReference(basket);
```

```
WithReference(cache)
    ConnectionStrings__cache="localhost:62354"
```

```
WithReference(apiservice)
    services__apiservice__http__0="http://localhost:5455"
    services__apiservice__https__0="https://localhost:7356"
```

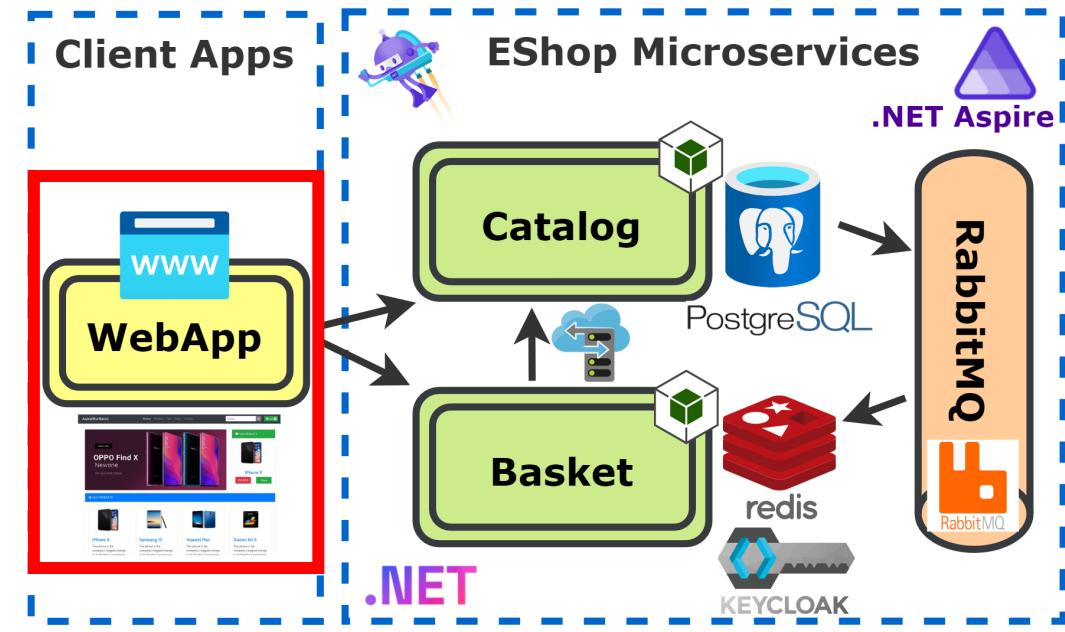


What Is Output Caching?

HTML output is cached, so repeated requests return the same content without re-running logic

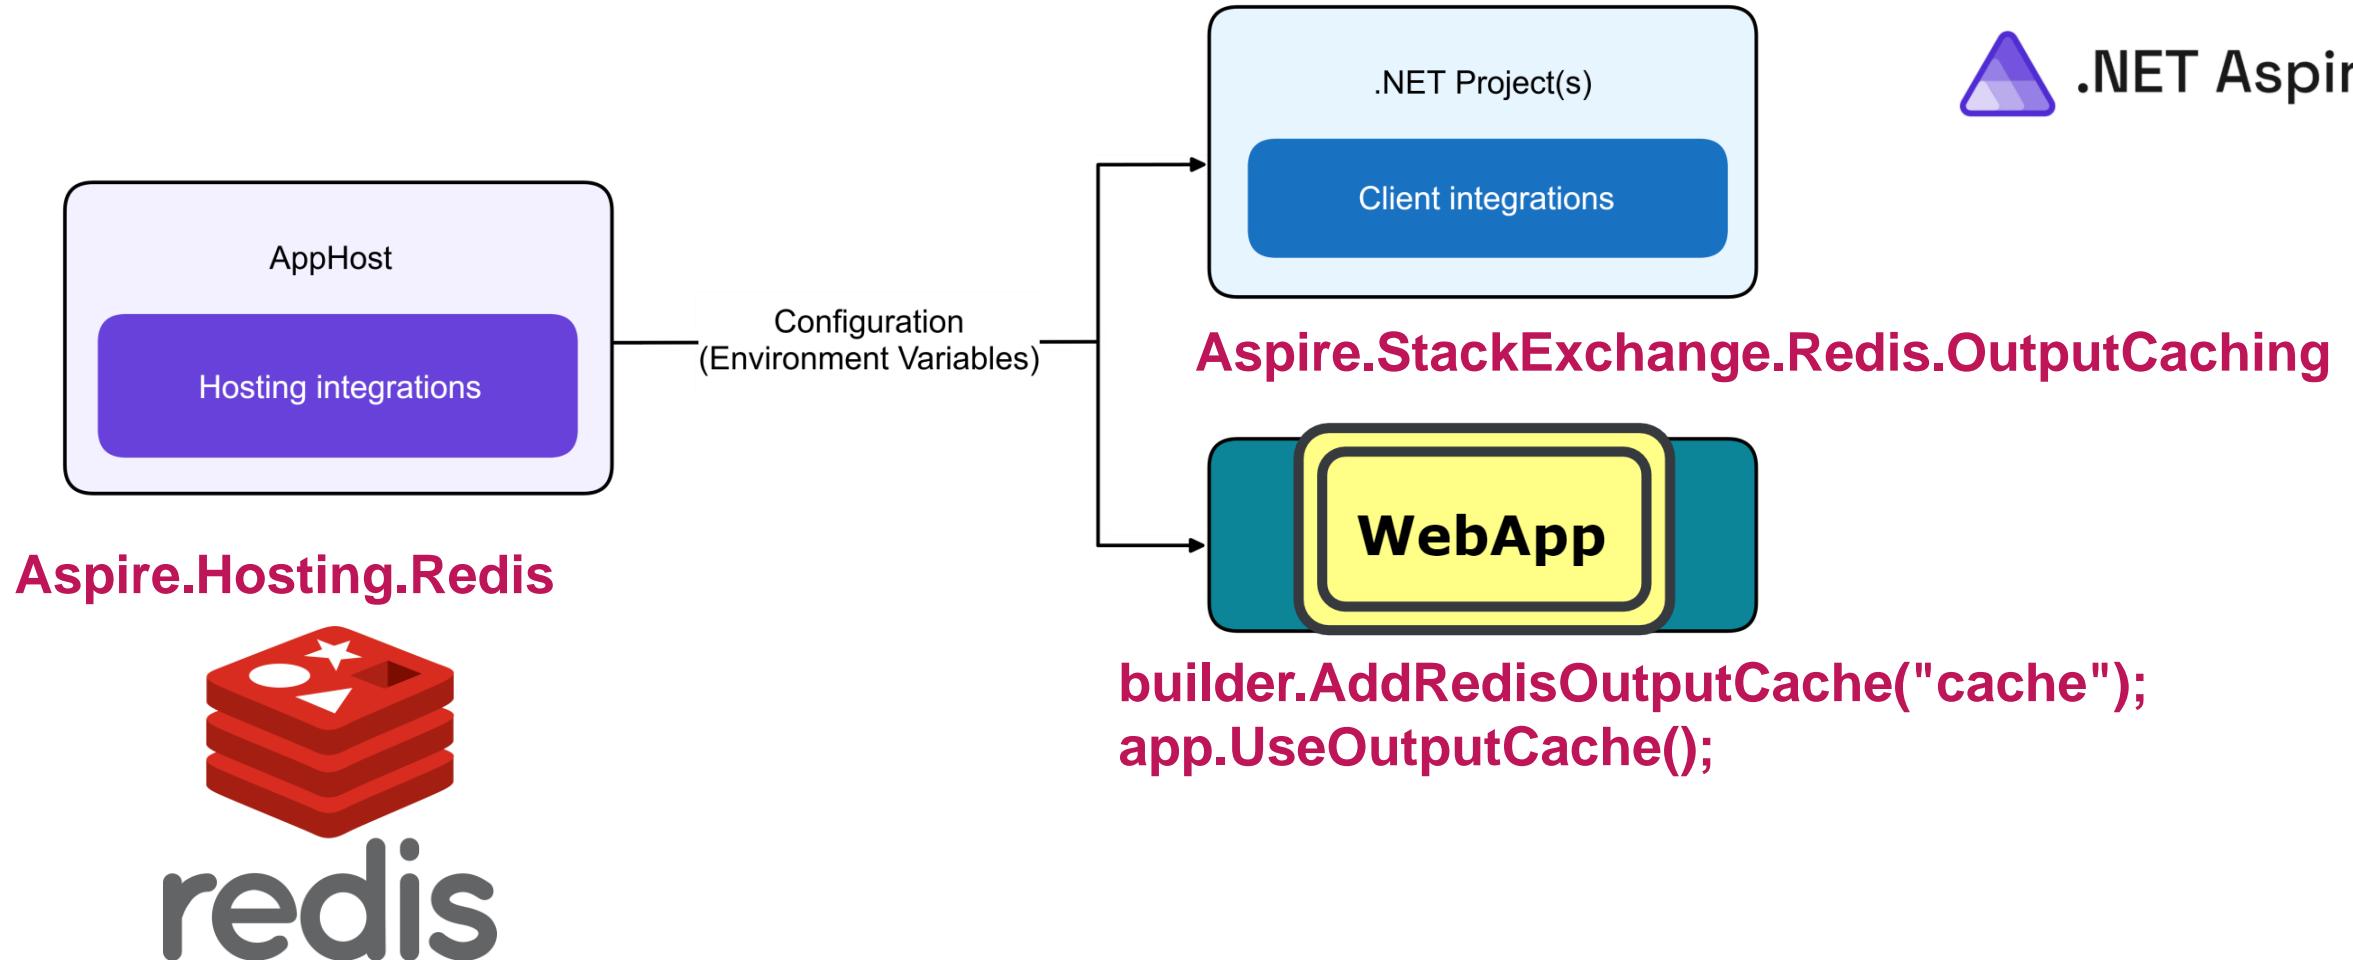
If data changes frequently, cache duration or invalidation must be considered

.NET Aspire integrates with Redis for distributed output caching



Secure Basket with Keycloak Authentication

Integrations



Deploy EShop Aspire Project to Azure Container Apps

What is ACA - Azure Container Apps ? How to Deploy ?

Data volumes don't work on ACA

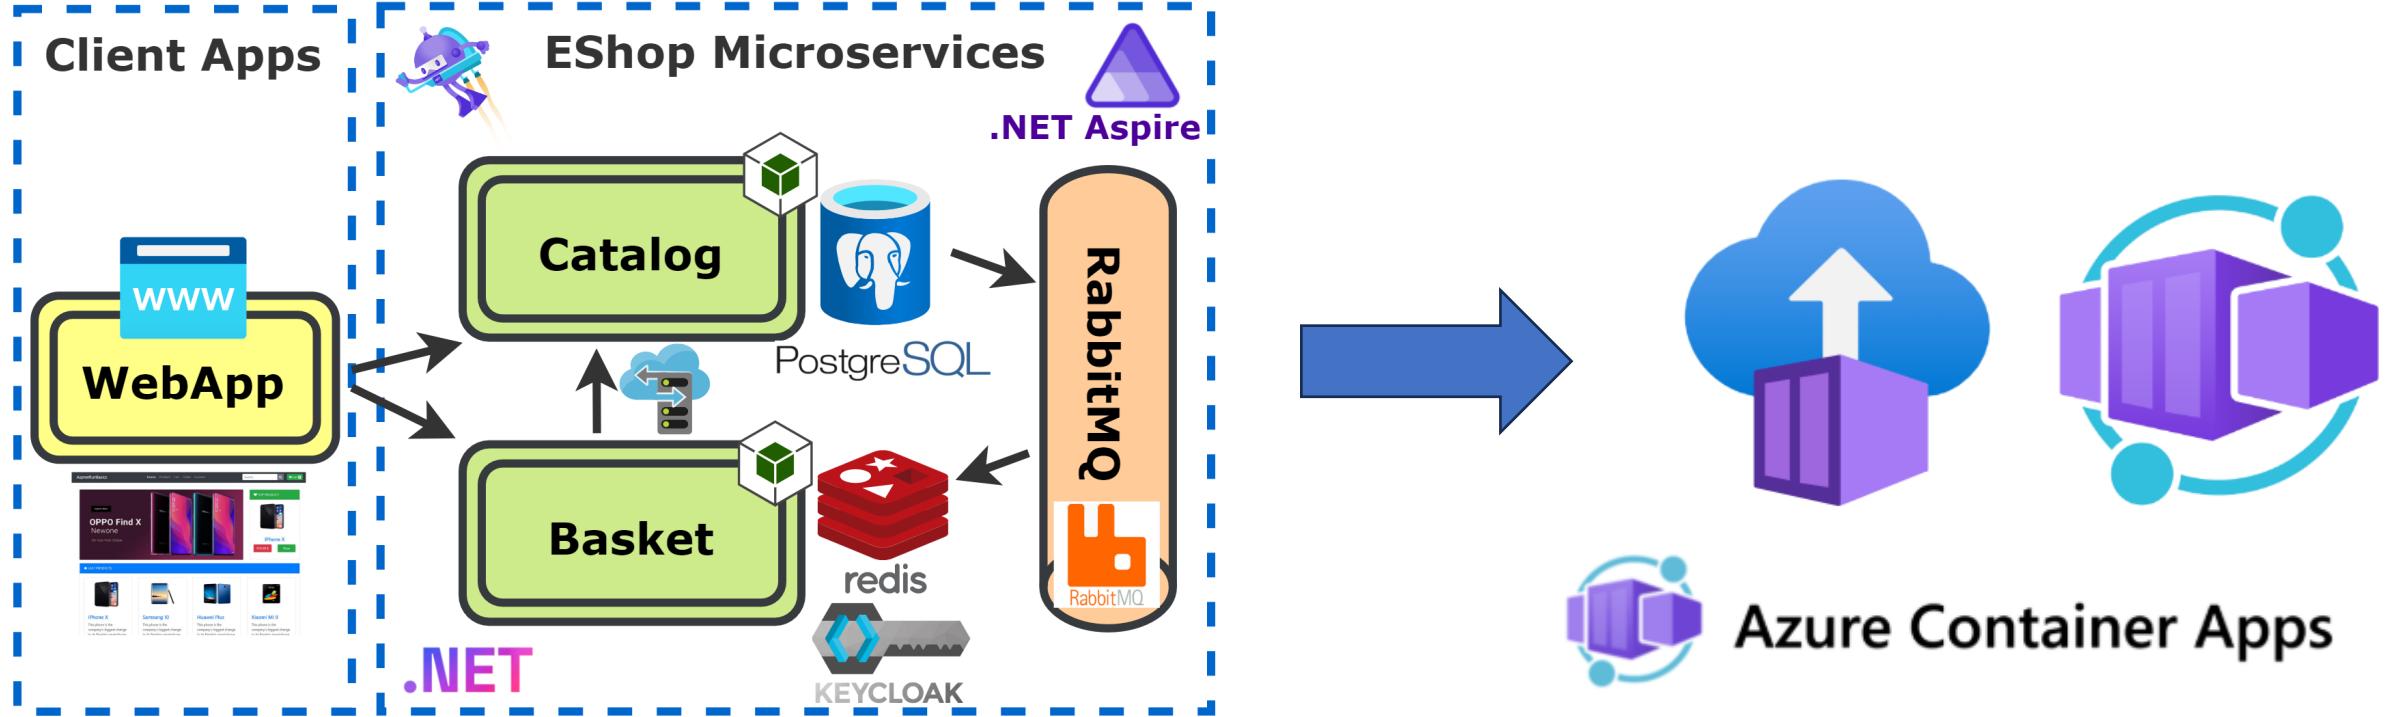
Deploy .NET Aspire App with azd commands to ACA

Test Successful Deployment to ACA

Mehmet Ozkaya



Deploy EShop Aspire Project to Azure Container Apps



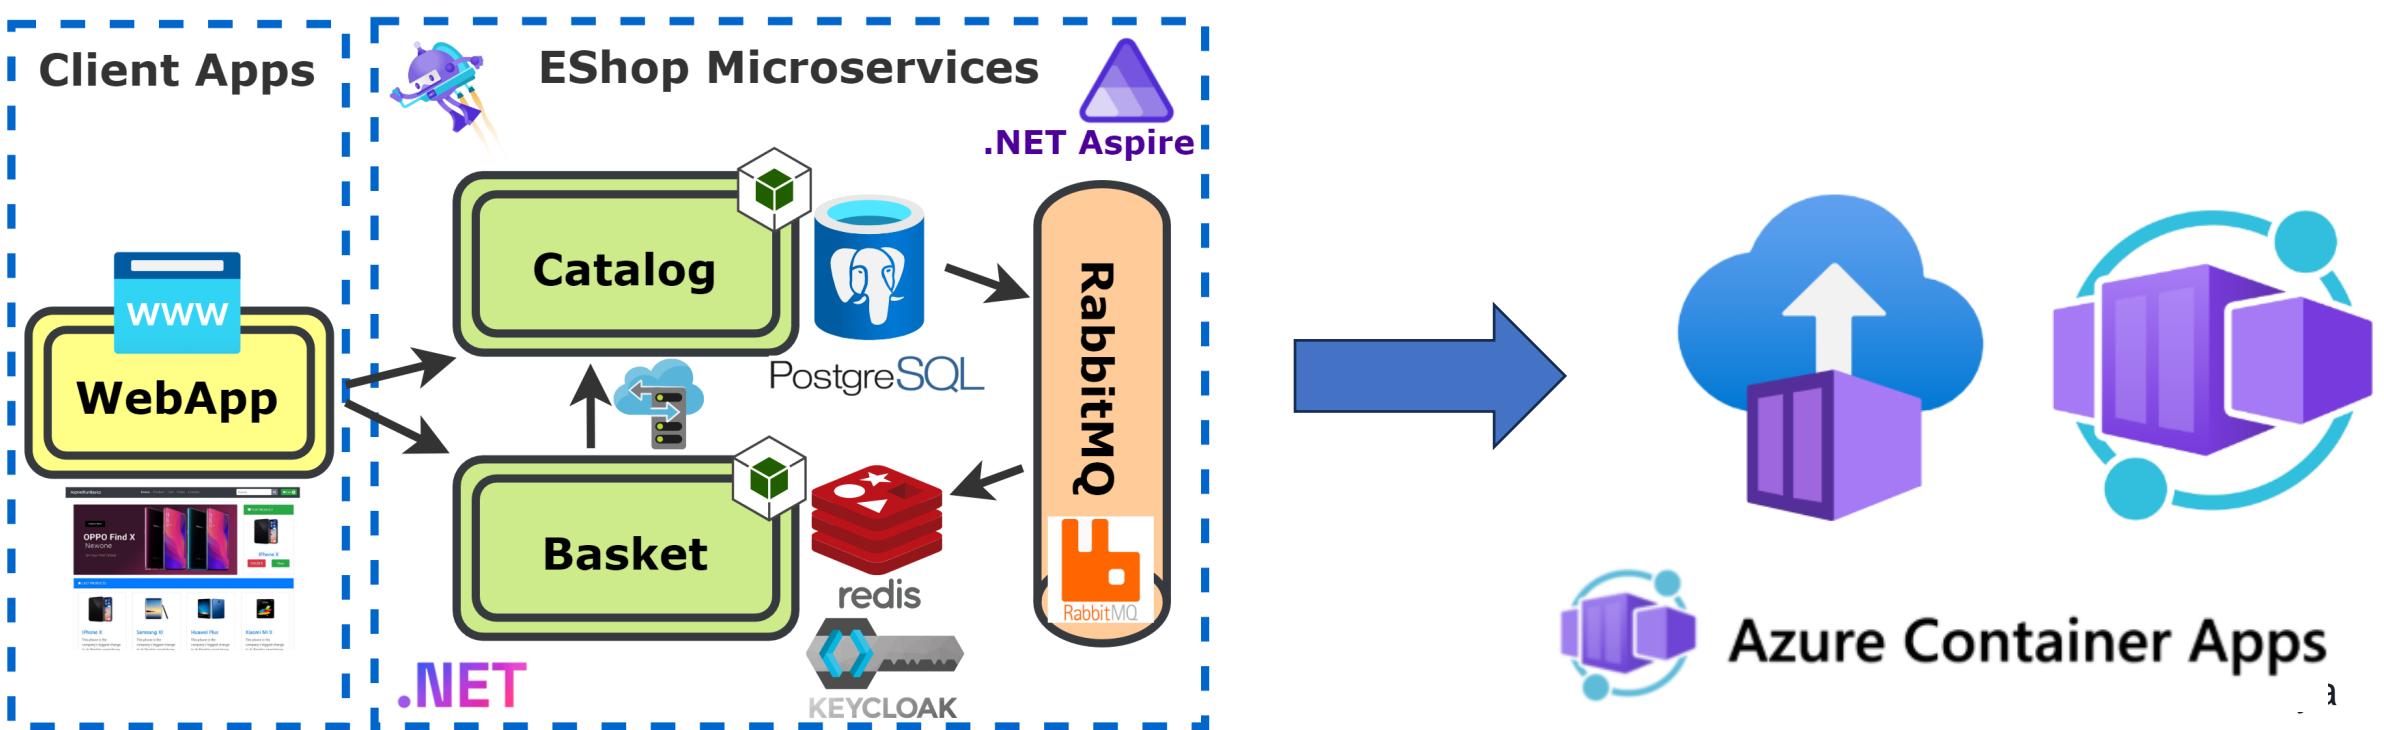
Deploy EShop Aspire Project to Azure

Azure Container Apps

Deploy a .NET Aspire project to Azure Container Apps

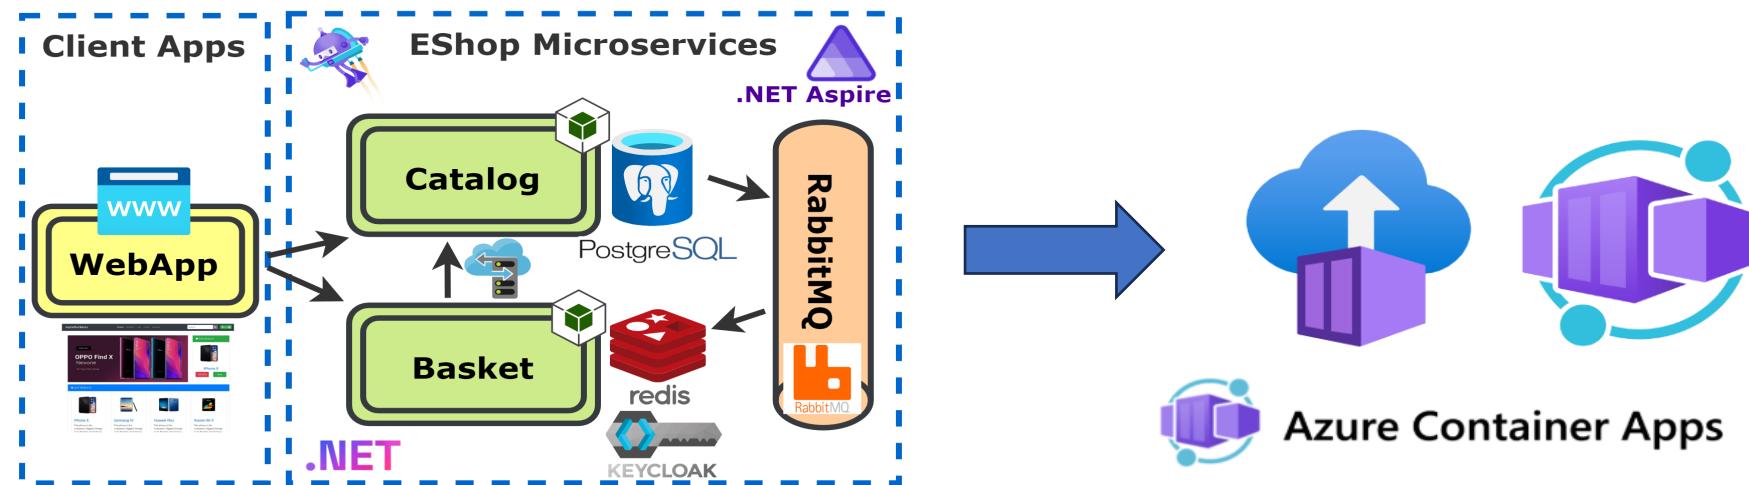
Deploy .NET Aspire projects with azd (Azure Developer CLI)

Azure Dashboard



Steps to Deploy EShop Aspire Project to Azure

1. Provision an Azure resource group and Container Registry
2. Publish the .NET Aspire projects as container images in Azure Container Registry
3. Provision a backing containers (postgres, redis) in Azure
4. Deploy the apps to an Azure Container Apps environment
5. View application console logs to troubleshoot application issues

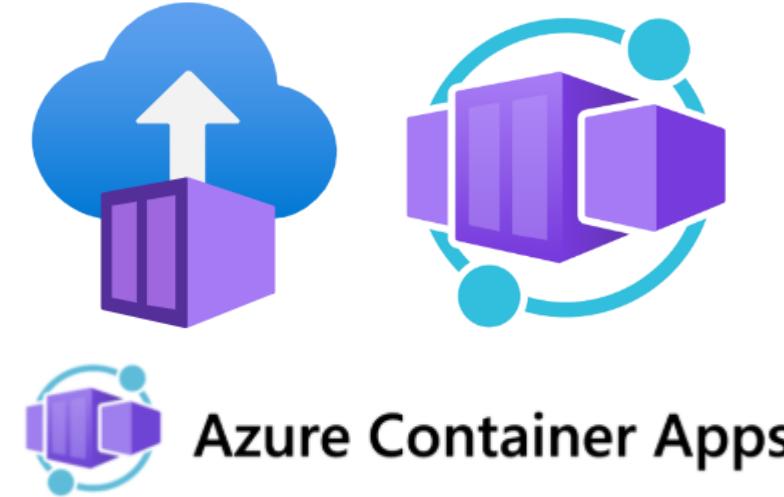


Azure Container Apps (ACA)

Microsoft's managed container hosting for microservices

Simplifies running containers in a serverless fashion

Perfect for .NET Aspire solutions needing easy scale & integrated logs



Key Benefits of Azure Container Apps

Automatic scaling and revision management

Dapr integration (optional) for sidecar patterns

KEDA for event-based scaling

Simple **container-based deployment** from the command line



Azure Container Apps

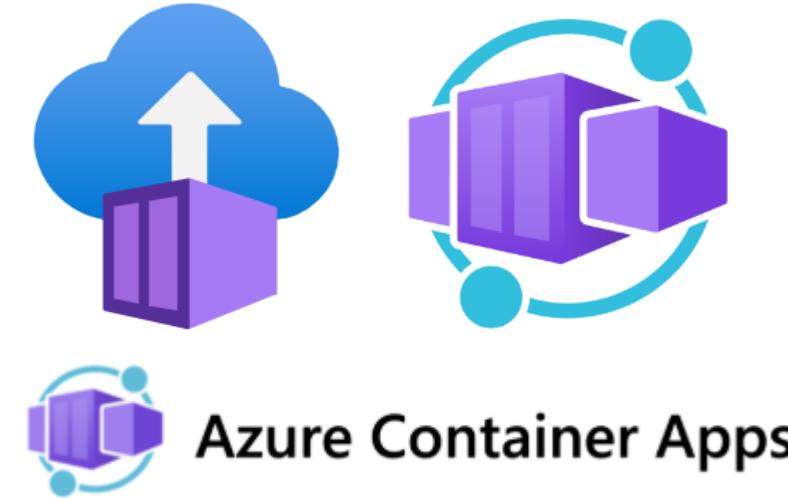
Deploying a .NET Aspire Project to ACA

Containerize each microservice automatically with
.NET Aspire

`azd init` to set up environment files

`azd up` to provision Azure resources

`azd down` to tear them down



Detailed azd commands

azd init

Creates .azure folder or config files for naming & region choices

Optionally picks an existing subscription

Typically done once per solution



Azure Container Apps

azd up

Builds containers, pushes to Azure Container Registry

Provisions Container Apps environment, secrets, and logs

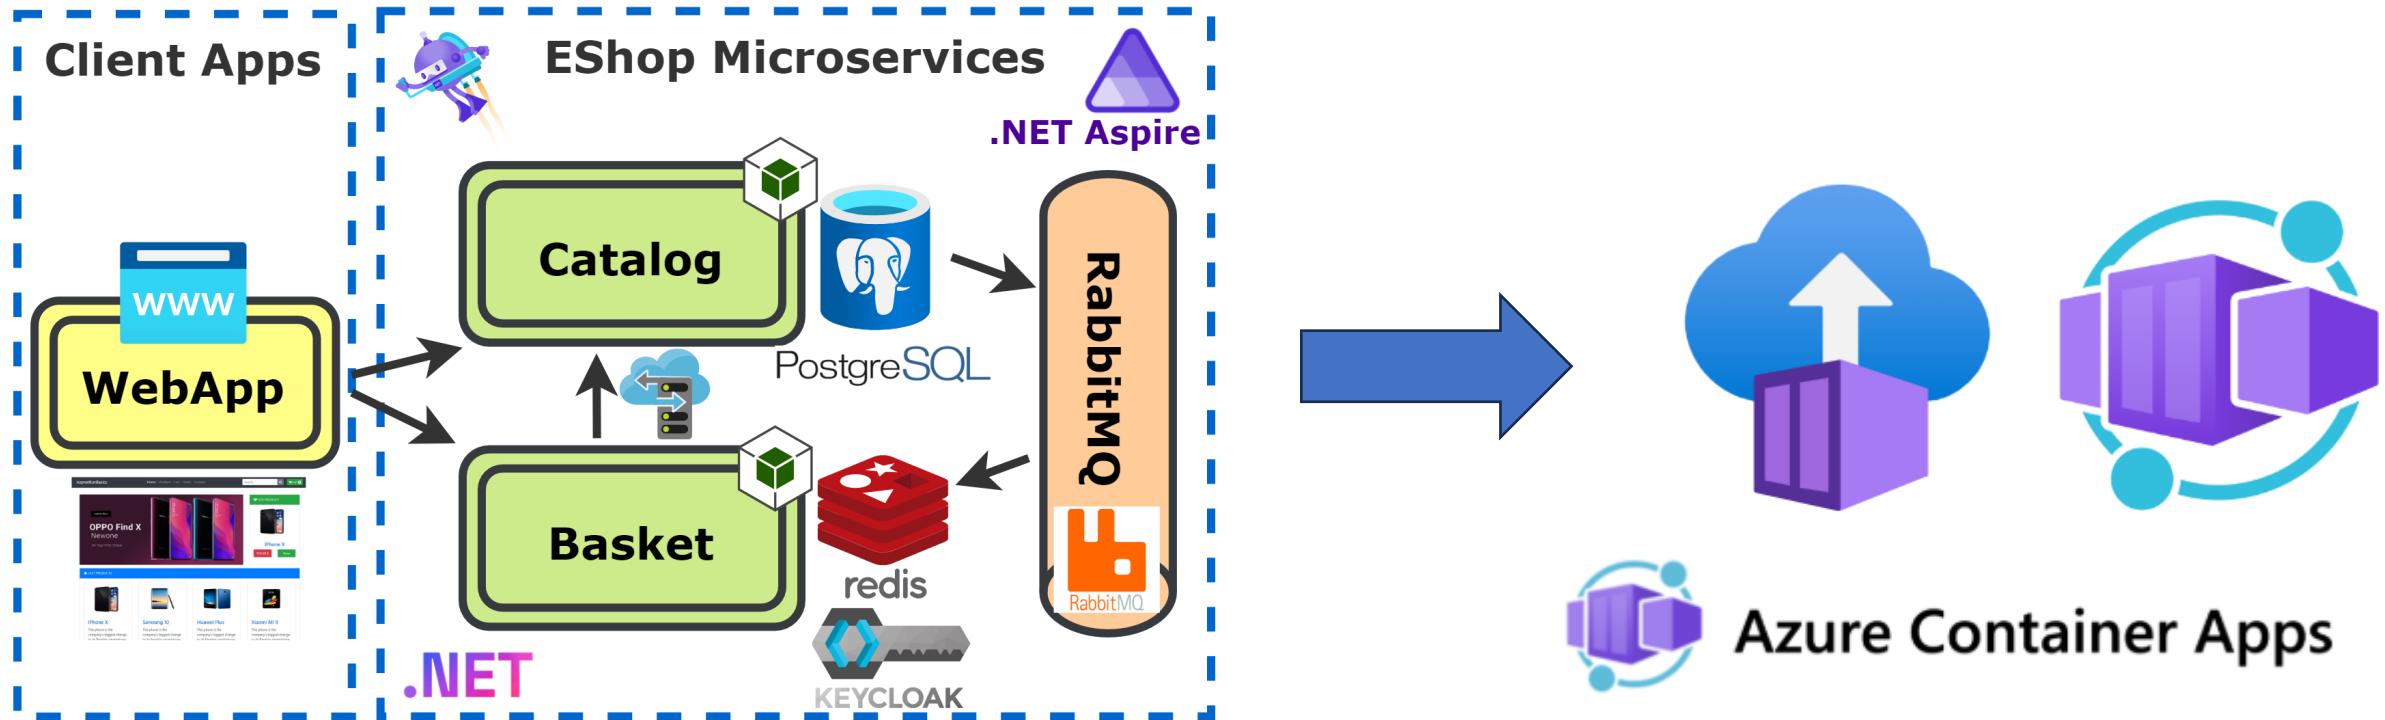
Deploys each microservice in your .NET Aspire solution

azd down

Frees up resource usage in your subscription

Removes container apps, registry, logs, secrets

Deploy EShop Aspire Project to Azure Container Apps

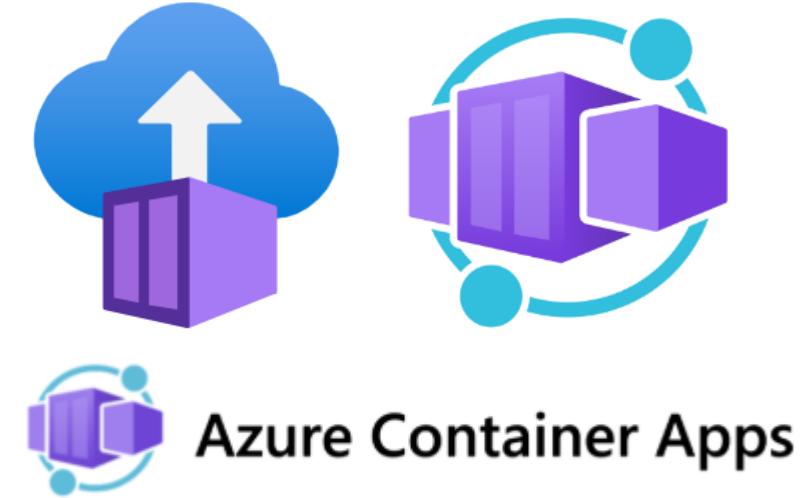


Data volumes don't work on ACA

ACA has Azure Storage limitations for container volumes

Some Linux containers (Postgres, etc.) can't mount volumes properly

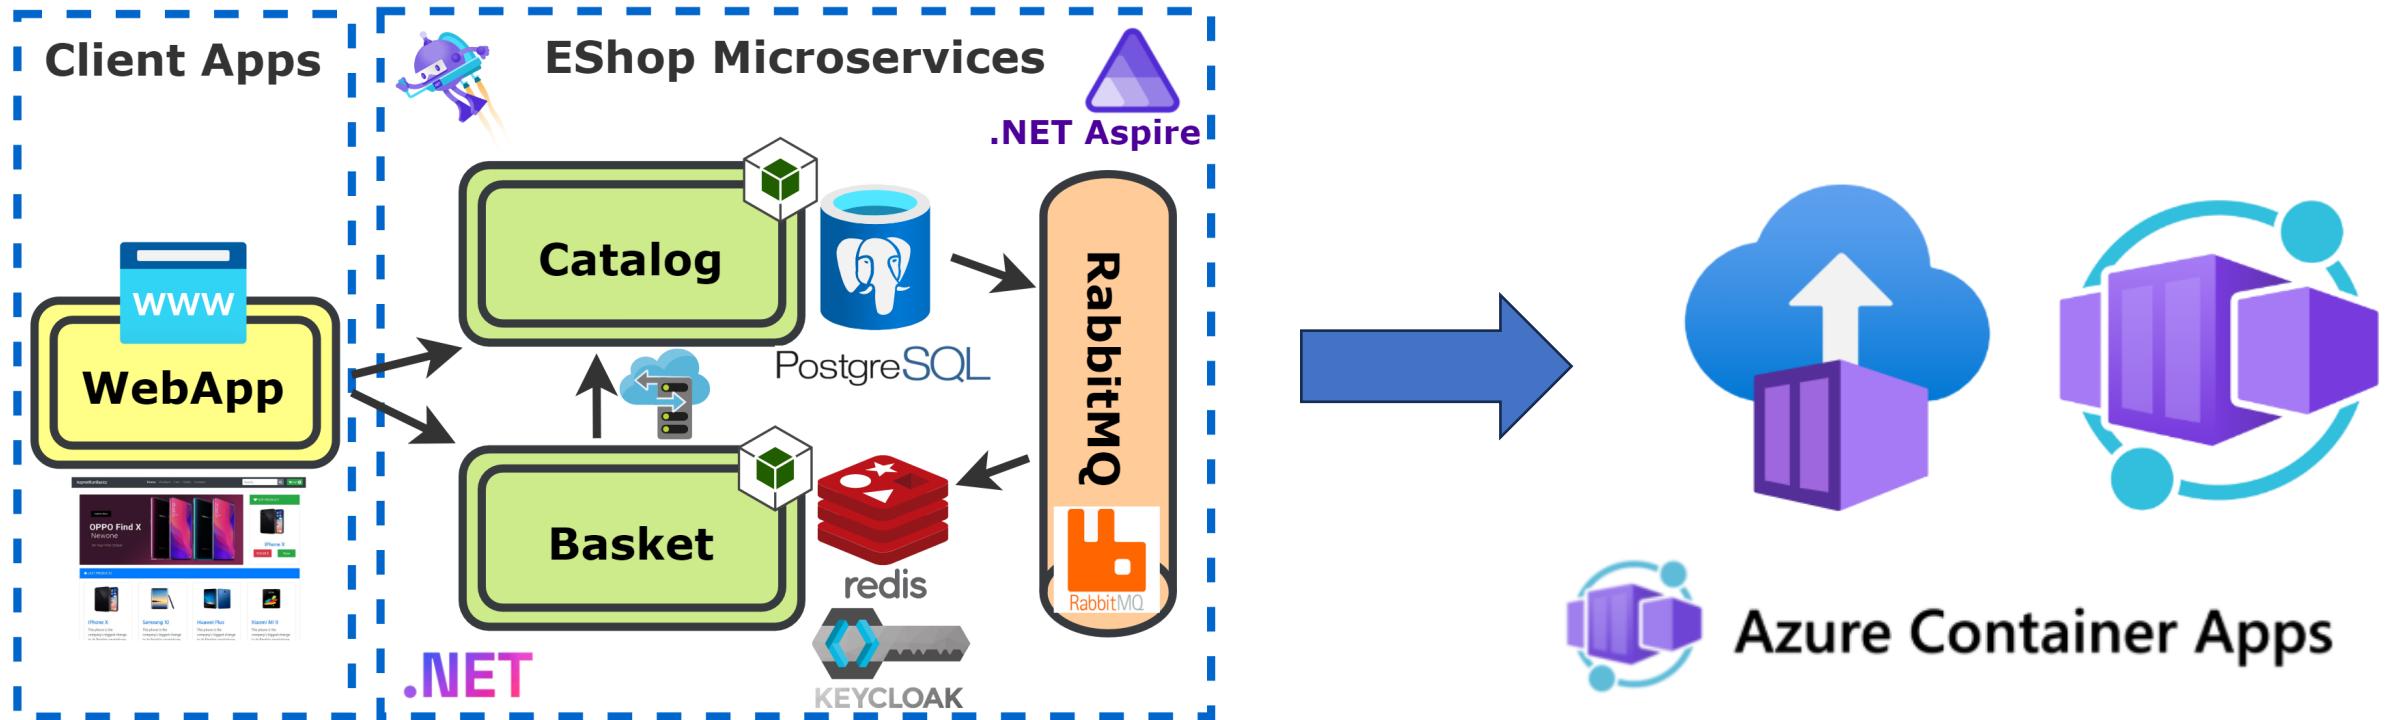
.NET Aspire approach: conditionally add volumes or switch to cloud services



[dotnet/aspire #6671 GitHub Issue](#)

[.NET Aspire Docs – Azure Postgres Resource](#)

Deploy EShop Aspire Project to Azure Container Apps



Prerequisites

An Azure subscription (e.g., Free Trial)

azd CLI installed on your machine

A .NET Aspire solution with container-based microservices

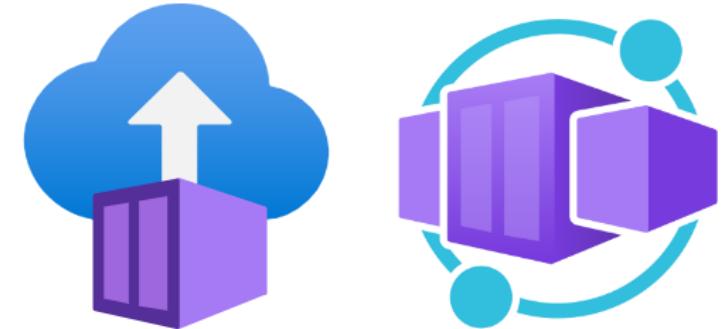
`azd version`

`azd auth login`

`azd init`

`azd up`

`azd down`



Azure Container Apps

.NET GenAI with Microsoft.Extensions.AI for Chat AI and Semantic Search

What are Large Language Models (LLMs) ?

Ollama – Run LLMs Locally

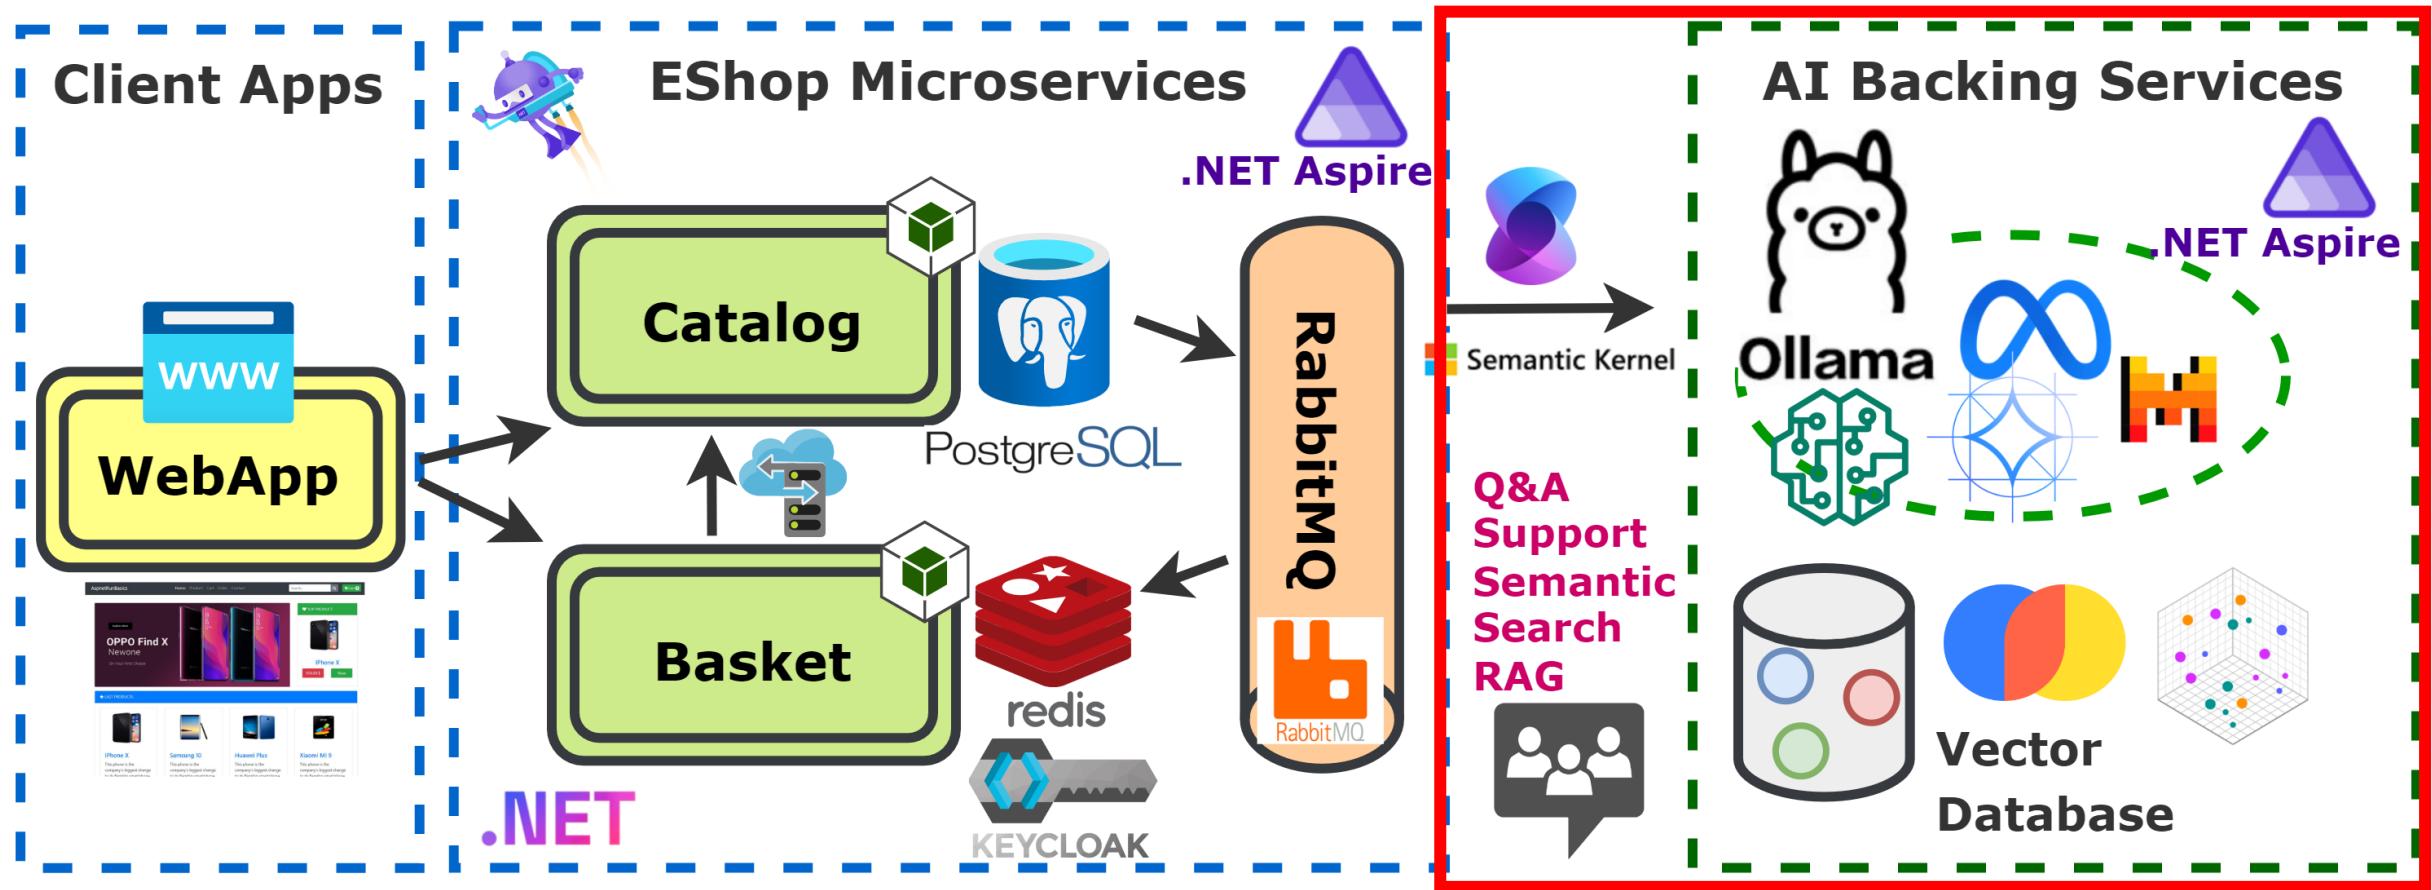
Using LLMs and VectorDBs as Cloud-Native Backing Services

Download Llama Model LLM inside Ollama with .NET Aspire

Mehmet Ozkaya



.NET GenAI with Microsoft.Extensions.AI for Chat AI and Semantic Search



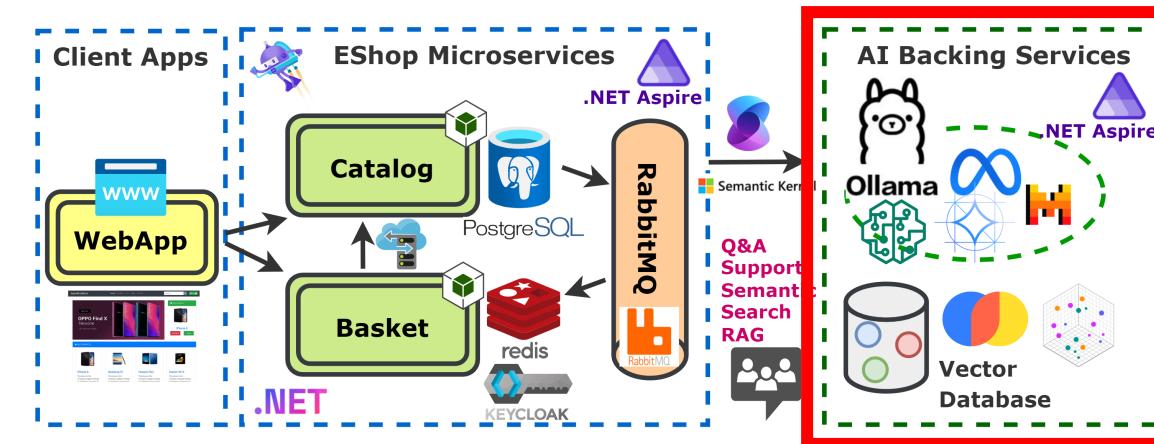
.NET GenAI with Microsoft.Extensions.AI for Chat AI

Ollama for Local LLM Inference, Integrate models like Llama3 or Phi3

Develop Customer Support Q&A chatbot

Vector Database (in-memory), Store embeddings for semantic search

Blazor Frontend for AI: Build pages that let users interact with the Q&A or advanced semantic product searches



What are Large Language Models (LLMs) ?

LLMs are advanced AI models trained on large datasets - massive amounts of text data from books, websites, research papers

They can understand, generate, and process natural language that mimics natural human communication

Heart of AI chatbots, content creation tools, translation services



Large Language Models (LLMs)

LLMs are machine learning models specialized in natural language

They operate on a vast number of parameters—billions or even trillions

Trained to perform tasks such as text generation, summarization, Q&A, translation, and more

Building intelligent language applications



Why Are LLMs Considered Large ?

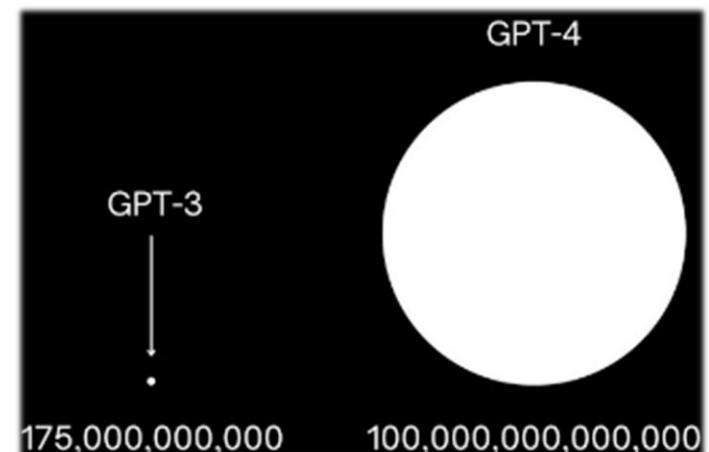
'Large' in LLM refers to both the size of the model: **the number of parameters and the training data**

LLMs contain **billions of parameters** (GPT-3 has 175 billion)

Training data: Trained on diverse text datasets (books, articles, web pages)

Larger models provide **better accuracy and context understanding**

More parameters = better performance but **higher computational costs**



Ollama – Run LLMs Locally

Ollama is a platform to **run large language models** (LLMs) or **small language models** (SLMs) **locally**

Offers **private, secure** AI solutions **without requiring the cloud**

Ideal for **developers** and **businesses** seeking **offline AI** capabilities w/ **privacy, low latency**, and **control** over AI models

Excellent choice for those looking to **integrate AI without depending** on constant **internet access**



Key Features of Ollama

Local Execution

Run LLMs and Small Language Models (SLMs) directly on your device

Pre-built Models

Includes optimized models for coding, chat, creative tasks, and more

Privacy-First

All data remains on your machine, protecting sensitive information

Customization

Allows model fine-tuning and adaptation for specific needs

Low Latency

Responses are quick without network dependency



Available Models on Ollama

Variety of Models: **llama, gemma, qwen, phi, mistral**

Code Generation Models-**codegemma,codellama**

AI assistants specialized in code generation

Creative Models-**llava**

Text-to-image, story generation, and poetry models

Domain-Specific Models-**medllama**

Finance, healthcare, and other industry-specific LLMs



Models

Filter by name...

llama3.2

Meta's Llama 3.2 goes small with 1B and 3B models.

Tools 1B 3B

1.5M Pulls 63 Tags Updated 3 weeks ago

llama3.1

Llama 3.1 is a new state-of-the-art model from Meta available in 8B, 70B and 405B parameter sizes.

Tools 8B 70B 405B

7M Pulls 94 Tags Updated 5 weeks ago

gemma2

Google Gemma 2 is a high-performing and efficient model available in three sizes: 2B, 9B, and 27B.

2B 9B 27B

1.6M Pulls 94 Tags Updated 5 weeks ago

Benefits and Use Cases for Ollama

Benefits

Offline Capability, Data Privacy, Cost Savings, Low Latency

Use Cases

Software Development, Customer Support, Education and Creative work



Models

Filter by name...

llama3.2

Meta's Llama 3.2 goes small with 1B and 3B models.

Tools 1B 3B

↓ 1.5M Pulls ⚡ 63 Tags ⏲ Updated 3 weeks ago

llama3.1

Llama 3.1 is a new state-of-the-art model from Meta available in 8B, 70B and 405B parameter sizes.

Tools 8B 70B 405B

↓ 7M Pulls ⚡ 94 Tags ⏲ Updated 5 weeks ago

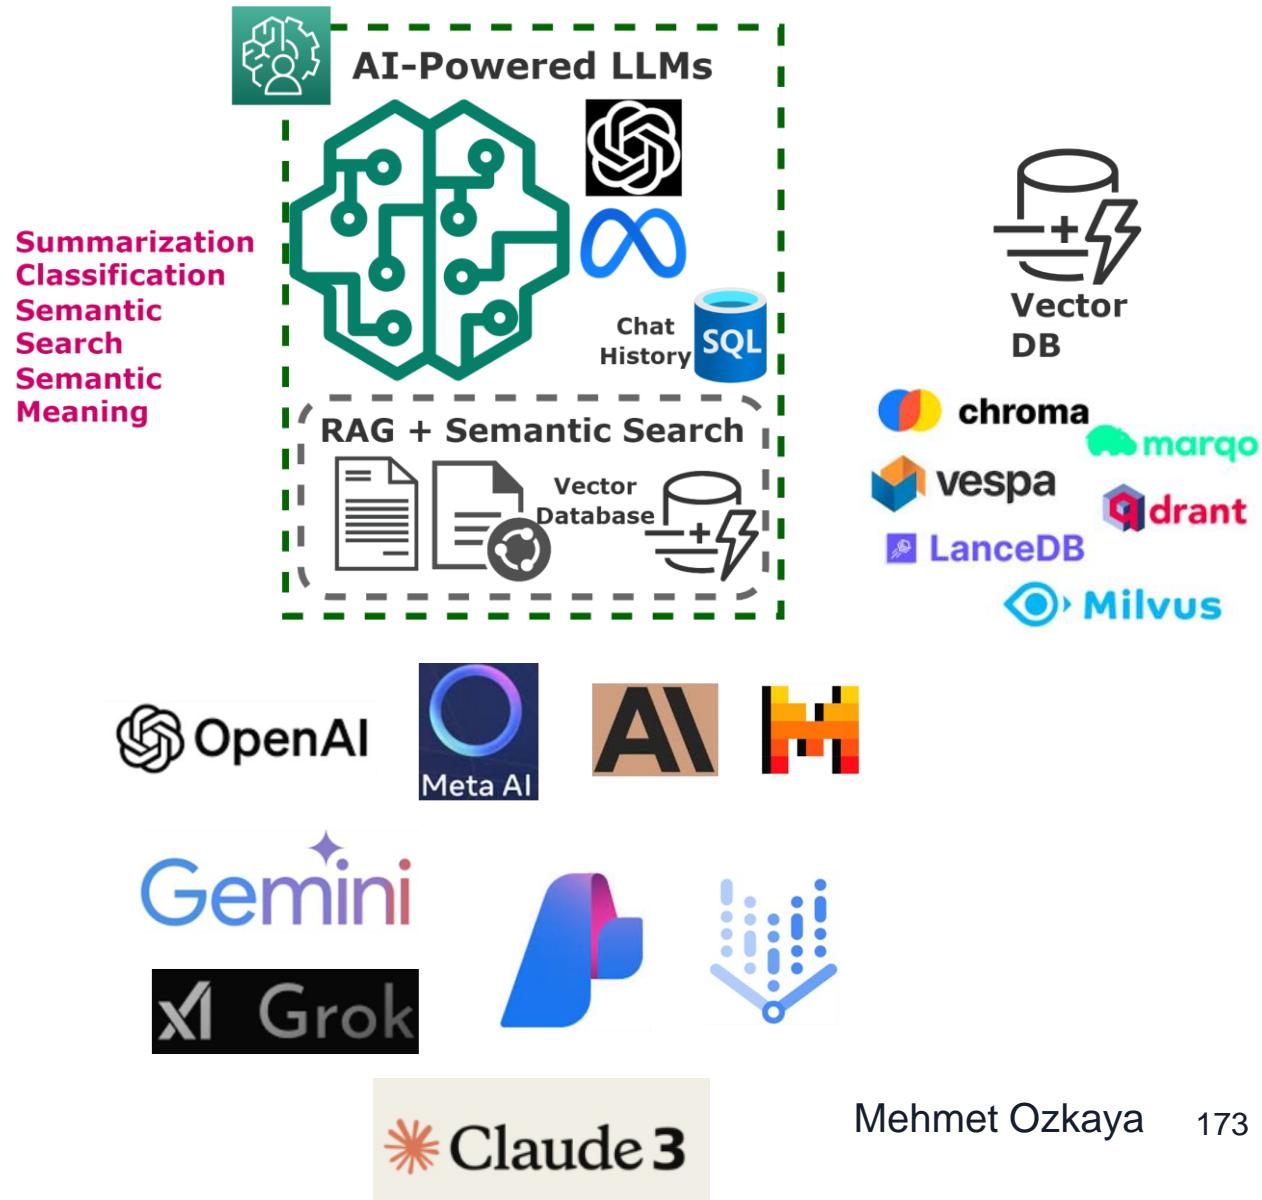
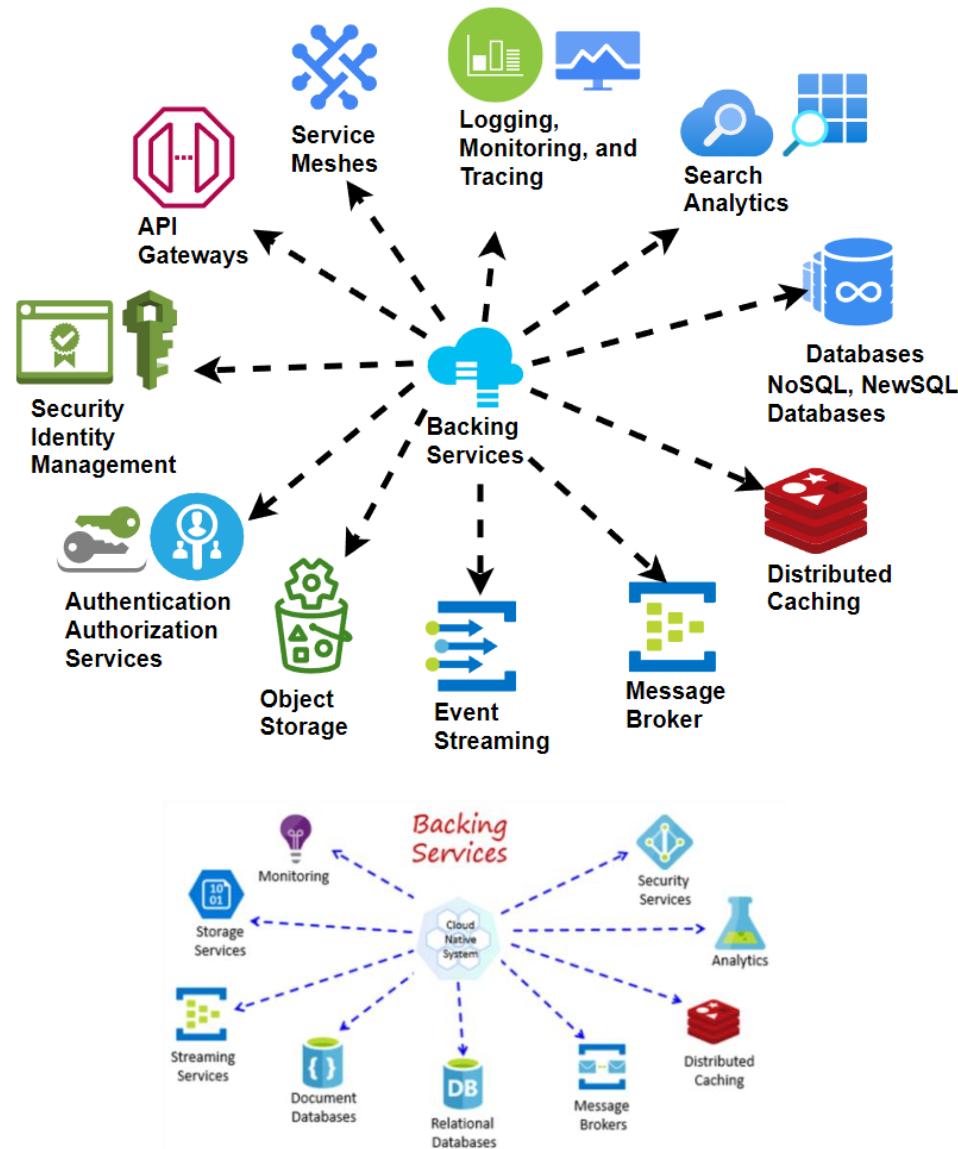
gemma2

Google Gemma 2 is a high-performing and efficient model available in three sizes: 2B, 9B, and 27B.

2B 9B 27B

↓ 1.6M Pulls ⚡ 94 Tags ⏲ Updated 5 weeks ago

The New Era of Cloud-Native Backing Services



Why Use LLMs and VectorDBs as Backing Services ?

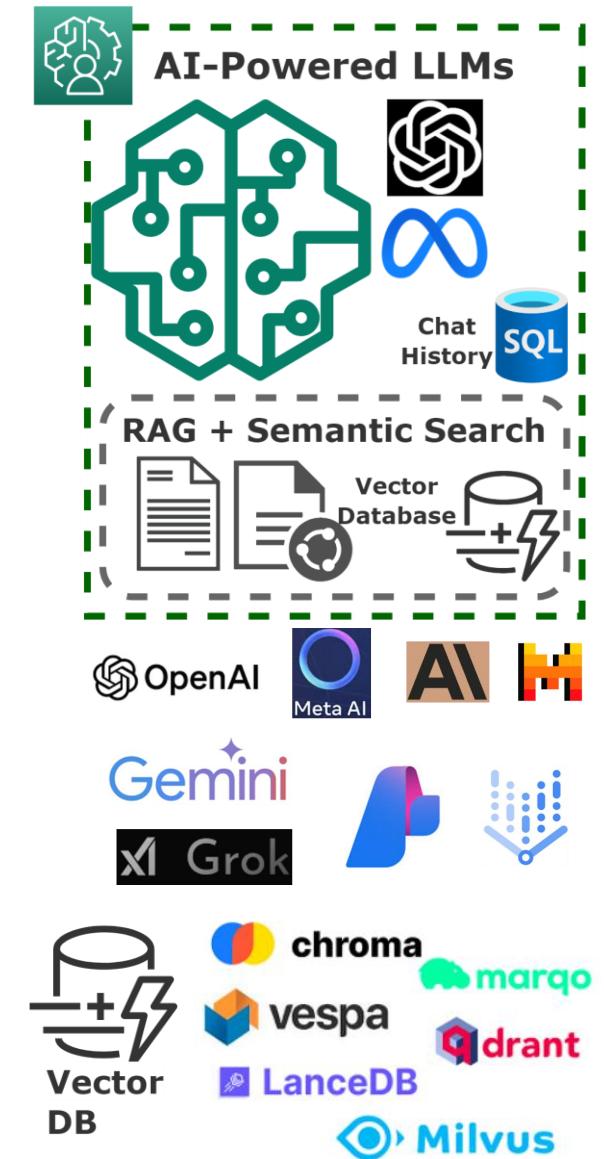
LLMs provide features like **summarization, classification, sentiment analysis** and chat-based Q&A

VectorDBs enable **similarity-based search and knowledge retrieval** to support **complex queries**

Leverage **pre-trained models** and **semantic search engines** to **integrate** powerful AI capabilities with minimal effort

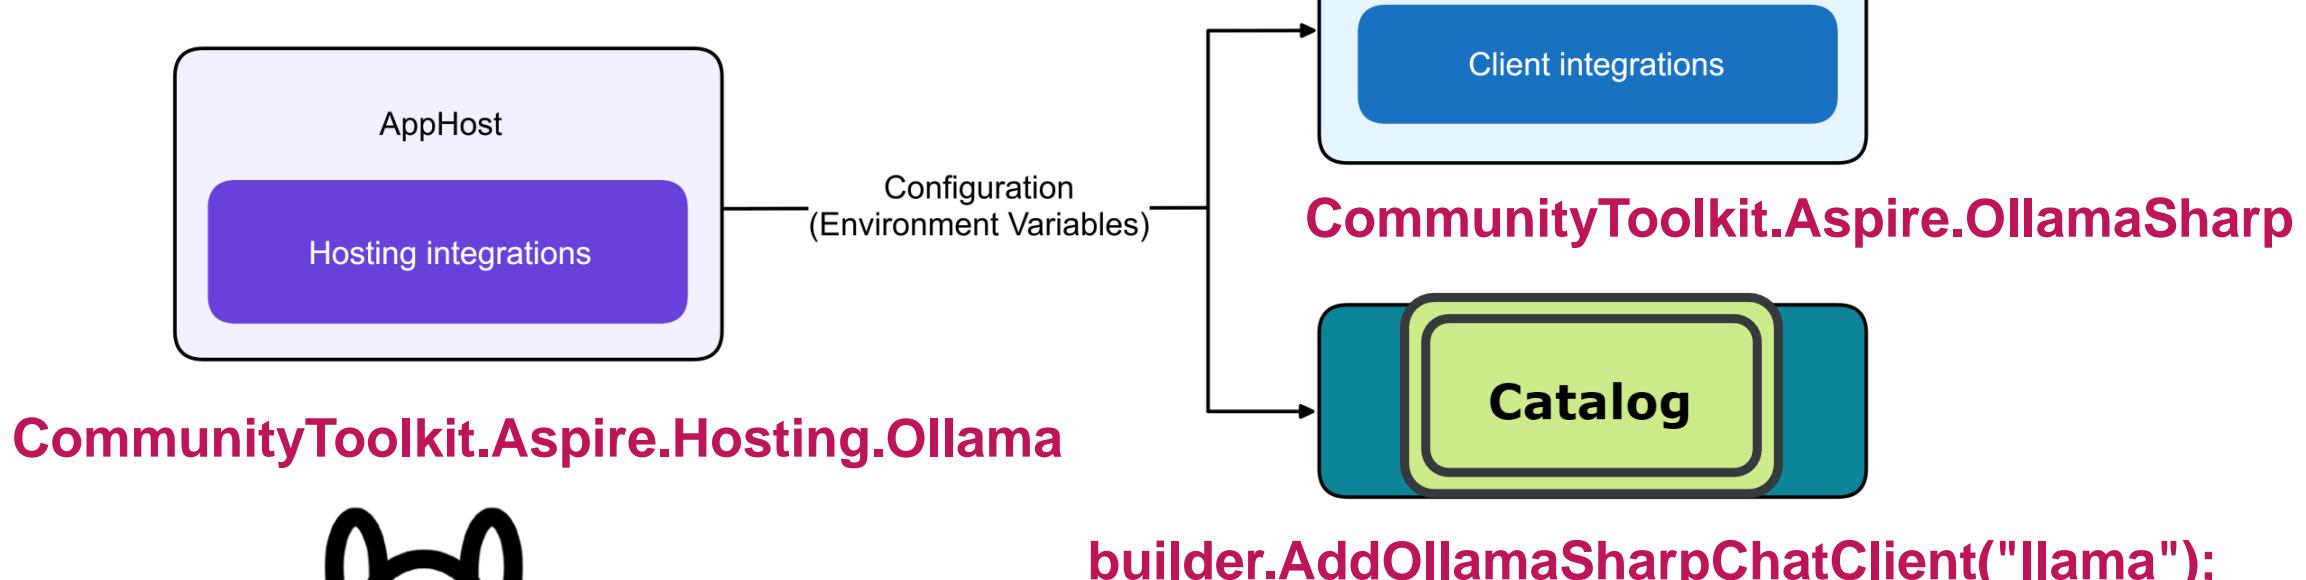
Automate **customer support**, real-time **recommendations, semantic search** for enterprise knowledge management

LLMs and VectorDBs as the **brain and memory** of your architecture



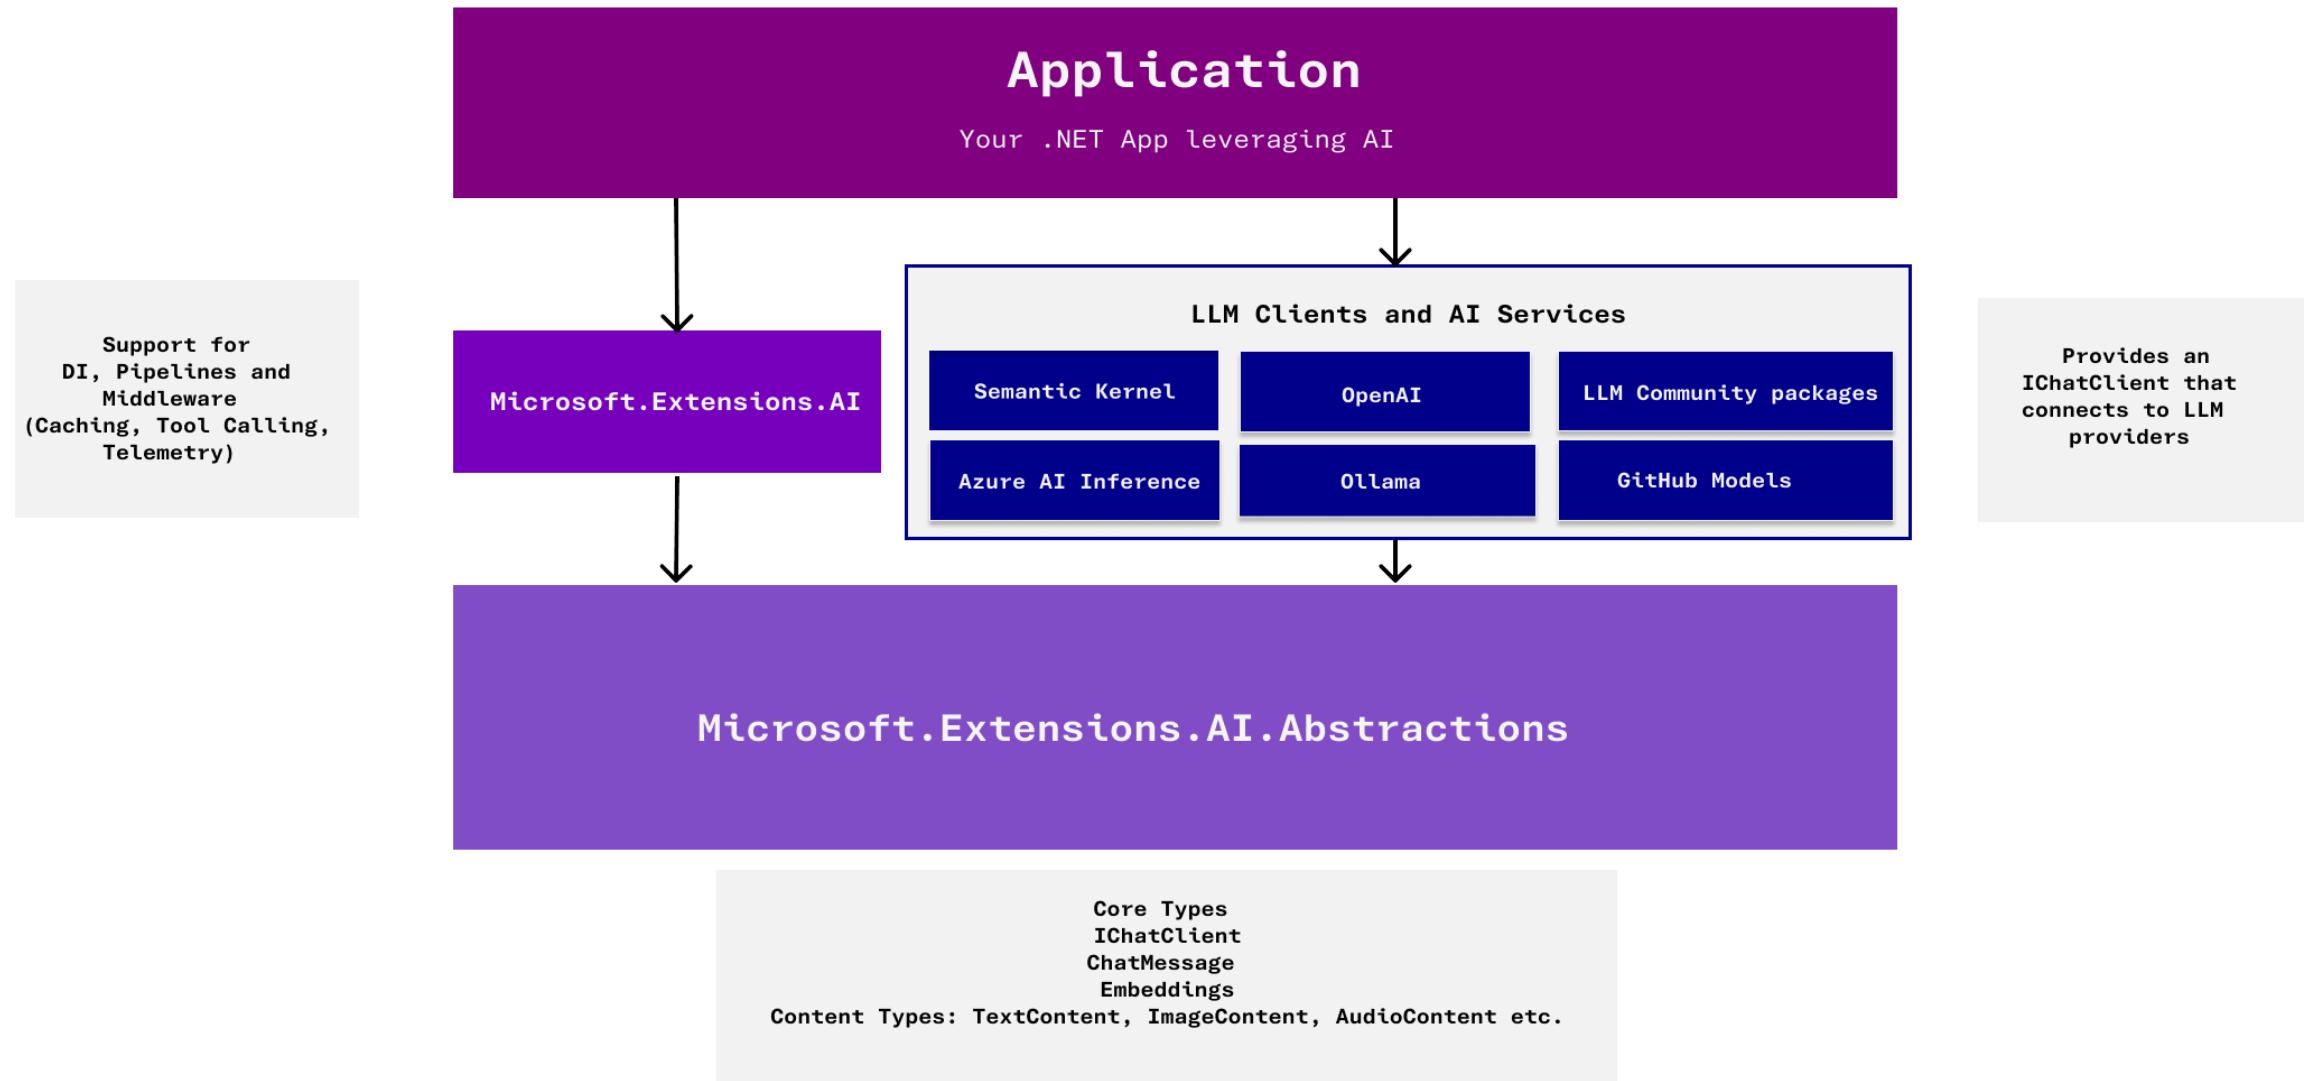
Ollama and Llama Model Integrations

Integrations



```
builder.AddOllamaSharpChatClient("llama");
```

Microsoft.Extensions.AI Structure

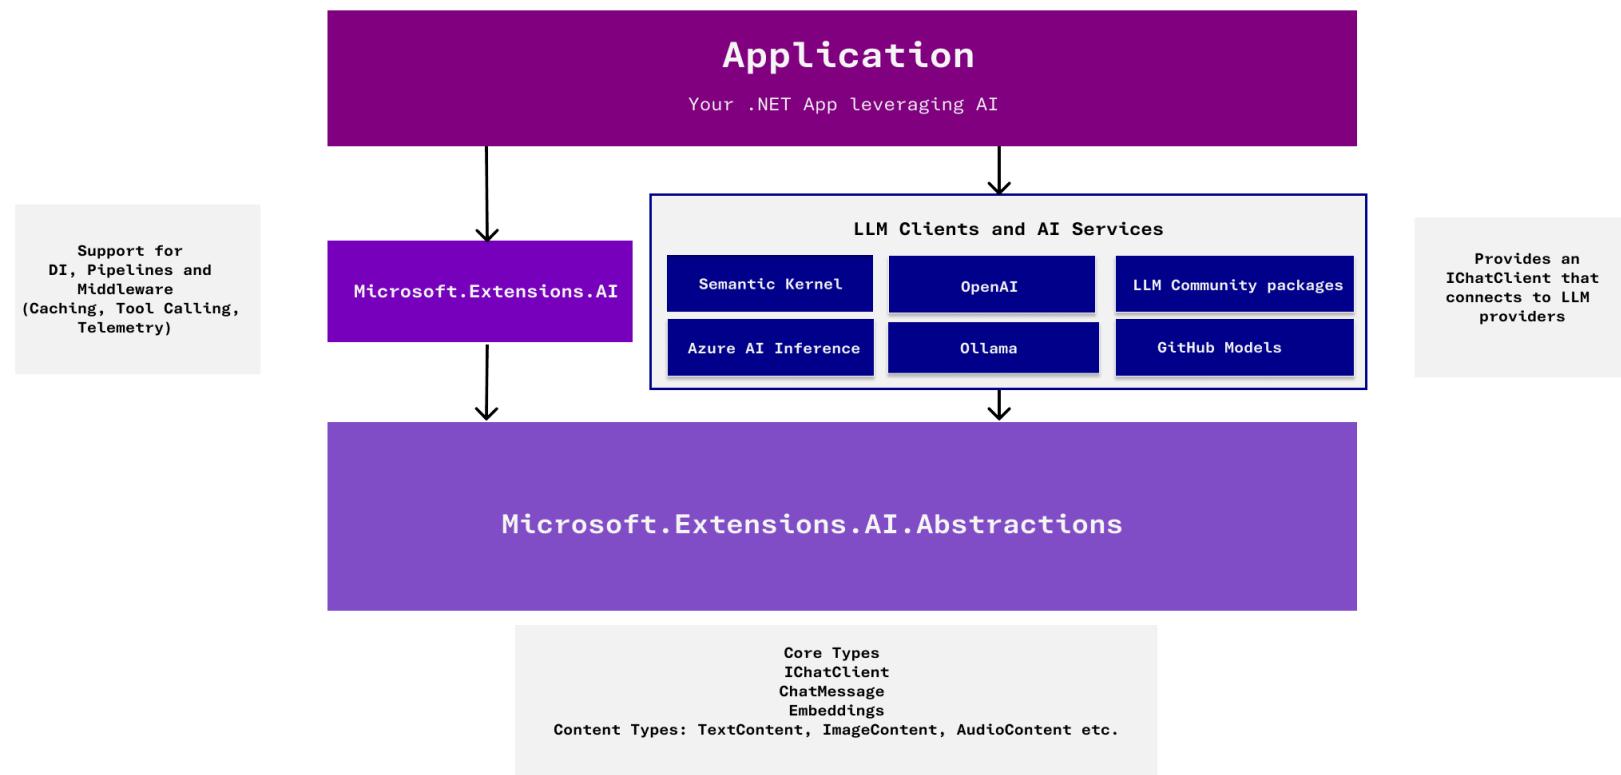


Work with Abstractions for Common AI Services

Chat: Conversational AI or ChatGPT-like interactions

Embeddings: Vector search & semantic understanding

Tool calling: Orchestrate AI requests with external services

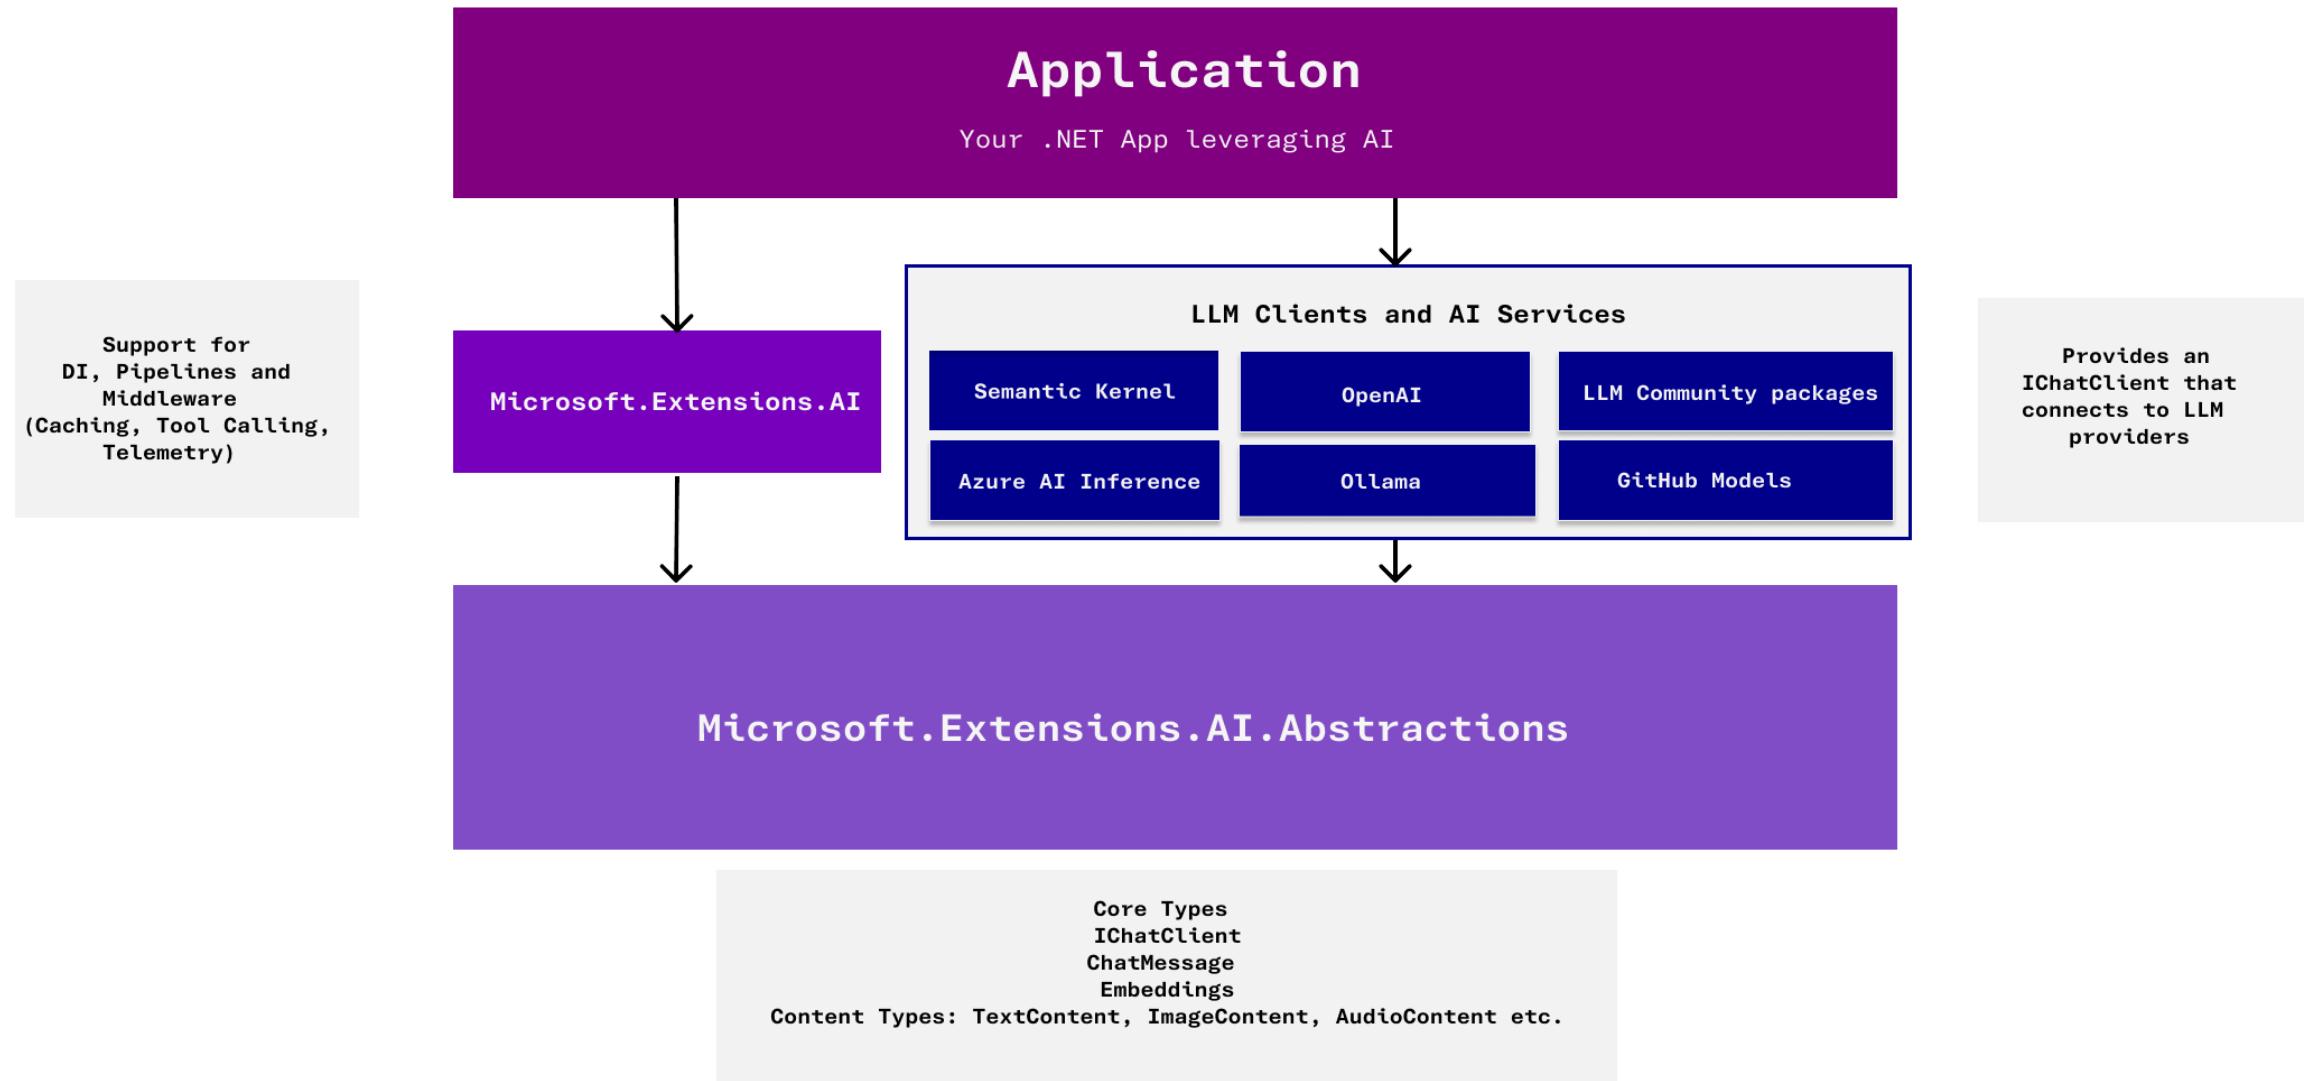


Example: IChatClient



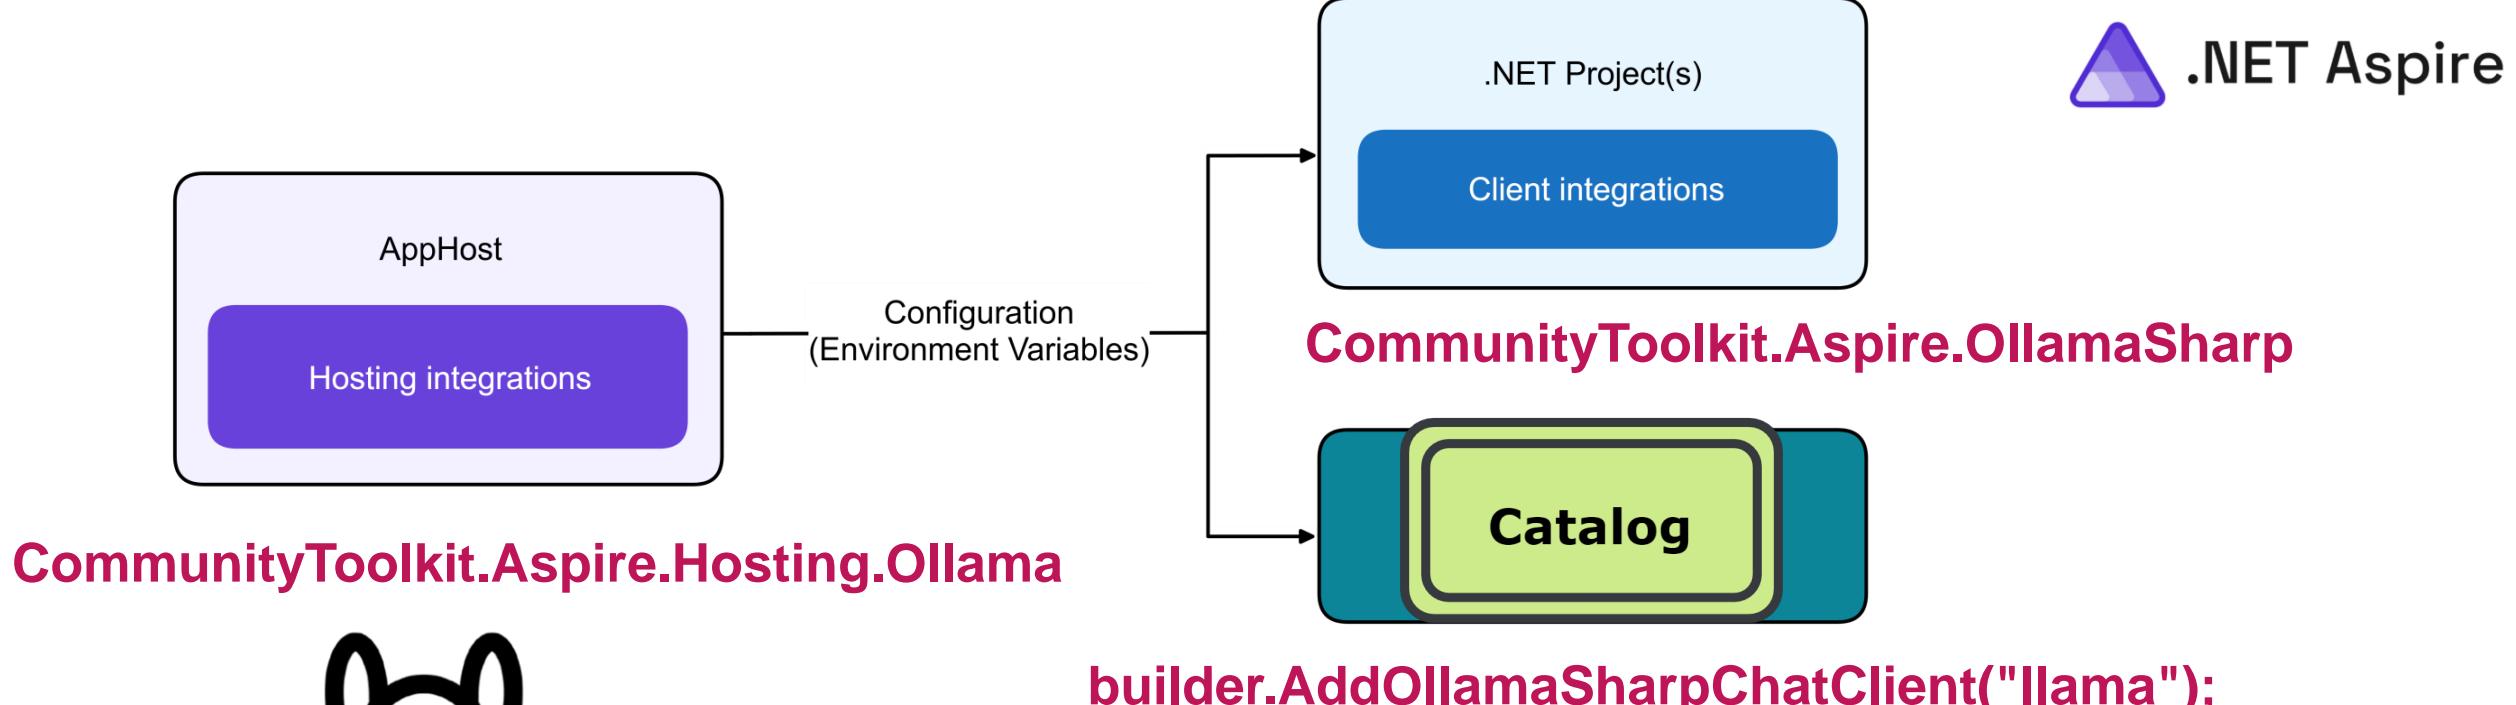
```
IChatClient client =  
    environment.IsDevelopment  
        ? new OllamaChatClient(...)  
        : new AzureAIIInferenceChatClient(...);  
  
var response = await client.GetResponseAsync(messages, options, cancellationToken);
```

Microsoft.Extensions.AI Structure

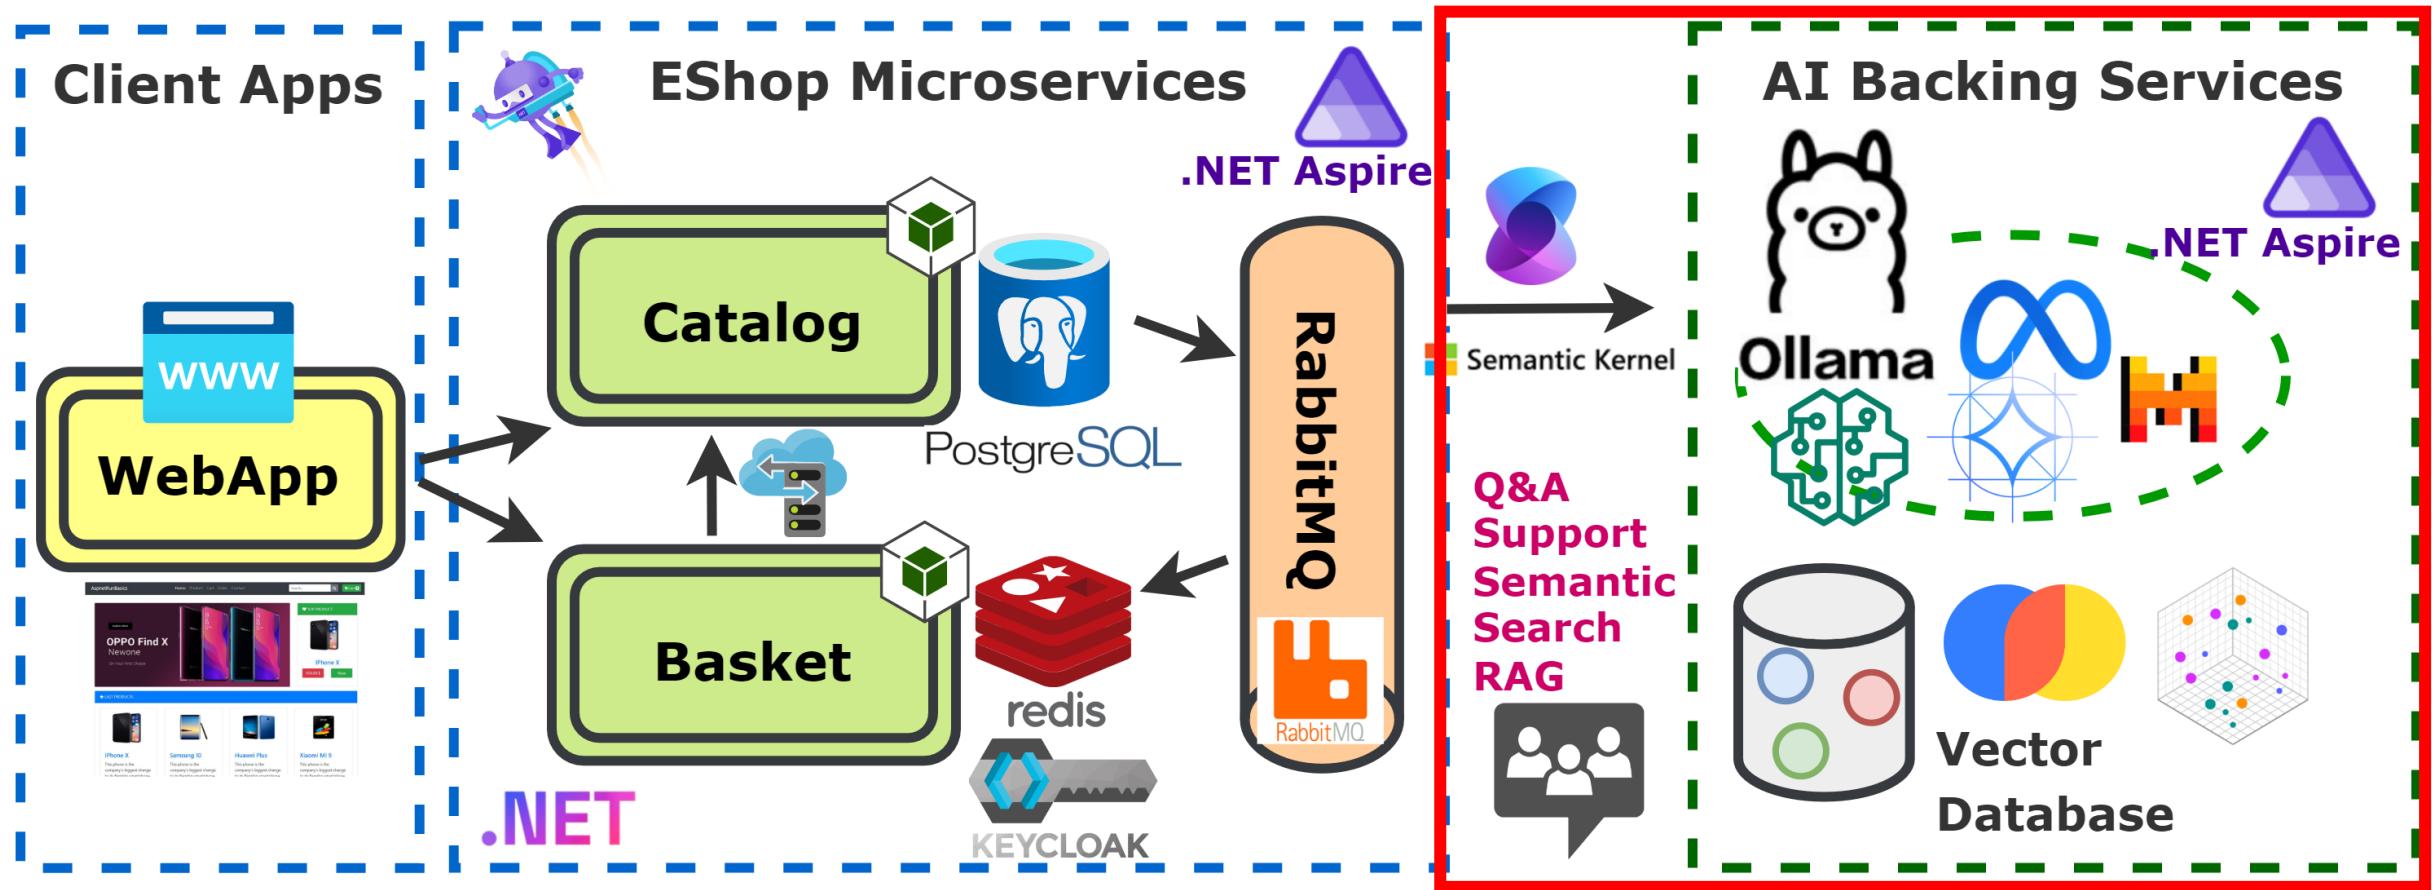


Ollama and Llama Model Integrations

Integrations



.NET GenAI with Microsoft.Extensions.AI for Chat AI and Semantic Search



Develop Blazor Web Application

Products

Here are some of our amazing outdoor products that you can purchase.

Image	Name	Description
	Hiking Poles	Ideal for camping and hiking trips
	Outdoor Rain Jacket	This product will keep you warm and dry in all weathers
	Survival Kit	A must-have for any outdoor adventurer
	Outdoor Backpack	This backpack is perfect for carrying all your outdoor essentials

Semantic Product Search with Vector Embeddings and Vector DB

What is a Vector Database?

What are Vectors and Vector Embeddings ?

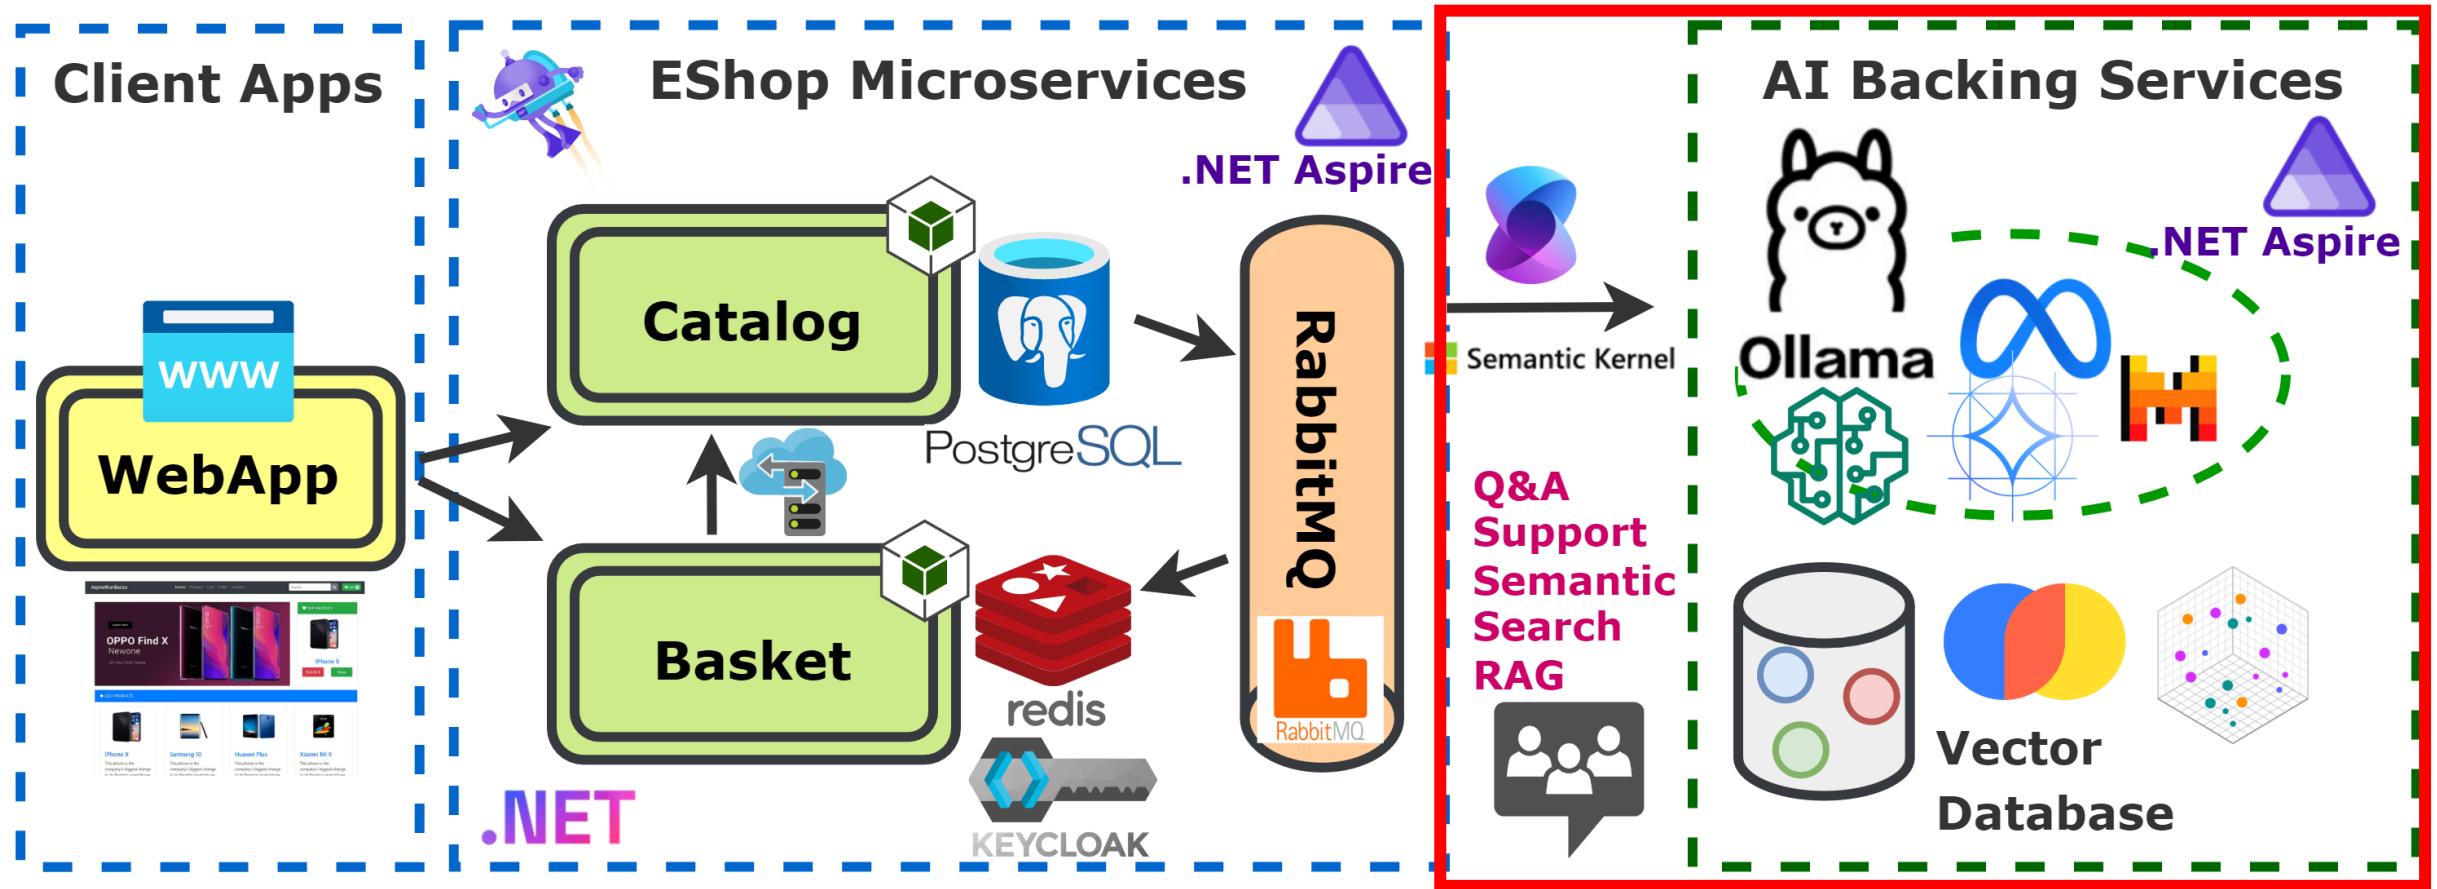
Hosting Integration for Ollama all-minilm Embeddings Model

Client Integration Packages with SemanticKernel and
Extensions.VectorData

Mehmet Ozkaya



Semantic Product Search with Vector Embeddings and Vector DB



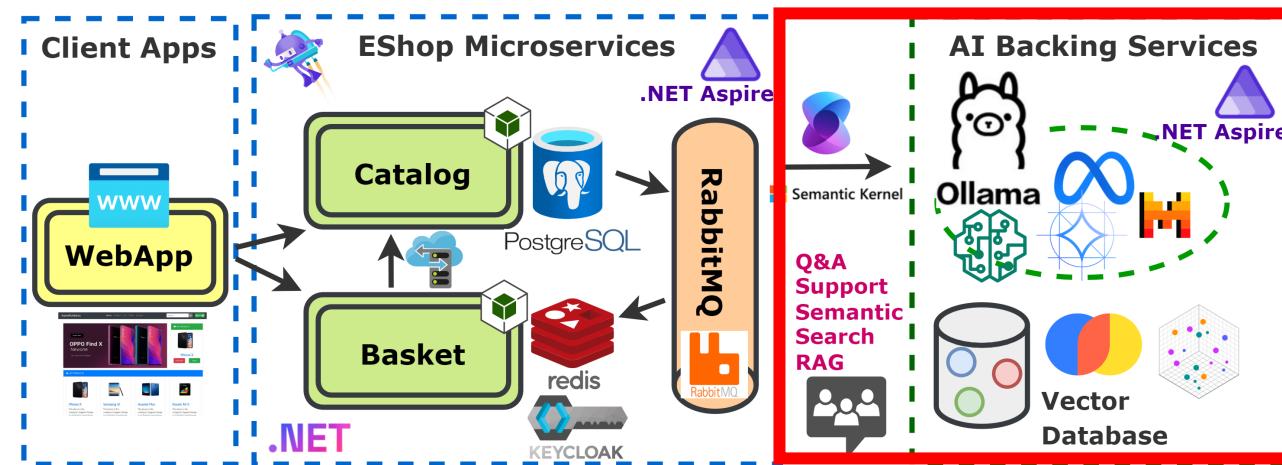
Semantic Product Search with Vector Embeddings and Vector DB

Hosting: Add a model (all-minilm) in Ollama for generating embeddings

Registration: Configure IEmbeddingGenerator to produce product embeddings

Vector Database: Store & Retrieve embeddings using a collection from Microsoft.Extensions.VectorData.Abstractions

Endpoints: We'll provide an API or method to transform a user query into an embedding



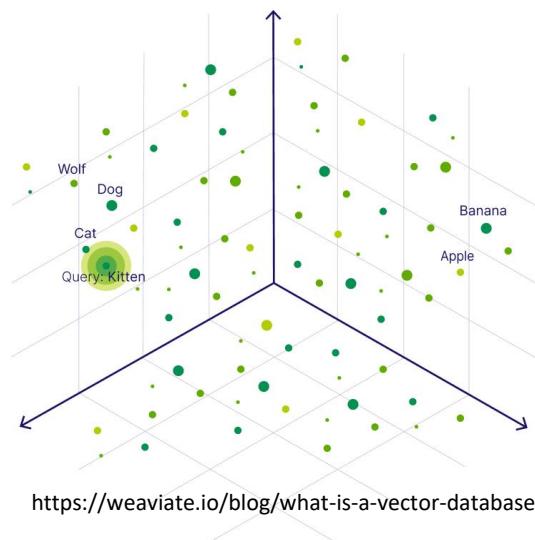
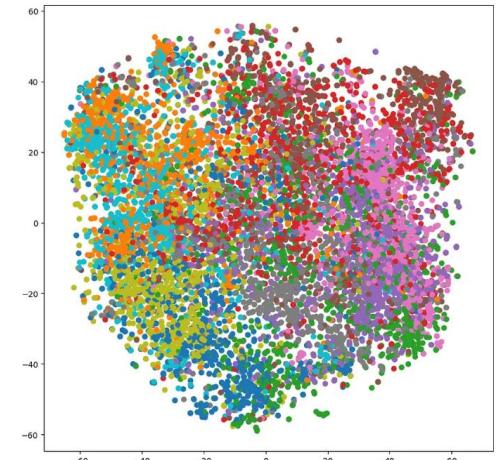
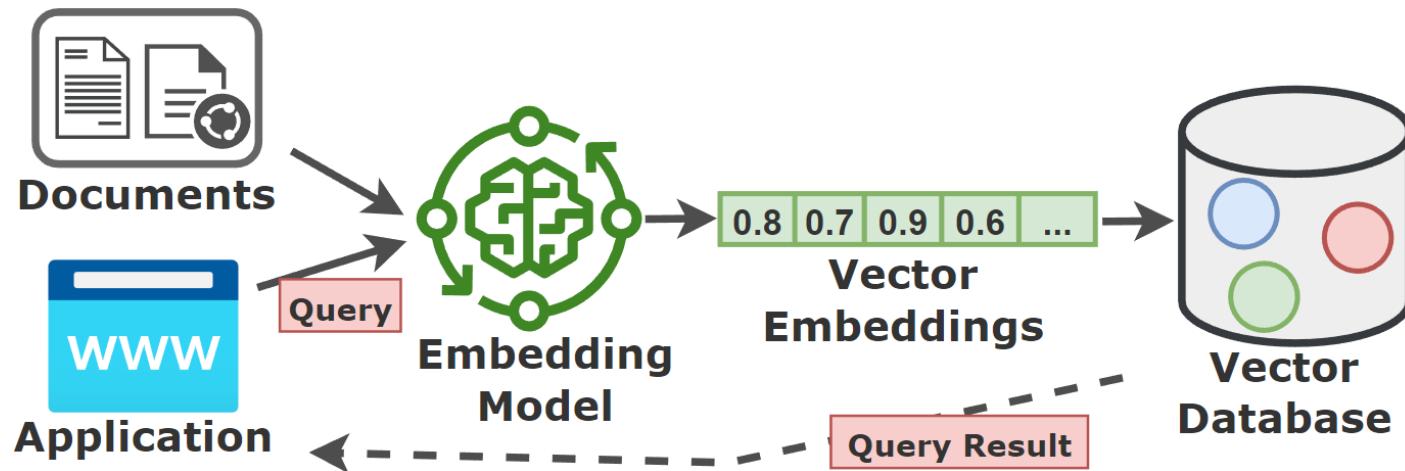
What is a Vector Database?

Vector Database is a **specialized database** designed to store, manage, and query **high-dimensional vectors**

Vectors are **numerical representations** that capture the **semantic meaning** of data

Indexes and **stores vector embeddings** for fast retrieval and **similarity search**

Enable **contextual searches based on the meaning** behind the data



<https://weaviate.io/blog/what-is-a-vector-database>

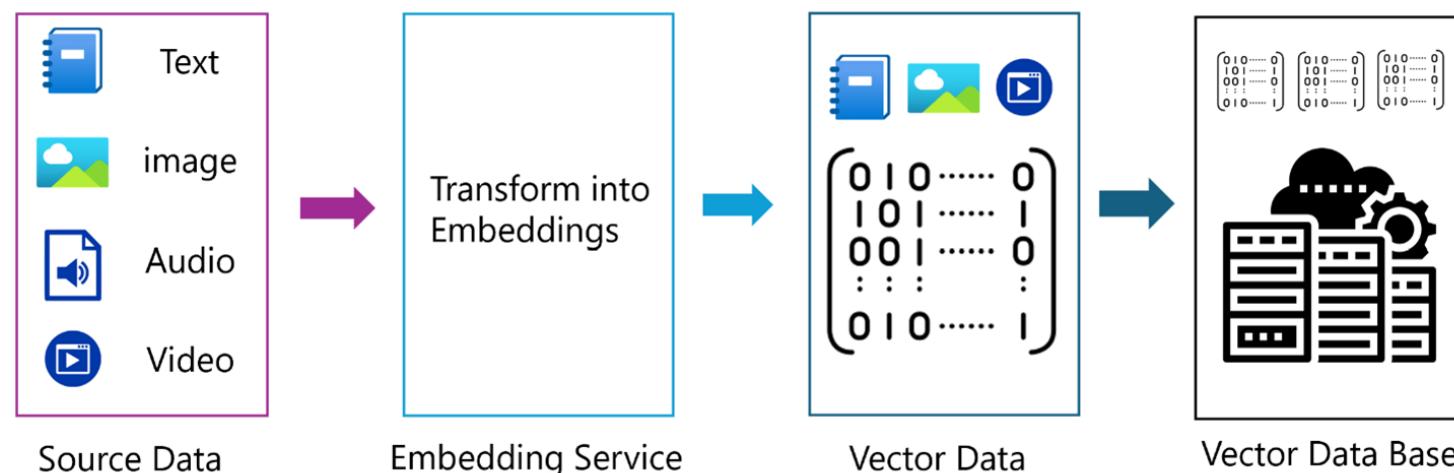
The Evolution from Traditional to Vector Databases

Traditional databases store data in **structured formats** like **rows** and **columns**

Fall short when we **need to perform semantic searches**—
searches **based on the meaning** of data rather exact words

Search for “laptop”: **Traditional Db** retrieve **exact word**, **Vector Db** understand that “notebook computer”

Vector Db store **unstructured data** like documents, images, audio, video, social media posts



CustomerID	FirstName	LastName	DateCreated	Cl
1	Homer	Simpson	13/06/2014 3:33:37 PM	
2	Peter	Griffin	13/06/2014 9:09:56 PM	
3	Stewie	Griffin	13/06/2014 9:16:07 PM	
4	Brian	Griffin	13/06/2014 9:16:36 PM	
5	Cosmo	Kramer	13/06/2014 9:16:41 PM	
6	Philip	Fry	13/06/2014 9:17:02 PM	
7	Amy	Wong	13/06/2014 9:22:05 PM	
8	Hubert J.	Farnsworth	13/06/2014 9:22:19 PM	
9	Marge	Simpson	13/06/2014 9:22:37 PM	
10	Bender	Rodriguez	13/06/2014 9:22:52 PM	
11	Turanga	Leela	13/06/2014 9:23:37 PM	
*	(New)		15/06/2014 9:00:01 PM	

Why Are Vectors Important?

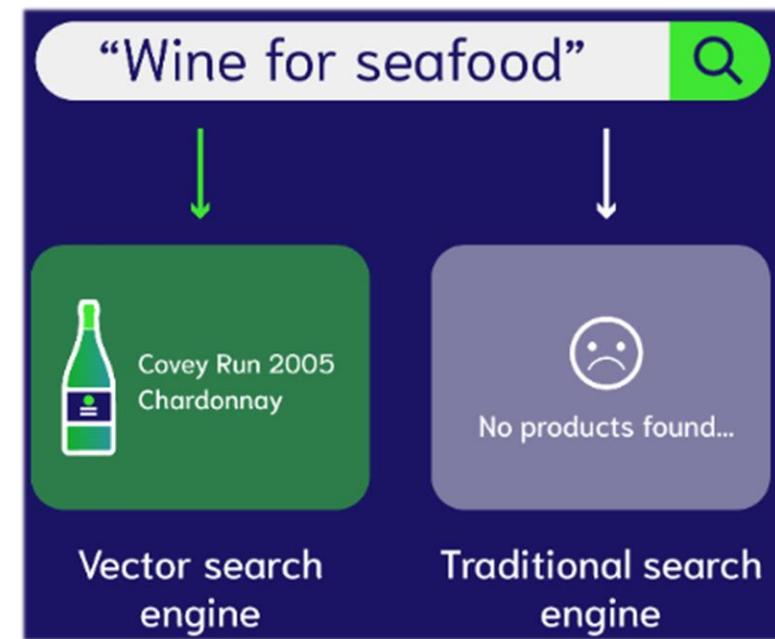
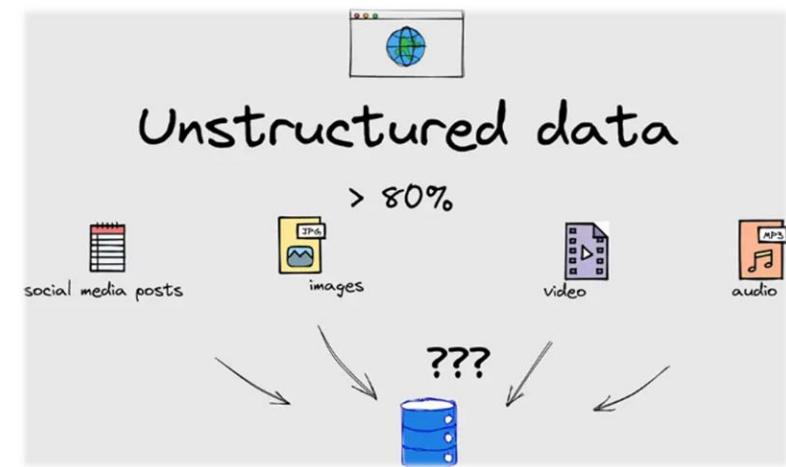
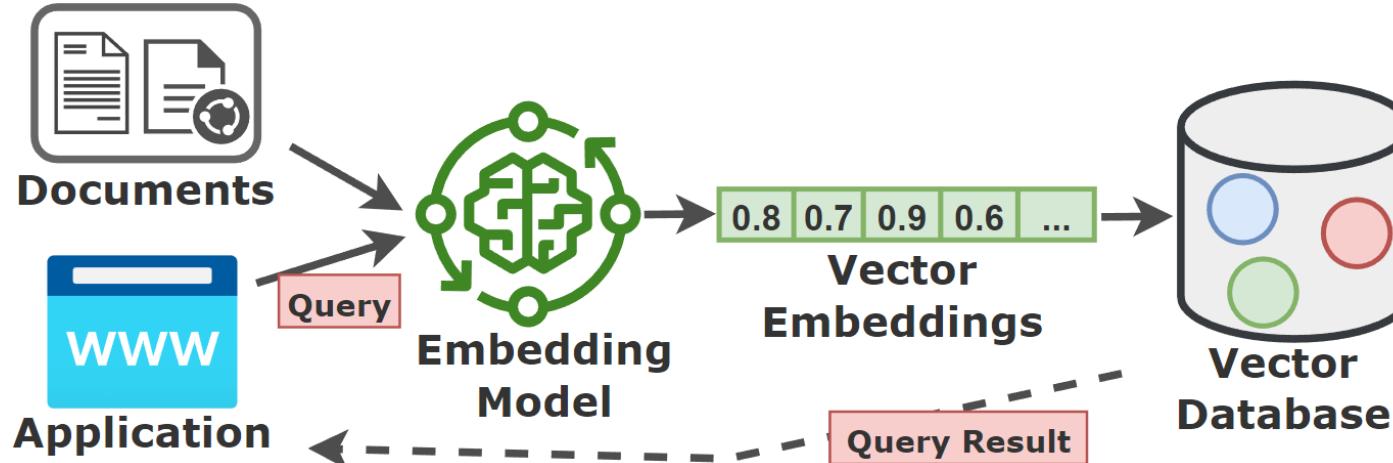
%80 of data on the internet is **Unstructured Data** (like documents, images, audio, video, social media posts)

Vectors **capture the meaning behind data**, not just its surface

Understand synonyms, paraphrases, and even nuanced relationships between data points

Search for “mobile phone” retrieve results for “smartphone”

Semantic search, personalized recommendations



<https://weaviate.io/blog/vector-embeddings-explained>

Core Features of Vector Databases

High-Dimensional Data Handling

Efficiently stores and queries data represented as vectors with hundreds or thousands of dimensions

Fast Similarity Search

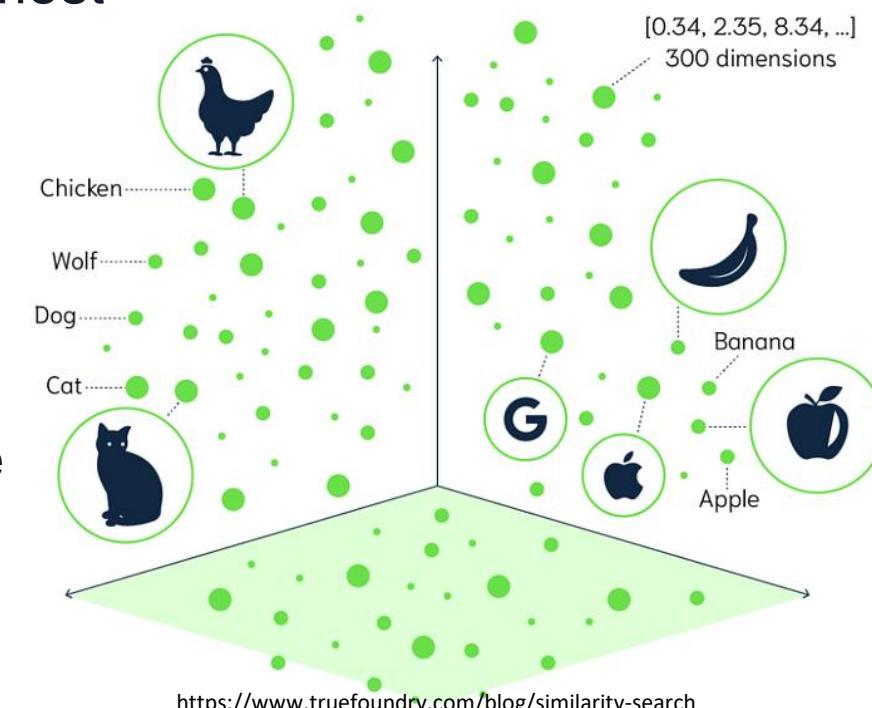
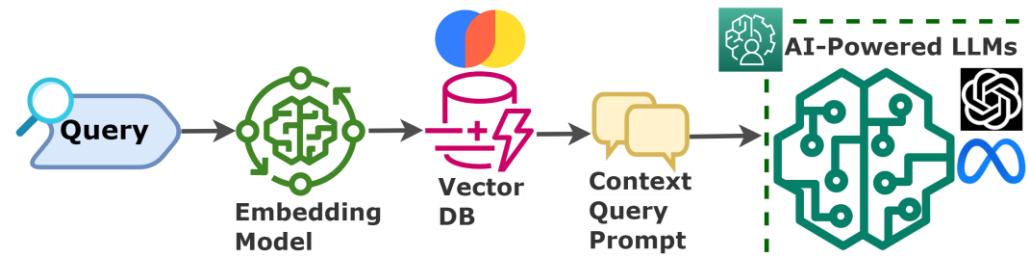
Approximate Nearest Neighbor (ANN) to quickly retrieve the most relevant vectors, even from large datasets

Scalability

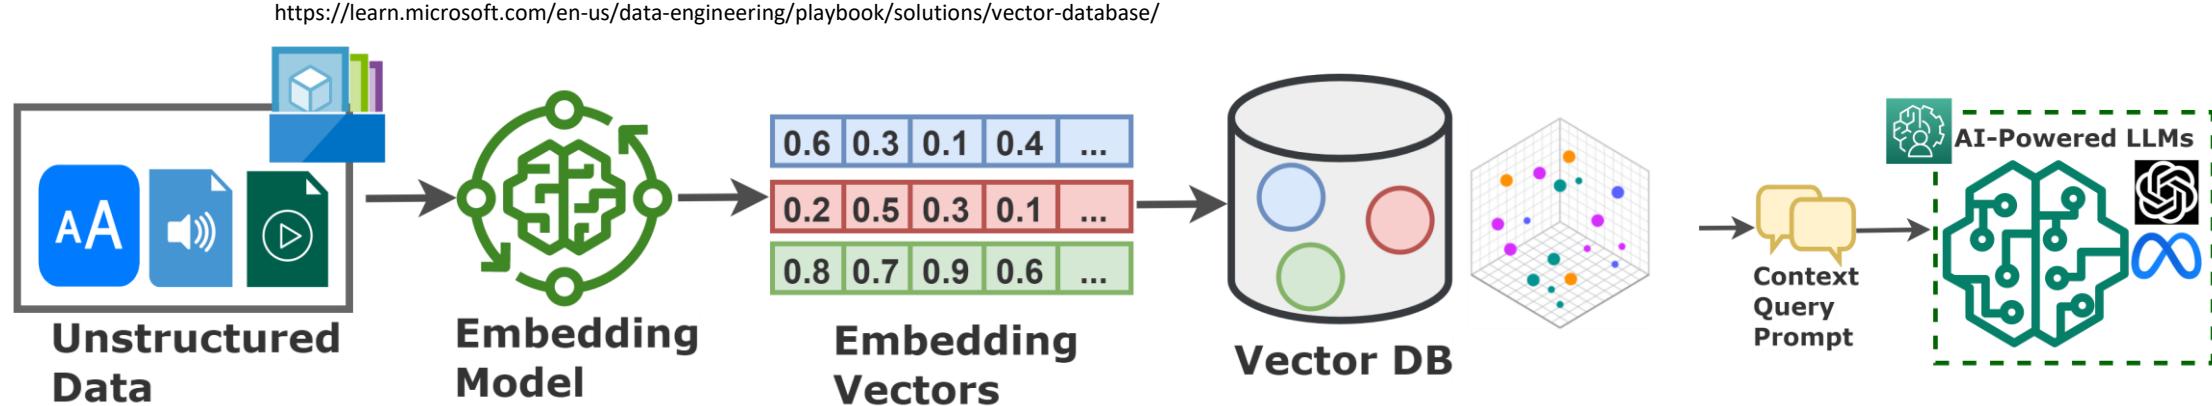
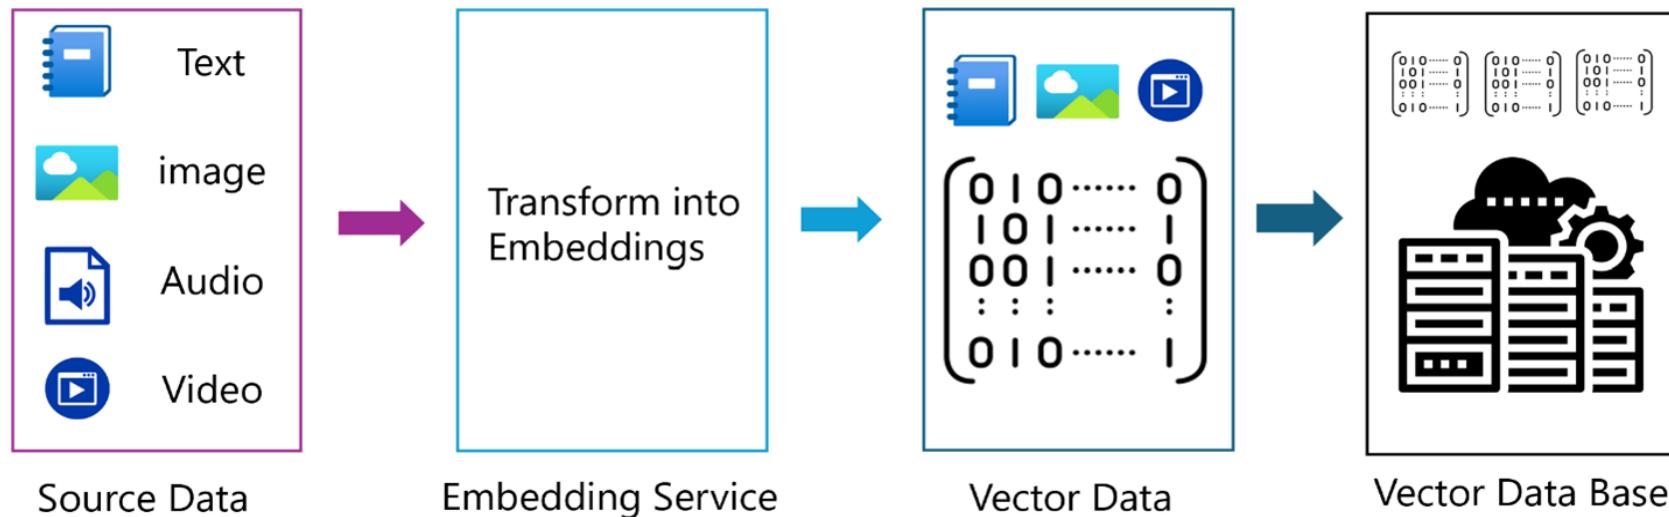
Supports millions or even billions of vectors, suitable for enterprise-scale applications

Integration with AI Models

Seamlessly works with embedding models from platforms like OpenAI, Hugging Face, or custom-trained models



What is a Vector Database? - Summary



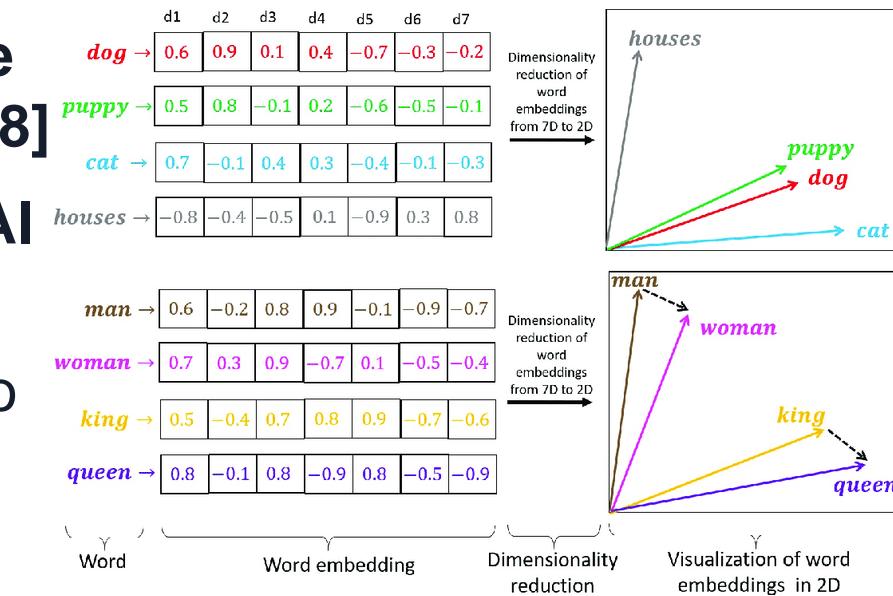
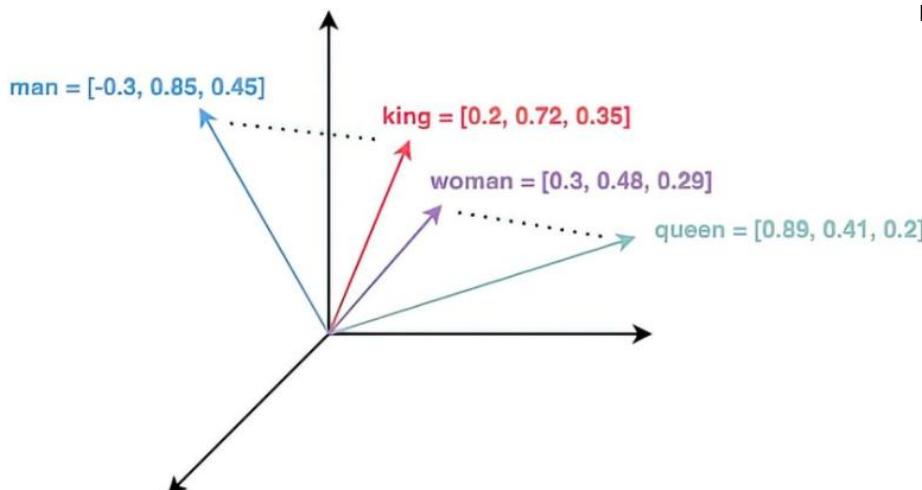
What is a Vector?

A **vector** is a **mathematical object** that has **both magnitude** and **direction**, represented as a list of numbers: [1.2, 3.4, -0.8]

Vectors are used to **represent complex data** in a way that **AI models** can **process, numerical summaries** of information

Vectors exist in **high-dimensional space**, which is **harder to visualize** but **incredibly powerful for computations**

Instead of saying “**it’s round, orange, and sweet**”, encode these features **numerically** in a **vector**: **[roundness: 0.9, color: 0.8, sweetness: 0.7]**



Rozado, David (2020). Word embeddings map words in a corpus of text to vector space.. PLOS ONE. Figure.
<https://doi.org/10.1371/journal.pone.0231189.g008>

What are Vector Embeddings?

Dense numerical representations of data, capture the semantic meaning of text, images, audio or other data types

Step 1: Input Data

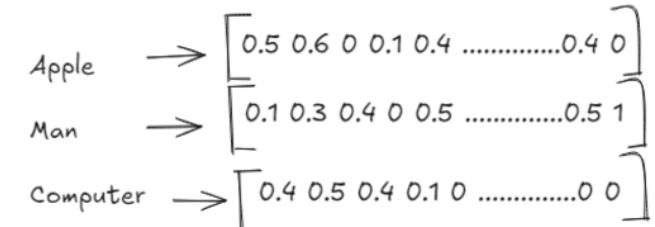
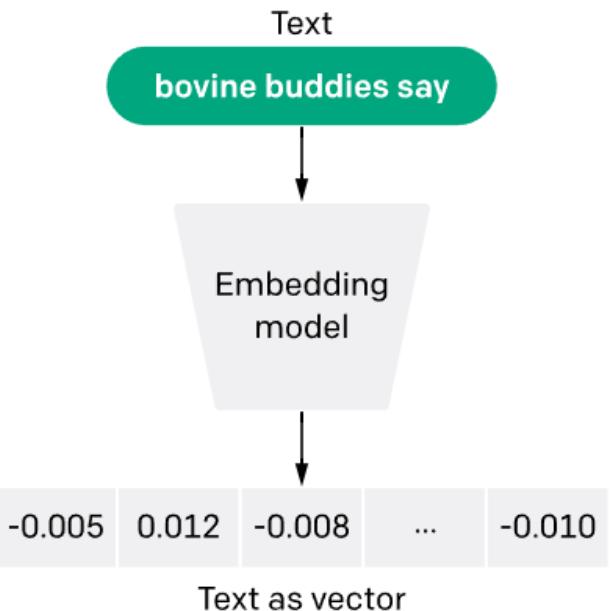
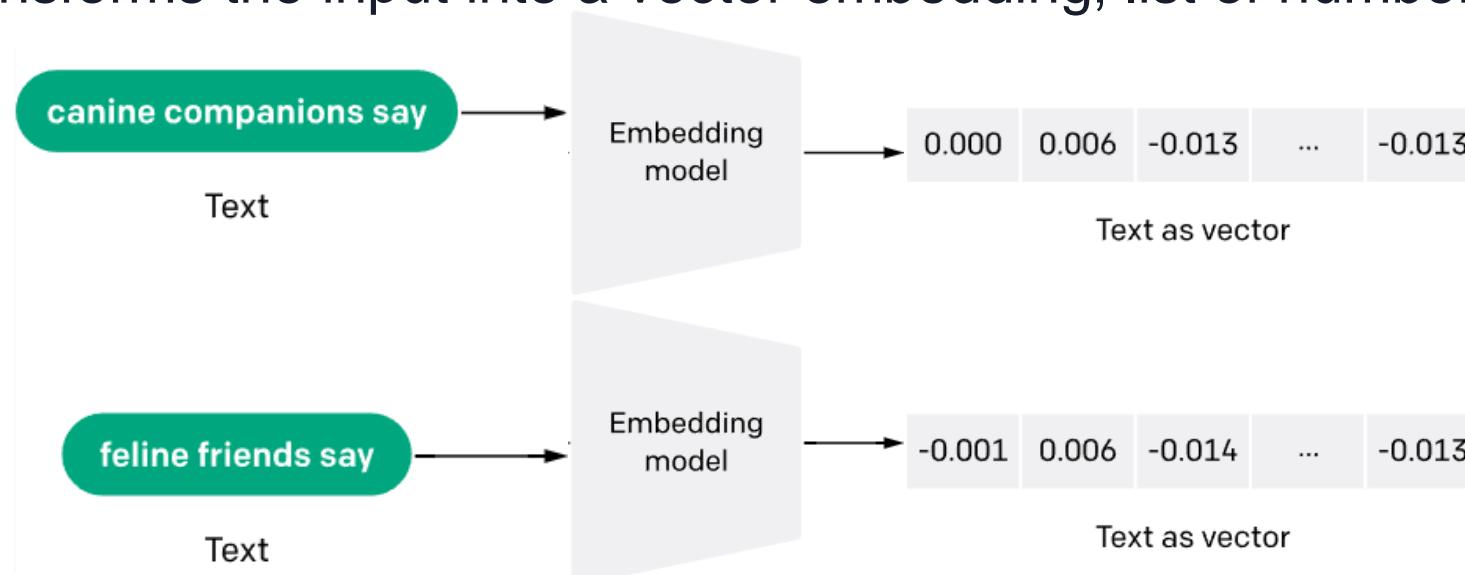
Start with raw data, like a sentence, an image, or a sound file

Step 2: Use an AI Embedding Model

Pass the data through an AI model, like a Transformer model

Step 3: Output Embedding

Transforms the input into a vector embedding, list of numbers



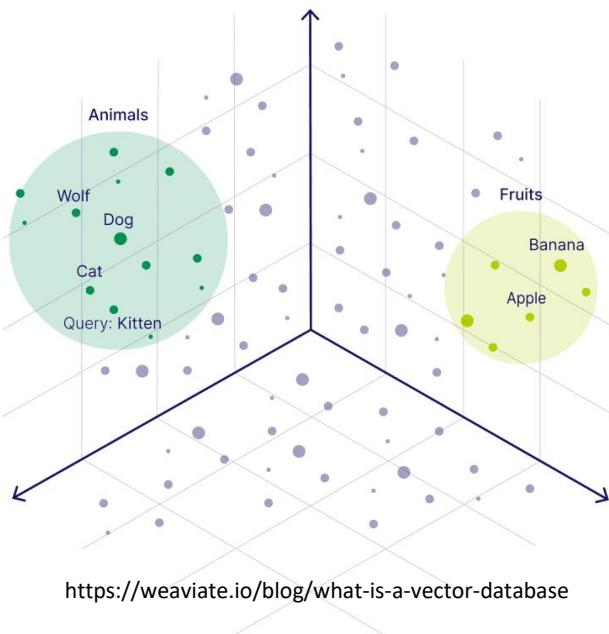
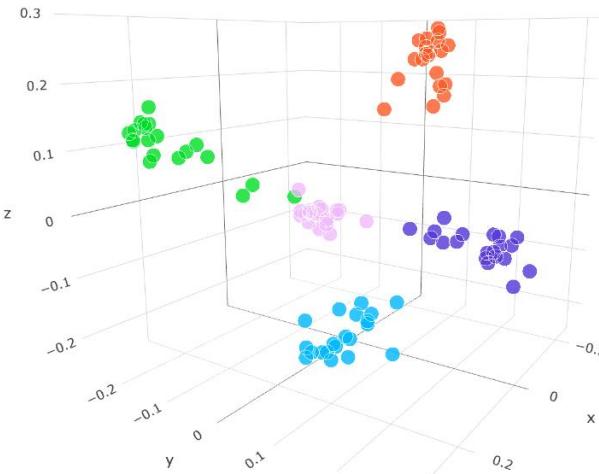
Why Are Embeddings Important?

Compare and analyze data based on **semantic meaning** rather than surface-level features

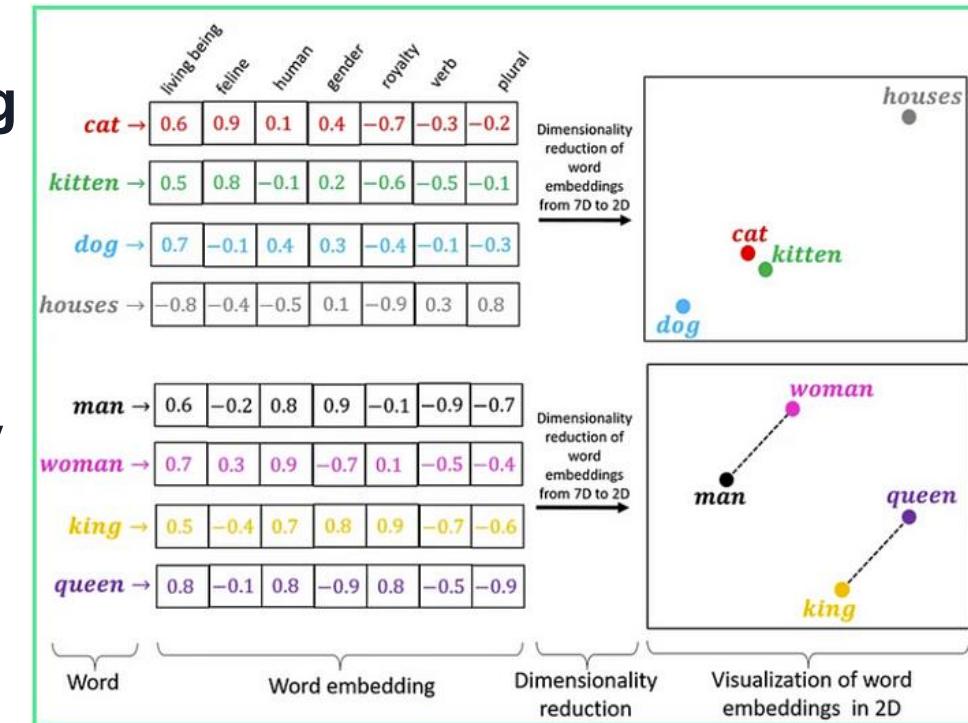
If two sentences have **similar meanings**, their embeddings will be close in high-dimensional space

Enable systems to **retrieve relevant information quickly**

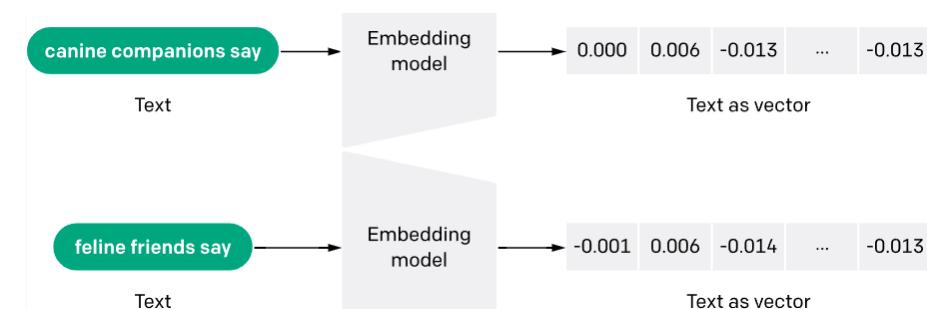
Same embedding techniques can be applied to text, images, and audio



<https://weaviate.io/blog/what-is-a-vector-database>

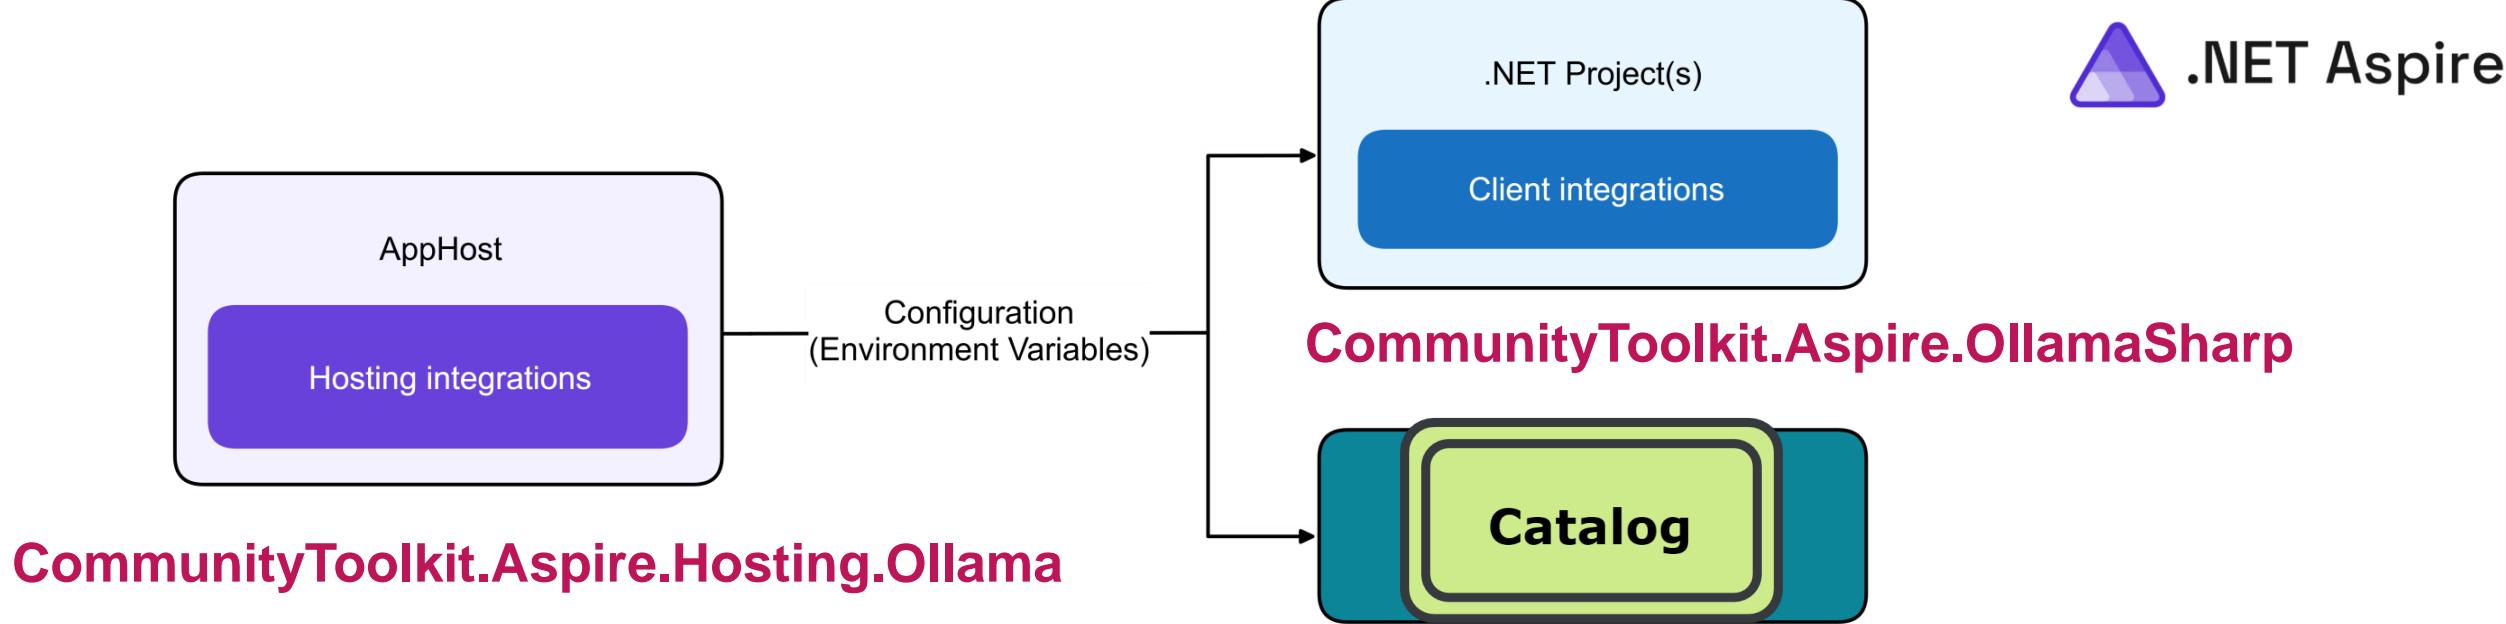


Rozado, David (2020). Word embeddings map words in a corpus of text to vector space.. PLOS ONE. Figure.
<https://doi.org/10.1371/journal.pone.0231189.g008>



Ollama and all-minilm Model Integrations

Integrations

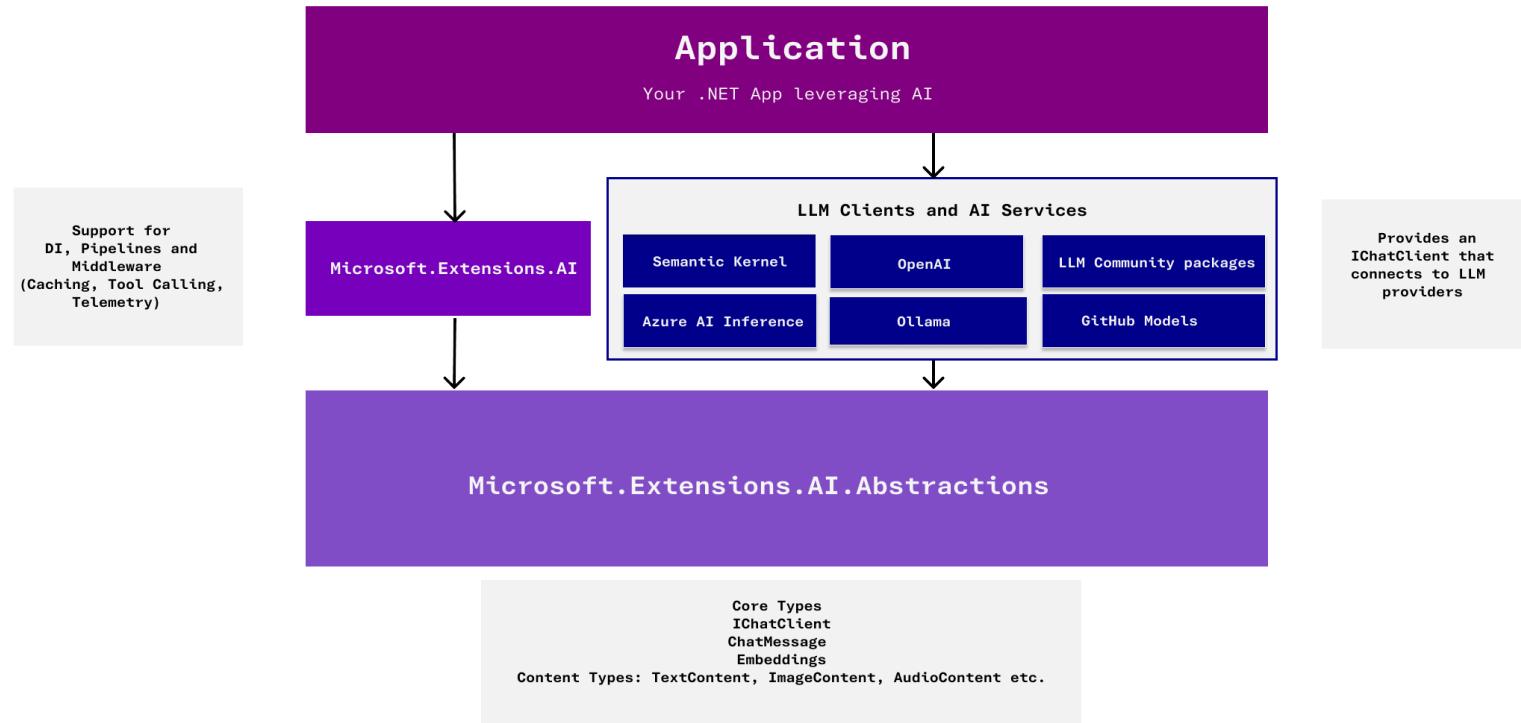


```
builder.AddOllamaSharpEmbeddingGenerator(..);  
builder.Services.AddInMemoryVectorStoreRecord  
Collection(..)
```

Microsoft.Extensions.VectorData

Abstraction layer manipulate vector data without coding directly against a specific database

Like Entity Framework for embeddings—create, read, update, delete on vector-based records in a consistent manner



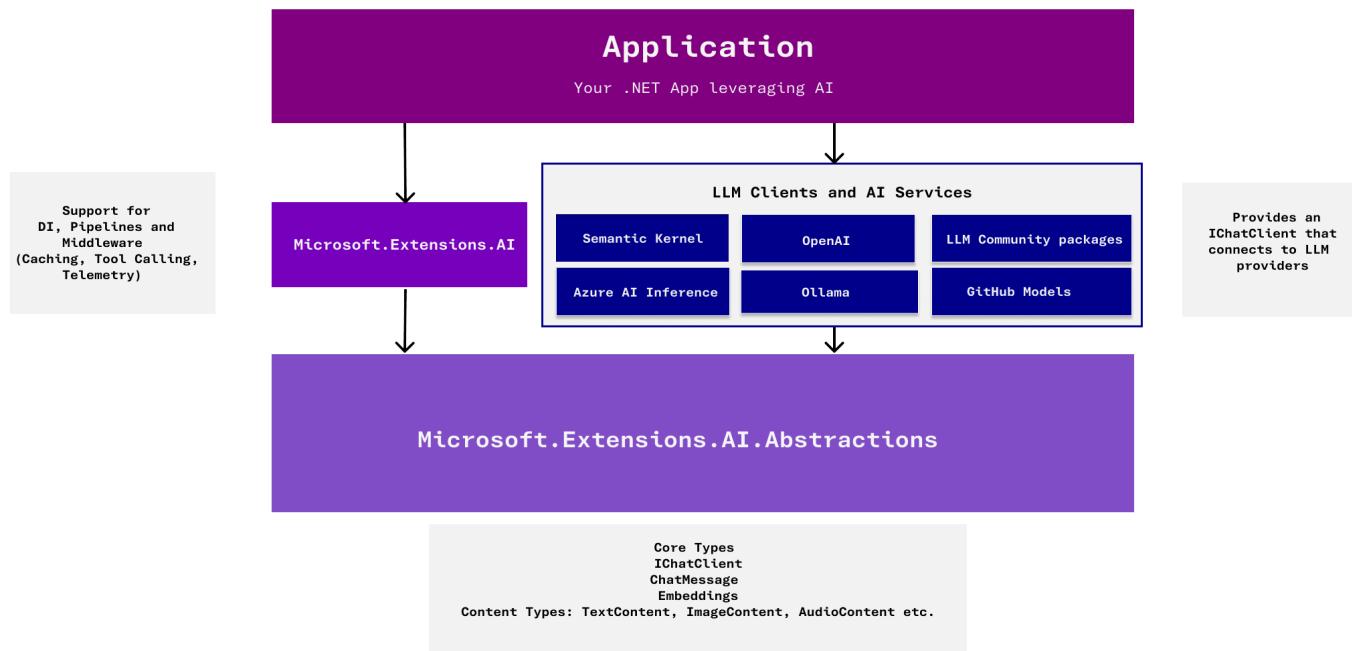
Getting Started with SemanticKernel Connectors

Use Semantic Kernel vector store connectors for `Extensions.VectorData`

Use the in-memory vector store implementation

Ollama reference implementation in `Microsoft.Extensions.AI`

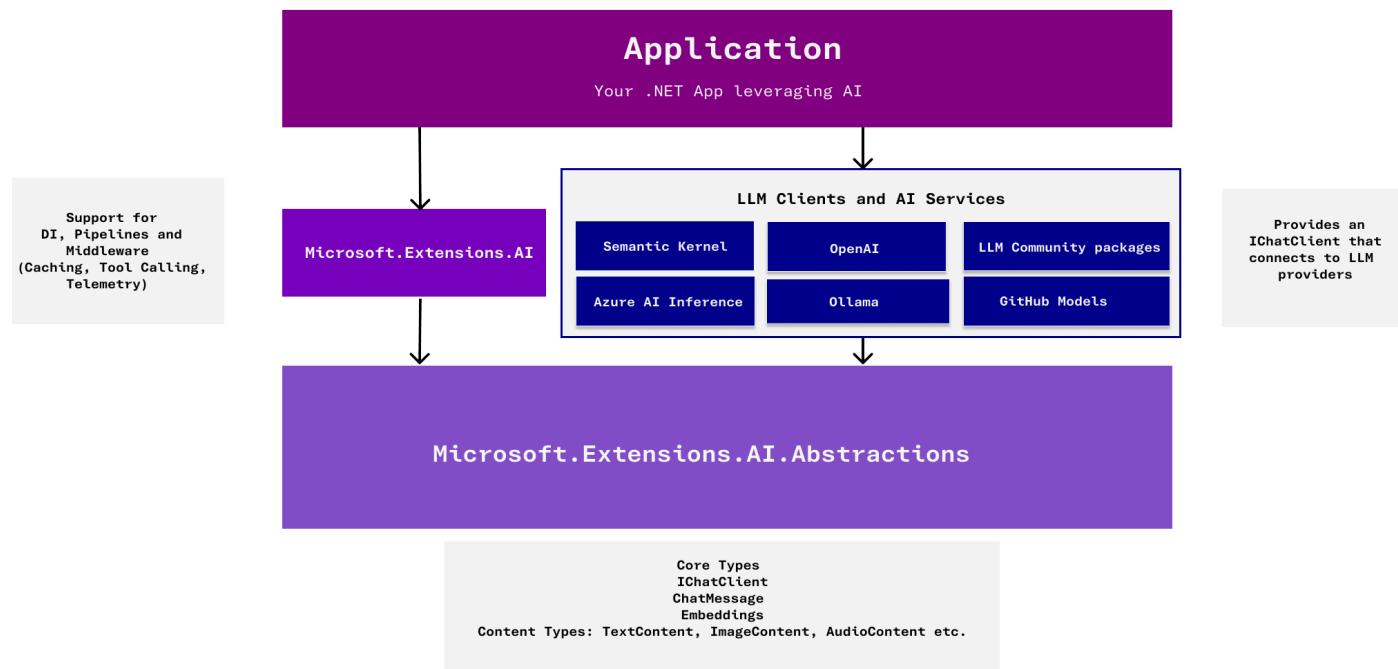
Reference `IRecordCollection<TRecord>` stays consistent across all connectors



Additional Packages & Setup

Microsoft.Extensions.VectorData.Abstractions

Microsoft.SemanticKernel.Connectors.InMemory



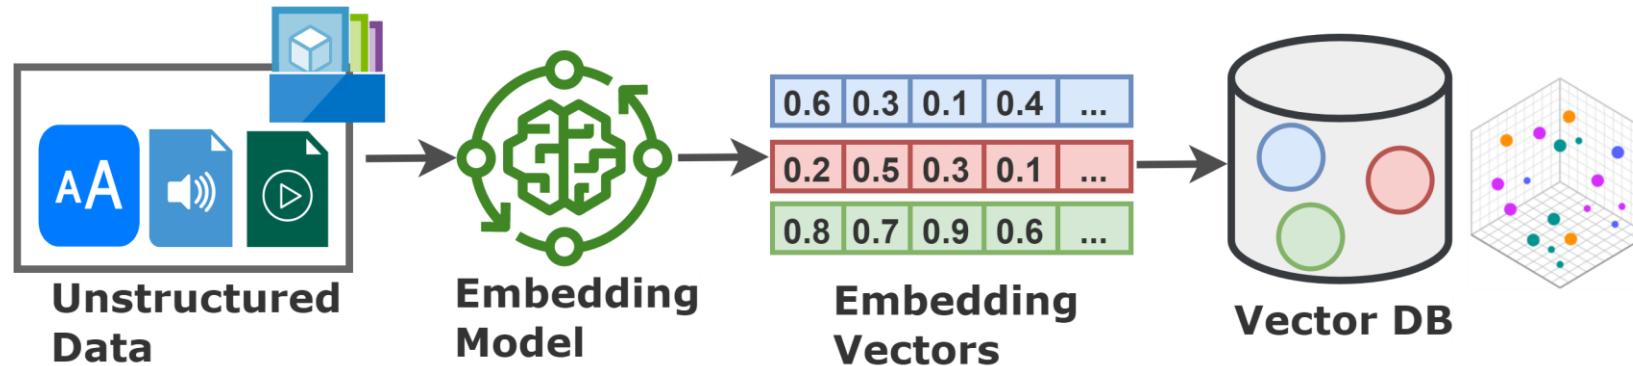
Product Vector Search

Models layer: The newly created `ProductVector` to store vector embeddings

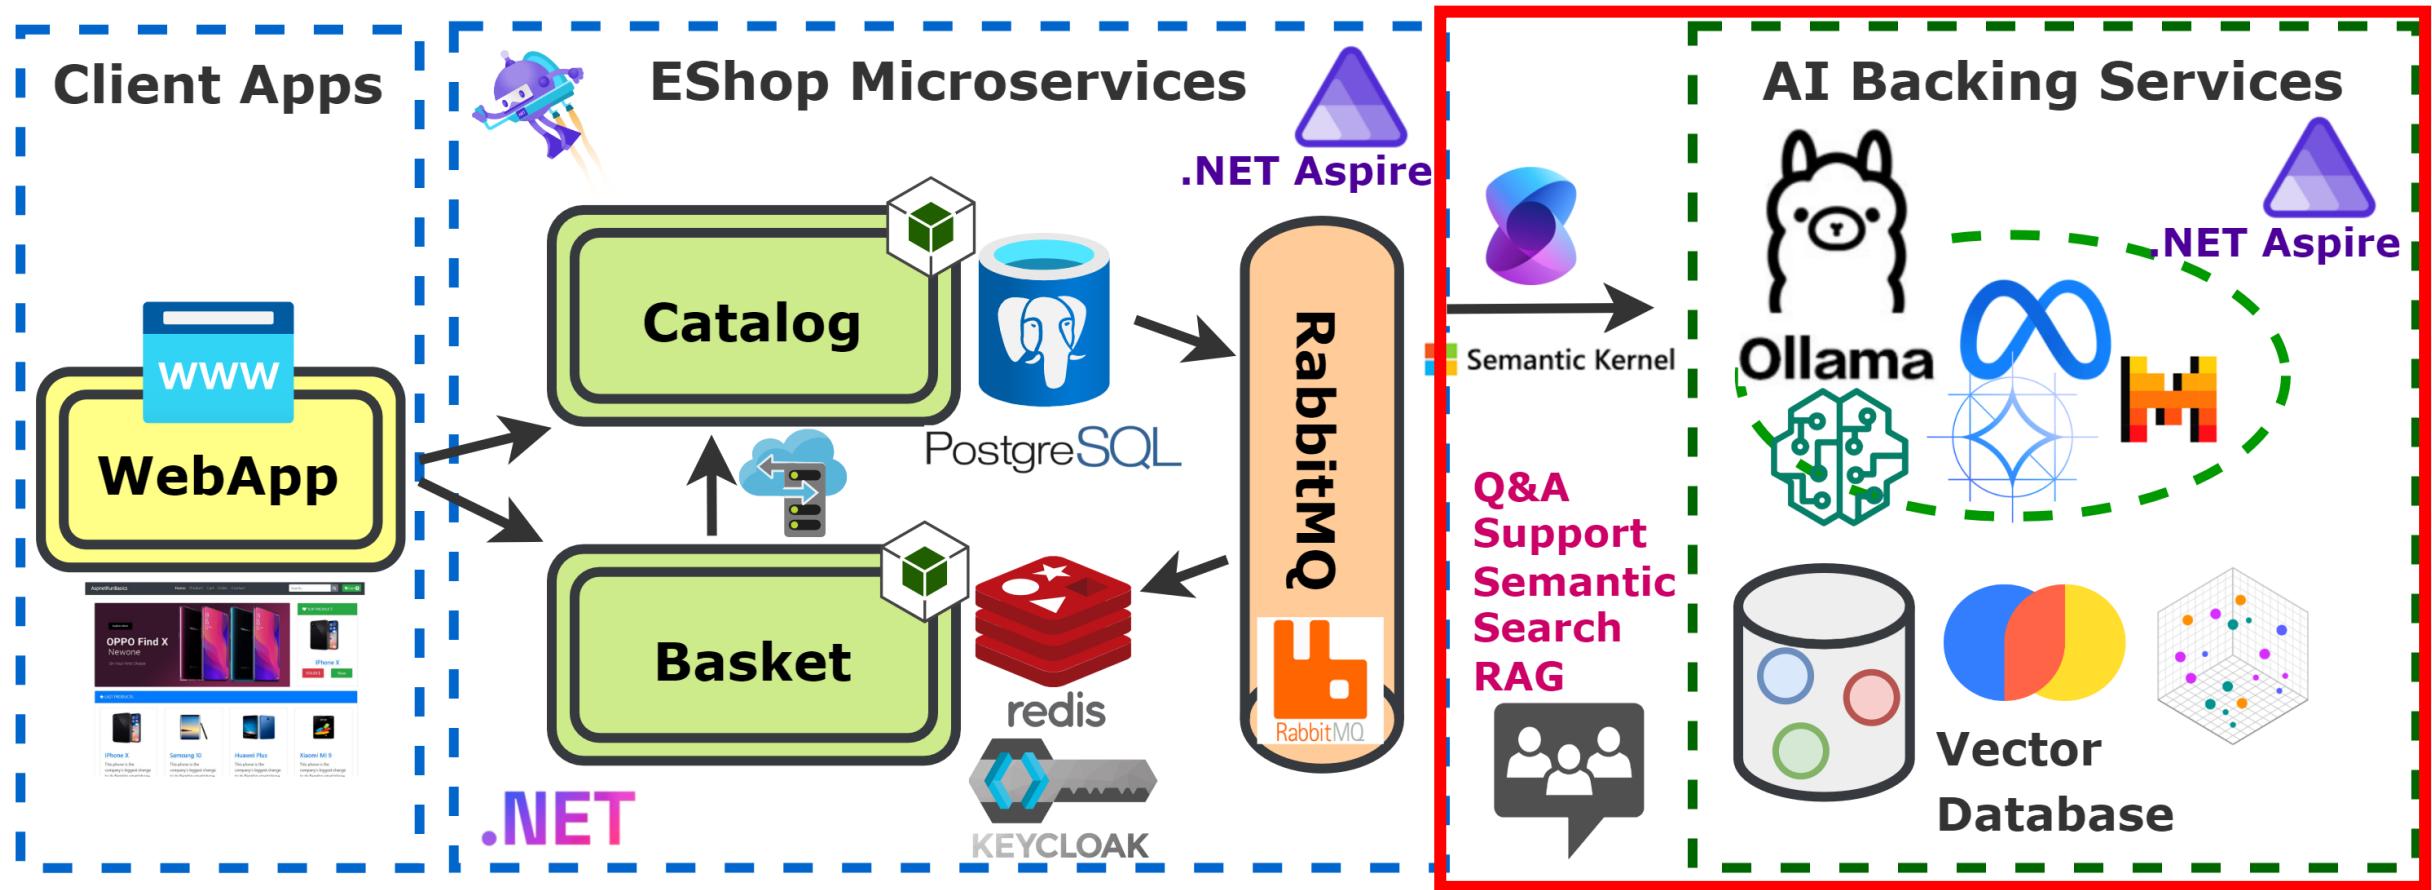
Services layer: Will handle the actual searching—keyword or semantic

Endpoints layer: Minimal API routes in the Catalog microservice

Frontend: A Blazor page with a checkbox toggling AI or normal search



.NET GenAI with Microsoft.Extensions.AI for Chat AI and Semantic Search



Develop Blazor Web Application

Products

Here are some of our amazing outdoor products that you can purchase.

Image	Name	Description
	Hiking Poles	Ideal for camping and hiking trips
	Outdoor Rain Jacket	This product will keep you warm and dry in all weathers
	Survival Kit	A must-have for any outdoor adventurer
	Outdoor Backpack	This backpack is perfect for carrying all your outdoor essentials

Thanks

Thank you so much for being with me on this journey.

Reviews and feedback is really encourage to me for pushing forward to create new courses like this.

Mehmet Ozkaya

